

DATA STRUCTURES PART II


Sid Chi-Kin Chau

[Lecture 3]





Recap from Previous Lecture

- Binary search tree
 - At most two children for each node
 - Left child node is smaller than its parent node
 - Right child node is greater than its parent node
 - Can support dynamic set operations
 - Search, Minimum, Maximum, Predecessor, Successor, Insert, and Delete




Worse-case Scenario

- Insert 6, 5, 4, 3, 2, 1 in an empty tree in order
- Depth of tree linear increases



Balanced Tree

- Balanced search tree
 - Belong to binary search tree
 - But with a height of $O(\log(n))$ guaranteed for n items
 - Height h = maximum number of **edges** from the root to a leaf
- Examples
 - Red-black tree
 - AVL tree
 - B-tree




Goals of This Lecture

- Red-black tree
 - Balanced tree by coloring the nodes
- AVL tree is similar to red-black tree
- B-tree
 - Balanced tree by keeping all leaves at same level
- Something else to learn
 - Applications of different data structures



Red-Black Tree

- Close to balanced tree
- Tree structure requires an extra one-bit color field in each node:
either **red** or **black**
 - Color is used to maintain balance in an approximate sense
- Node:
 - Key
 - Color
 - Left
 - Right
 - Parent
- Note: Leave nodes are NULL nodes (with no child)




Red-Black Properties

- The red-black properties:
 1. Every node is either **red** or **black**
 2. (a) Root and (b) leaves (NULL node) are **black**
 - Note: this means every “real” node has 2 children
 3. If a node is **red**, both children are **black**
 - Note: cannot have 2 consecutive reds on a path
 4. Every path from node to descendent leaf contains the same number of **black** nodes
- “Black-height” is the number of **black** nodes on a path to a leaf
 - Red-black property #4 ensures the black-height is independent of any leaf




Example

- The red-black properties:
 1. Every node is either **red** or **black**
 2. Root and leaves (NULL node) are **black**
 3. If a node is **red**, both children are **black**
 4. Every path from node to descendent leaf contains the same number of **black** nodes




Example

- The red-black properties:
 1. Every node is either **red** or **black**
 2. Root and leaves (NULL node) are **black**
 3. If a node is **red**, both children are **black**
 4. Every path from node to descendent leaf contains the same number of **black** nodes
- Add a new node with $x.\text{key} = 8$




Example

- The red-black properties:
 1. Every node is either **red** or **black**
 2. Root and leaves (NULL node) are **black**
 3. If a node is **red**, both children are **black**
 4. Every path from node to descendent leaf contains the same number of **black** nodes
- Add a new node with $x.\text{key} = 8$
 - Insert x as in binary search tree
 - Set $x.\text{color} \leftarrow$ **red**, satisfying all red-black properties



Example

- The red-black properties:
 1. Every node is either **red** or **black**
 2. Root and leaves (NULL node) are **black**
 3. If a node is **red**, both children are **black**
 4. Every path from node to descendent leaf contains the same number of **black** nodes
- Add a new node with $y.key = 11$
 - Insert y as in binary search tree
 - Set $y.color \leftarrow ???$
 - Cannot satisfy all red-black properties




Red-Black Tree Insertion

- Insertion: the basic idea
 - Insert the new node as in binary search tree
 - Color the new node **red**
 - Only red-black property #3 may be violated
 - It will be violated if the new node's parent is **red**
 - If so, move violation up the tree until a place is found where it can be fixed




Example

- Insertion:
 - Insert the new node as in binary search tree
 - Color the new node **red**
- Add a new node with $y.key = 11$
 - Insert y as in binary search tree
 - Set $y.color \leftarrow ???$




Example

- Insertion:
 - Insert the new node as in binary search tree
 - Color the new node **red**
- Add a new node with $y.key = 11$
 - Insert y as in binary search tree
 - Set $y.color \leftarrow$ **red**
 - Recolor the tree




Example

- Insertion:
 - Insert the new node as in binary search tree
 - Color the new node **red**
- Add a new node with z.key = 10
 - Insert z as in binary search tree
 - Set z.color \leftarrow **red**




Example

- Insertion:
 - Insert the new node as in binary search tree
 - Color the new node **red**
- Add a new node with z.key = 10
 - Insert z as in binary search tree
 - Set z.color \leftarrow ??
 - Tree is too unbalanced
 - Need to restructure
 - See how to deal with it next ...




Tree Rotation

- Our basic operation for changing tree structure is called rotation
 - Rotation preserves inorder key ordering
 - How would tree rotation actually work?



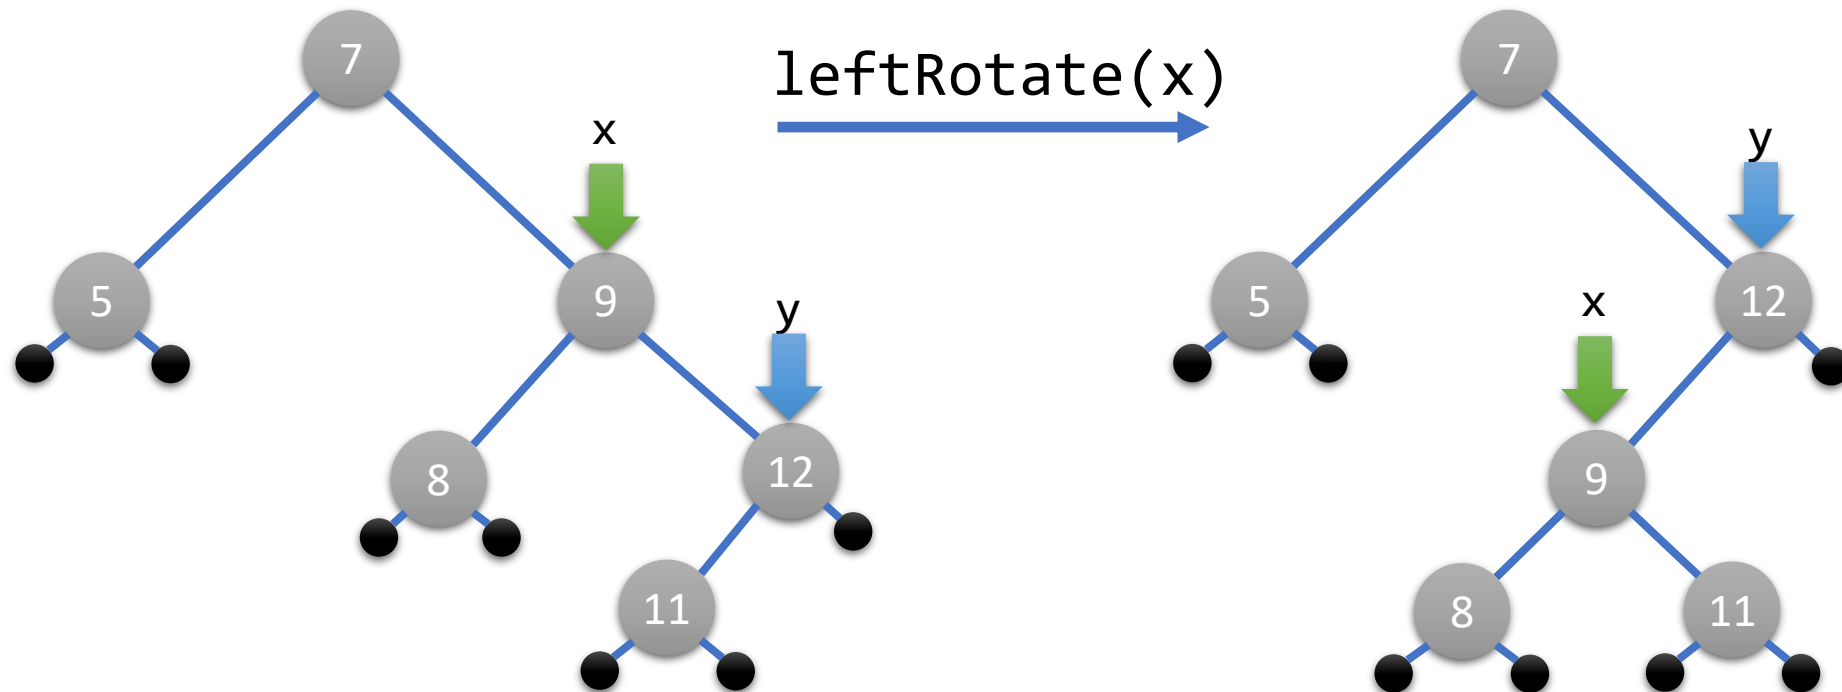
Tree Rotation

- Need a couple of pointer manipulations
 - x keeps its left child
 - y keeps its right child
 - x's right child becomes y's left child
 - x's and y's parents change







Example

- Left Rotate at node 9



Three Possible Cases




- Suppose that the parent of the new node is the left child of its grandparent
 - “Uncle” is the other child of the grand parent
- Cases:
 1. New node z’s uncle is **red**
 2. New node z is right child and its uncle is **black**
 3. New node z is left child and its uncle is **black**



Mirrored Possible Cases

- Cases:


1. New node z's (right) uncle is **red**
2. New node z is right child and its uncle is **black**
3. New node z is left child and its uncle is **black**
4. New node z's (left) uncle is **red**
5. New node z is left child and its uncle is **black**
6. New node z is right child and its uncle is **black**



Red-Black Tree Insertion: Case 1

- Case 1: Uncle is red
 - Solution: Recoloring
 - Assume all subtree Δ 's have equal black-height
 - Recolor parent, uncle y and grandparent to satisfy the property that all paths have equal black height
 - Note: If c is root, then reset the color to be black at the end of insertion


```
// Case 1  
If (y.color = RED) Then  
  z.parent.color ← BLACK  
  y.color ← BLACK  
  z.parent.parent.color ← RED
```



Red-Black Tree Insertion: Case 1

- Case 1: Uncle is red
 - Solution: Recoloring
 - New node z can be either left or right child


```
// Case 1  
If (y.color = RED) Then  
    z.parent.color ← BLACK  
    y.color ← BLACK  
    z.parent.parent.color ← RED
```



Red-Black Tree Insertion: Case 2

- Case 2: Uncle is black
 - New node z is right child
 - Solution: Transformation
 - Transform to case 3 via a left-rotation


```
// Case 2
If (z = z.parent.right) Then
    z ← z.parent
    leftRotate(z)
// continue with Case 3
```



Red-Black Tree Insertion: Case 3


- Case 3: Uncle is black
 - New node z is left child
 - Solution: Rotation
 - Recolor and right rotate

```
// continue with Case 3  
z.parent.color ← BLACK  
z.parent.parent.color ← RED  
rightRotate(z.parent.parent)
```




Example

- Let us return to the example
- Add a new node with z.key = 10
 - Which Case is it?
 - It is Case 3
 - Recolor and rotate right



Example

- Let us return to the example
- Add a new node with z.key = 10
 - Which Case is it?
 - It is Case 3
 - Recolor and rotate right



Red-Black Tree Insertion: Cases 4-6

- Cases:
 1. New node z's (right) uncle is **red**
 2. New node z is right child and its uncle is **black**
 3. New node z is left child and its uncle is **black**
 4. New node z's (left) uncle is **red**
 5. New node z is left child and its uncle is **black**
 6. New node z is right child and its uncle is **black**
- Cases 1-3 hold if z's parent is a left child of its grandparent
- If z's parent is a right child of its grandparent, Cases 4-6 are symmetric to Cases 1-3
 - Swap left for right, and vice versa



Red-Black Tree Insertion

RB-Insert[T,z]

// Call BST-Insert
BST-Insert[T,z]

z.left \leftarrow null; z.right \leftarrow null
z.color \leftarrow RED

// Call RB-InsertFixup to
// recolor and restructure tree
RB-InsertFixup[T,z]

RB-InsertFixup[T,z]

```
While z.parent.color = RED Do
  If z.parent = z.parent.parent.left Then
    y  $\leftarrow$  z.parent.parent.right
    // Case 1
    If (y.color = RED) Then
      z.parent.color  $\leftarrow$  BLACK
      y.color  $\leftarrow$  BLACK
      z.parent.parent.color  $\leftarrow$  RED
      // Continue While loop
      z  $\leftarrow$  z.parent.parent
    Else
      If (z = z.parent.right) Then
        // Case 2
        z  $\leftarrow$  z.parent
        leftRotate(z)
        // continue with case 3
        z.parent.color  $\leftarrow$  BLACK
        z.parent.parent.color  $\leftarrow$  RED
        rightRotate(z.parent.parent)
      Else
        // Cases 4-6 with "right" & "left" exchanged
        T.root.color  $\leftarrow$  BLACK
```


Exercise



- Consider the following operations to red-black tree
 - Insert 12, 43, 34, 11, 44, 1
- What is the tree height of the final tree?
- How many **red** nodes are in the final tree?




Red-Black Tree Deletion

- x is the to-be-deleted node (disregarding the colors):


- Cases:

1. x has only one child

- Just remove x and replace x by its child
- If the red-black property #3 is now broken, recolor y in black keeping the black-height, since x was definitely **black** (as its parent is **red**)



2. x has two children




Red-Black Tree Deletion

NOT in exam!

- Cases:

2. x has two children, x's successor is z

- Replace the successor z by left child w
- Remove x and replace x by its successor z
- But if x is **red** and z is **black**, then there will be an extra black node in z's new left child




- See Chapter 14 of “Introduction to Algorithms” (by Cormen, Leiserson and Rivest)



Black-height

- The red-black properties:

1. Every node is either **red** or **black**
2. Root and leaves (NULL node) are **black**
3. If a node is **red**, both children are **black**
4. Every path from node to descendent leaf contains the same number of **black** nodes




- Let black-height ($BH(x)$) be the number of black nodes from x to a leaf
 - $BH(x)$ is independent of any leaf node
- Note: No adjacent nodes on any path on red-black tree can be both **red**
 - Otherwise, it violates red-black property #3



Quick Fact

- Note that Black-height \leq tree height
- Let the “internal” nodes be non-leaf nodes
- Consider the extreme case:
 - All nodes are **black** (ignoring all **red** nodes)
 - The tree is perfectly balanced
 - Check: satisfy all red-black properties
 - Total number of **black** internal nodes is
$$2^0 + 2^1 + 2^2 + \dots + 2^{BH(\text{root})-1} = 2^{BH(\text{root})} - 1$$
- Hence,
 - Total number of internal nodes $\geq 2^{BH(\text{root})} - 1$
 - Since adding **red** internal nodes back will not affect the black-height



Red-Black Tree: Worst-case Running Time

- What is the minimum black-height of the root with height h ?
 - Answer: $BH(\text{root}) \geq h/2$
 - Because the number of **black** nodes from the root to a leaf is at least $h/2$
 - Otherwise, two **red** nodes will be adjacent to each other

Theorem: A red-black tree with n internal nodes has height $h \leq 2 \log(n + 1)$

Proof: Note that $n \geq 2^{BH(\text{root})} - 1$

Since $BH(\text{root}) \geq h/2$, we obtain

$$n \geq 2^{BH(\text{root})} - 1 \geq 2^{h/2} - 1$$

$$\Rightarrow n \geq 2^{h/2} - 1$$

Then, $\log(n + 1) \geq h/2 \Rightarrow h \leq 2 \log(n + 1)$

Therefore, $h = O(\log(n))$

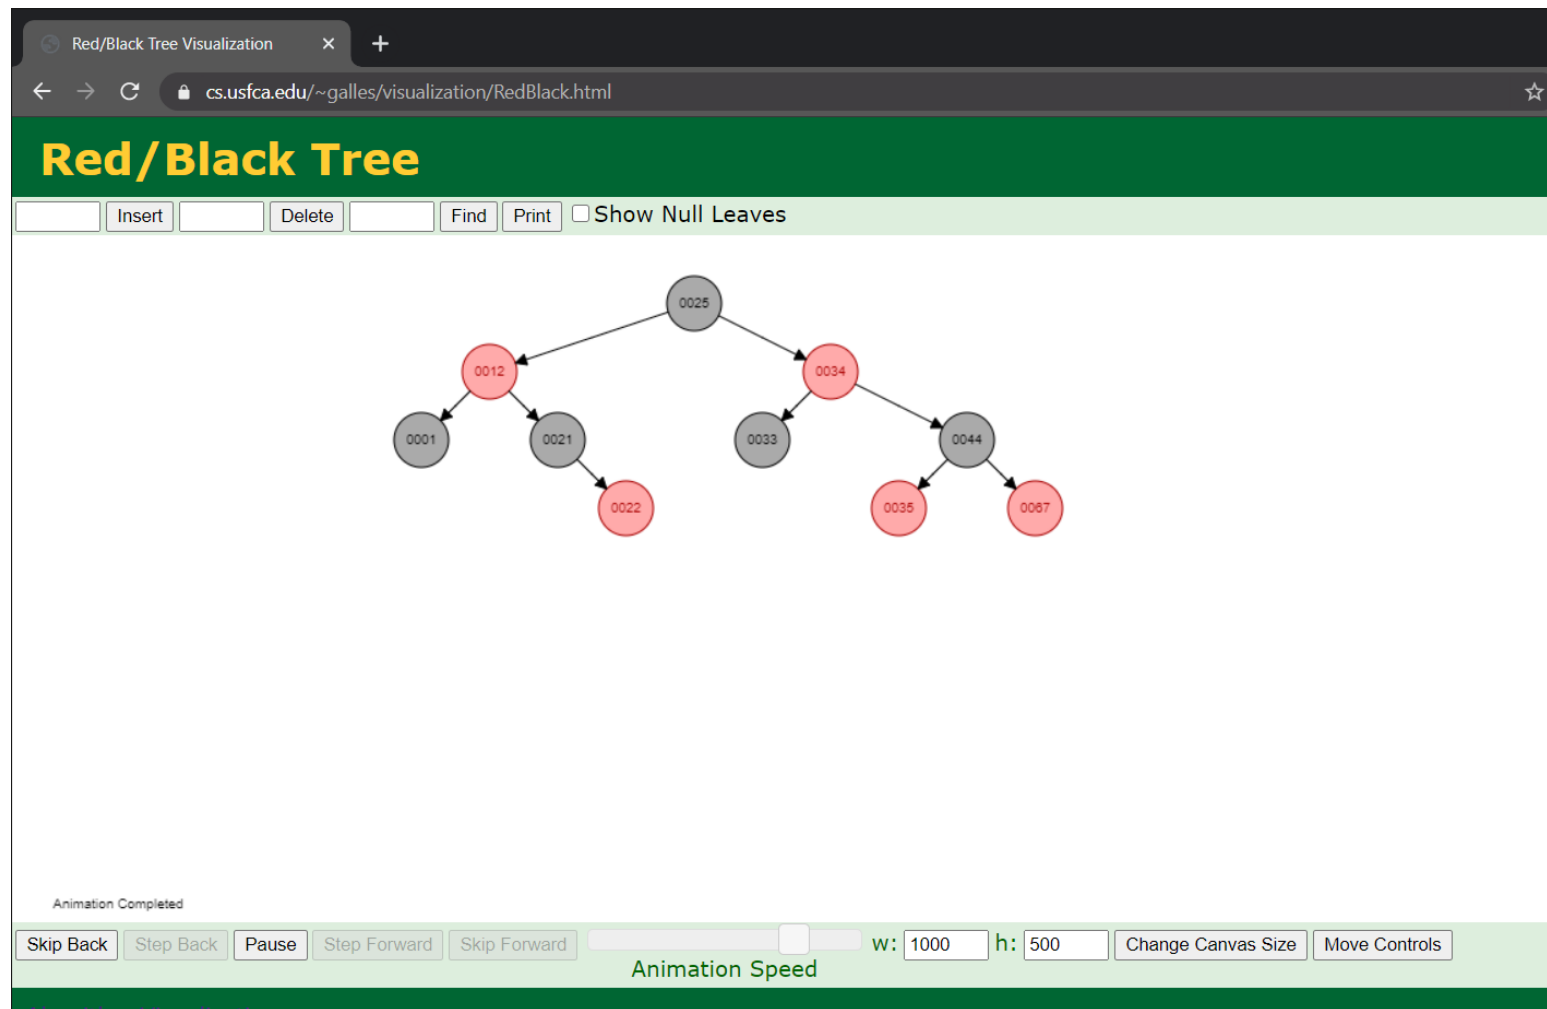


Red-Black Tree: Worst-case Running Time

- So we have shown that a red-black tree has $O(\log(n))$ height
- Corollary: These operations take $O(\log(n))$ time
 - Minimum, Maximum
 - Successor, Predecessor
 - Search
- Insert and Delete
 - Will also take $O(\log(n))$ time
 - But will need to take special care since they modify tree structure



Demo



<https://www.cs.usfca.edu/~galles/visualization/RedBlack.html>



Implement a Red-Black Tree


```
public class RBTree <T extends Comparable<T>> {  
  
    private void insert(Node<T> x) { ... }  
  
    public void rotateLeft(Node<T> x) { ... }  
  
    public void rotateRight(Node<T> x) { ... }  
  
    public Node<T> search(T key) {  
    }  
}
```

```
public class Node<T> {  
  
    Colour colour; // Node colour  
    T key;          // Node value  
    Node<T> parent; // Parent node  
    Node<T> left, right; // Child nodes  
}
```



Summary

- Balanced search tree
 - Red-black tree
 - Red-black properties
 - Insertion
 - Black-height



Reference

- Visualizations
 - <https://www.cs.usfca.edu/~galles/visualization/RedBlack.html>
- Reference:
 - Chapter 14 in Introduction to Algorithms (by Thomas H. Cormen, Charles E. Leiserson, and Ronald L. Rivest)

