

DATA STRUCTURES PART III


Sid Chi-Kin Chau

[Lecture 4]




Recap from Previous Lecture

- Balanced search tree
 - Belong to binary search tree
 - But with a height of $O(\log(n))$ guaranteed for n items
 - Height h = maximum number of **edges** from the root to a leaf
- Examples
 - Red-black tree
 - AVL tree
 - B-tree



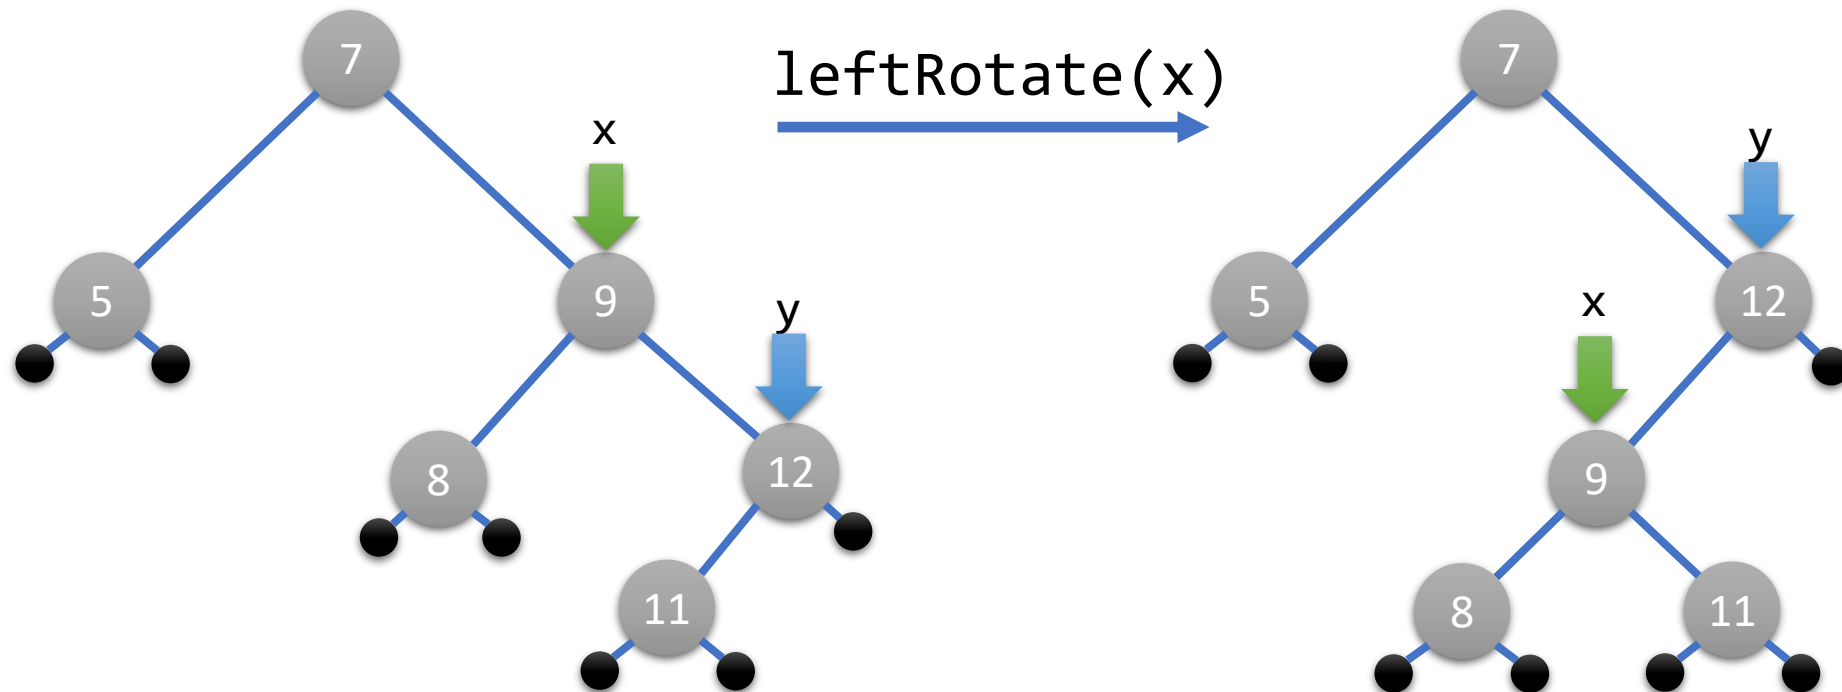
Tree Rotation

- Our basic operation for changing tree structure is called rotation
 - Rotation preserves inorder key ordering
 - How would tree rotation actually work?



Example

- Left Rotate at node 9




AVL Tree

- AVL tree is binary search tree that has a height difference between with left and right subtrees of a node is at most 1
 - In BST, left subtree keys are less than the root and right subtree keys are greater than the root
 - Each node has a height of the subtree rooted at that node
- Balance:
 - Define the balance of a node n
 - $b(n) = \text{Height with only Right Subtree} - \text{Height with only Left Subtree}$
 - Legal values for AVL tree: $b(n)$ in $\{-1, 0, 1\}$
 - Otherwise, $b(n) > 1$ or $b(n) < -1$, then it is an unbalanced AVL tree





AVL Tree Example




AVL Tree Insertion

- Adding a new node cannot make an unbalanced parent
 - Assume the parent p 's balance in $b(p)$ in $\{-1, 0, 1\}$ before insertion
 - The parent will have
 - $b(p) = 0$ if already has a child,
 - $b(p) = 1$ if a new left child, or $b(p) = -1$ if a new right child
 - Otherwise it would not be our parent or the parent would have been unbalanced already



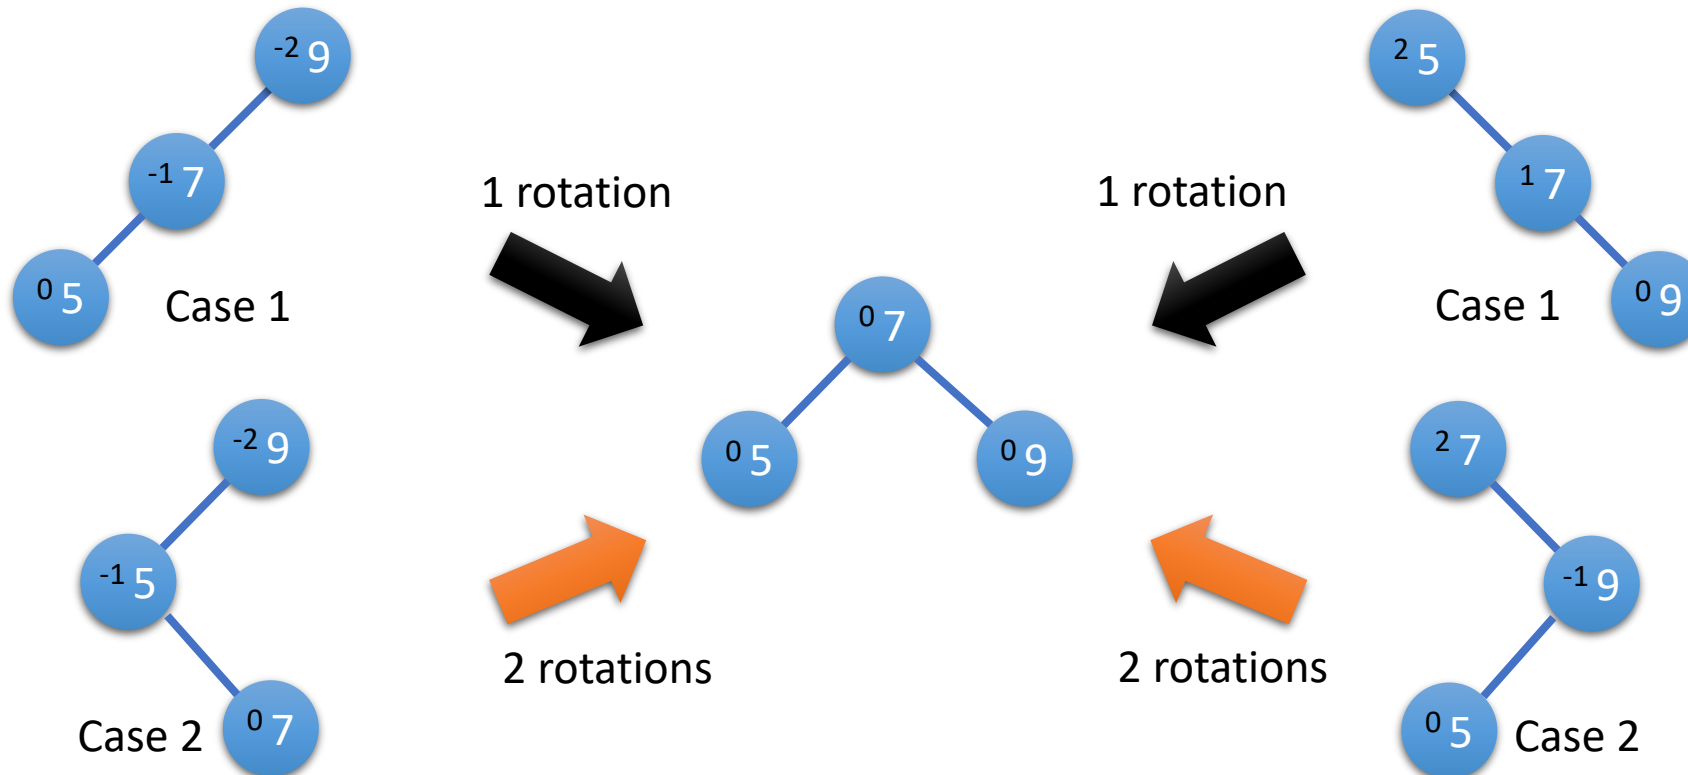
AVL Tree Insertion

- Even if the parent is not unbalanced, it is possible that the grandparent becomes unbalanced
- Hence, we need a way to re-balance the grandparent



AVL Tree Insertion

- Rotations are required to balance AVL tree
 - Case 1 (left-left or right-right): Requires 1 rotation
 - Case 2 (left-right or right-left): Requires 2 rotations (first convert to Case 1)



AVL Tree Insertion

- Insert[n]
 - If empty tree, then
 - Set n as root, $b(n) = 0$. Return
 - Else insert n (by walking the tree to a leaf p and inserting the new node as its child), set $b(n) = 0$, and look at its parent p
 - If $b(p)$ was -1, then $b(p) = 0$. Return
 - If $b(p)$ was +1, then $b(p) = 0$. Return
 - If $b(p)$ was 0, then update $b(p)$ and call Insert-fix[p, n]



AVL Tree Insertion

- Basic idea:
 - Work up ancestor chain updating balances of the ancestor chain or fix a node that is unbalanced
- Insert-fix[p, n]
 - Precondition: p and n are balanced: $b(p), b(n) \in \{-1, 0, 1\}$
 - Postcondition: g, p, and n are balanced: $b(g), b(p), b(n) \in \{-1, 0, 1\}$
 - If p is null or p.parent is null, return
 - Let $g \leftarrow p.parent$

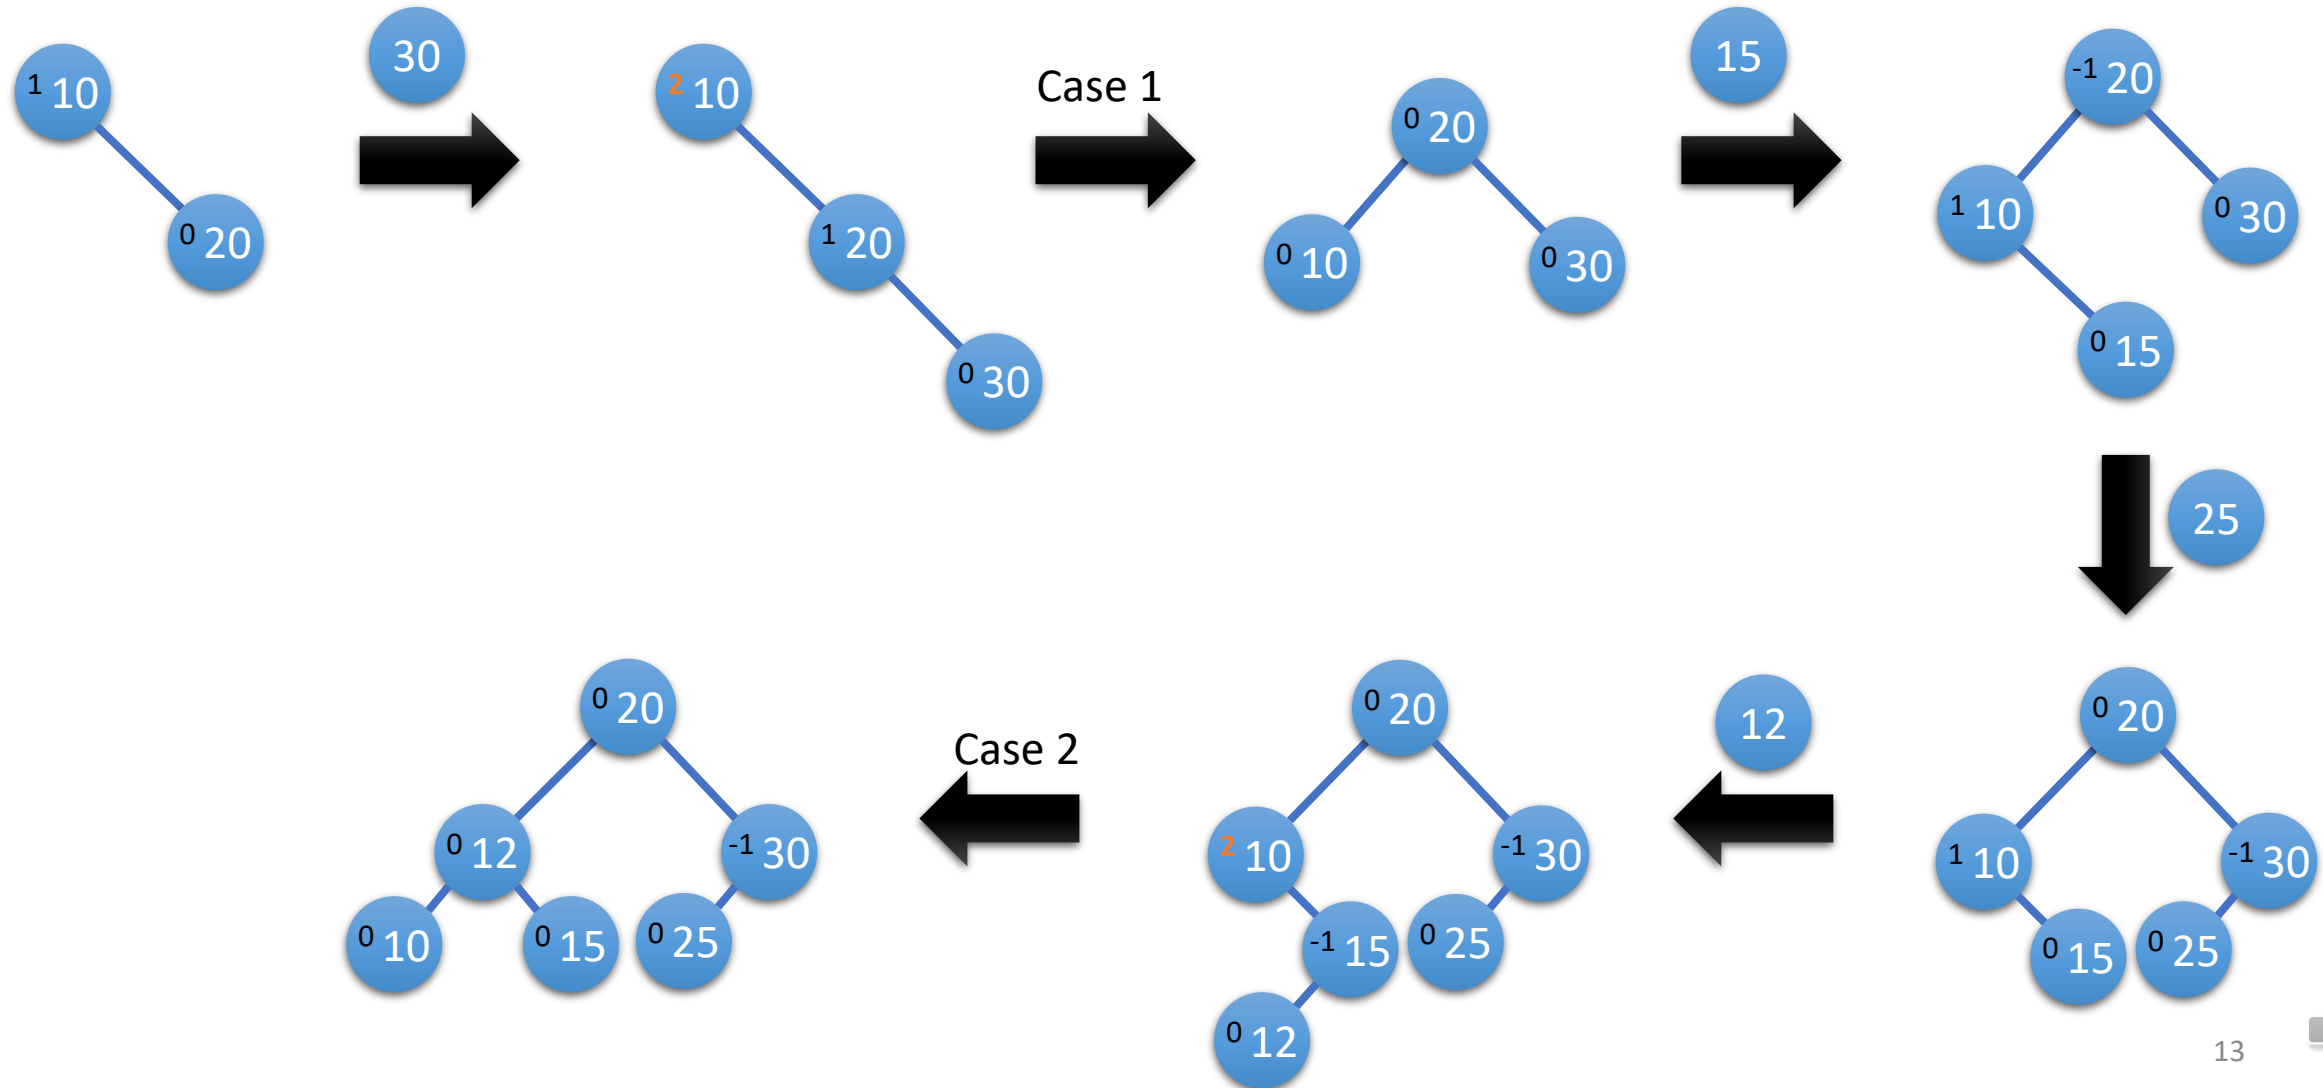


AVL Tree Insertion


- `Insert-fix[p, n]`
 - Assume p is left child of g // If p is right child, swap left/right, +/-
 - Update $b(g)$ // Update g 's balance to new accurate value for now
- Case a: if $b(g) = 0$, return
- Case b: if $b(g) = -1$, `Insert-fix[g, p]` // Fix recursively
- Case c: if $b(g) = -2$
 - If Case 1, then `rotateRight(g)`;
 - If Case 2, then `rotateLeft(p)`; `rotateRight(g)`;



Example



Example

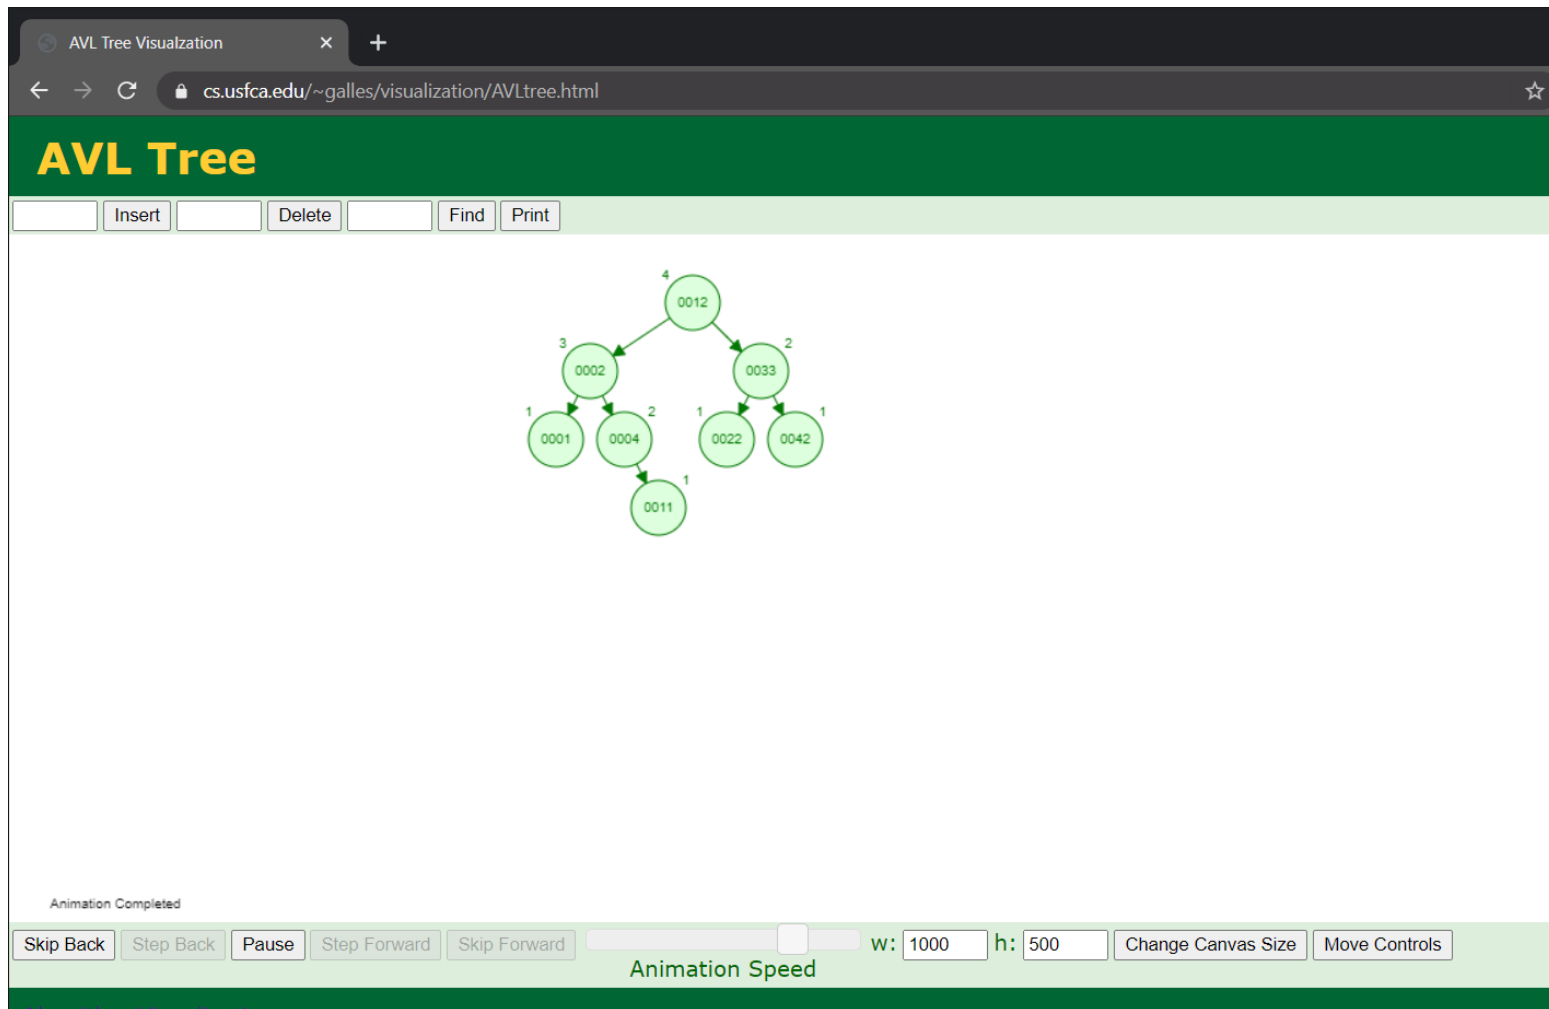


AVL Tree Deletion

- Deletion operations may also require rebalancing via rotations
- The key idea is to update the balance of the nodes on the ancestor pathway
- If an ancestor gets out of balance then perform rotations to rebalance
 - Unlike insertion, rotations during deletion does not mean you are done, but need to continue recursively to the root
- More cases in AVL tree deletion



Demo



<https://www.cs.usfca.edu/~galles/visualization/AVLtree.html>




B-Tree


- Generalization of binary search trees
 - Not binary trees (i.e., more than two children per node)
 - The “B” stands for Bayer (the inventor)
- Designed for searching data stored on block-oriented devices
 - Map tree nodes into blocks
- Balanced tree
 - By keeping all leaves at the same level
- Each node can hold multiple items and have multiple children
- B-tree grows and shrinks from the root
 - Binary search tree only grows downward and also shrink from downward



B-Tree Structure



Example



B-Tree Properties


- T is a B-tree of degree $m (\geq 3)$ if it satisfies the following properties:
 1. All leaves of T appear on the same level of T
 2. Every node of T has at most m children
 3. Every node of T, except for the root and the leaves, has at least $m/2$ children
 4. The root of T is either a leaf or has at least two children
 5. An internal node with k children stores $k-1$ keys, and a leaf stores between $m/2-1$ and $m-1$ keys. The keys $v.\text{key}_i$, where $1 \leq i \leq k-1$, of a node v are maintained in sorted order, $v.\text{key}_1 \leq \dots \leq v.\text{key}_{k-1}$
 6. If v is an internal node with k children, the $k-1$ keys of v separate the range of keys stored in the subtrees rooted at the children of v . If x_i is any key stored in the subtree rooted at the i -th child, the following holds:

$$x_1 \leq v.\text{key}_1 \leq x_2 \leq \dots \leq v.\text{key}_{k-2} \leq x_{k-1} \leq v.\text{key}_{k-1} \leq x_k$$



Searching B-Tree

- Determine the range at each node
- Follow the pointer of the corresponding range as in binary search tree



Organization of B-Tree


- Each non-leave node can have a variable number of children
- Must all be in a specific key range
- Number of child nodes typically vary between $m/2$ and m (with $m-1$ keys)
 - Split nodes that would otherwise have contained $m + 1$ children (after insertion)
 - Merge nodes that contain less than $m/2$ children (after deletion)




B-Tree Insertion

- Consider B-tree where each node can have up to 2 keys (3 children)


- Add 1:




- Add 2:




- Add 3:



Split




- Add 4:




B-Tree Insertion


- Add 5:



- Add 6:

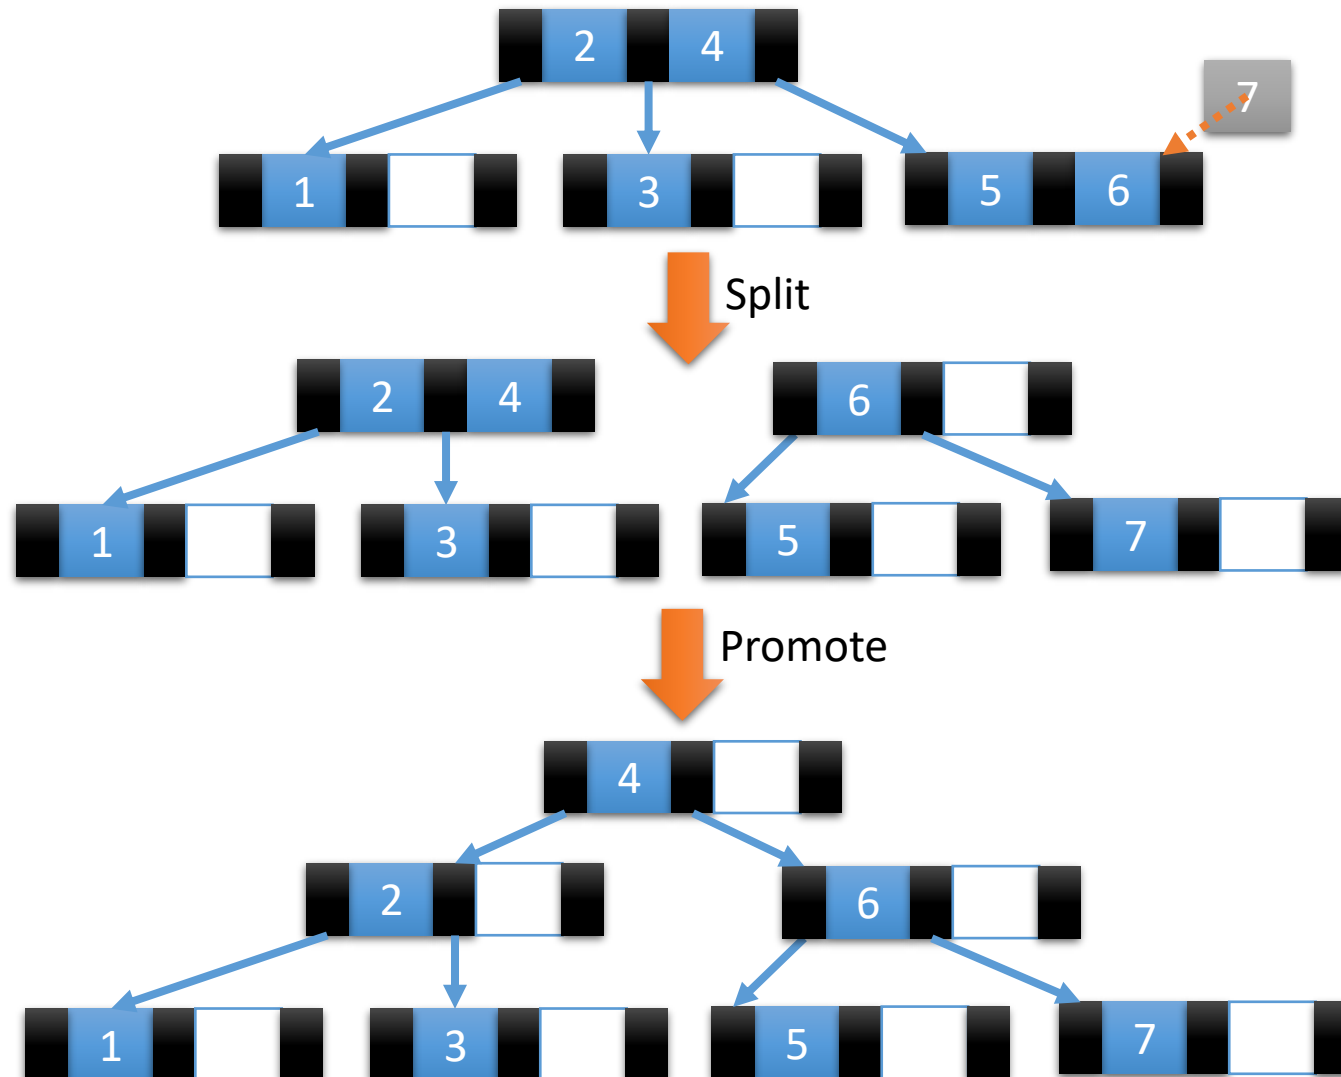


- Add 7:



B-Tree Insertion


- Add 7:



B-Tree Insertion Basic Operations


- Split:

- When trying to add to a full node
- Split node at the central value



- Promote:

- Must insert root of split node higher up
- May require a new split



B-Tree Insertion (Basic Idea)

- Find target leaf L
- If L has less than $m - 2$ entries :
 - Add the entry
- Else
 - Repeat
 - Allocate new leaf L'
 - Pick the $m/2$ highest keys of L and move them to L'
 - Insert highest key of L and corresponding address leaf into the parent
 - If the parent is full
 - Split it and add the middle key to its parent node
 - Until a parent is found that is not full



Demo

B-Tree Visualization

cs.usfca.edu/~galles/visualization/BTree.html

B-Trees

Insert Delete Find Print Clear

☒ Max. Degree = 3 ☐ Preemptive Split / Merge (Even max degree only)

☐ Max. Degree = 4

☐ Max. Degree = 5

☐ Max. Degree = 6

☐ Max. Degree = 7

```
graph TD; 0021[0021] --> 0006[0006]; 0021 --> 0023[0023]; 0006 --> 0001[0001]; 0006 --> 0002[0002]; 0006 --> 0009[0009]; 0023 --> 0022[0022]; 0023 --> 0025[0025]; 0023 --> 0042[0042];
```



Summary

- AVL tree
 - Keeping every node's balance in $\{-1, 0, 1\}$
 - Rotations are required during insertion and deletion
 - AVL tree is almost balanced tree
- B-tree
 - B-tree properties
 - Insertion



Reference

- Visualizations
 - <https://www.cs.usfca.edu/~galles/visualization/AVLtree.html>
 - <https://www.cs.usfca.edu/~galles/visualization/BTree.html>

