

BENCHMARKING & PERFORMANCE

Sid Chi-Kin¹ Chau

[Lecture 8]



Benchmarking

- Benchmarking
 - Compare the performance of different algorithms and systems
 - Evaluate practical performance with real-world input data
 - Optimize best practice, improve implementation and plan resource allocation
 - Collect and analyze practical performance data
 - Provide assurance and confidence before practical deployment
- Benchmarking trials
 - Construct a suite of independent trials for which the algorithm is executed
 - Trials are executed and milli/nanosecond-level timings are taken before and after the algorithm is executed
 - Eliminate inconsistent measurements



Benchmarking

- Performance measurements may be different in a different time, even with same code and implementation
 - Computer background processes may affect practical performance
 - Eliminate outlier performance data
- In Java, the system garbage collector may affect the performance
 - The system garbage collector is invoked immediately prior to the trial
 - Call `System.gc()`
 - Although this cannot guarantee that the garbage collector does not execute during the trial, it may reduce the impact



Benchmarking in Java

- Example of benchmarking simple summation

```
public class Trial {  
    public static void main (String[] args) {  
        for (long len = 1000000; len <= 5000000; len += 1000000) {  
            for (int i = 0; i < 30; i++) {  
                System.gc(); //Invoke garbage collector  
                long start = System.currentTimeMillis();  
  
                // Simple summation to be timed  
                long sum = 0;  
                for (int x = 1; x <= len; x++) { sum += x; }  
  
                long end = System.currentTimeMillis();  
                // Output runtime  
                System.out.println("trial:" + len + " runtime:" + (end - start));  
            }  
        }  
    }  
}
```



Benchmarking in Java


- Instead of millisecond-level timers, nanosecond timers could be used
- In Java, invoke `System.nanoTime()`

```
for (long len = 1000000; len <= 5000000; len += 1000000) {  
    for (int i = 0; i < 30; i++) {  
        System.gc(); //Invoke garbage collector  
        long start = System.nanoTime(); //Nanosecond timer  
  
        // Simple summation to be timed  
        long sum = 0;  
        for (int x = 1; x <= len; x++) { sum += x; }  
  
        long end = System.nanoTime();  
        // Output runtime  
        System.out.println("trial:" + len + " runtime:" + (end - start));  
    }  
}
```




Benchmarking Data Structures

- Benchmark binary search tree, red-black tree and AVL tree
 - Consider worst-case (i.e., highly unbalanced tree) and average-case (i.e., random input sequences)
 - Which one of the tree data structures is the best practically?
 - BST is surprisingly efficient. Why?



Aspects of Performance Analysis

- Does your software perform as what you expect?
 - Does it complete fast?
 - Does it work well with more inputs?
 - Does it break?
- Metrics of performance:
 - Latency, throughput, memory size
 - Network bandwidth
 - Concurrency
- Performance analysis
 - Best case
 - Average case
 - Worst case

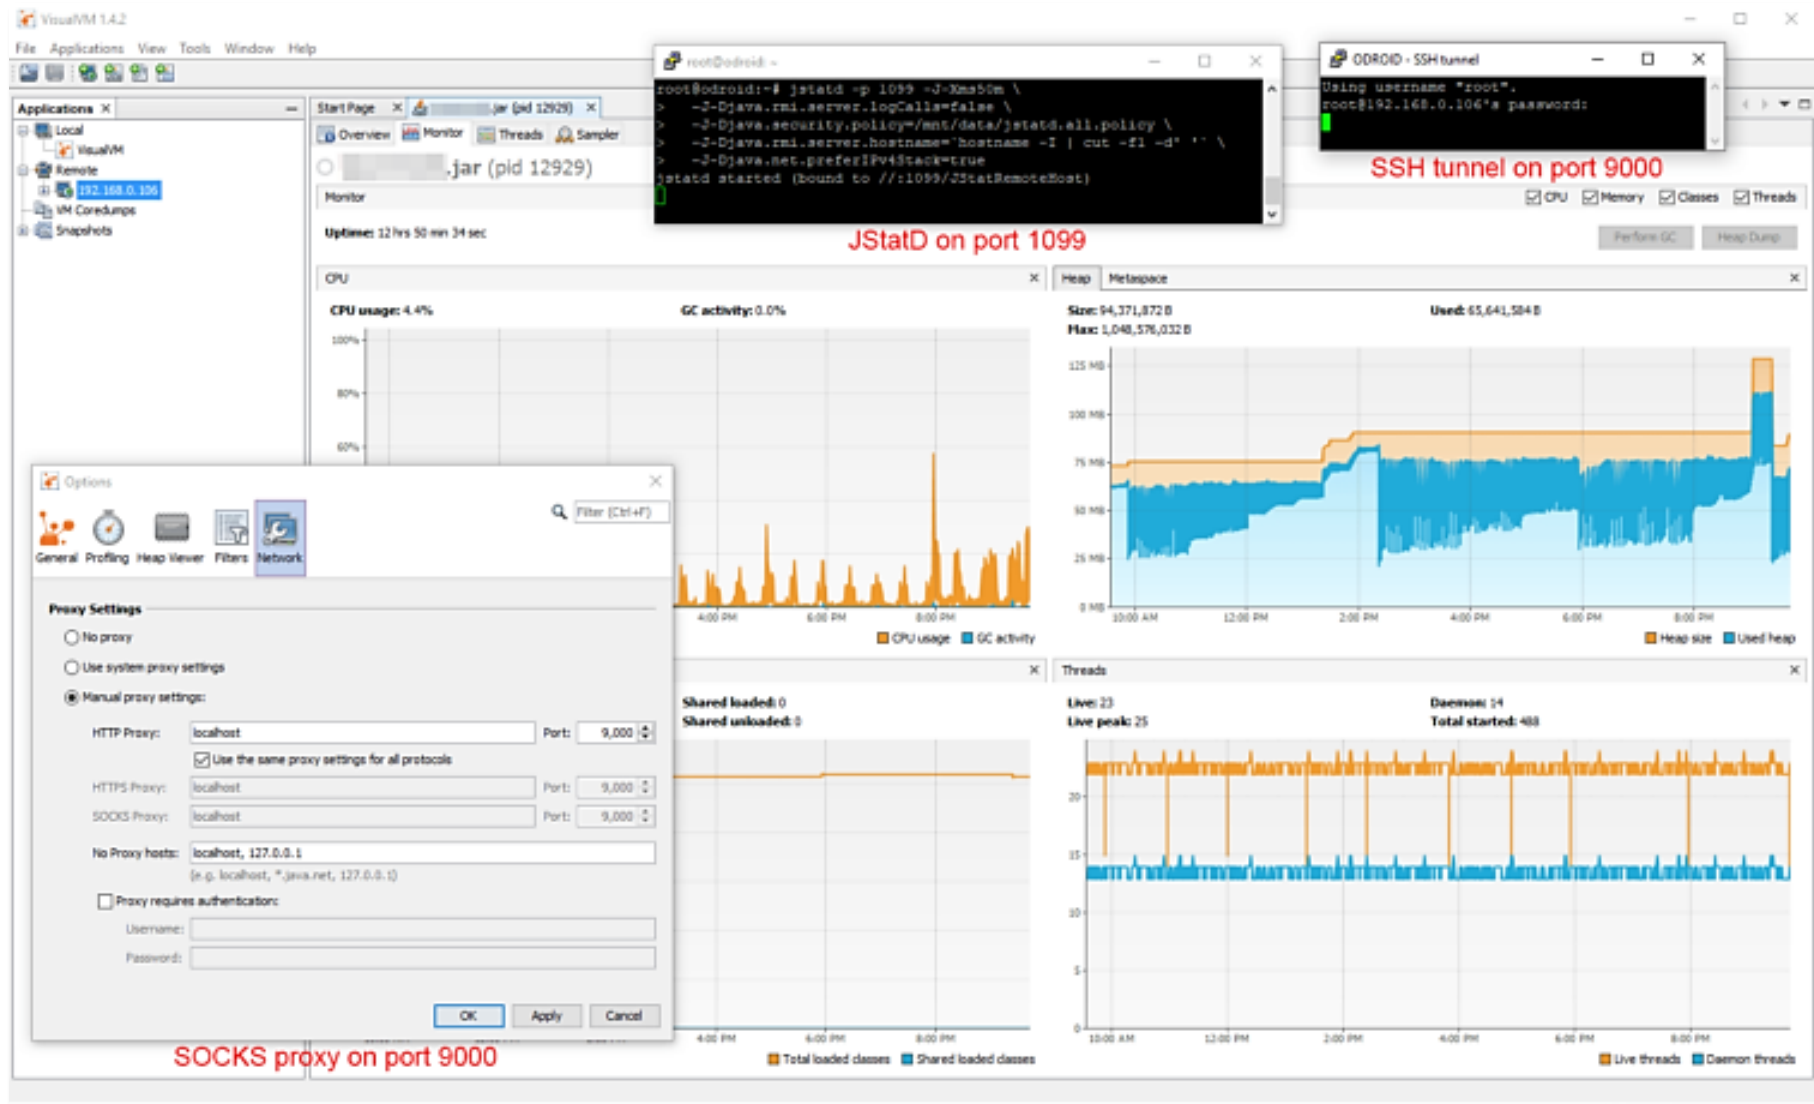


Performance Profiling

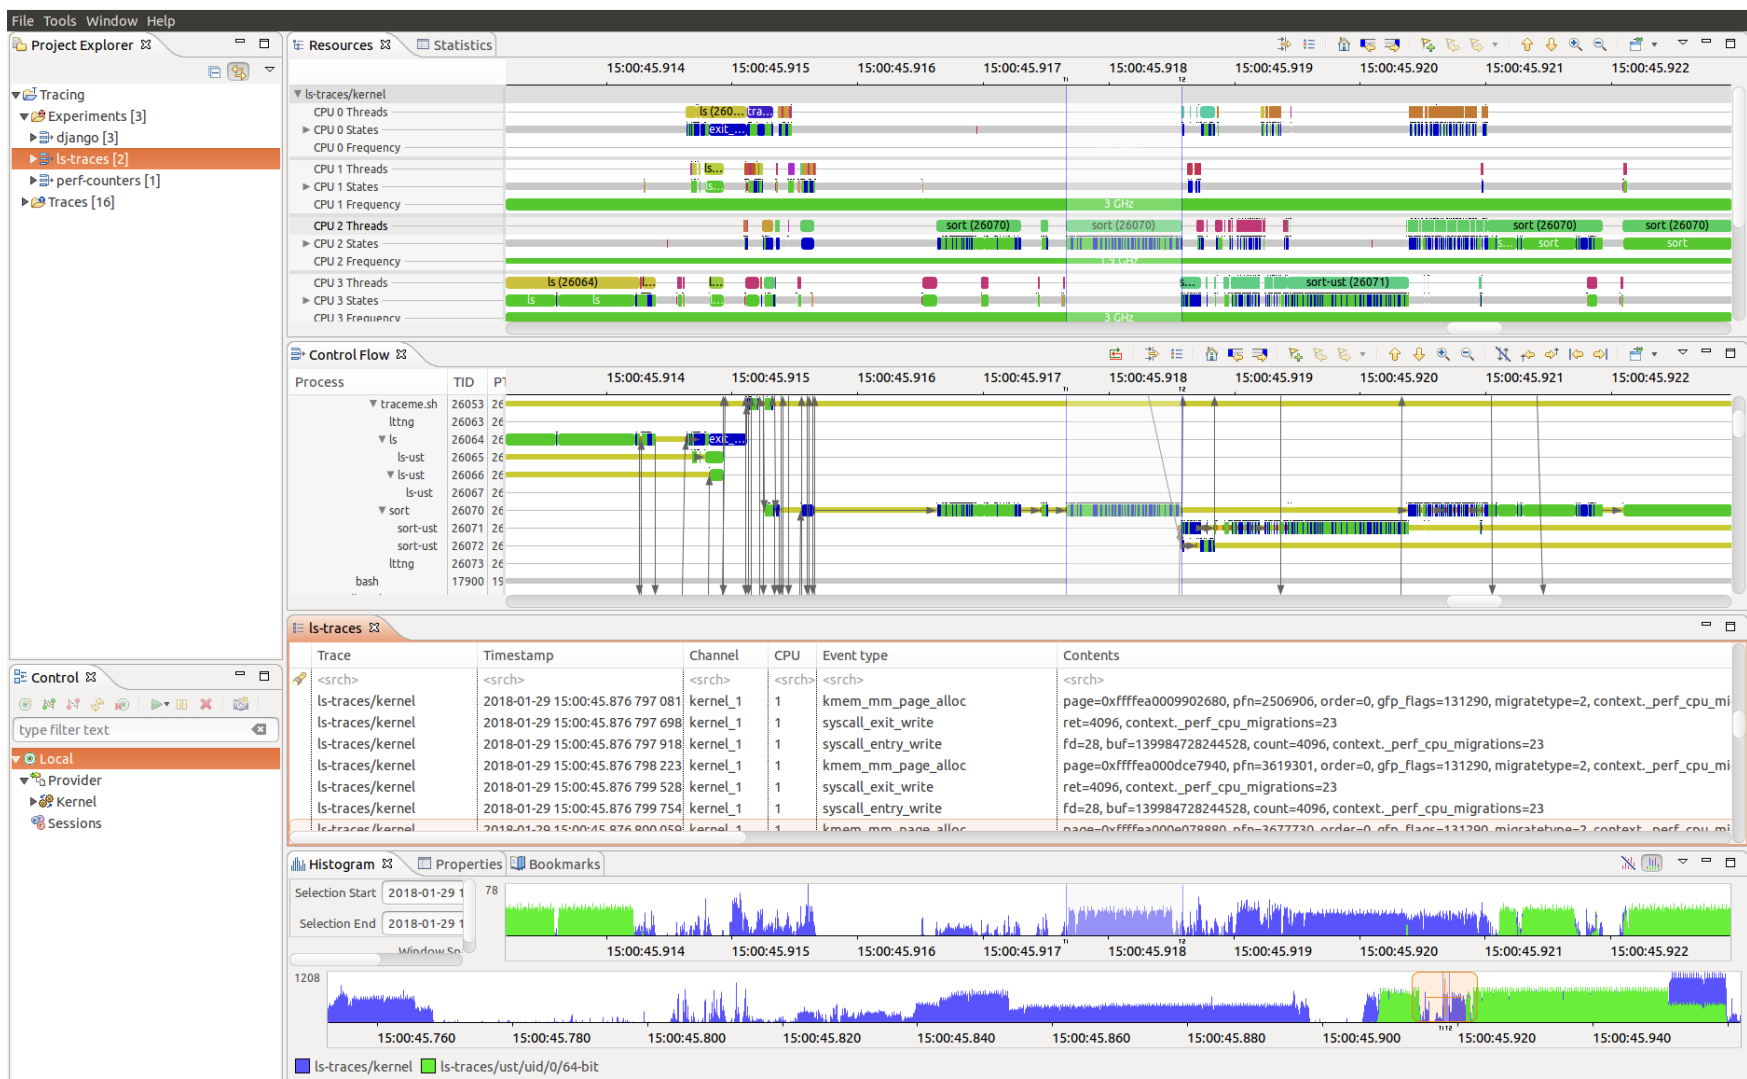
- Tools that provide a visual interface for detailed information about the runtime operations of a program
 - Understand how your program utilizes different computing resources
 - e.g. memory, CPU, GPU, hard disk, network
 - Identify the bottleneck of your program
 - Optimize program implementation
 - Locate potential bugs in your program
- Example:
 - JConsole
 - VisualVM
 - Eclipse Trace Compass



VisualVM



Trace Compass



Performance Evaluation by Simulation

- How to obtain practical performance evaluation of algorithms and systems considering realistic inputs?
 - *Real-world deployment*
 - Setup a small-scale deployment for performance evaluation
 - Expensive or only small-scale; Cannot obtain prior insights
 - *Modeling and analysis*
 - Mathematical reasoning of performance
 - Only apply to simple systems; modeling needs simplifying assumptions
 - *Simulation*
 - Generate artificial inputs to estimate real-world performance
 - A balance between realism and efficiency



Performance Evaluation by Simulation

- Simulation is cost-effective and does not require many simplifying assumptions, which is a viable approach for performance evaluation
- Dynamic systems
 - Systems (which are controlled by certain algorithms) change with time and respond to random inputs, e.g., a system playing Tetris
 - Sample path is the evolution of states in a dynamic system
 - Computational construction of sample paths is a major part of simulation
- Discrete-event simulations
 - Some sample paths are characterized by finite events
 - Construct a random generation model for the discrete events



Motivating Question of Performance

- You have a program X with two component parts A and B
 - Each of which takes 10 minutes. What is the latency of X?
 - Latency is the time from the beginning to the end to complete a job
- Suppose that you can speedup part B by a factor of 5
 - What is the latency now?
 - What is the overall speedup?
 - If A and B are sequential, then Amdahl's Law provides an answer

CPU Processing Time



Amdahl's Law

- How much extra performance can you have if you speed up some part of your program?
- Notations:
 - S is the overall performance gain
 - k is the speed-up factor
 - α is the portion of speed-up

- $$T_{new} = \overset{\text{Unimproved part}}{(1 - \alpha)T_{old}} + \overset{\text{Improved part}}{\alpha \frac{T_{old}}{k}} = T_{old} \left((1 - \alpha) + \frac{\alpha}{k} \right)$$

- $$S = \frac{T_{old}}{T_{new}} = \frac{1}{(1 - \alpha) + \frac{\alpha}{k}}$$



Example

- Your program has one very slow procedure that consumes 70% of the total time. Next, you improve it by a factor of 2
- What is the performance gain in the overall latency?

- $\alpha = 0.7$ (70%)

- $k = 2$

- $$S = \frac{T_{old}}{T_{new}} = \frac{1}{(1-\alpha) + \frac{\alpha}{k}} = \frac{1}{(1-0.7) + \frac{0.7}{2}} = 1.538$$

CPU Processing Time



Example

- Floating point instructions could be improved by 2x. Only 15% of instructions are floating point
- What is the performance gain in the overall latency?

- $\alpha = 0.15$ (15%)

- $k = 2$


- $$S = \frac{T_{old}}{T_{new}} = \frac{1}{(1-\alpha) + \frac{\alpha}{k}} = \frac{1}{(1-0.15) + \frac{0.15}{2}} = 1.081$$

CPU Processing Time



Amdahl's Law

S (performance gain)




k (speed-up factor)



Application to Parallel Processing


- Divide the program into sequential part, $1-P$, and parallel part, P
- Assume there are N processors then the improvement of the parallelizable part is N
- Based on Amdahl's law, the performance gain from N processors is:

$$S = \frac{1}{(1-P) + \frac{P}{N}}$$



Limit as $N \rightarrow \infty$

- The performance gain from a very large number of processors is
- $$S = \lim_{N \rightarrow \infty} \frac{1}{(1-P) + \frac{P}{N}} = \frac{1}{1-P}$$
- Fundamental limitation of performance gain from parallelization (diminishing returns); adding more CPUs may not improve performance
 - Neglects other potential bottlenecks, e.g., memory bandwidth and I/O bandwidth



Reference

- Amdahl's Law paper: "Validity of the single processor approach to achieving large scale computing capabilities"
 - <https://inst.eecs.berkeley.edu/~n252/paper/Amdahl.pdf>
- Java and Parallel Programming
 - <https://www.oracle.com/technical-resources/articles/java/fork-join.html>

