

Image Filtering

Disclaimer: Many of the slides used here are obtained from online resources (including many open lecture materials given at MIT, etc.) without due acknowledgement. They are used here for the sole purpose of classroom teaching. All the credits and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any, any purposes other than helping you study the course of ENGN4528/6528.

Announcement: Student Rep. & Lab Assign.& Lab Session

- Student Representatives
 - Duan Dao
 - Maitreyi Singh
- Lab Assignment One Released
 - Due date: due 11:59pm, 28th March 2021 (end of week 5)
- Lab Sessions for Lab Assignment One starts this week. Please attend the lab session for getting support.
 - Lab session: week 03 & week 04, no lab session at week 05
- About closing the discussion forum on wattle and move all discussions to Piazza.

Announcement: Practice questions, tutorial

- Introduce practice questions, type of questions for exam
 - Our tutor, Huiyu Gao will explain answers to these questions during the lab session.
- Tutorial: Mar. 10th, about Image processing by our tutor: Jiayu Yang

Announcement

- Mathematics Reference Book: Linear algebra, probability
 - <https://mml-book.github.io/book/mml-book.pdf>
 - eg. CDF (cumulative distribution function, see chapter 6.2 discrete and continuous distribution function)
- Lecture Series by Prof. Shree Nayar, Professor of Computer Science in the School of Engineering at Columbia University, Senior researcher in computer vision with high profile.
 - <https://fpcv.cs.columbia.edu/>

Outline

- Continue: Linear filter
 - Gaussian filter
 - Edge detection
- non-linear filter
 - median filter
 - bilateral filter

Images as functions

- We can think of an **image** as a function, f , from \mathbf{R}^2 to \mathbf{R} :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
 - $f: [a, b] \times [c, d] \rightarrow [0, 1]$

Image Noise

Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise



Gaussian noise

Moving Average In 2D

$$F[x, y]$$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |

Moving Average In 2D

$F[x, y]$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$G[x, y]$

| | | | | | | | | | | |
|---|----|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | |
| | | | | | | | | | | |
| 0 | 10 | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |

Moving Average In 2D

$F[x, y]$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$G[x, y]$

| | | | | | | | | | | |
|--|--|--|---|----|----|--|--|--|--|--|
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | 0 | 10 | 20 | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |

Moving Average In 2D

$F[x, y]$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$G[x, y]$

| | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Moving Average In 2D

$$F[x, y]$$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | | |
|--|--|--|---|----|----|----|----|--|--|--|
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | 0 | 10 | 20 | 30 | 30 | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |
| | | | | | | | | | | |

Moving Average In 2D

$$F[x, y]$$

| | | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | |
|--|----|----|----|----|----|----|----|----|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | 30 | 30 | 30 | 20 | 10 | |
| | 0 | 20 | 40 | 60 | 60 | 60 | 40 | 20 | |
| | 0 | 30 | 60 | 90 | 90 | 90 | 60 | 30 | |
| | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 | |
| | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 | |
| | 0 | 20 | 30 | 50 | 50 | 60 | 40 | 20 | |
| | 10 | 20 | 30 | 30 | 30 | 30 | 20 | 10 | |
| | 10 | 10 | 10 | 0 | 0 | 0 | 0 | 0 | |
| | | | | | | | | | |

Correlation filtering

Say the averaging window size is $(2k+1) \times (2k+1)$:

$$G[x, y] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \sum_{u=-k}^k \sum_{v=-k}^k F[x+u, y+v]$$

Loop over all pixels in neighborhood around image pixel $F[x,y]$

Now we can generalize to allow **different weights** depending on neighboring pixel's relative position:

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] \underbrace{F[x+u, y+v]}_{\text{Non-uniform weights}}$$

Correlation filtering

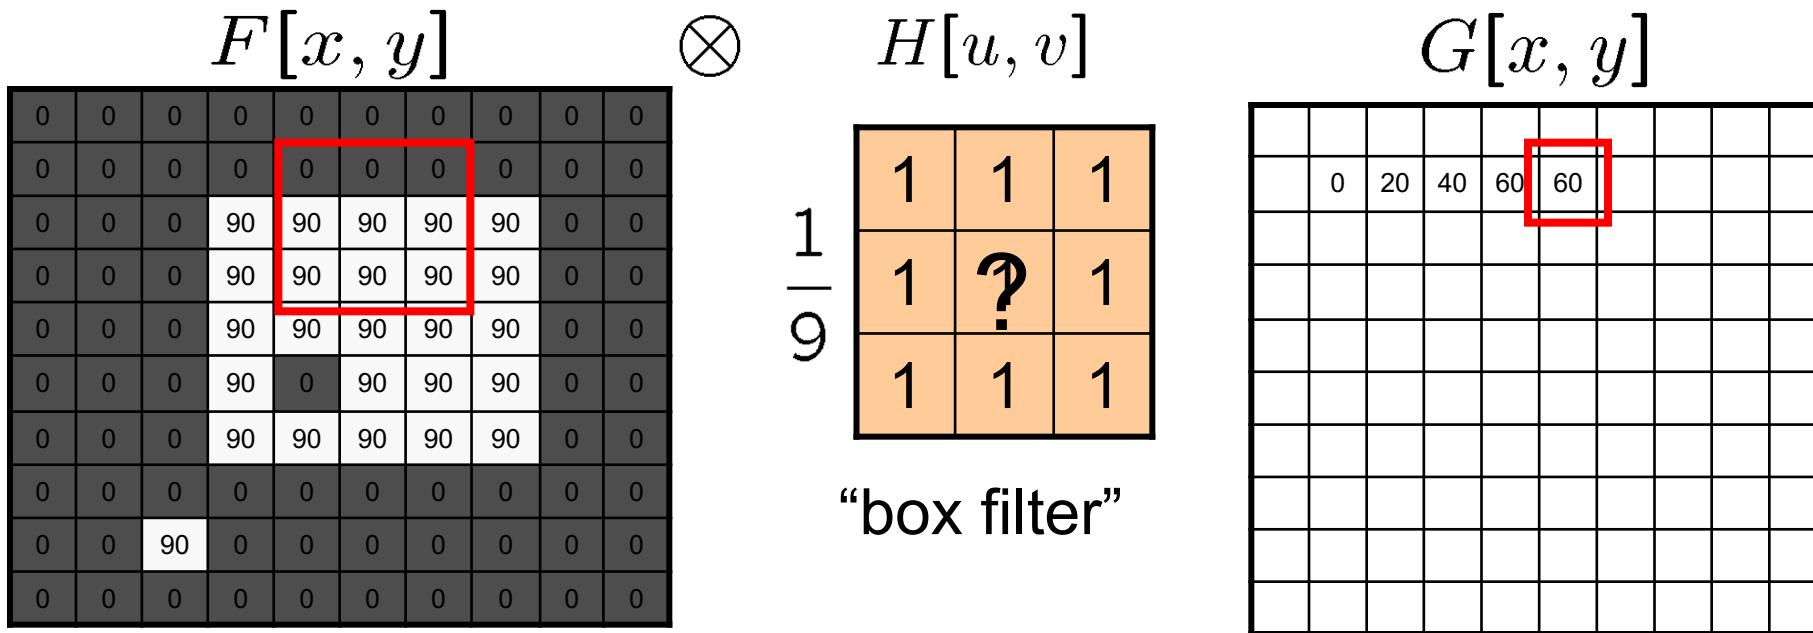
$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[x + u, y + v]$$

This is called cross-correlation, denoted $G = H \otimes F$

- Filtering an image: replace each pixel with a linear combination of its neighbors.
- The filter “**kernel**” or “**mask**” $H [u, v]$ is the prescription for the weights in the linear combination.

Averaging filter

- What values belong in the kernel H for the moving average example?



$$G = H \otimes F$$

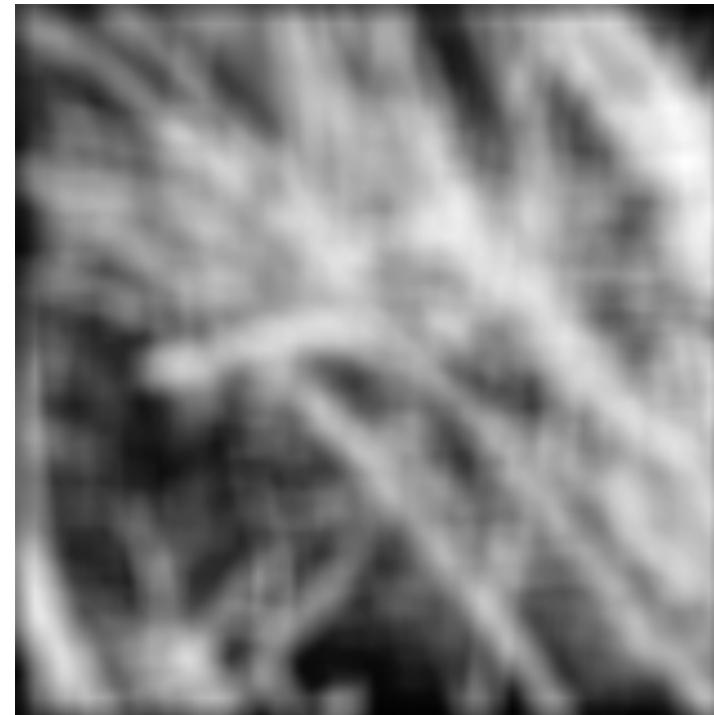
Smoothing by averaging



depicts box filter:
white = high value, black = low value



original



filtered

We can see the ‘Block’ effect on the filtered image.

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

| | | | | | | | | | | | |
|---|---|---|----|----|----|----|----|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

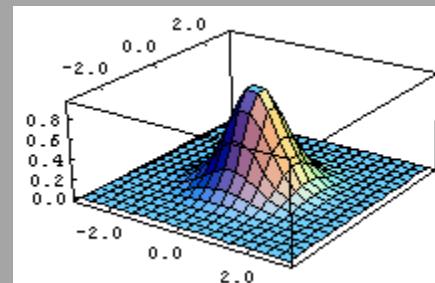
$F[x, y]$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

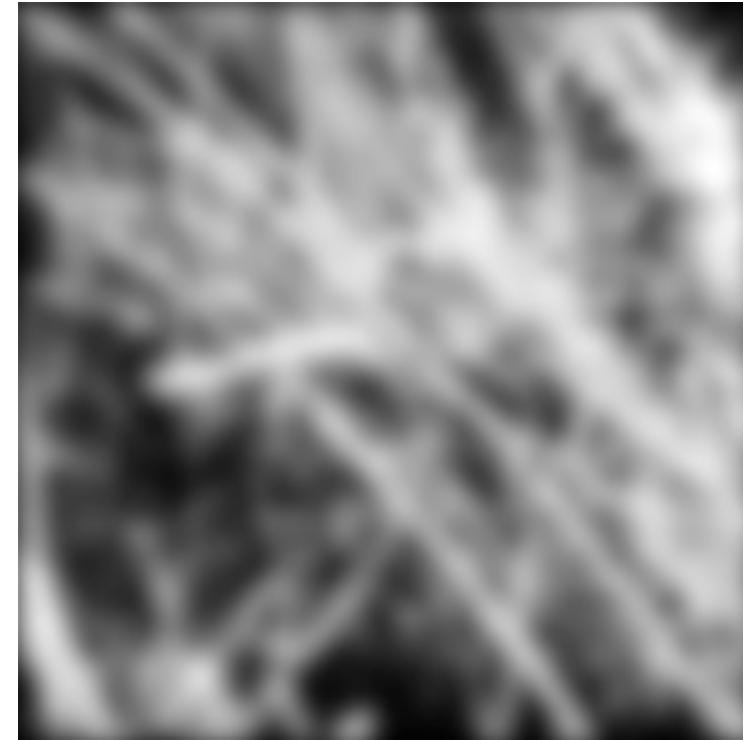
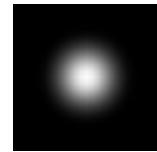
$H[u, v]$

- This kernel is an approximation of a 2d Gaussian function:

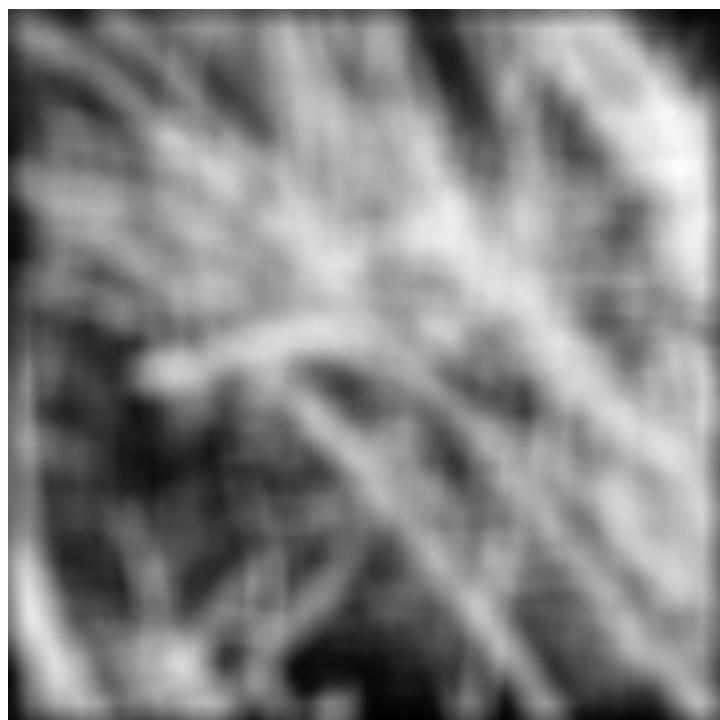
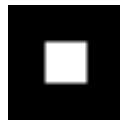
$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



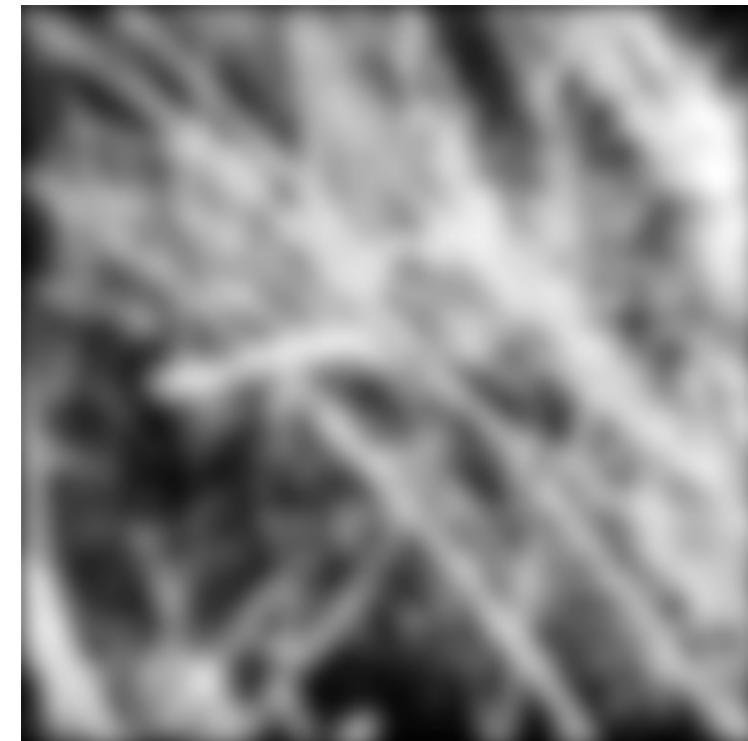
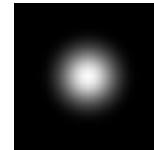
Smoothing with a Gaussian



Compare box filter with Gaussian filter effects



Smoothed image by box filter

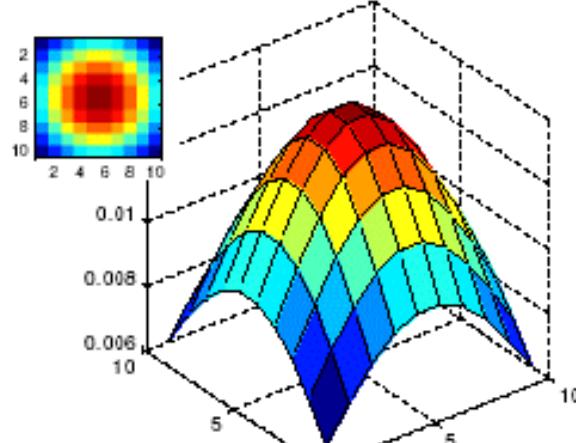
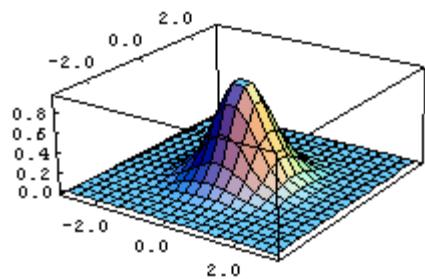


Smoothed image by Gaussian filter

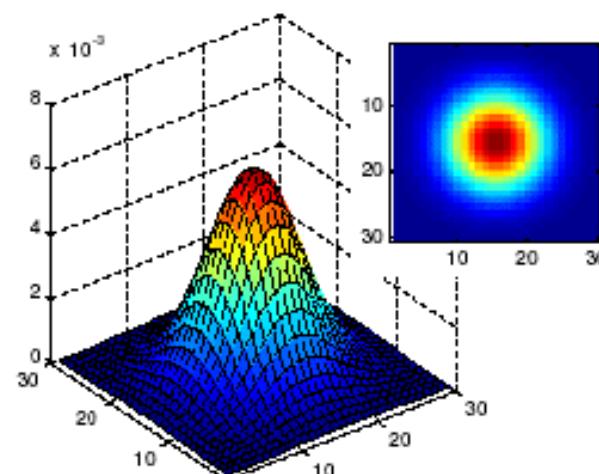
Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



$\sigma = 5$ with
10 x 10
kernel

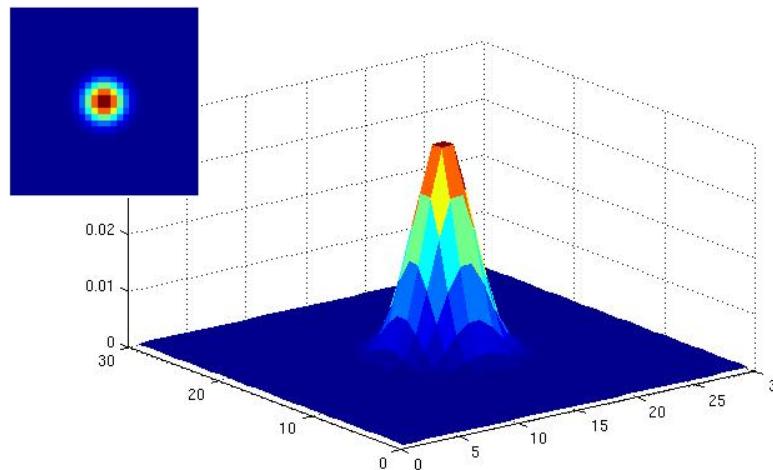


$\sigma = 5$ with
30 x 30
kernel

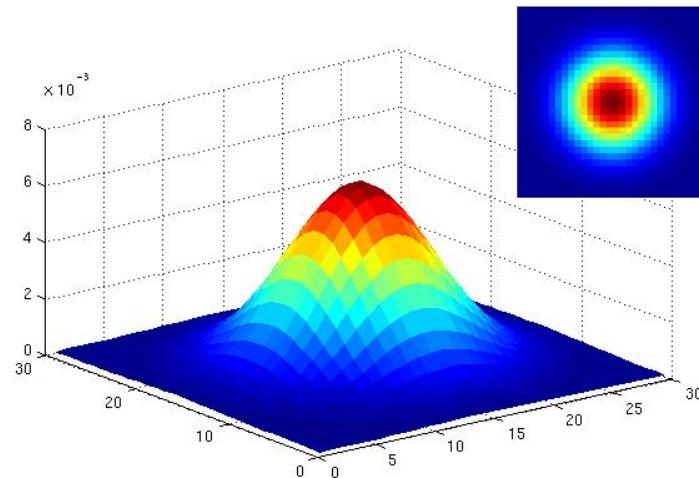
• Slide credit:
Kristen Grauman

Gaussian filters

- What parameters matter here?
- **Variance** of the Gaussian function: determines extent of smoothing



$\sigma = 2$ with
30 x 30
kernel



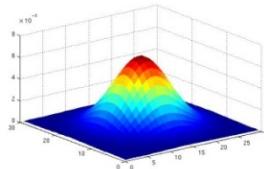
$\sigma = 5$ with
30 x 30
kernel

• Slide credit:
Kristen Grauman

Matlab

```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```

```
>> mesh(h);
```



```
>> imagesc(h);
```



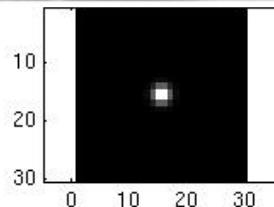
```
>> outim = imfilter(im, h); % correlation  
>> imshow(outim);
```



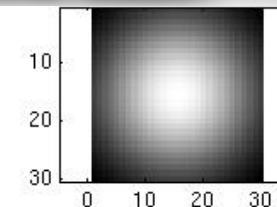
outim

Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

•Slide credit:
Kristen Grauman

Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[x - u, y - v]$$

$$G = H \star F$$



*Notation for
convolution
operator*

Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$G(i, j) = \sum_{u=0}^M \sum_{v=0}^N F(u, v) H(i - u, j - v)$$

$$G = F \star H$$



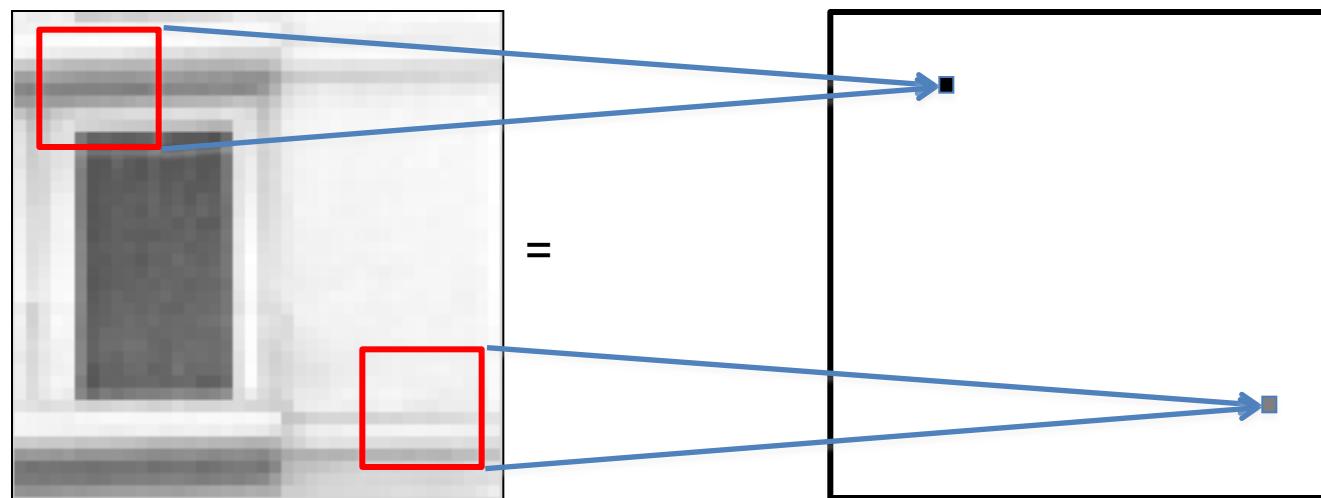
*Notation for
convolution
operator*

Discrete linear system



In vision, many times, we are interested in operations that are spatially invariant.
For a linear spatially invariant system:

$$g[m, n] = h \star f = \sum_{k, l} h[m - k, n - l]f[k, l]$$



Properties of convolution

- **Shift invariant:**

- Operator behaves the same everywhere, i.e. the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood.

- **Superposition:**

- $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$

Properties of convolution

- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Distributes over addition

$$f * (g + h) = (f * g) + (f * h)$$

- Scalars factor out

$$kf * g = f * kg = k(f * g)$$

- Identity:

$$\text{unit impulse } e = [..., 0, 0, 1, 0, 0, ...]. \quad f * e = f$$

Matlab's Filter vs. Convolution

- 2D correlation filtering

$h = \text{filter}$ $g = \text{image}$

- `f=filter2(h,g);` or
`f=imfilter(h,g);`

$$f[m, n] = \sum_{k, l} h[k, l] g[m + k, n + l]$$

- 2D convolution

- `f=conv2(h,g);`

$$f[m, n] = \sum_{k, l} h[k, l] g[m - k, n - l]$$

Separability

- In some cases, the filter is separable, and we can factor into two steps:
 - Convolve all rows with a 1D filter
 - Convolve all columns with a 1D filter

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix} = \begin{bmatrix} 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \end{bmatrix} \circ \begin{bmatrix} 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ 0 & 0 & 0 \end{bmatrix}$$

Separability example

- 2D convolution (center location only)(65)

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

h f

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

- The filter is factored into a product of 1D filters.

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

- Perform convolutions along rows.

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 11 & & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$

Source: K. Grauman

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 11 & & \\ \hline 18 & & \\ \hline 18 & & \\ \hline \end{array}$$

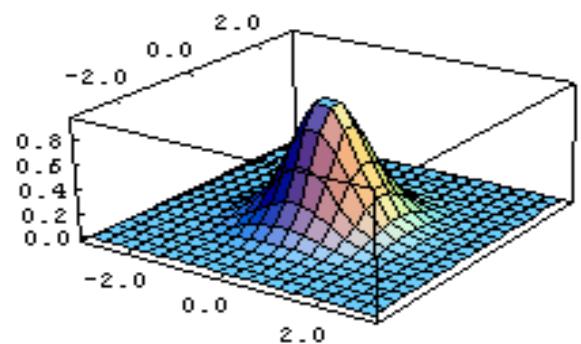
- Perform convolutions along columns.

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 11 & & \\ \hline 18 & & \\ \hline 18 & & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \quad 65$$

Source: K. Grauman

This kernel is an approximation of
a 2d Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



Properties of Gaussian filter

- Rotational symmetry treats features of all orientations equally (isotropy).
- Convolution with self gives another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into the product of two 1D Gaussians

Separability of the Gaussian filter

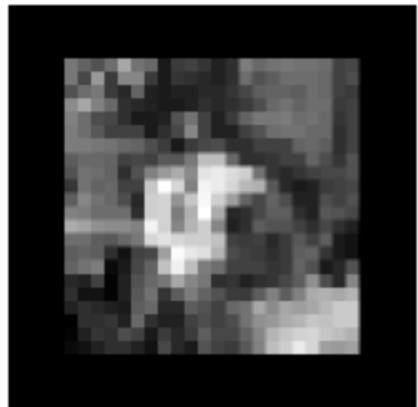
$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

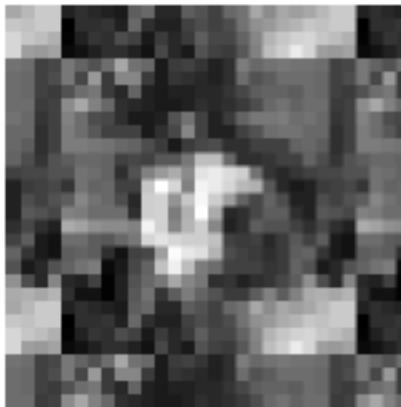
In this case, the two functions are the (identical) 1D Gaussian

- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
- What if the kernel is separable?
 - $O(n^2 m)$

How to handle borders



zero



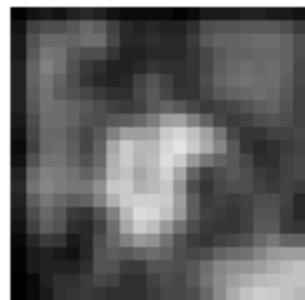
wrap



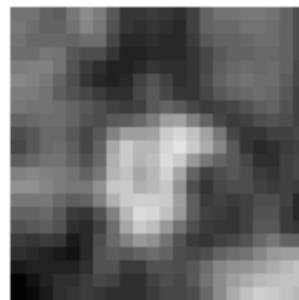
clamp



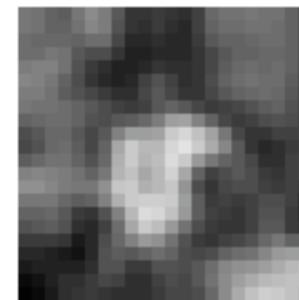
mirror



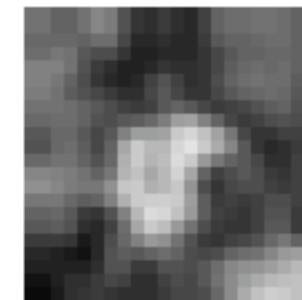
blurred: zero



normalized zero



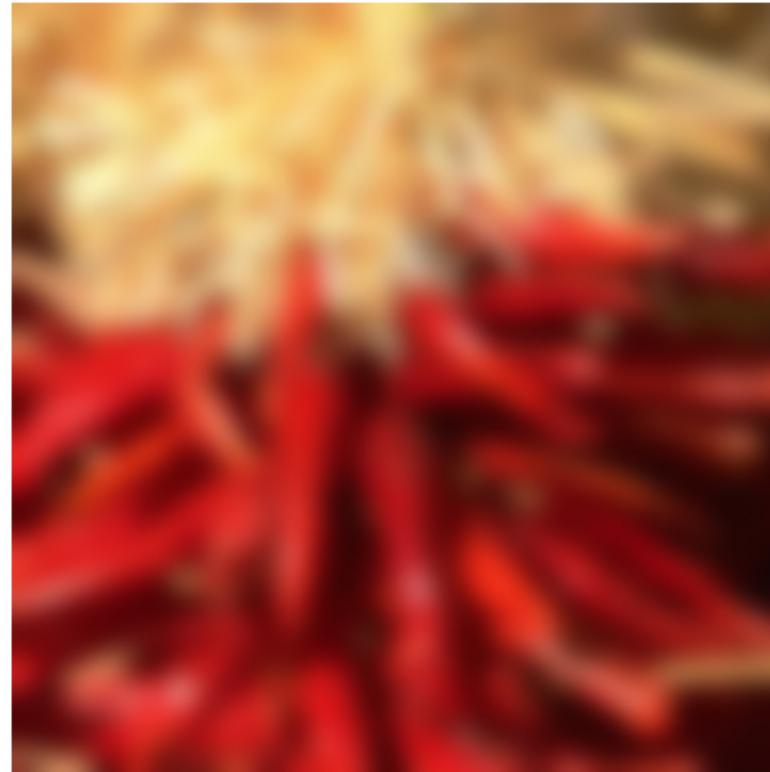
clamp



mirror

From Szeliski, Computer Vision, 2010

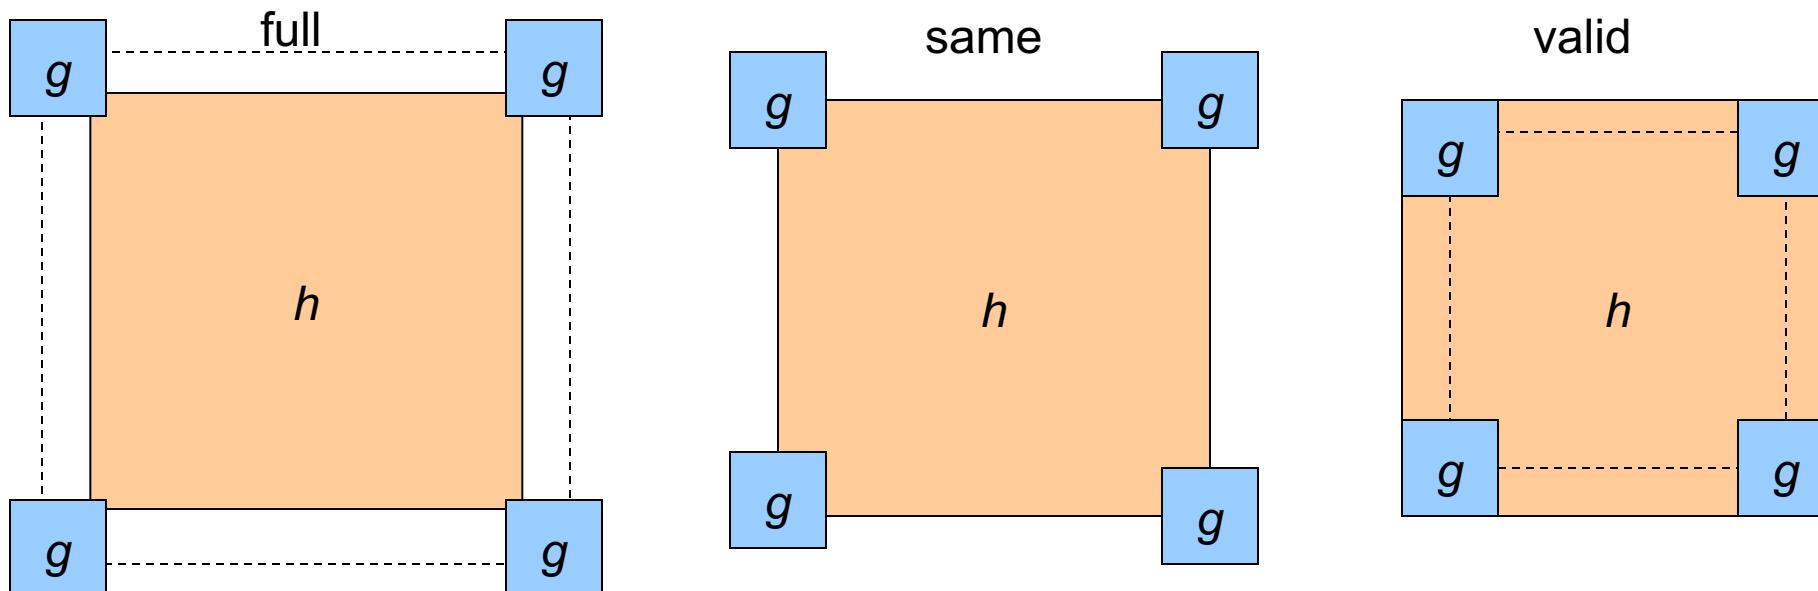
Example: Border handling



Source: S. Marschner

Matlab's border handling

- What is the size of the output?
- MATLAB: `filter2(g, h, shape)`
 - *shape* = 'full': output size is sum of sizes of *h* and *g*
 - *shape* = 'same': output size is same as *h*
 - *shape* = 'valid': output size is difference of sizes of *h* and *g*



Discrete Filtering

- Linear filter: Weighted sum of pixels over rectangular neighborhood—*kernel* defines weights
- Think of kernel as template being matched by **correlation** (Matlab: `imfilter`, `filter2`)
- **Convolution:** Correlation with kernel rotated 180° (Matlab: `conv2`)

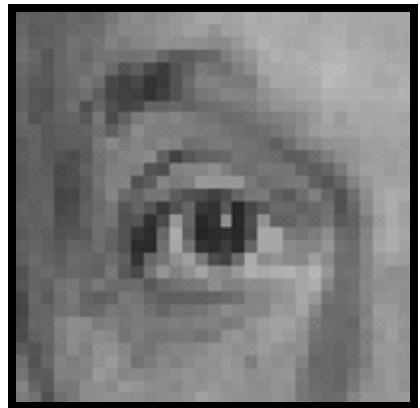
| | | |
|---|----|----|
| 1 | -1 | -1 |
| 1 | 2 | -1 |
| 1 | 1 | 1 |

Practice with linear filters

A taxonomy of useful filters

- Impulse, Shifts,
- Blur
 - Box filter
 - Gaussian filter
 - Bilateral filter
 - Asymmetrical filter: motion blur
- Edges
 - [-1 1]
 - Derivative filter
 - Derivative of a gaussian
 - Oriented filters
 - Gabor filter
 - Quadrature filters: phase and magnitude.
 - Elongated edges: filling gaps...

Practice with linear filters

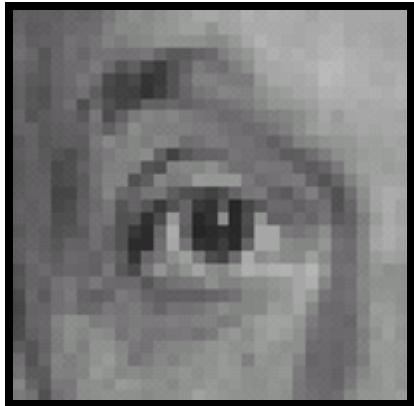


Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

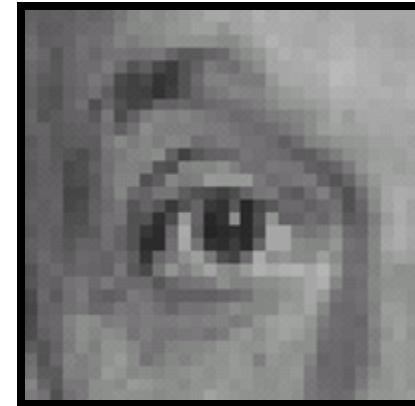
?

Practice with linear filters



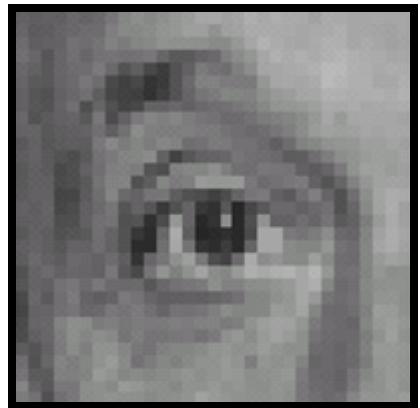
Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |



Filtered
(no change)

Practice with linear filters

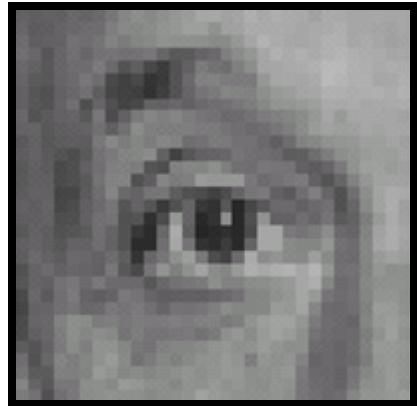


Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |

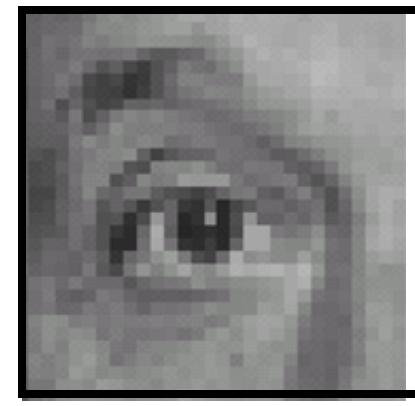
?

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |



Shifted left
By 1 pixel

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 0 |
| 0 | 0 | 0 |

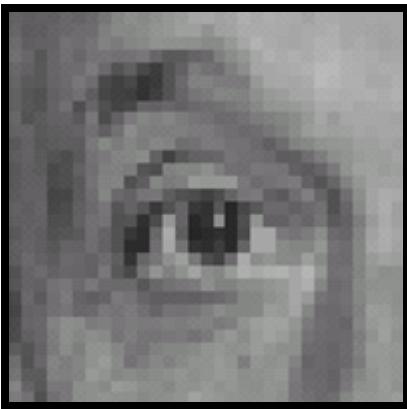
-

| | | | |
|---------------|---|---|---|
| $\frac{1}{9}$ | 1 | 1 | 1 |
| 1 | 1 | 1 | 1 |
| 1 | 1 | 1 | 1 |

?

(Note that filter sums to 1)

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 0 |
| 0 | 0 | 0 |

-

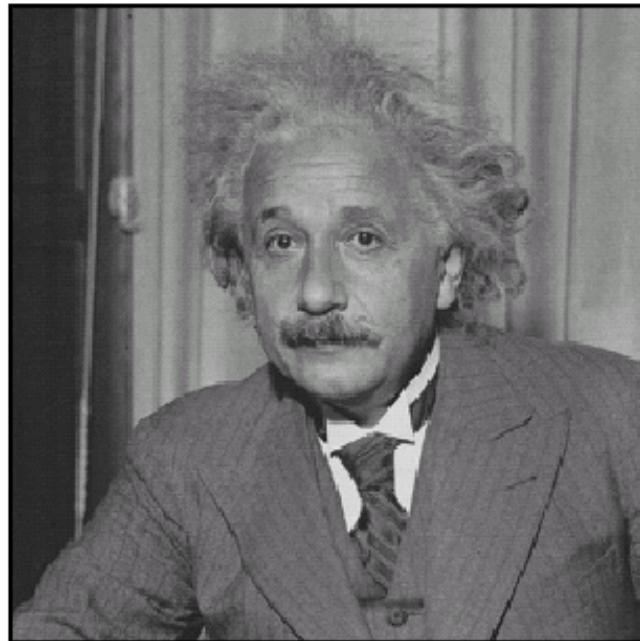
| | | | |
|---------------|---|---|---|
| $\frac{1}{9}$ | 1 | 1 | 1 |
| 1 | 1 | 1 | 1 |
| 1 | 1 | 1 | 1 |



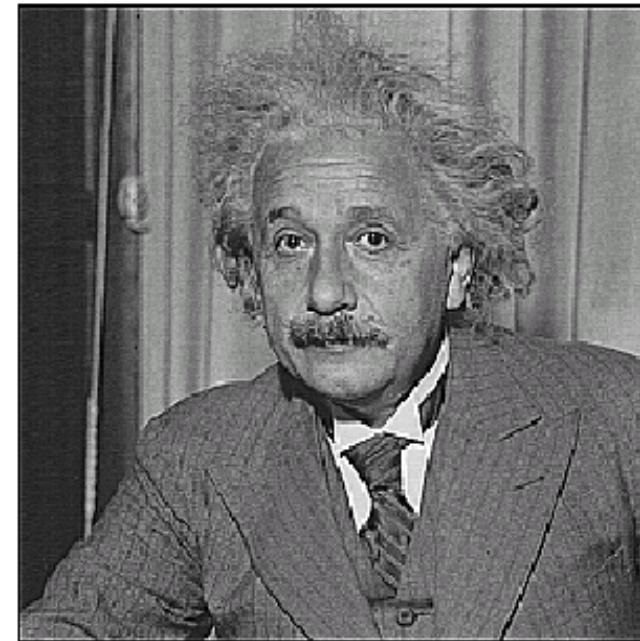
Sharpening filter

- Accentuates differences with local average

Sharpening

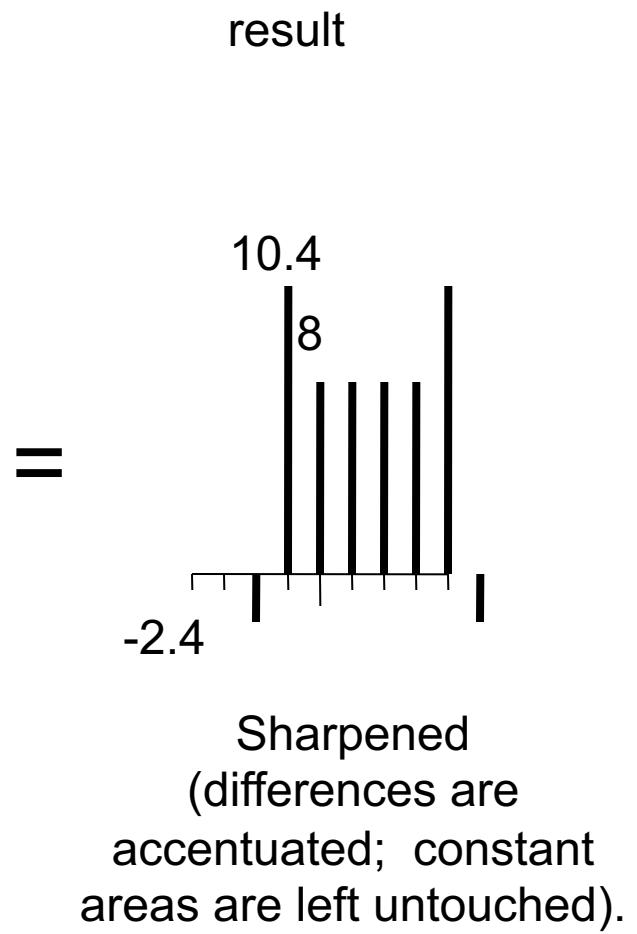
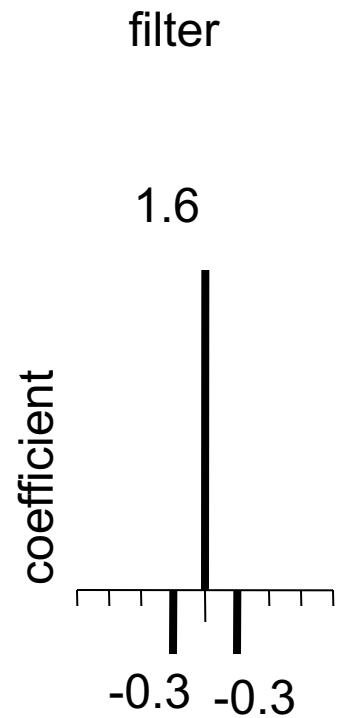
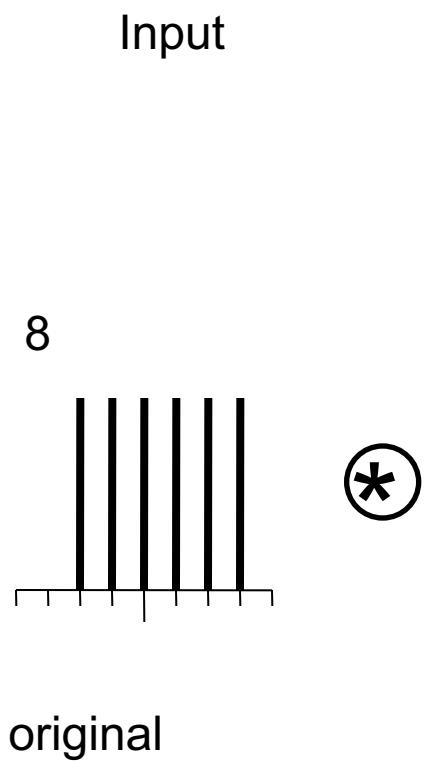


before



after

1-D sharpening example



$[-1 \ 1]$

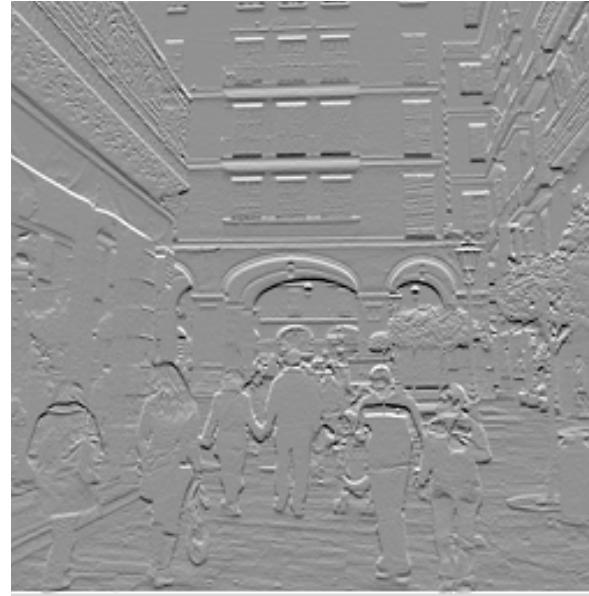
$$\frac{\partial \mathbf{I}}{\partial x} \simeq \mathbf{I}(x, y) - \mathbf{I}(x - 1, y)$$

 \otimes $[-1, 1]$ $h[m,n]$ 

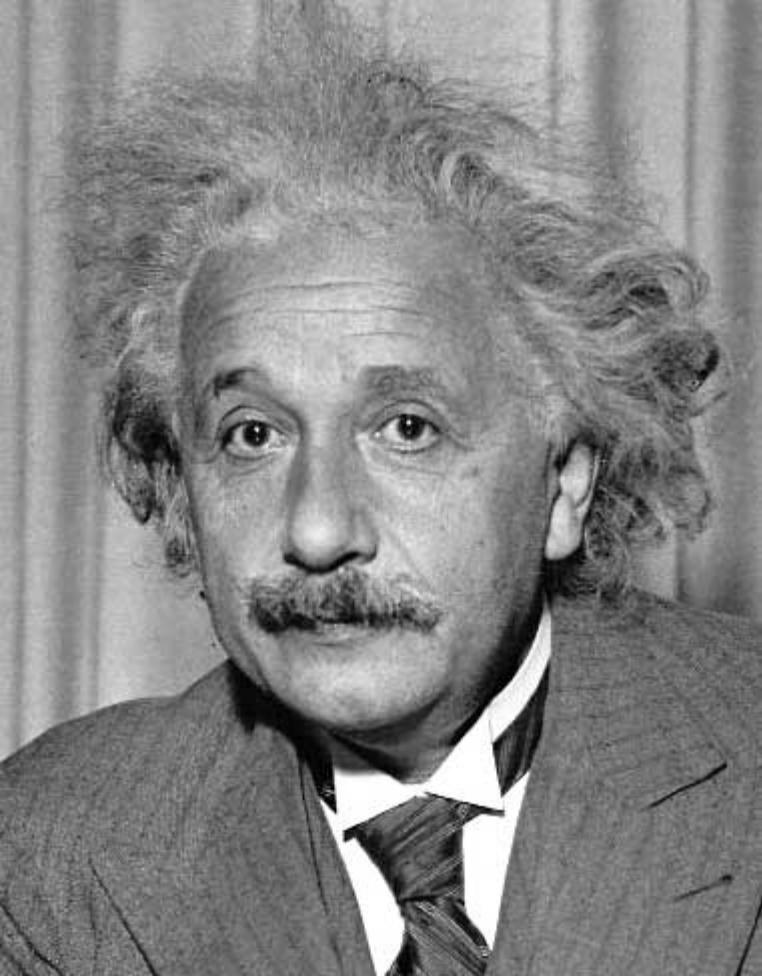
$$[-1 \ 1]^\top$$

 \otimes

$$[-1, 1]^\top = \\ h[m,n]$$



Derivative filters



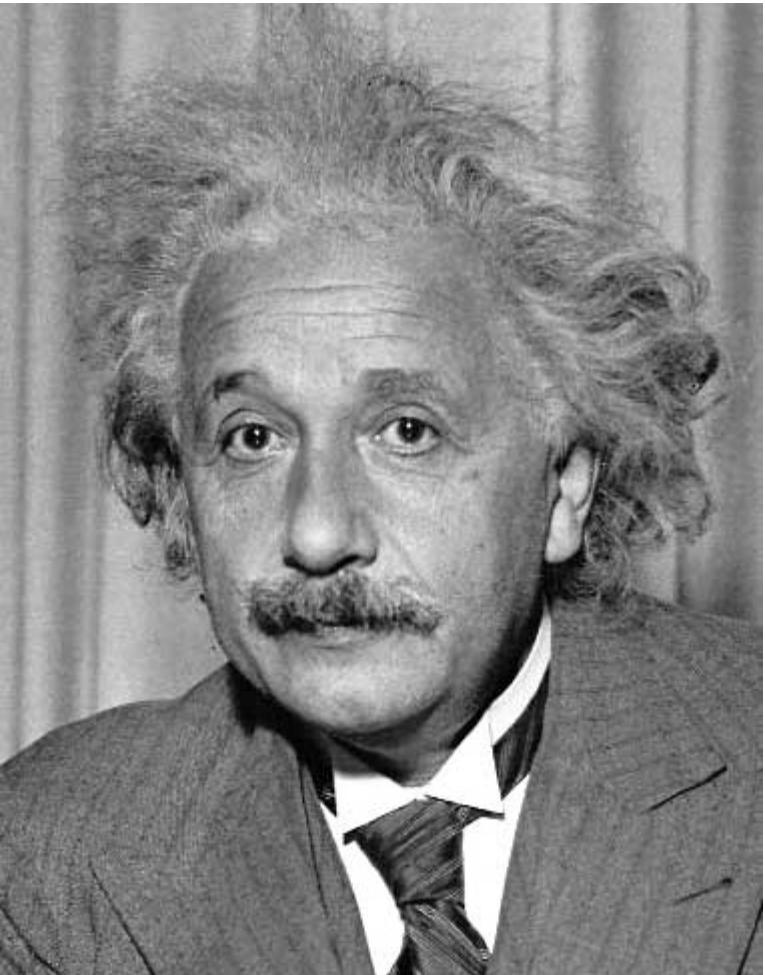
| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

Sobel



Vertical Edge
(absolute value)

Derivative filters



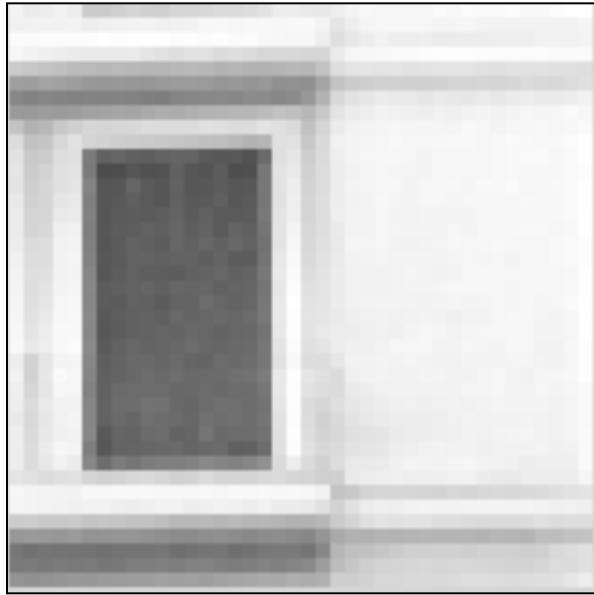
| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |

Sobel



Horizontal Edge
(absolute value)

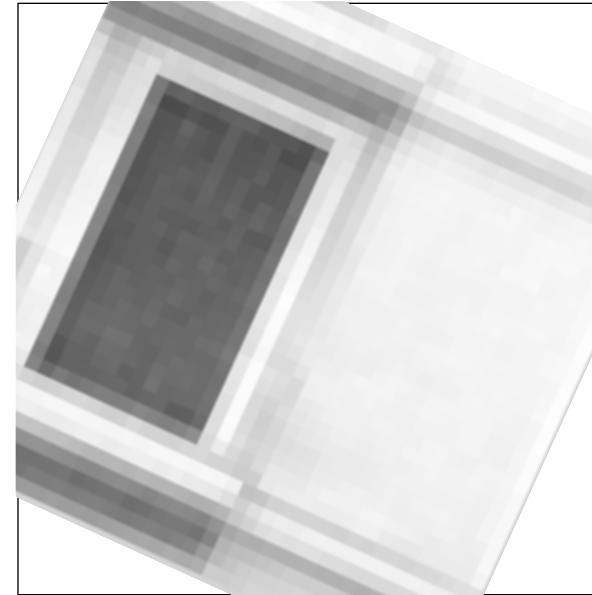
Image rotation



$g[m,n]$

\otimes ? =

$h[m,n]$

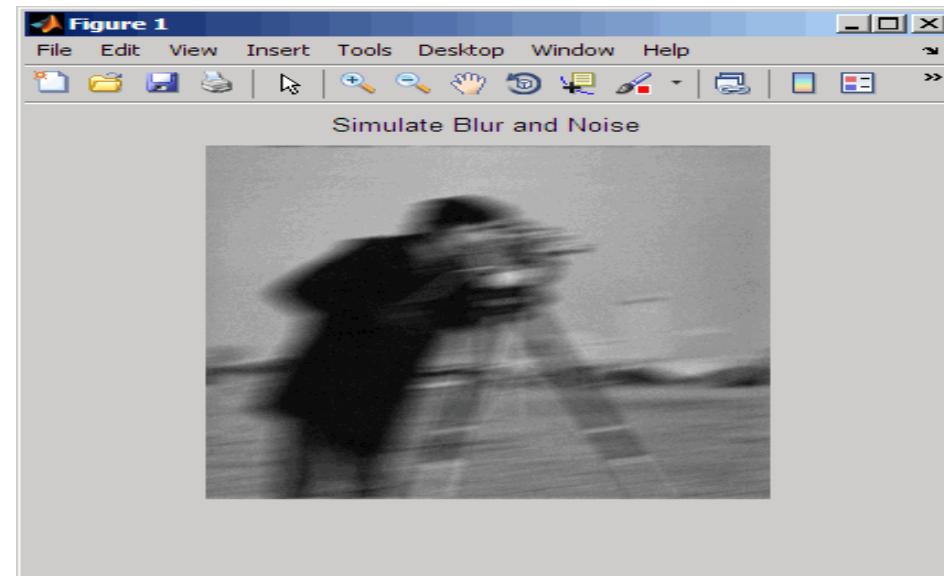


$f[m,n]$

It is linear, but not a spatially invariant operation. There is not convolution.

How could we synthesize motion blur?

```
theta = 30; len = 20;  
fil = imrotate(ones(1, len), theta, 'bilinear');  
fil = fil / sum(fil(:));  
figure(2), imshow(imfilter(im, fil));
```



Motion blur filter



$g[m,n]$

\otimes



=

$h[m,n]$



$f[m,n]$

Motion blur filter



$g[m,n]$

\otimes

$h[m,n]$

=



$f[m,n]$

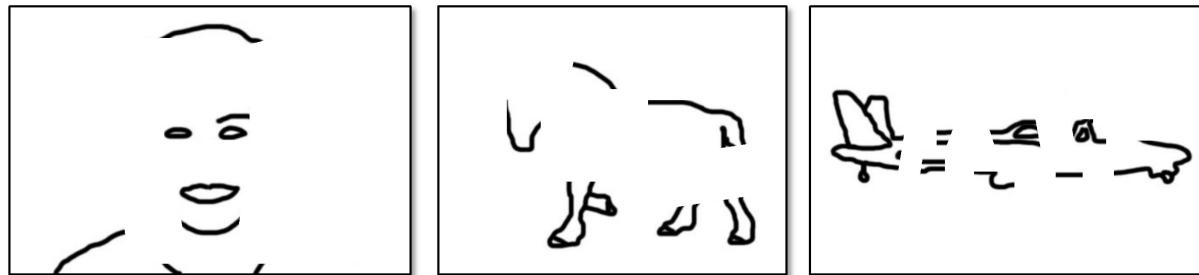
Summary

- Linear filters and convolution useful for
 - Enhancing images (smoothing, removing noise)
 - Box filter
 - Gaussian filter
 - Separable filters more efficient

Edge Detection

Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**



•Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

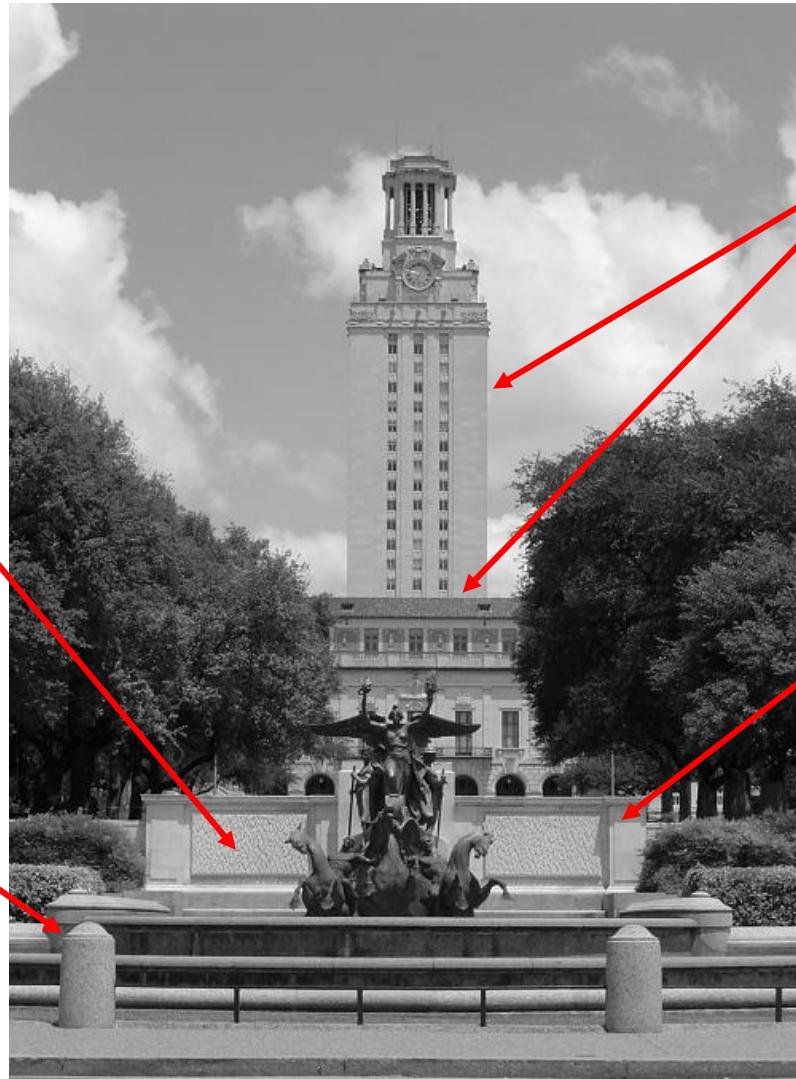
What causes an edge?

- Reflectance change:
appearance
information, texture

- Change in surface
orientation: shape

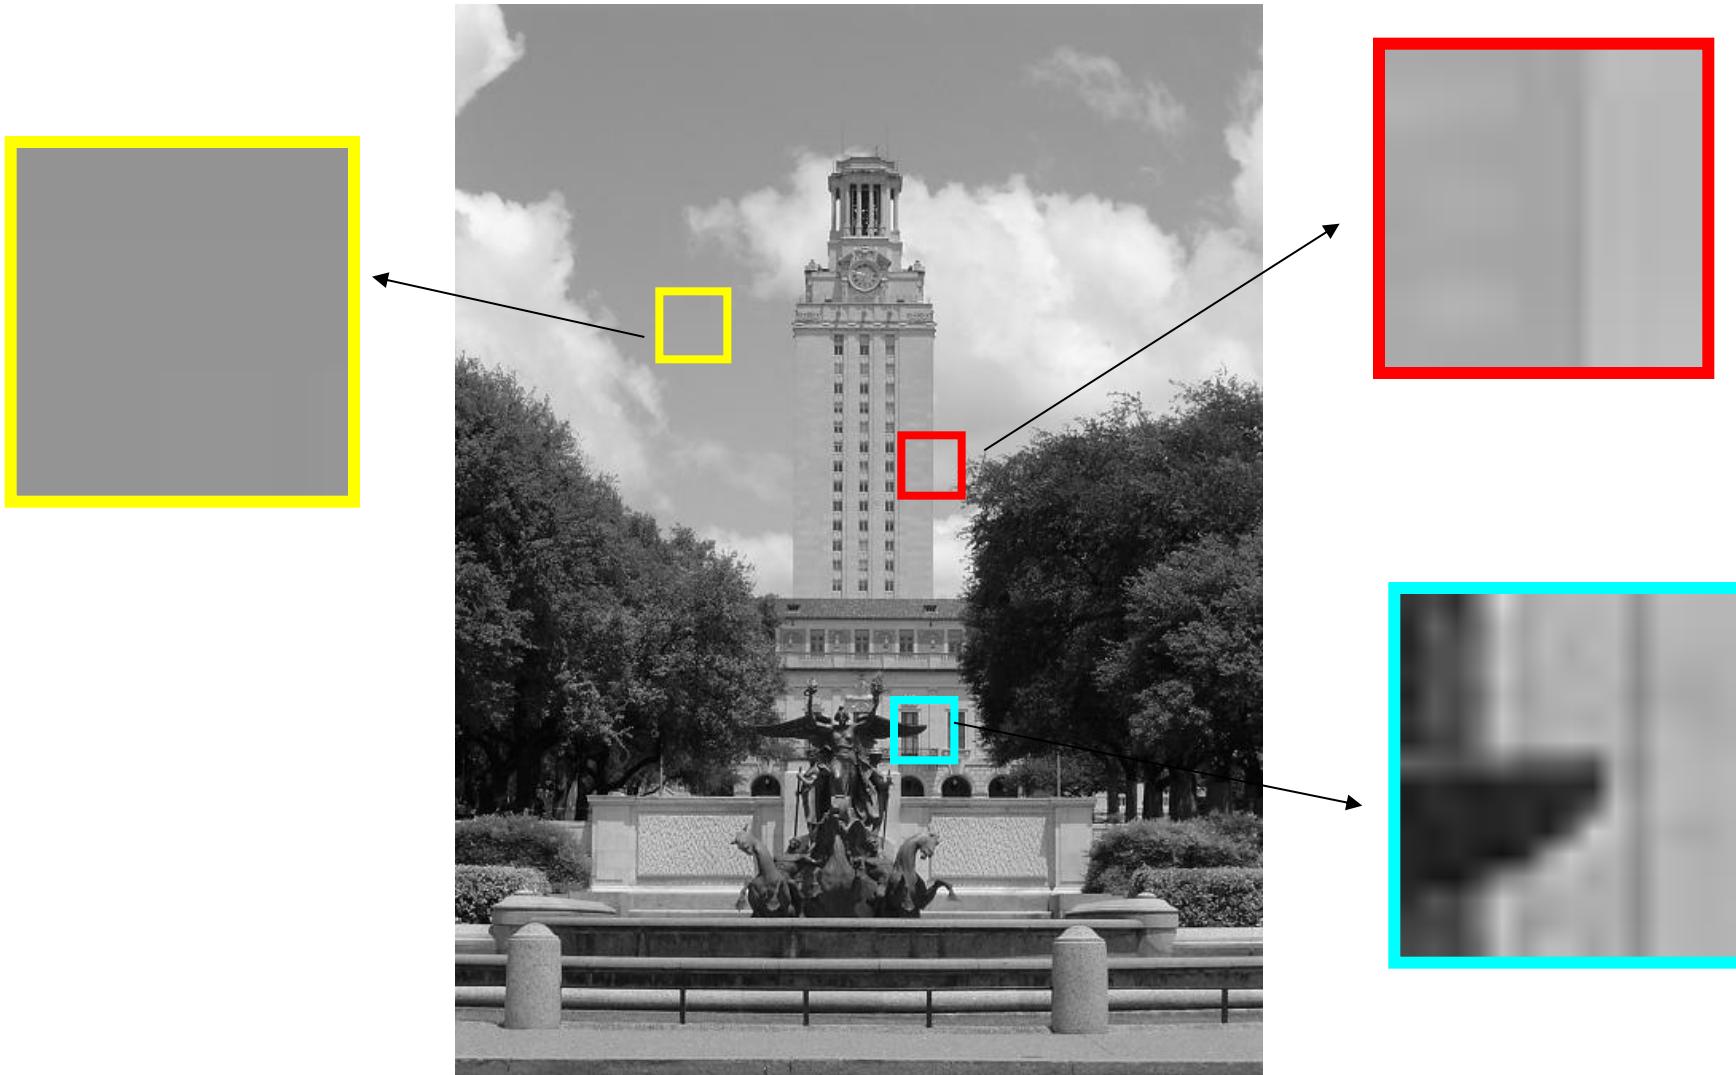
- Depth discontinuity:
object boundary

- Cast shadows



• Slide credit:
Kristen Grauman

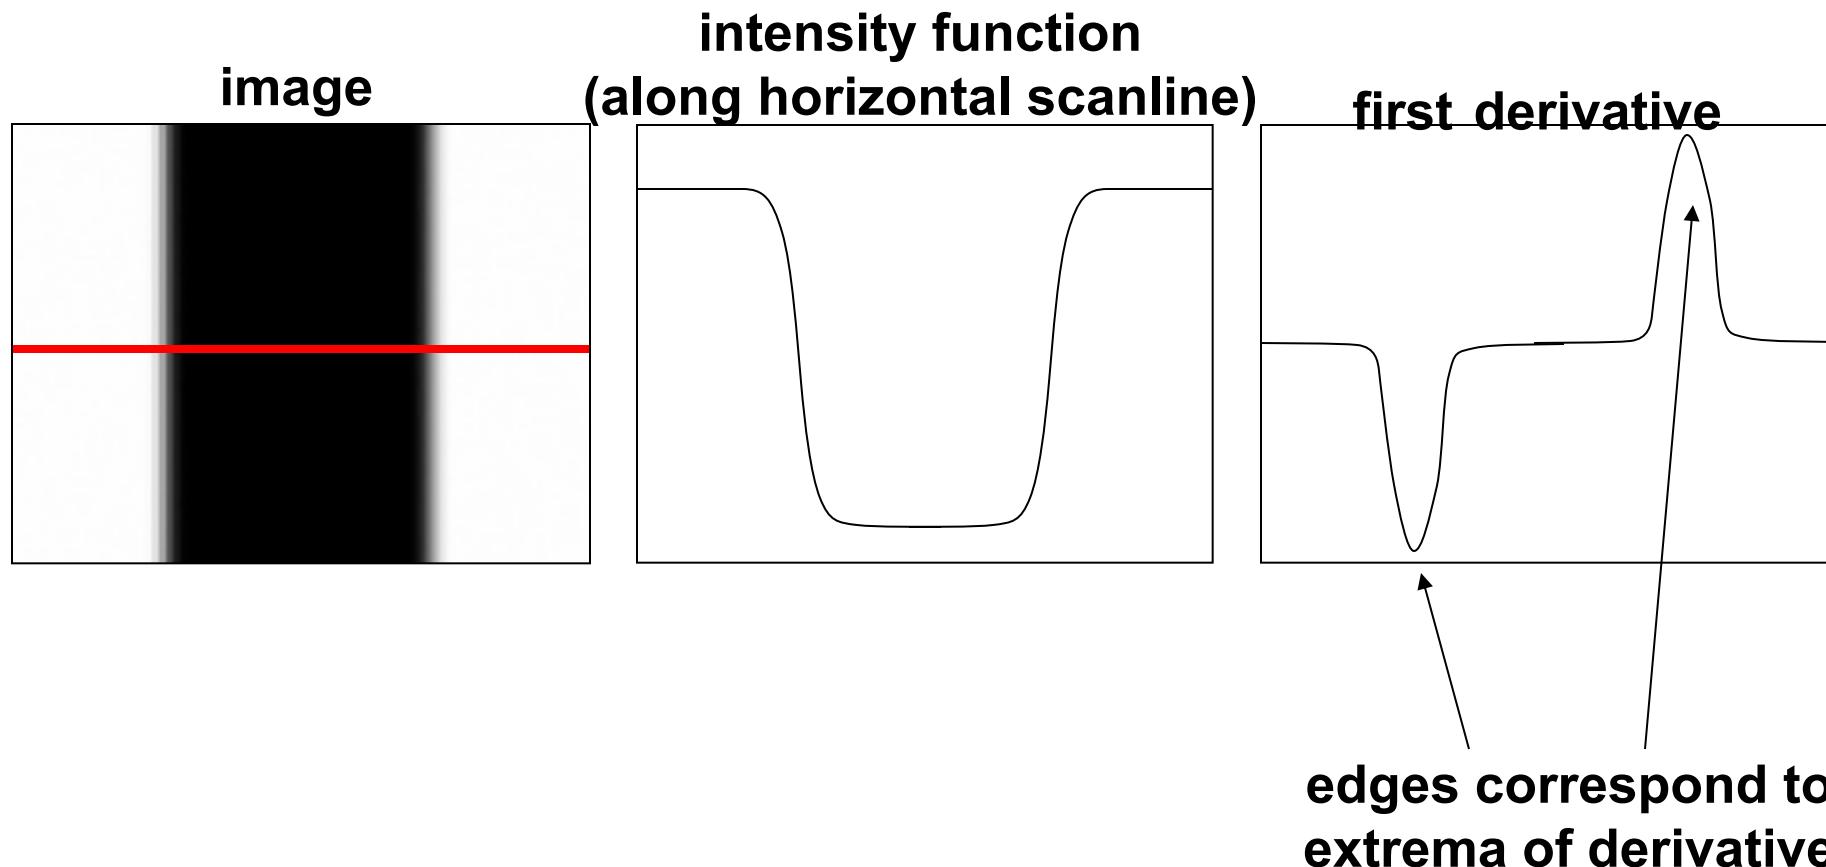
Edges/gradients and invariance



•Slide credit:
Kristen Grauman

Derivatives and edges

- An edge is a place of rapid change in the image intensity function.



Methods

- Using first order derivative: *Look for extrema*
 - Sobel operator : `edge(I, 'sobel');`
 - Prewitt, Roberts,...
 - Derivative of Gaussian
- Using second order derivative: *Look for zero-crossings*
 - Laplacian : isotropic
 - Second directional derivative
 - LOG/DOG (Laplace of Gaussian/Difference of Gaussians)

$$\nabla^2 \mathbf{I}$$

Common edge detectors

- Basic gradient edge detectors
 - **Sobel**, Prewitt, Gaussian derivative ...
- **LoG/Marr** edge detector
- **Canny** edge detector

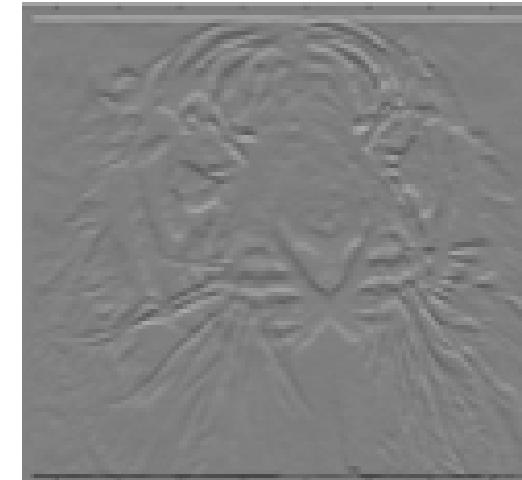
Gradient Edge Detectors

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
>> My = fspecial('sobel');  
>> outim = imfilter(double(im), My);  
>> imagesc(outim);  
>> colormap gray;
```



Eg: Convolute an image with a Sobel kernel

K

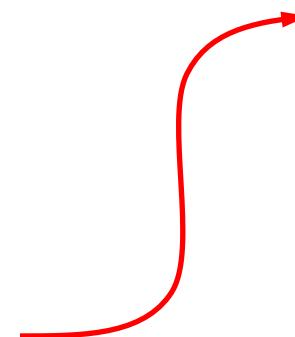
| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

•Rotate
↓

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

I

| | | | |
|---|---|---|---|
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |

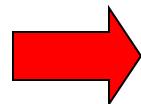


Step

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

| | | | |
|---|---|---|---|
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |

| | | | | |
|----|---|---|---|---|
| -1 | 0 | 1 | | |
| -2 | 0 | 0 | 2 | 2 |
| -1 | 0 | 0 | 2 | 2 |
| 0 | 0 | 1 | 2 | |
| 0 | 0 | 2 | 2 | |



| | | | |
|---|--|--|--|
| 0 | | | |
| | | | |
| | | | |
| | | | |

I

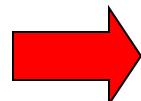
I'

Step

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

| | | | |
|---|---|---|---|
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |

| | | | |
|----|---|---|---|
| -1 | 0 | 1 | |
| 0 | 0 | 4 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |



| | | | |
|---|---|--|--|
| 0 | 6 | | |
| | | | |
| | | | |
| | | | |

I

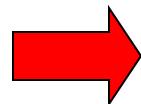
I'

Step

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

| | | | |
|---|---|---|---|
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |

| | | | |
|---|----|---|---|
| | -1 | 0 | 1 |
| 0 | 0 | 0 | 4 |
| 0 | 0 | 0 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |



| | | | |
|---|---|---|--|
| 0 | 6 | 6 | |
| | | | |
| | | | |
| | | | |

I

I'

Step

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

| | | | |
|---|---|---|---|
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |
| 0 | 0 | 2 | 2 |

| | | | | |
|---|---|----|---|---|
| 0 | 0 | -1 | 0 | 1 |
| 0 | 0 | -4 | 0 | 2 |
| 0 | 0 | -2 | 0 | 1 |
| 0 | 0 | 2 | 2 | |
| 0 | 0 | 2 | 2 | |

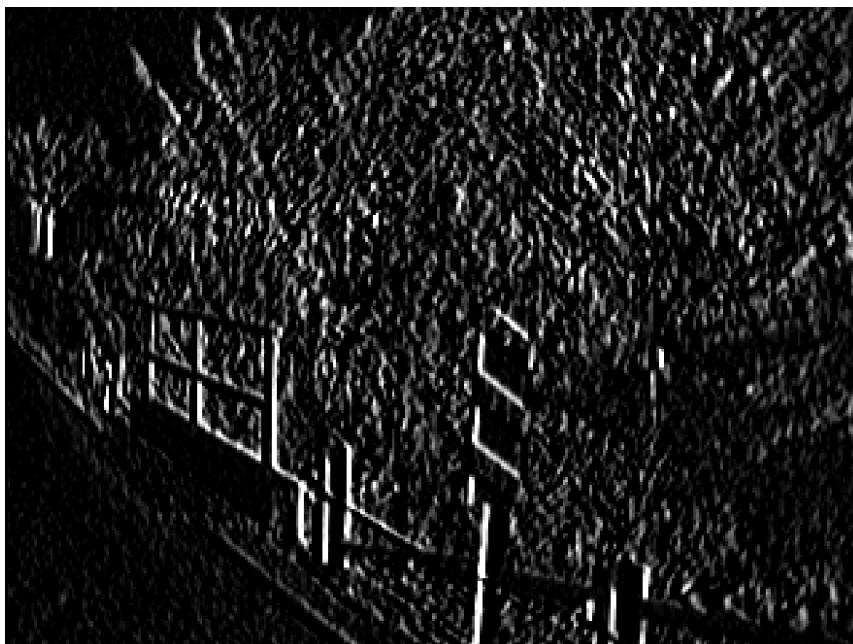
| | | | |
|---|---|---|----|
| 0 | 6 | 6 | -6 |
| | | | |
| | | | |
| | | | |

border
artifact

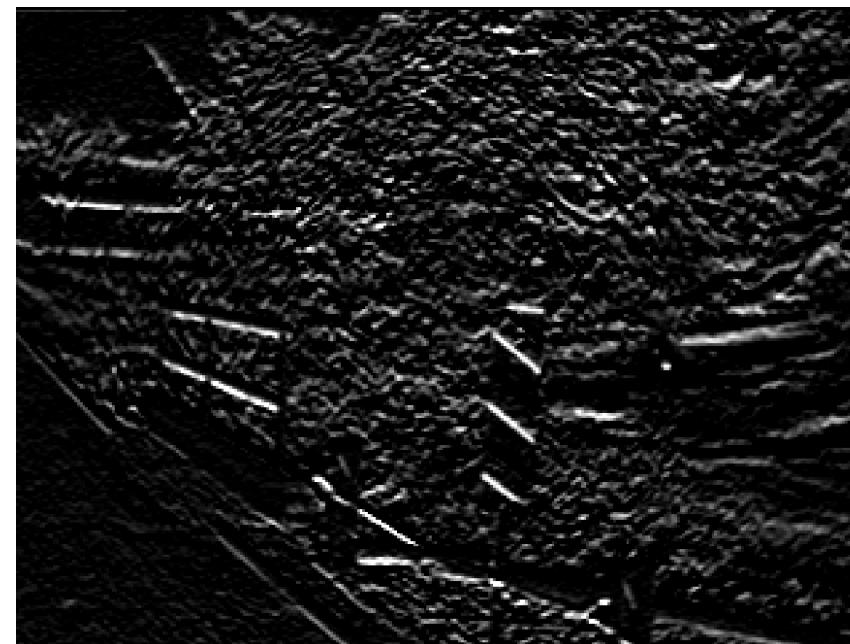
I

I'

Sobel edge detection: example



- Vertical edges



- Horizontal edges

Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

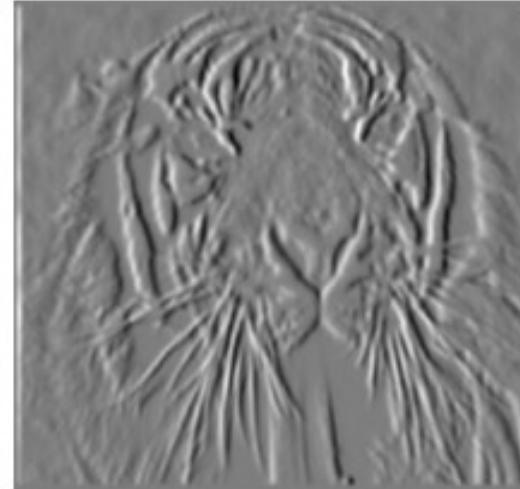
To implement above as convolution, what would be the associated filter?

Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

| | |
|----|---|
| -1 | 1 |
|----|---|

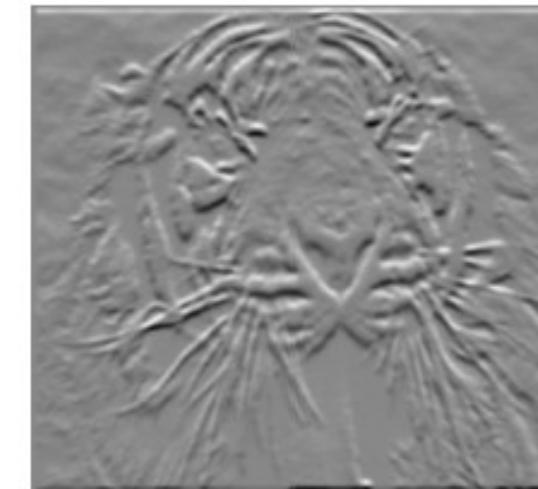


$$\frac{\partial f(x, y)}{\partial y}$$

-1 or 1

| | |
|----|---|
| -1 | 1 |
|----|---|

| | |
|---|----|
| 1 | -1 |
|---|----|



- Which shows changes with respect to x?

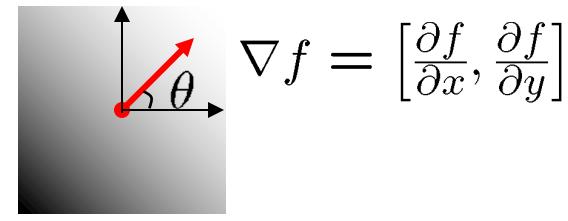
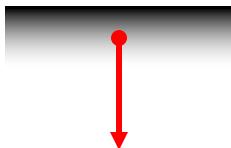
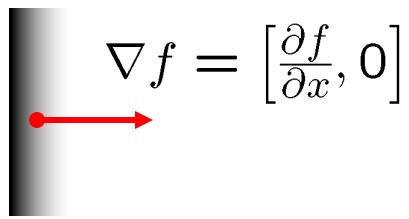
- (showing filters for correlation)

Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity

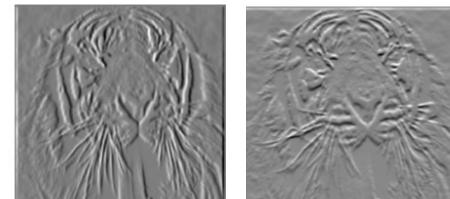


- The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- The **edge strength** is given by the gradient magnitude

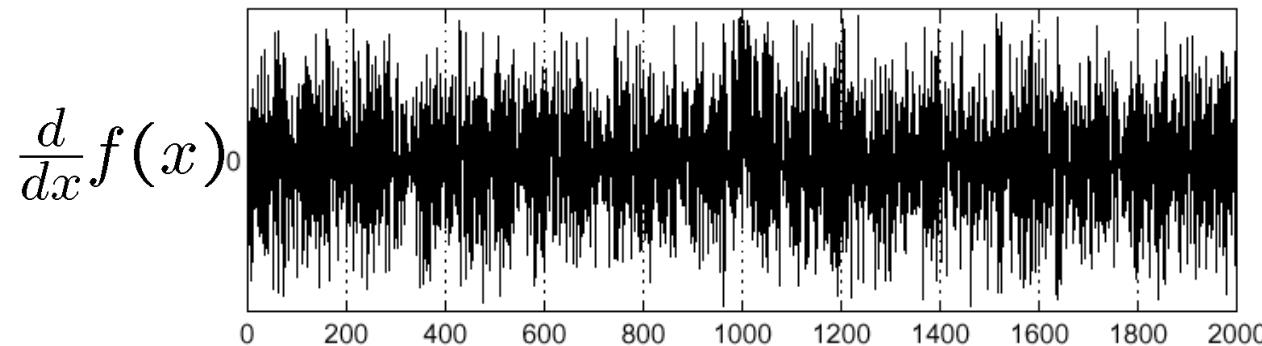
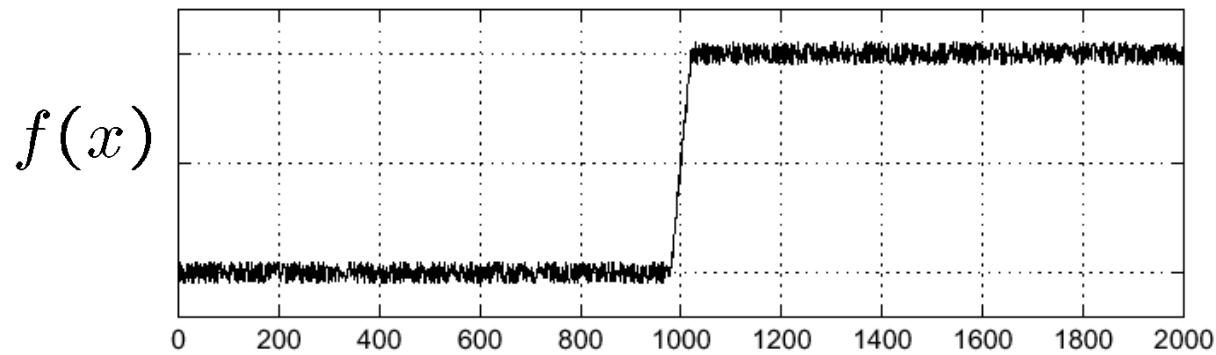
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



Effects of noise

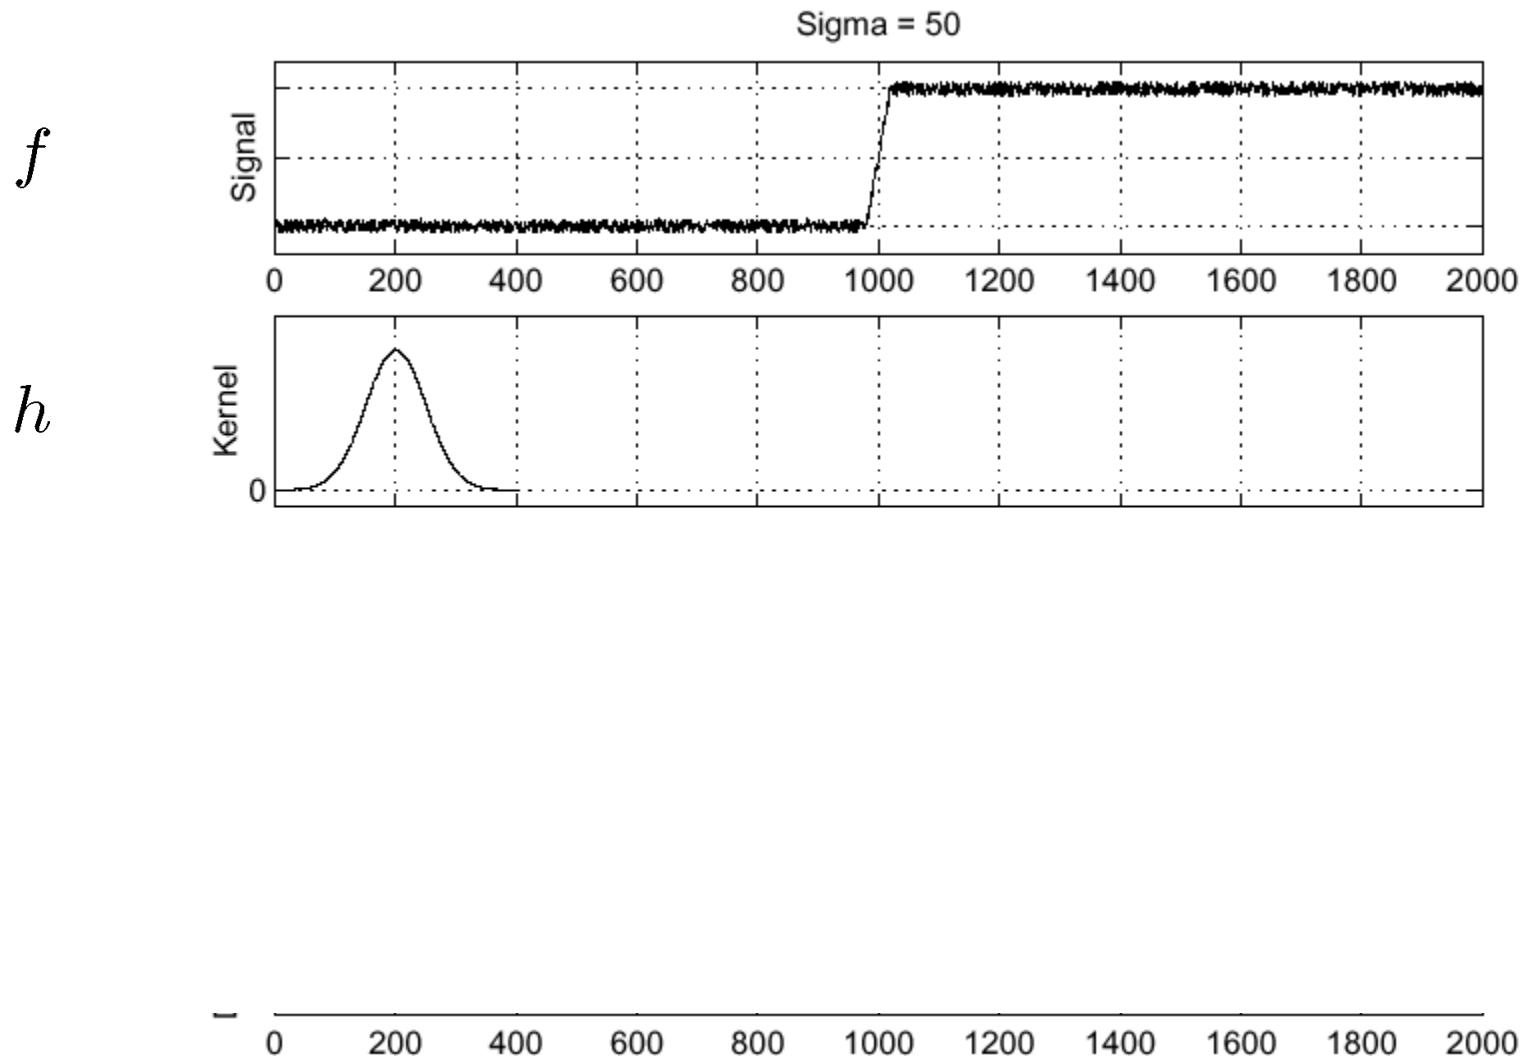
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



- Where is the edge?

Solution: smooth first



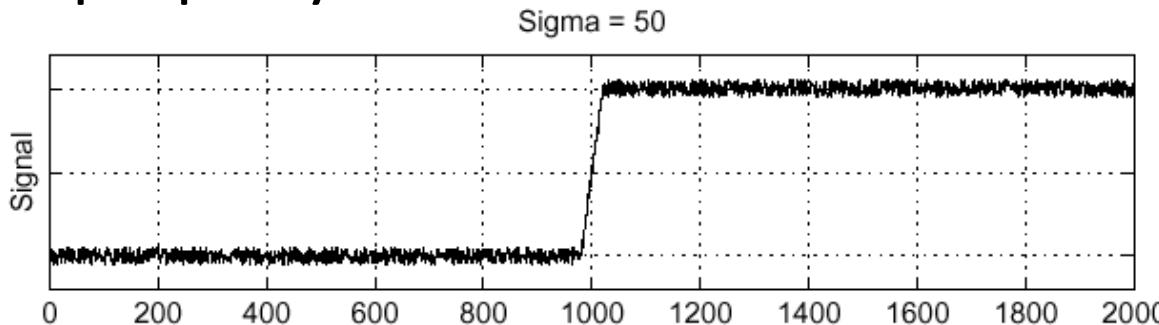
- Where is the edge?
- Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution

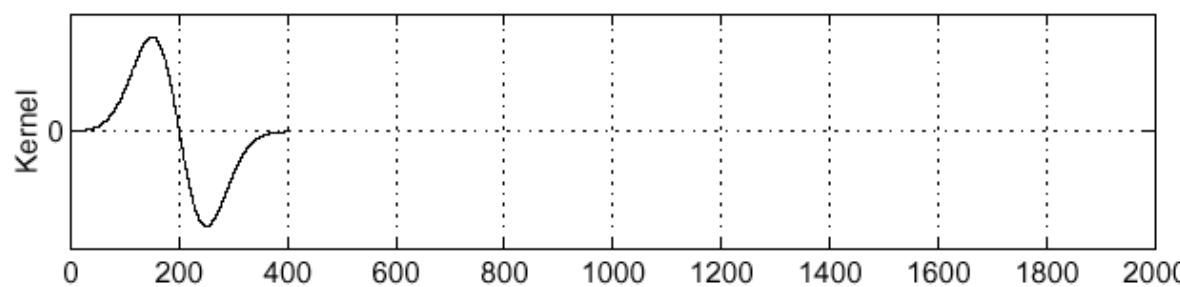
$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

Differentiation property of convolution.

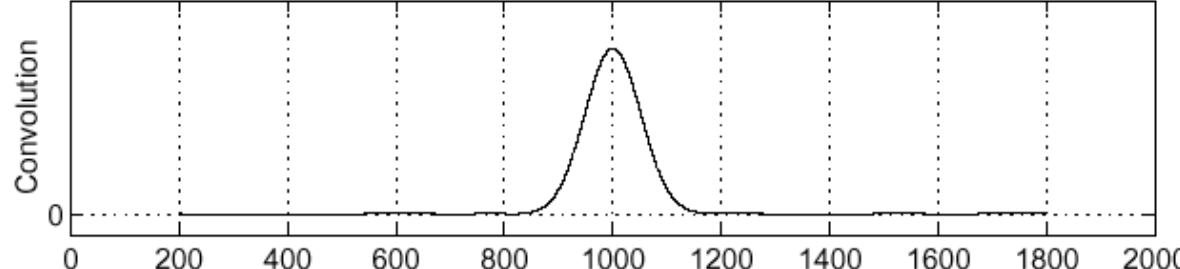
f



$\frac{\partial}{\partial x}h$



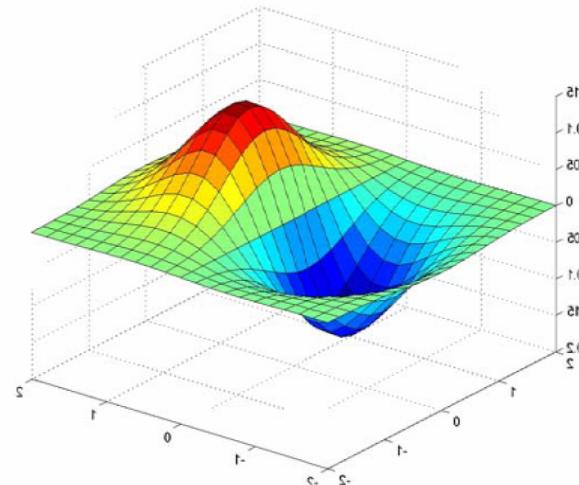
$(\frac{\partial}{\partial x}h) \star f$



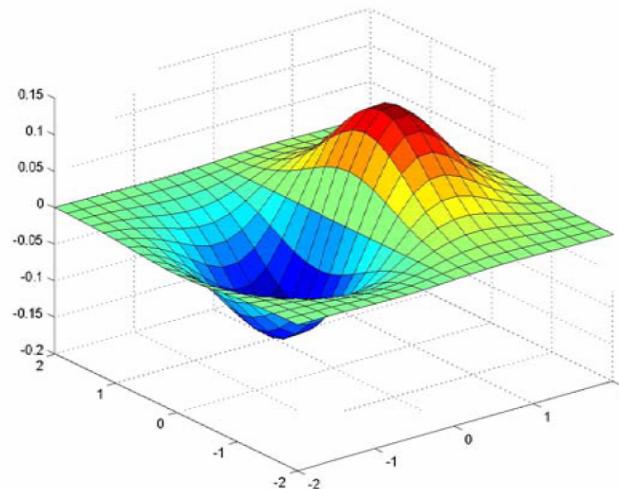
Derivative of Gaussian filters

$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

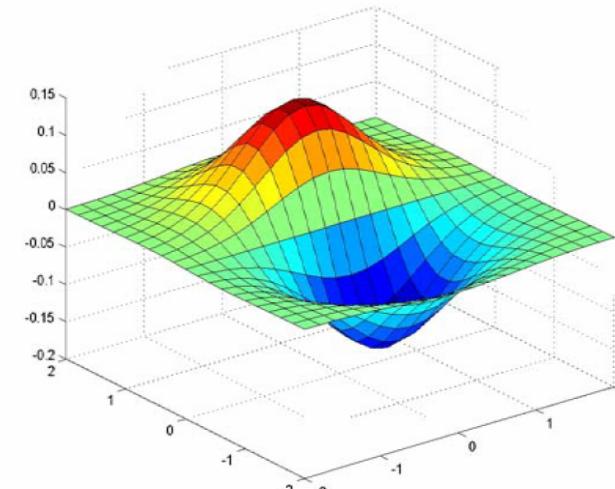
$$\begin{bmatrix} \bullet & 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ \bullet & 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ \bullet & 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ \bullet & 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ \bullet & 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \quad \otimes \quad \begin{bmatrix} 1 & -1 \end{bmatrix}$$



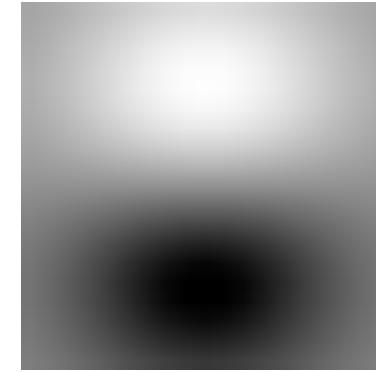
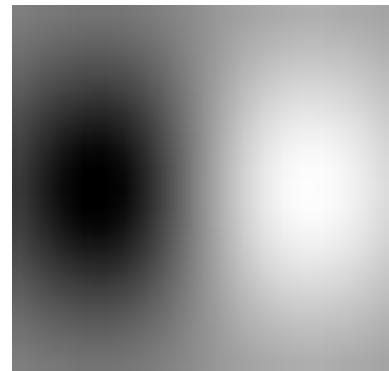
Derivative of Gaussian filters



•x-direction



•y-direction



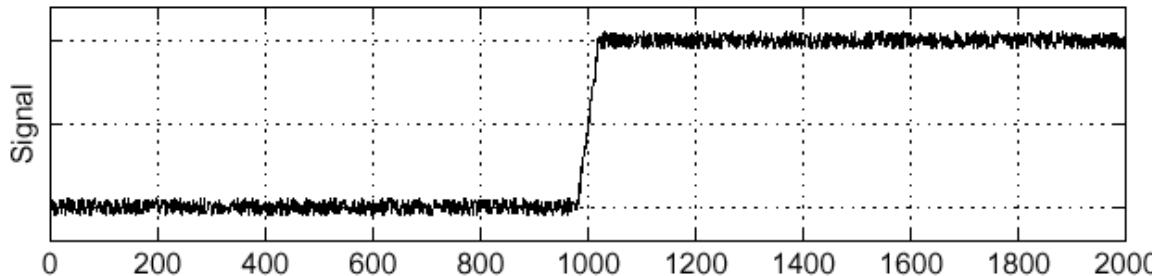
Laplacian of Gaussian

Consider

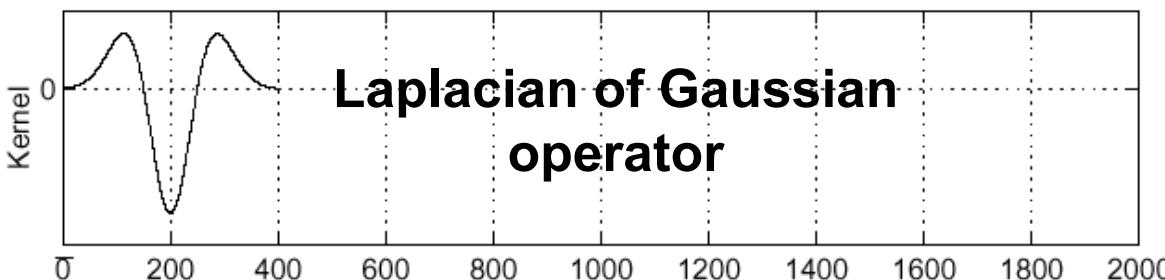
$$\frac{\partial^2}{\partial x^2}(h \star f)$$

Sigma = 50

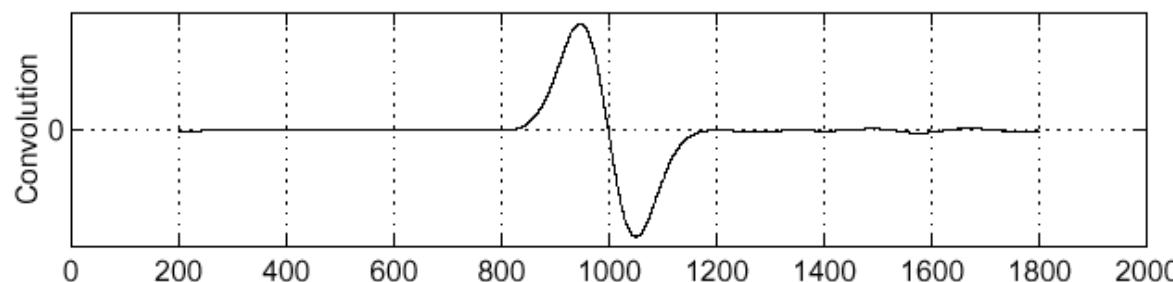
f



$$\frac{\partial^2}{\partial x^2} h$$



$$(\frac{\partial^2}{\partial x^2} h) \star f$$

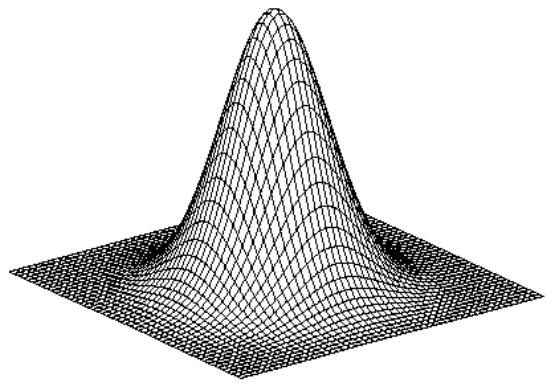


- Where is the edge?

- Zero-crossings of bottom graph

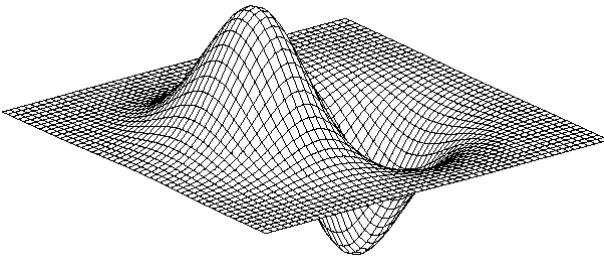
• Slide credit: Steve Seitz

2D edge detection filters



Gaussian

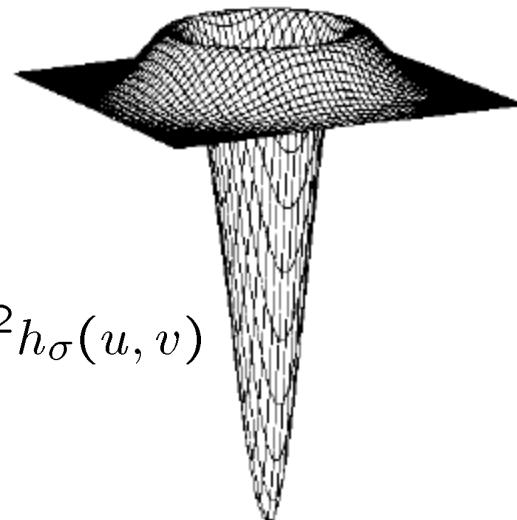
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



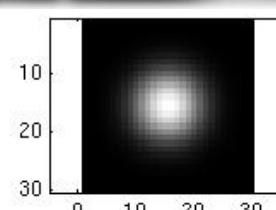
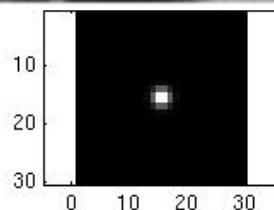
$$\nabla^2 h_\sigma(u, v)$$

- ∇^2 is the Laplacian operator:

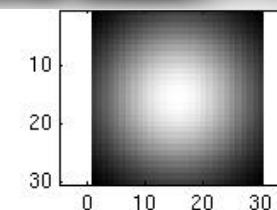
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Smoothing with a Gaussian

- Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



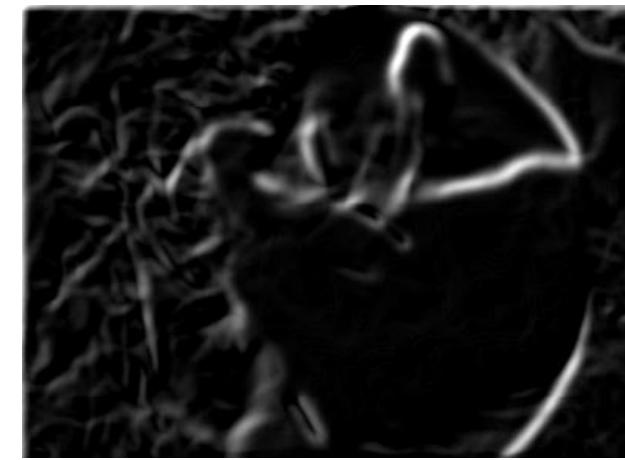
• ...



Effect of σ on derivatives



$\sigma = 1$ pixel

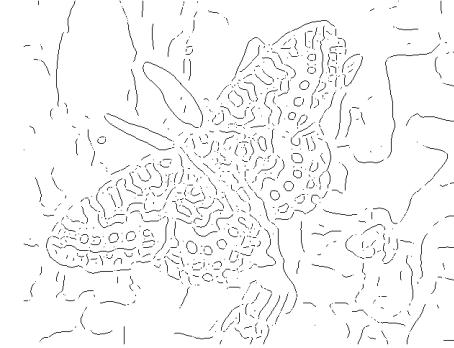


$\sigma = 3$ pixels

- The apparent structures differ depending on Gaussian's scale parameter.
- Larger values: larger scale edges detected
- Smaller values: finer features detected



Gradients -> edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from filter output are actually edges vs. noise

- Threshold, Thin

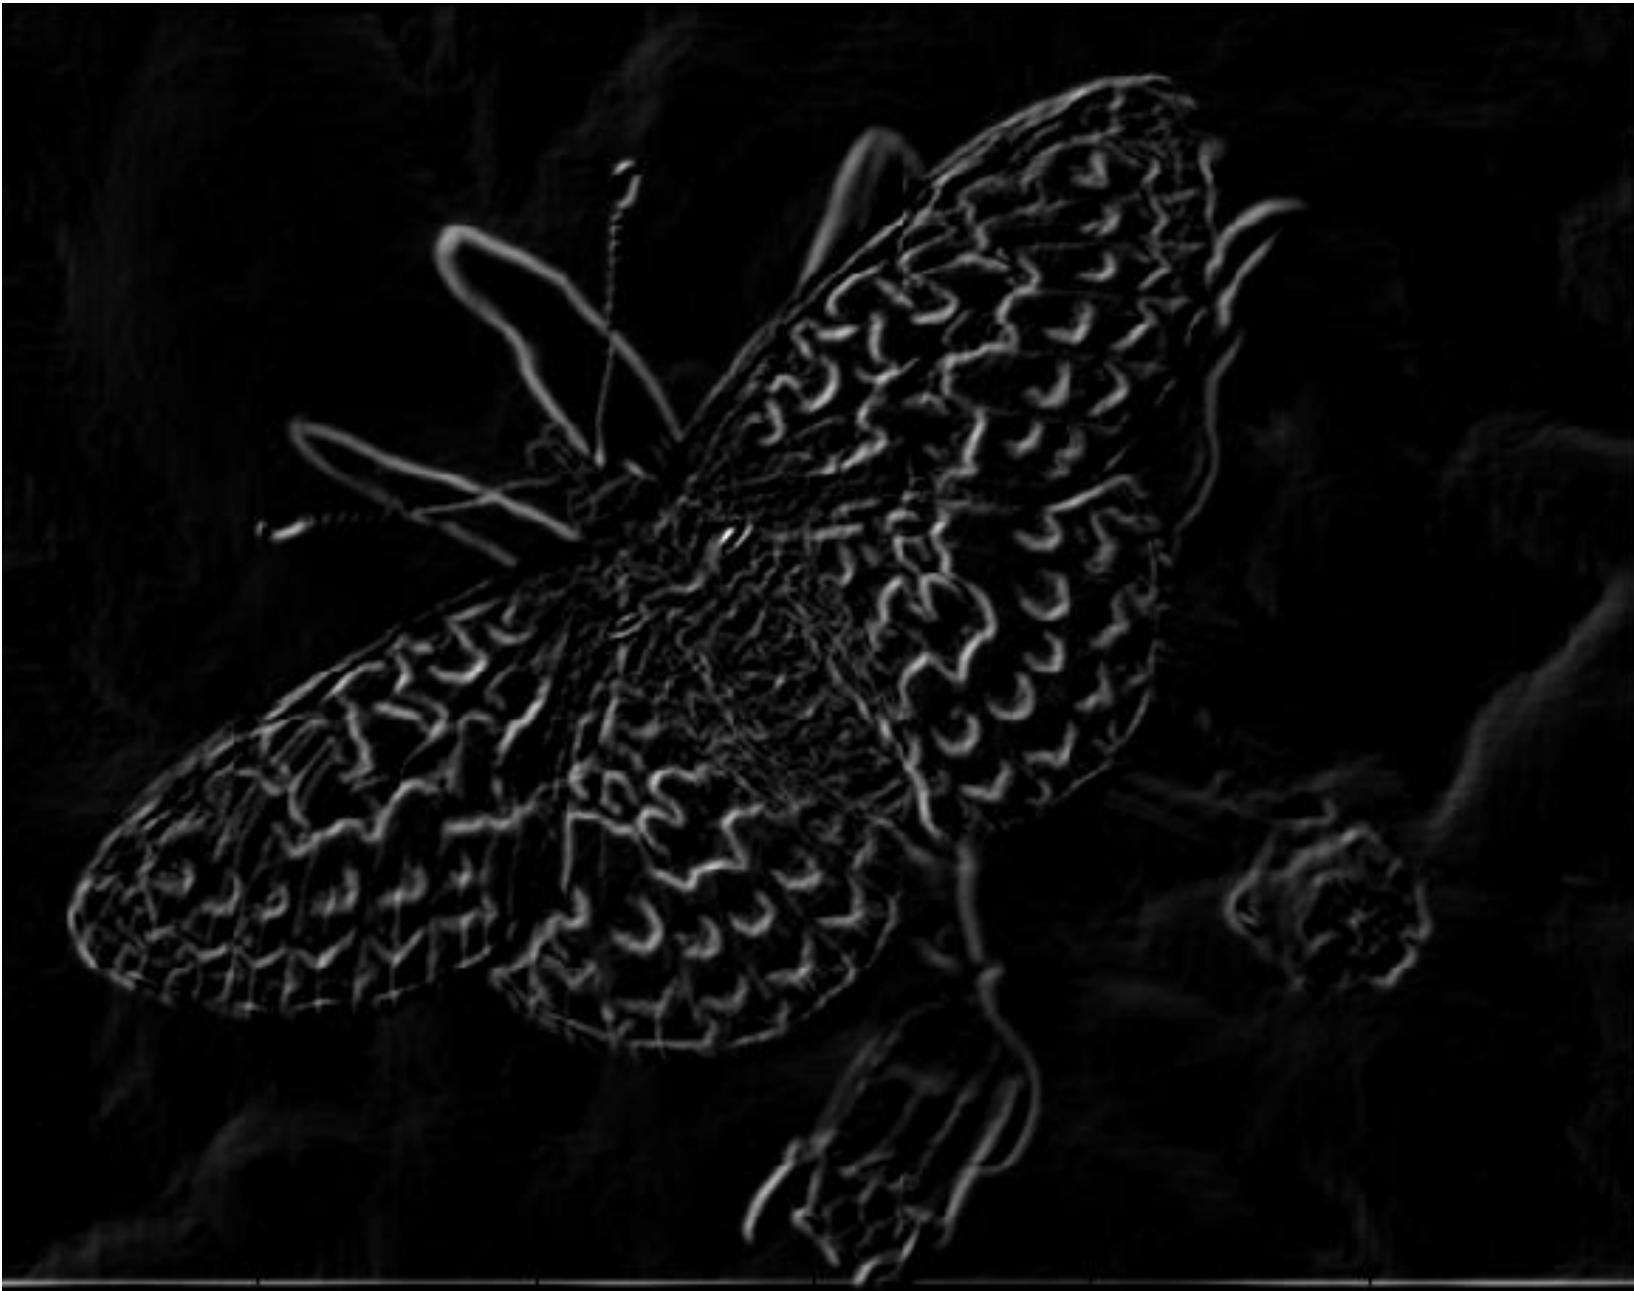
Thresholding

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`

The Canny edge detector



original image (Lena)

The Canny edge detector



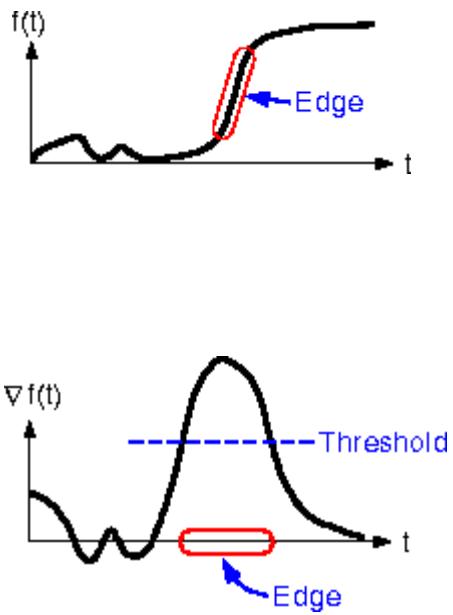
- norm of the gradient

The Canny edge detector



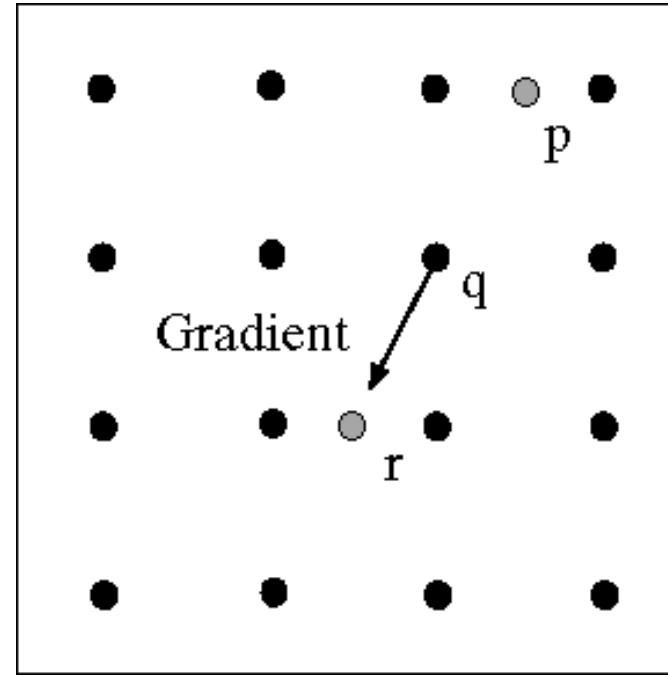
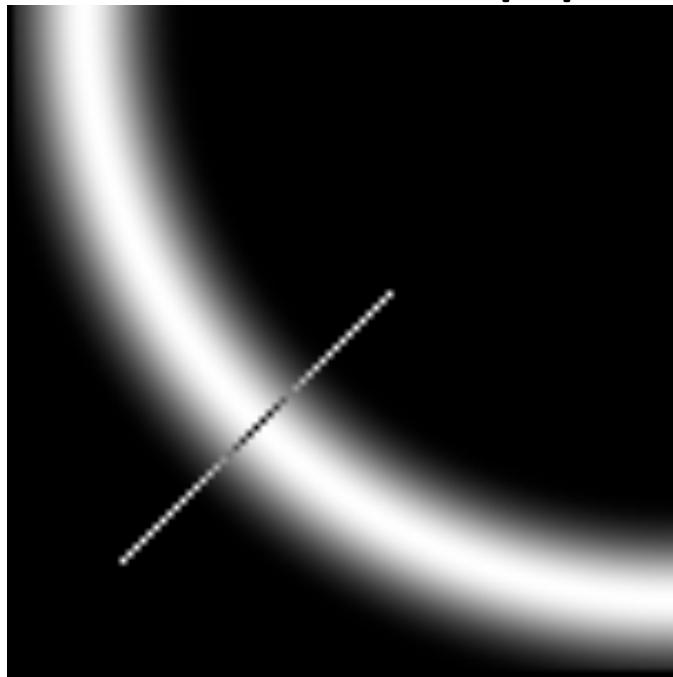
- thresholding

The Canny edge detector



- How to turn these thick regions of the gradient into curves?

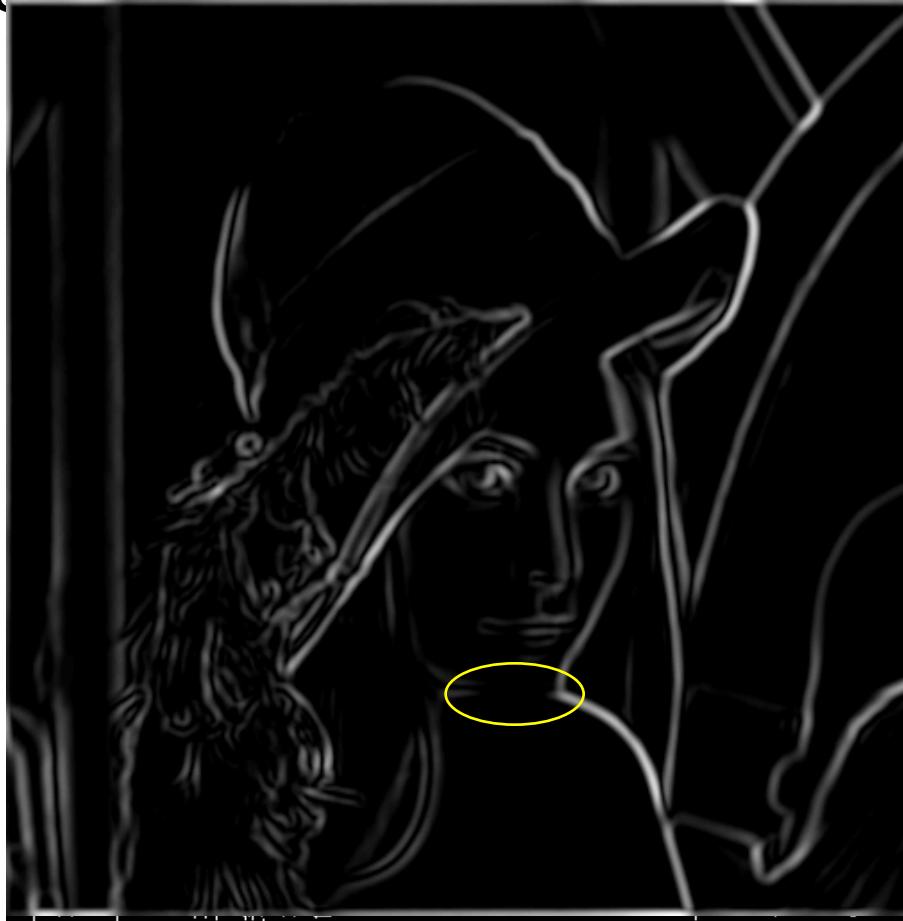
Non-maximum suppression



Check if pixel is local maximum along gradient direction, select single max across width of the edge

- requires checking interpolated pixels p and r

The Canny edge detector



- thinning
- (non-maximum suppression)

- Problem:
pixels along
this edge
didn't survive
the
thresholding

Hysteresis thresholding

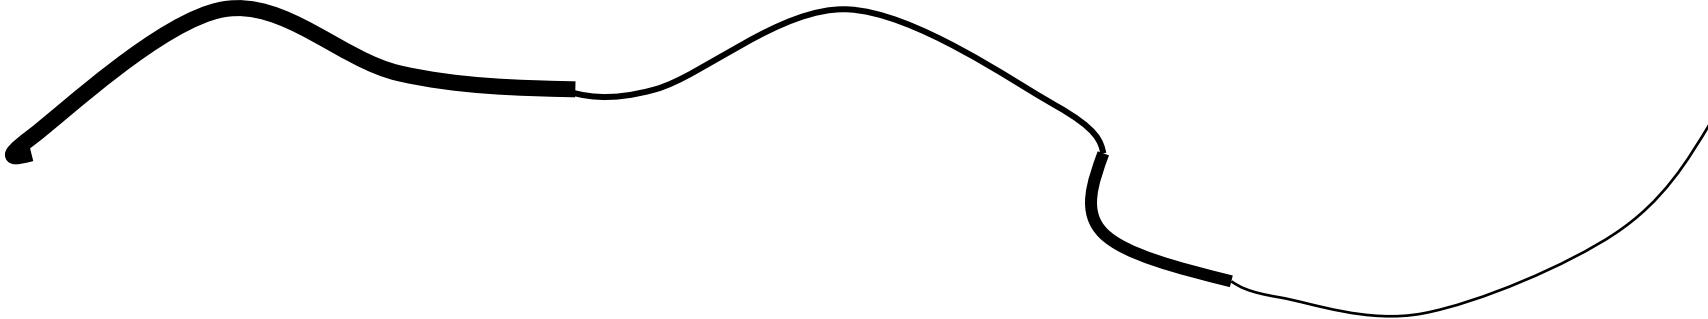
- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



•Credit: James Hays

Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



Final Canny Edges



Credit: James

Recap: Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`

Mask properties

- Smoothing

- Values positive
- Sum to 1 → constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components; “low-pass” filter

- Derivatives

- Opposite signs used to get high response in regions of high contrast
- Sum to 0 → no response in constant regions
- High absolute value at points of high contrast

Non-linear Filter

- Median Filter
- Bilinear Filter

Nonlinear Image Filtering

Disclaimer: Many of the slides used here are obtained from online resources (including many open lecture materials given at MIT, etc.) without due acknowledgement. They are used here for the sole purpose of classroom teaching. All the credits and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any, any purposes other than helping you study the course of ENGN4528/6528.

Median Filter (a nonlinear filter)

Median filter

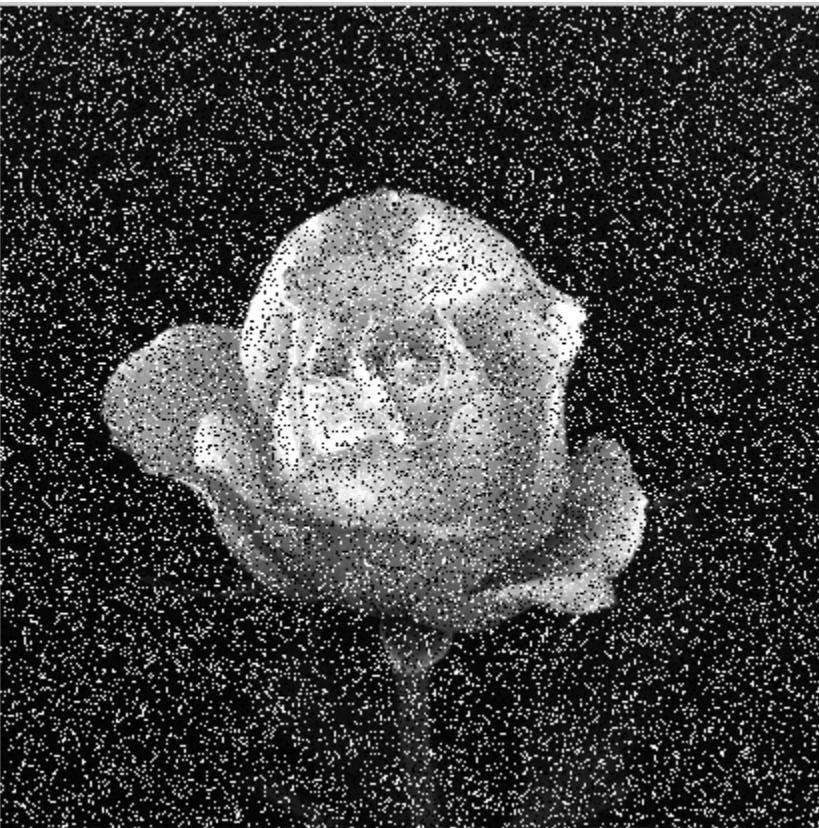
- The median filter is a sliding-window based spatial filter.
- It replaces the value of the center pixel with the median of the intensity values in the neighborhood of that pixel.
- Median filtering is a nonlinear operation often used in image processing to reduce "salt and pepper" noise. A median filter is more effective than convolution when the goal is to simultaneously reduce noise and preserve edges.
- **Median filters** are particularly effective in the presence of *impulse noise*, also called '**salt – and – pepper**' noise because of its appearance as white and black dots superimposed on an image.
- For every pixel, a 3x3 neighborhood with the pixel as center is considered. In **median filtering**, the value of the pixel is replaced by the median of the pixel values in the 3x3 neighborhood.

Questions

What is the value of the yellow box after 3x3 median filtering?

| | | | | |
|---|---|---|---|---|
| 1 | 2 | 2 | 2 | 3 |
| 2 | 5 | 2 | 4 | 3 |
| 2 | 2 | 2 | 2 | 3 |
| 4 | 3 | 7 | 5 | 3 |
| 3 | 3 | 3 | 1 | 1 |

Example: Median filter result



Bilateral filter (a nonlinear filter)

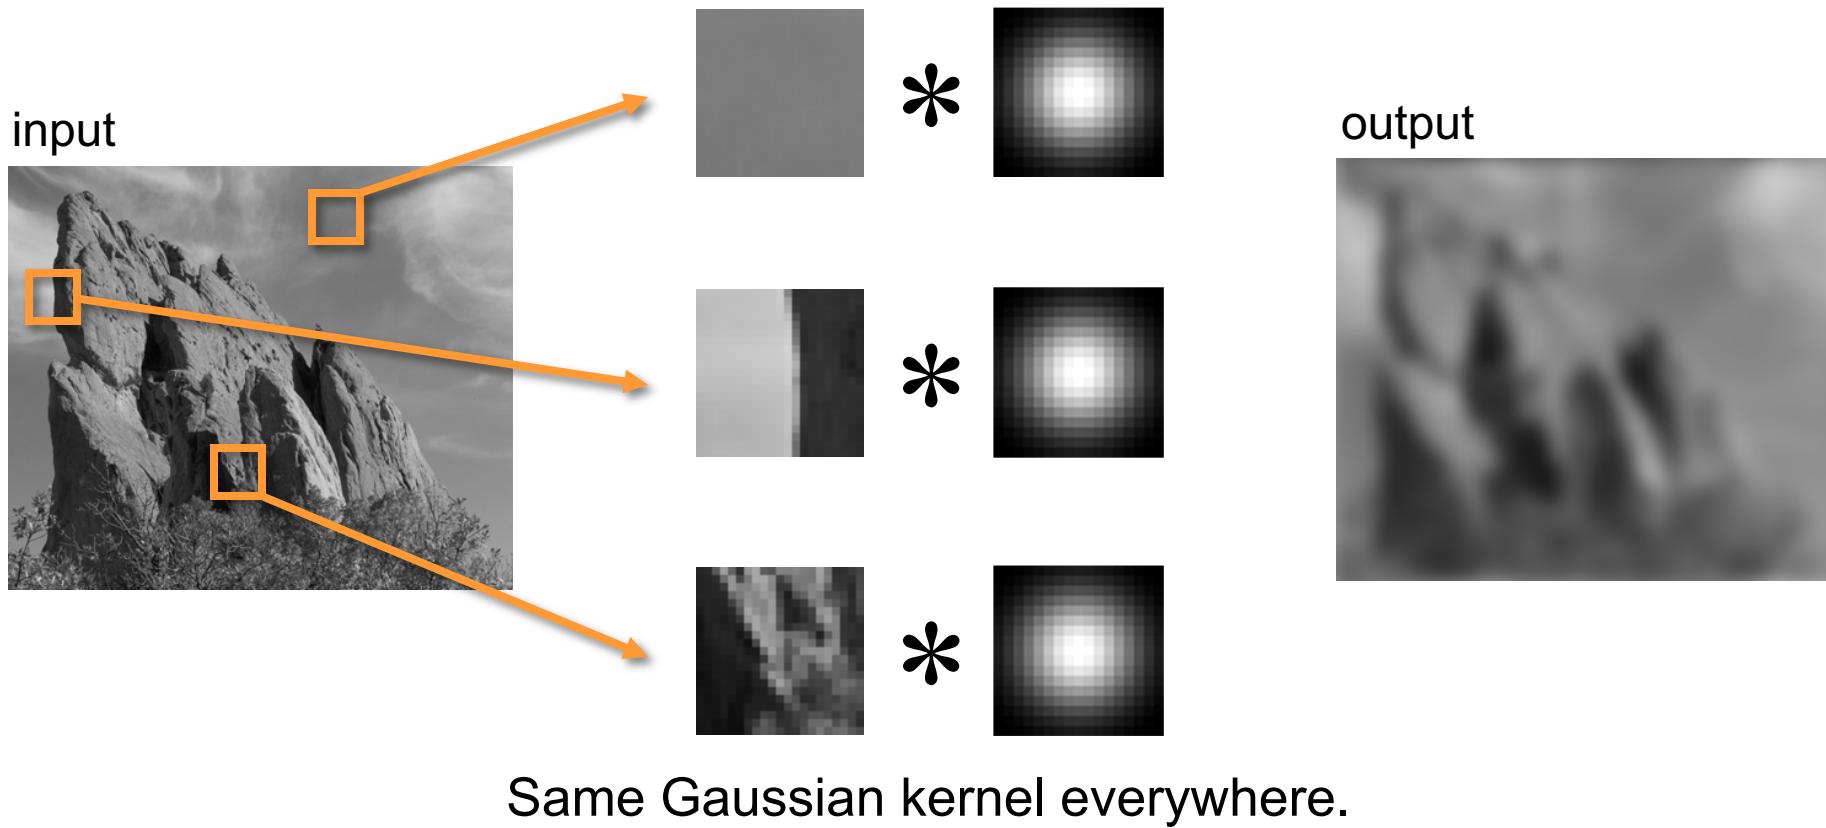
How to smooth an image without causing excessive edge-blur
?

Bilateral Filter (Tomasi et al.1998)

- It smooths regions while preserving edges



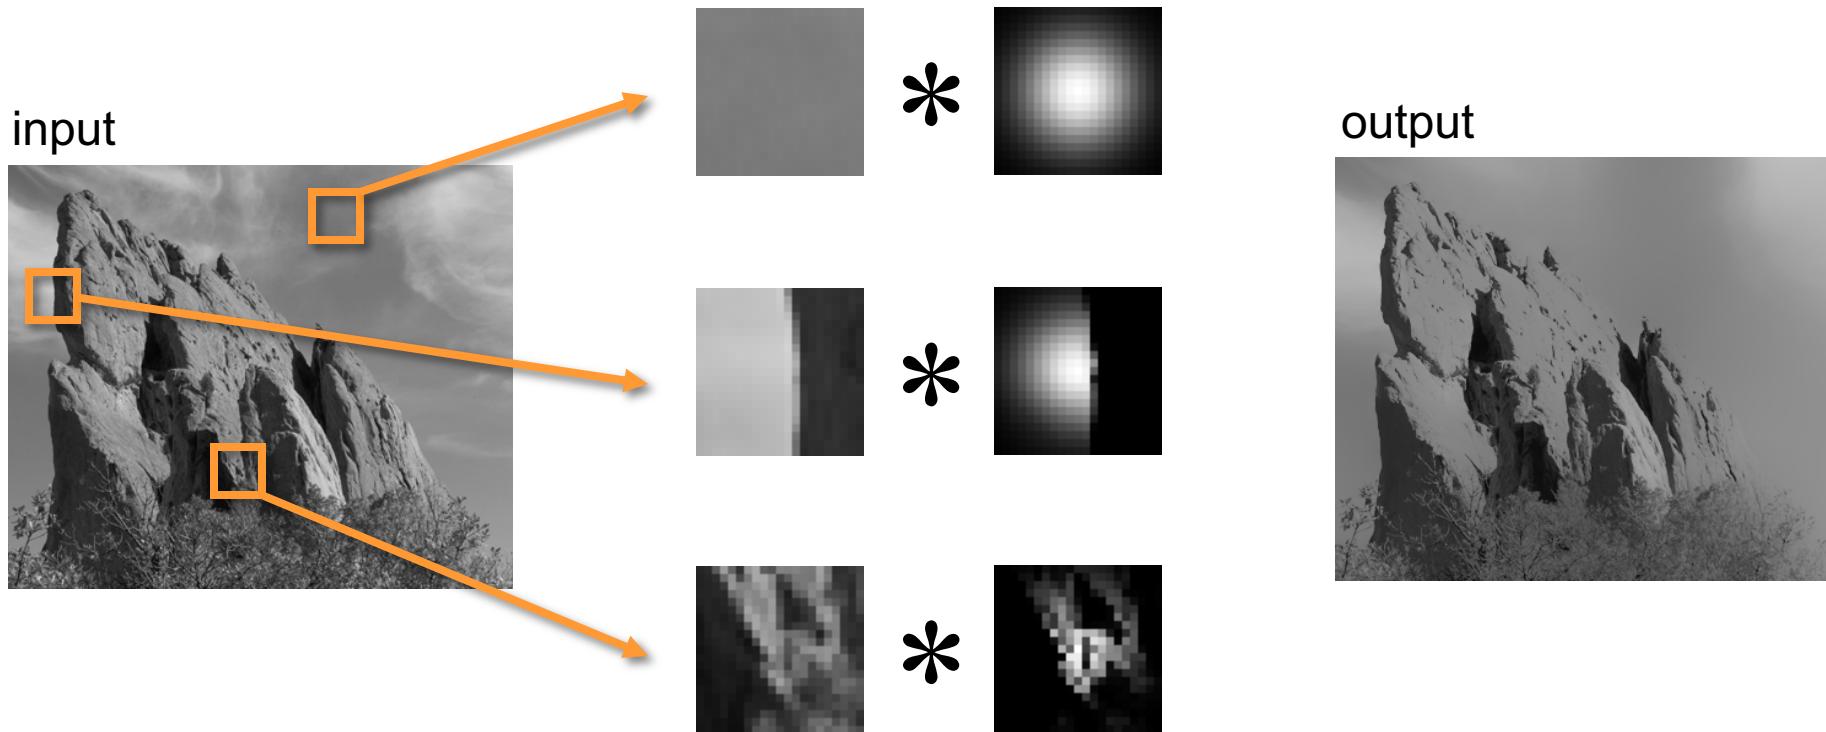
Blur Comes from Averaging across Edges



Bilateral Filter

No Averaging across Edges

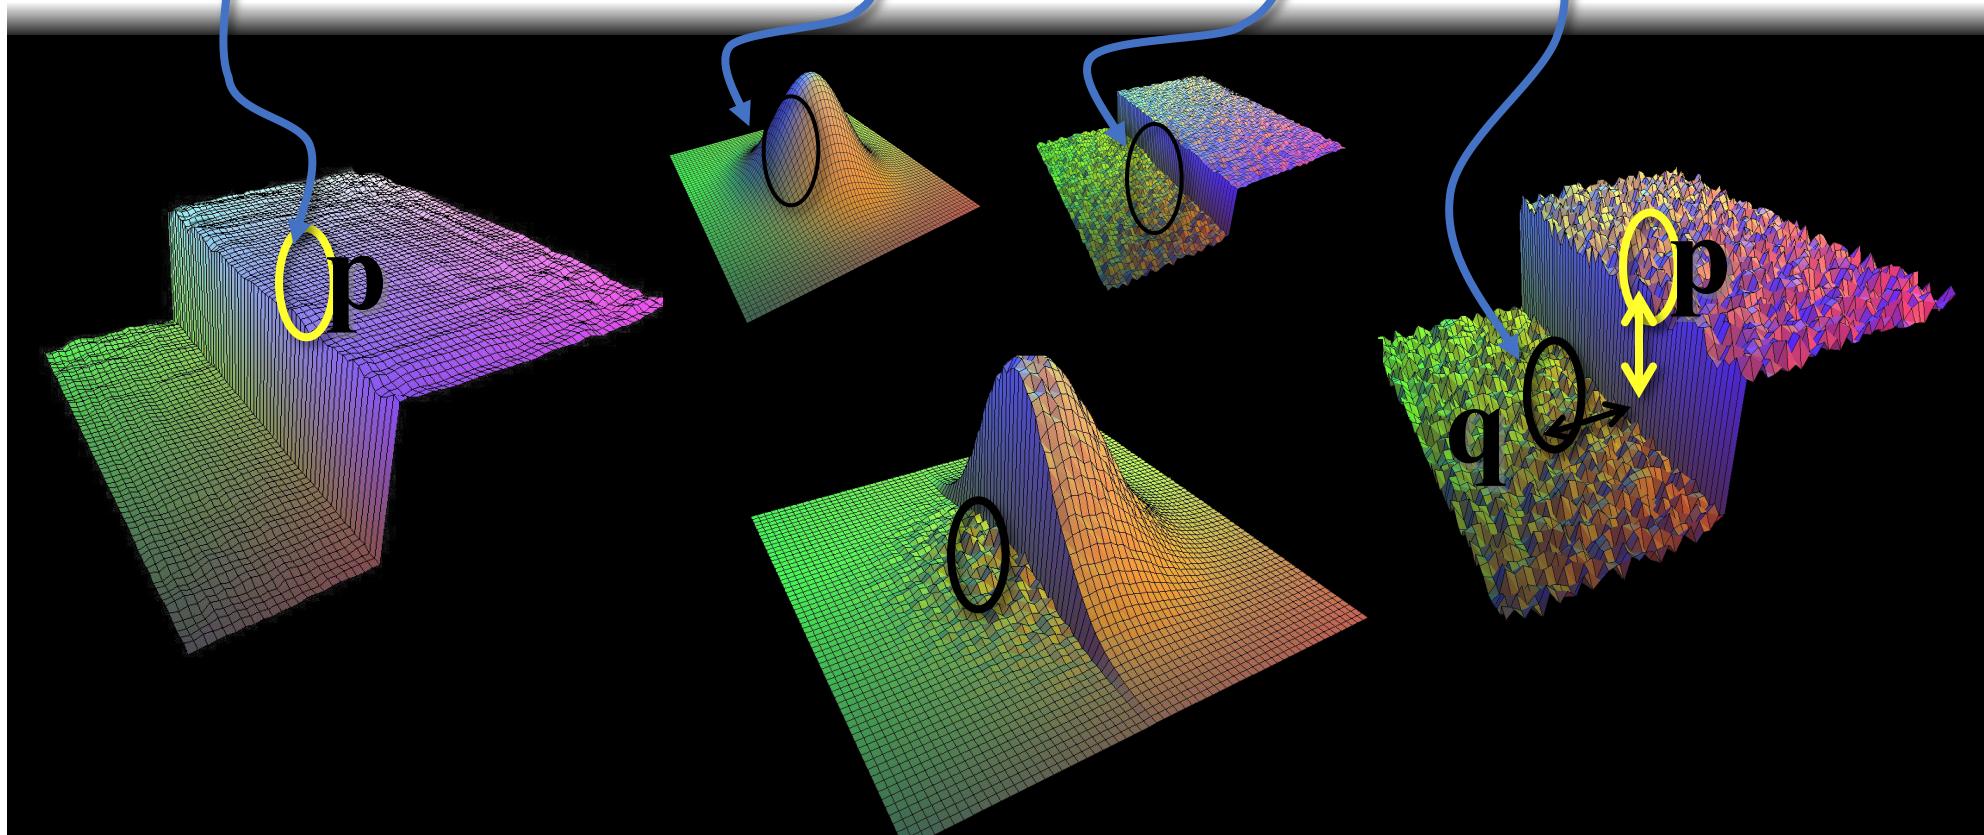
[Aurich 95, Smith 97, Tomasi 98]



The kernel shape depends on the image content.

Bilateral Filter on a Height Field

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$



reproduced
from [Durand 02]

Bilateral Filter on a Height Field

- Domain filter: G_{σ_s}
- Range filter: G_{σ_r}

Varying the Space Parameter



input

$\sigma_s = 2$



$\sigma_s = 6$



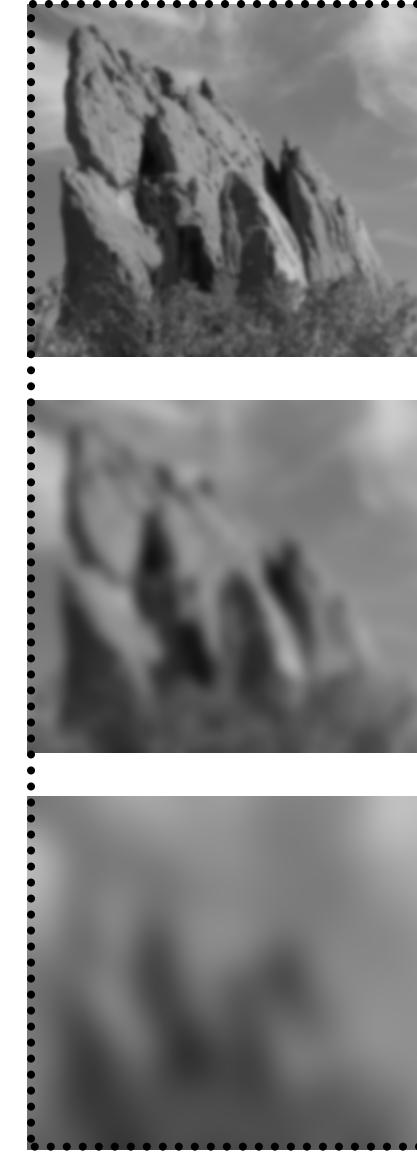
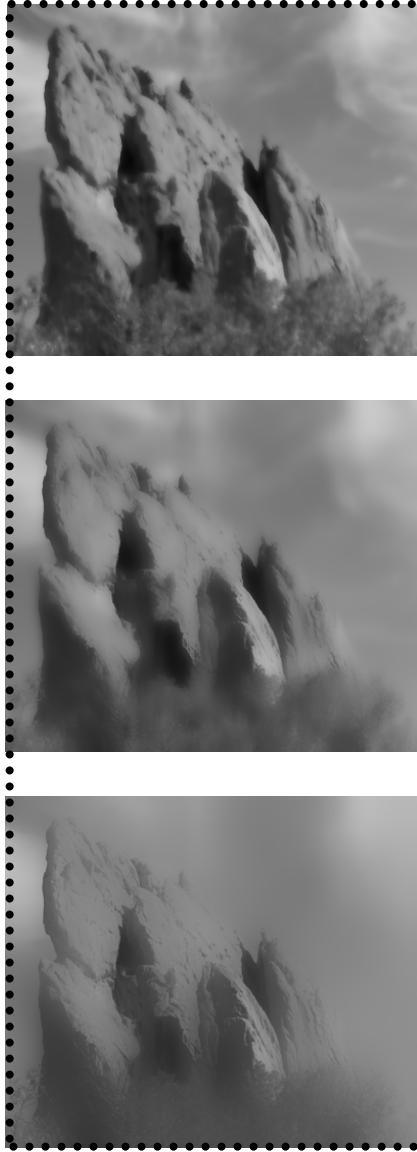
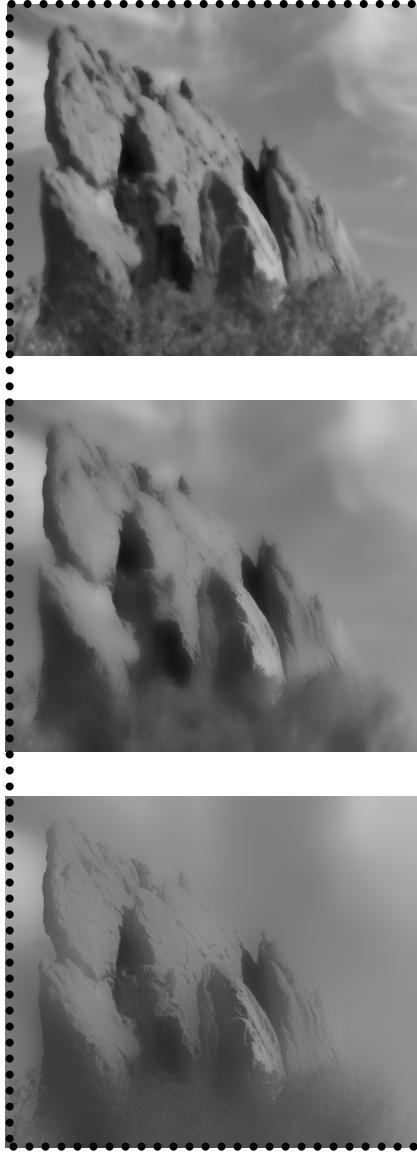
$\sigma_s = 18$



$\sigma_r = 0.1$

$\sigma_r = 0.25$

$\sigma_r = \infty$
(Gaussian blur)



input



$$\sigma_s = 2$$



$$\sigma_s = 6$$



$\sigma_s = 18$ 

Readings

- Computer Vision: Algorithms and Applications, Chapter 3.1, 3.2 and 3.3