

Face

Disclaimer: Many of the slides used here are obtained from online resources (including many open lecture materials) without specific acknowledgements. They are used here for the sole purpose of classroom learning. All copyrights belong to the original authors or publishers. You should not copy it, redistribute it, put it online, or use it for any purposes other than help you learning
ENGN4528/6528.

Announcements

- Survey (47 students participated in the survey)
 - (Lab-session-week-1, Lab-session-week-2). 9
 - (Lab-session-week-2, Lab-session-week-3). 38
- Lab-Assignment-2 has been released.
Due date: May 2nd, 23:59pm
- Tutorial on Lab assignment 2: Wed. 4:00pm

Announcement

- Mid-term Quiz

Unit Codes

You can search for multiple units by putting a comma between your unit codes.

SEARCH : FINAL TIMETABLE

Displaying records 1 to 1 of 1

Exam Code	Exam Title	Date	Time	Writing Time (minutes)	Reading Time (minutes)	Venue	Exam Conditions
ENGN6528_Semester 1	Computer Vision	Tuesday 20/04/2021	12:00pm	90	15	Online Exam - Wattle	Any material ,

Displaying records 1 to 1 of 1

- **90 minutes** in total including reading, downloading question paper, and uploading results.
- Questions are expected to be finished in one hour.
- **Open book**, download and upload papers through wattle. **You need to work on questions by yourself.**
- We are going to have multiple versions of the quiz papers. Students will be randomly assigned to different groups for exam.

Announcements

- Type of questions (see sampled exam paper)
 - Basic Calculation
 - Basic Concept
 - Algorithm design

Announcements

- We will provide instructions two weeks before the exam.

ENGN6528_Sem1_2021

- Participants
- Grades
- General
- Links for Lectures and Labs
- Exams
- Course Supp reading
- Lectures-S1-2021
- Tutorials-S1-2021
- Practice Exercises-S1-2021
- Lab-Assignments-S1-2021
- Lecture Notes

%%%%%
Survey for the lab arrangement.

Please fill in this lab time preference survey

Exams

Sample Exam Paper 337.7KB PDF document Uploaded 20/03/21, 18:10

Course Supp reading

Proof of the Convolution of Two Gaussians 108.3KB PDF document Uploaded 17/03/21, 23:50

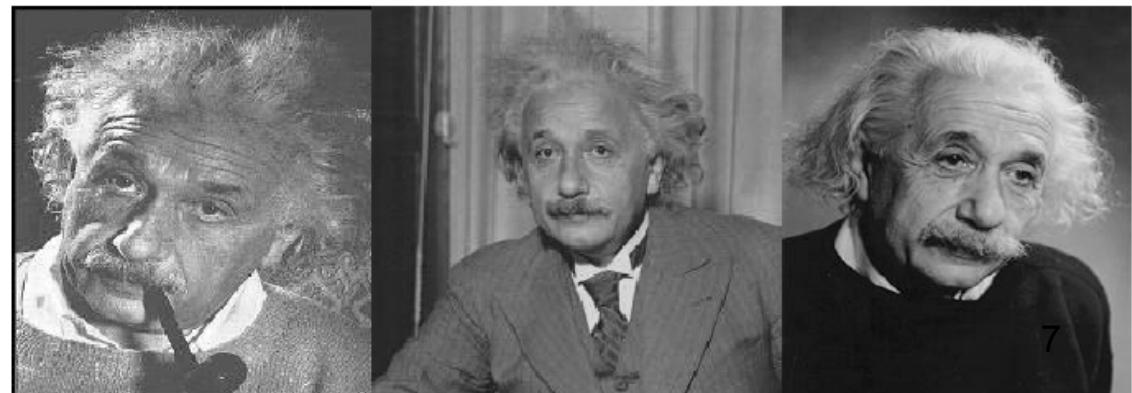
5

Review

- Face detection

Natural Face Recognition Ability

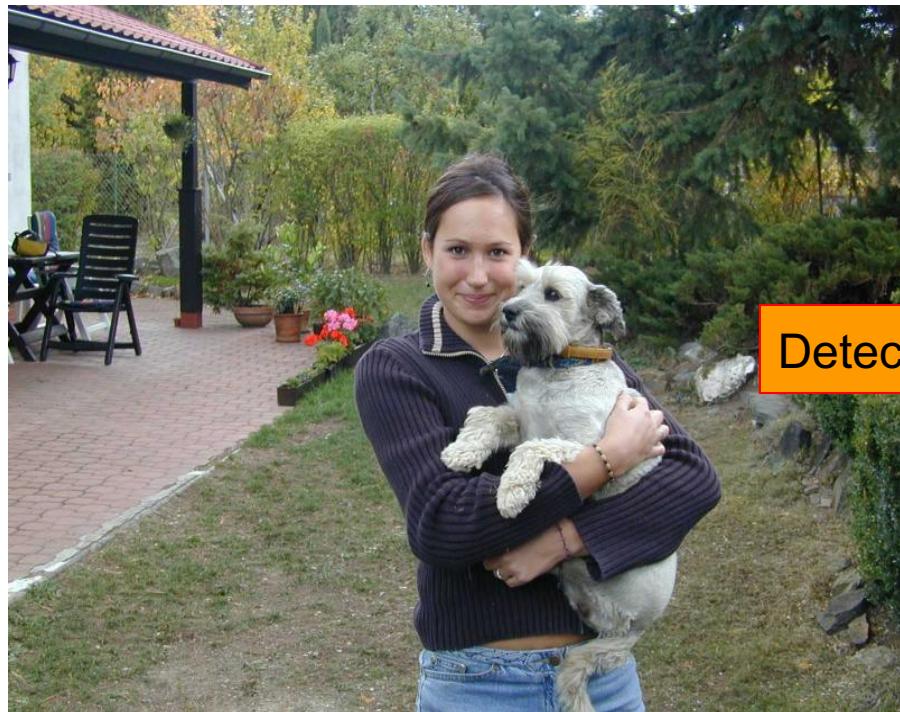
- Humans perform face recognition almost instantaneously
- Highly invariant to pose, scale, rotation, lighting changes
- Can handle partial occlusions and changes in age
- And we can do all this for faces of several thousand individuals



Today's lecture

- PCA and EigenFace representation [Turk & Pentland]
- Viola-Jones face detection: [Viola & Jones]

Face detection versus recognition



Detection



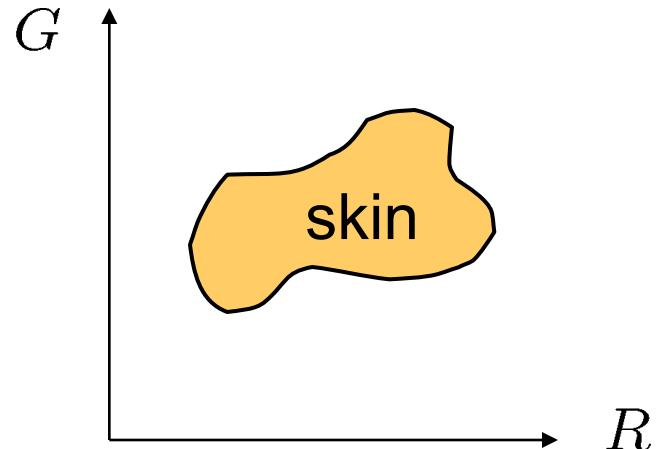
Recognition

“Sally”

Common face detection methods

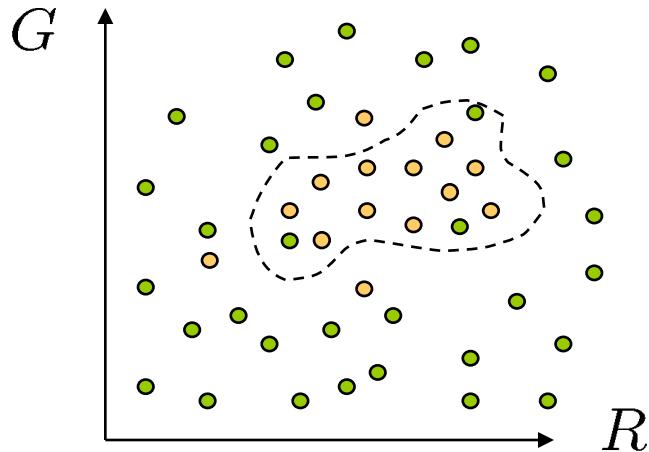
- Colour segmentation
- Template matching
- Eigen Face: Face space representation
- Viola-Jones Adaboost face detection(Part I)
- SIFT BoW detection
- Deep-Net detection
- ...

Skin color detection



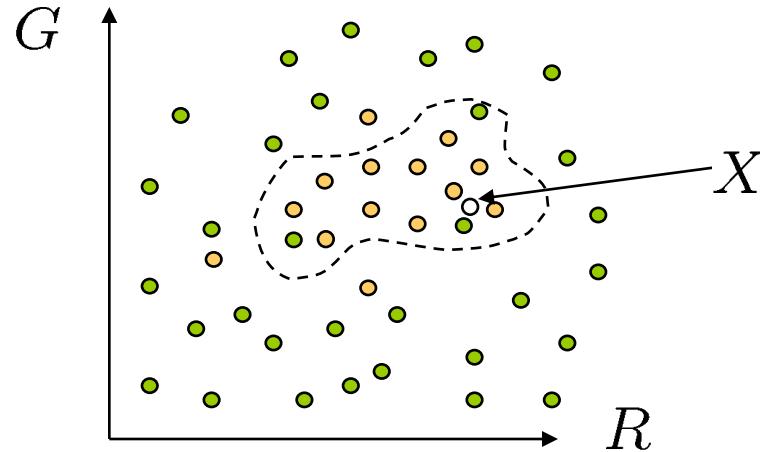
- Skin pixels have a distinctive range of colors
 - Corresponds to region(s) in RGB color space
- Skin classifier
 - A pixel $X = (R, G, B)$ is skin if it is in the skin (color) region
 - How to find this region?

Skin colour detection



- **Learn the skin region from examples**
 - Manually label skin/non pixels in one or more “training images”
 - Plot the training data in RGB space
 - skin pixels shown in orange, non-skin pixels shown in green
 - some skin pixels may be outside the region, non-skin pixels inside.

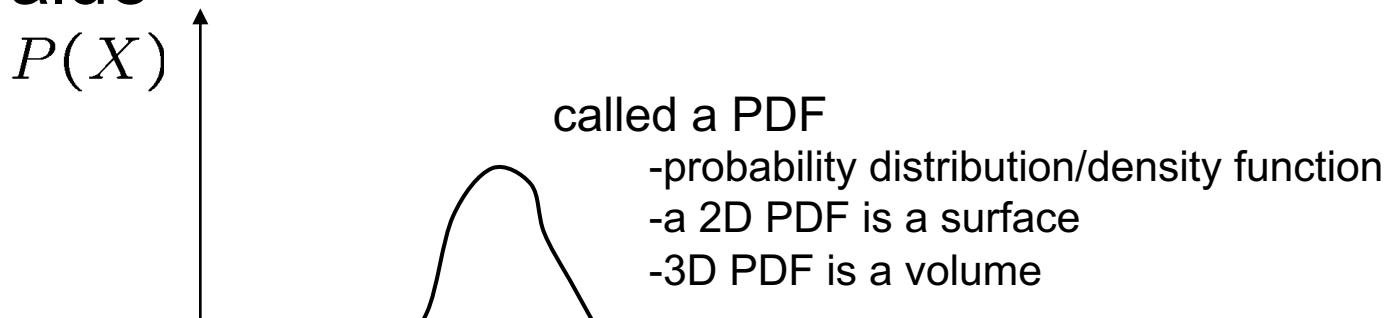
Skin classifier



- Given $X = (R, G, B)$: how to determine if it is skin or not?
 - Nearest neighbor
 - find labeled pixel closest to X
 - Find plane/curve that separates the two classes
 - popular approach: Support Vector Machines (SVM)
 - Data modeling
 - fit a probability density/distribution model to each class

Probability

- X is a random variable
- $P(X)$ is the probability that X achieves a certain value



$$0 \leq P(X) \leq 1$$

$$\int_{-\infty}^{\infty} P(X)dX = 1$$

continuous X

$$\sum P(X) = 1$$

discrete X

Bayes' Law

$$P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)}$$

- In the context of our problem:

$$P(\text{skin}|R) = \frac{P(R|\text{skin}) P(\text{skin})}{P(R)}$$

↑
what we want
(posterior)

what we can measure
(likelihood) domain knowledge
(prior)

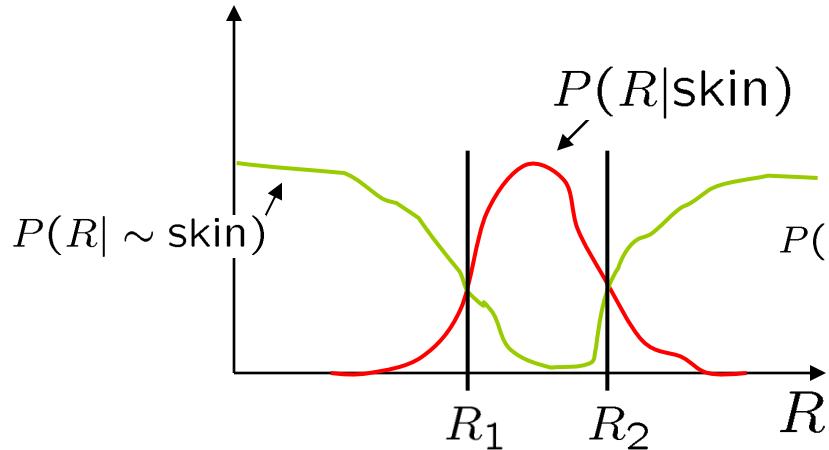
normalization term

$$P(R) = P(R|\text{skin})P(\text{skin}) + P(R|\sim \text{skin})P(\sim \text{skin})$$

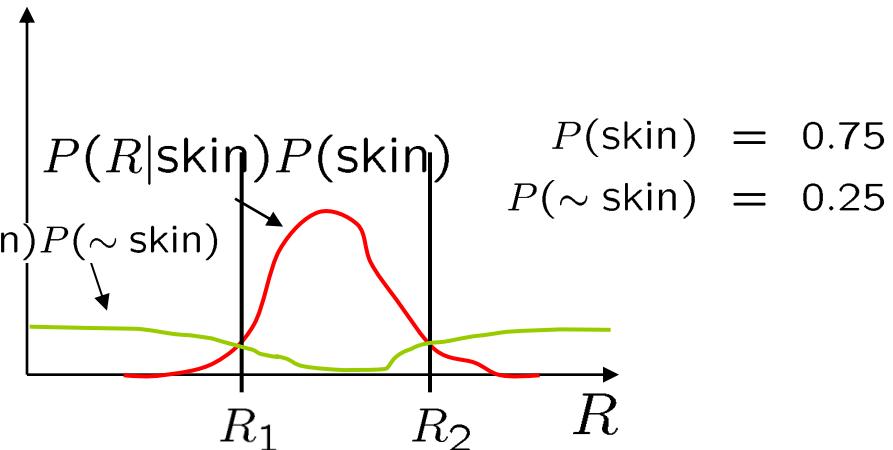
What can we use for the prior $P(\text{skin})$?

- Domain knowledge:
 - $P(\text{skin})$ may be larger if we know the image contains a person
 - For a portrait, $P(\text{skin})$ may be higher for pixels in the center
- Learn the prior from the training set. How?
 - $P(\text{skin})$ is proportion of skin pixels in training set

Bayesian estimation



likelihood



posterior (unnormalized)

- Bayesian estimation
 - Goal is to choose the label (skin or \sim skin) that maximizes the posterior \leftrightarrow minimizes probability of misclassification
 - this is called **Maximum A Posteriori (MAP) estimation**

Template Matching

Objects can be represented by
storing sample images or "templates"



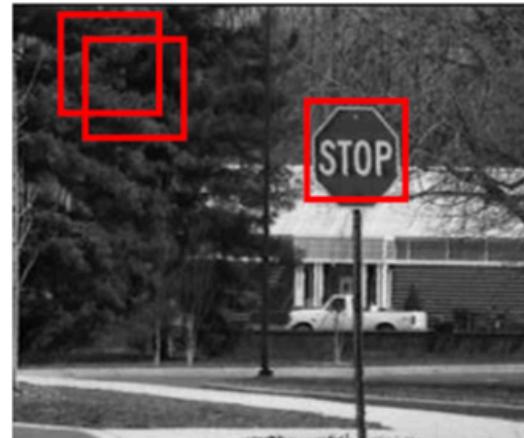
Stop sign template

Hypotheses from Template Matching

- Place the template at every location on the given image.

- Compare the pixel values in the template with the pixel values in the underlying region of the image.

- If a "good" match is found, announce that the object is present in the image.



- Possible measures are: SSD, SAD, Cross-correlation, Normalized Cross-correlation, max difference, etc.

Limitations of Template Matching

- If the object appears scaled, rotated, or skewed on the image, the match will not be good.



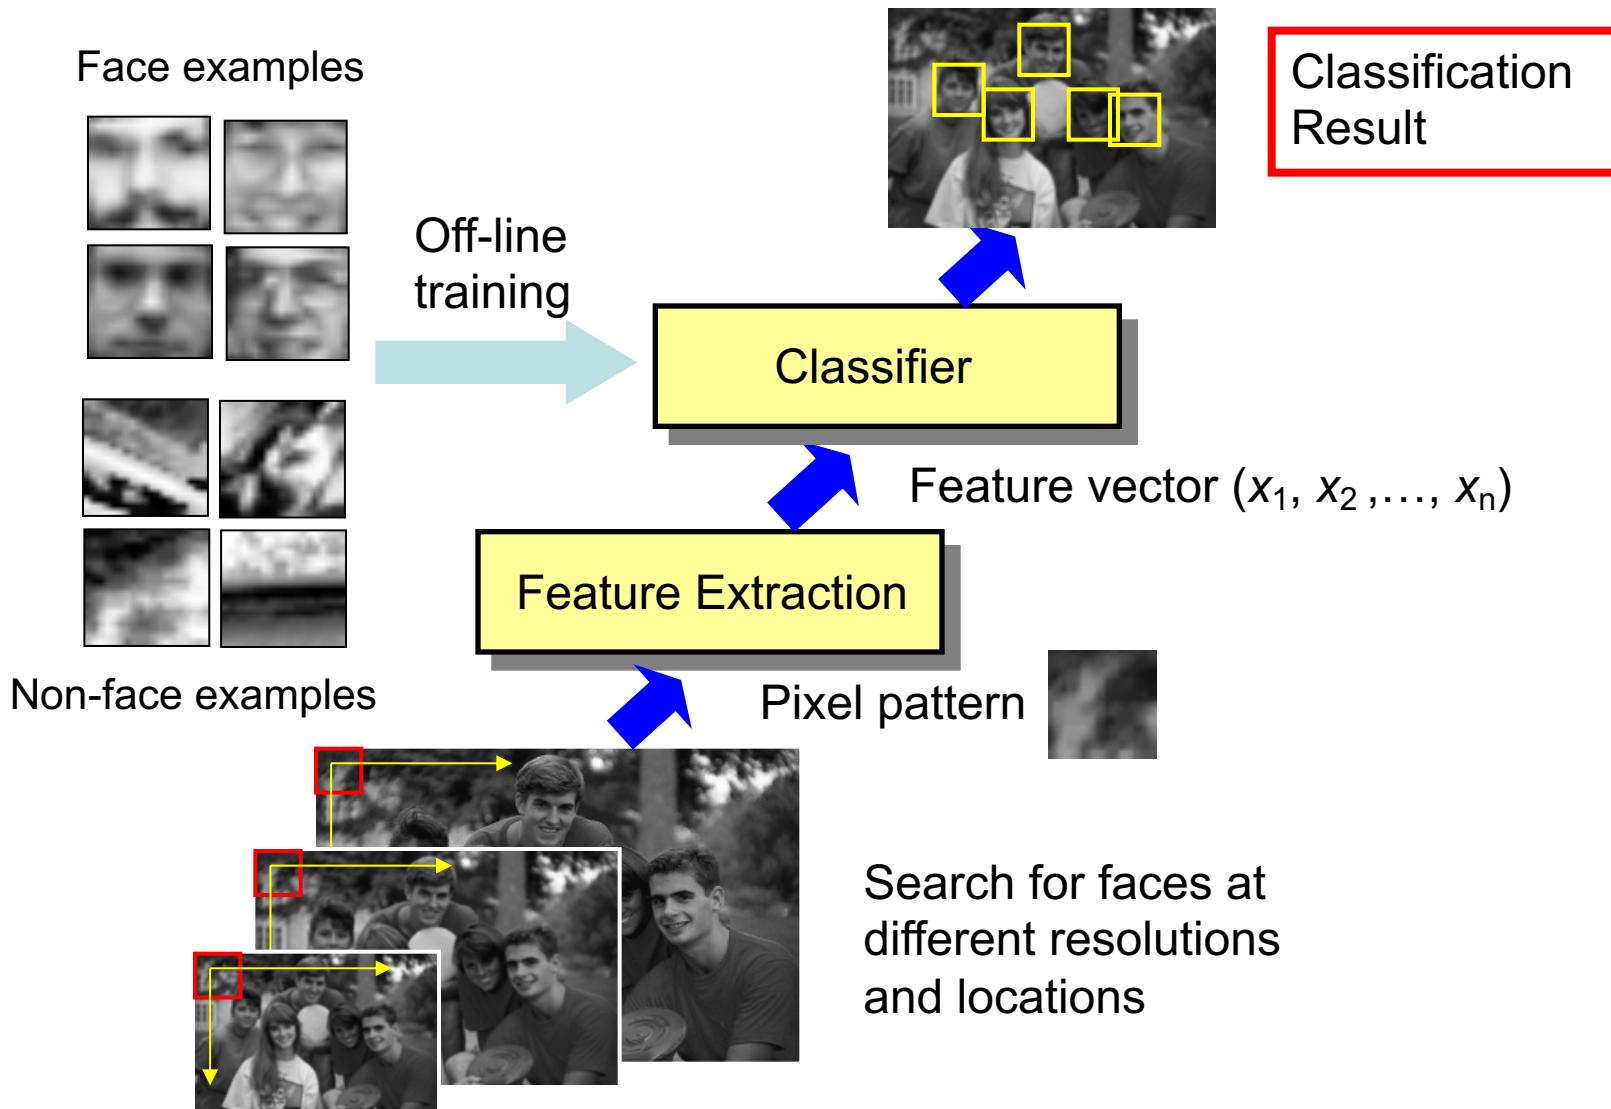
Solution:

- Search for the template and possible transformations of the template:



Not very efficient! (but doable ...)

Face Detection – basic scheme



Face Detection – basic scheme

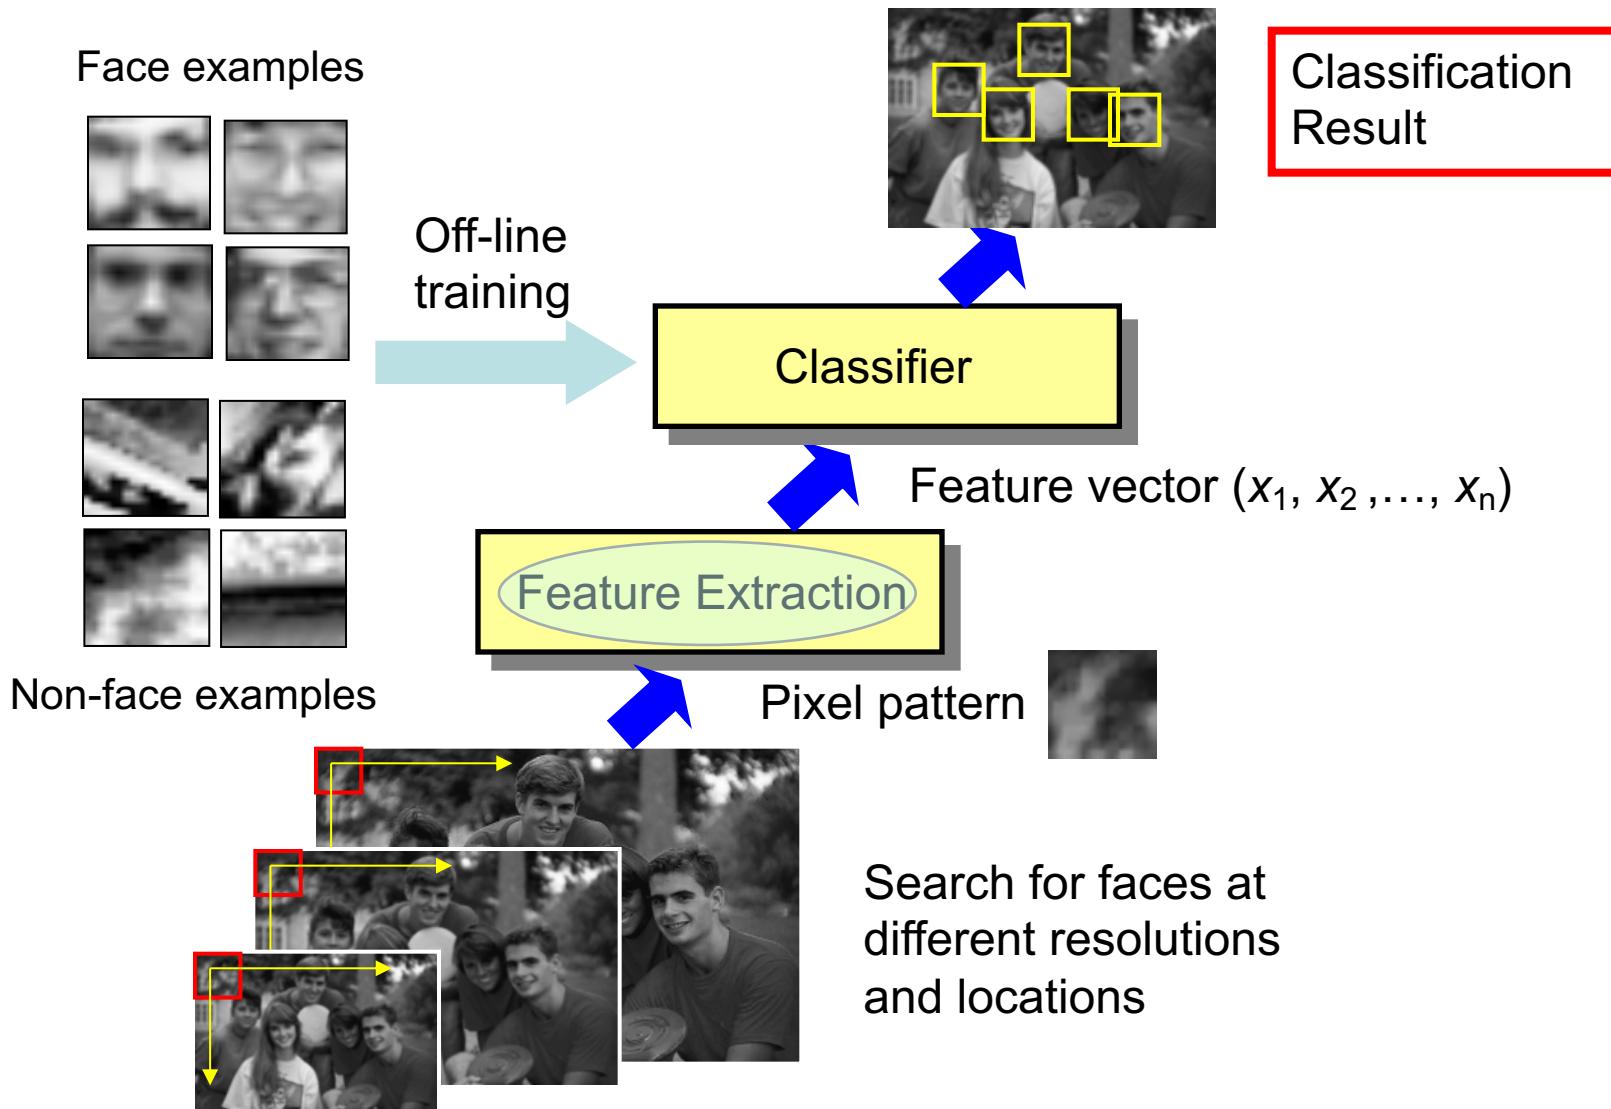
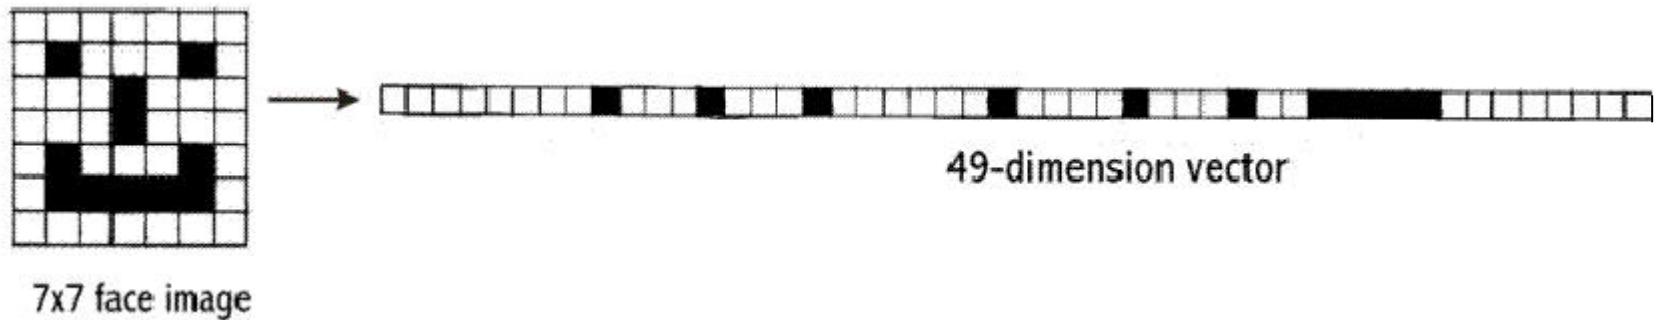


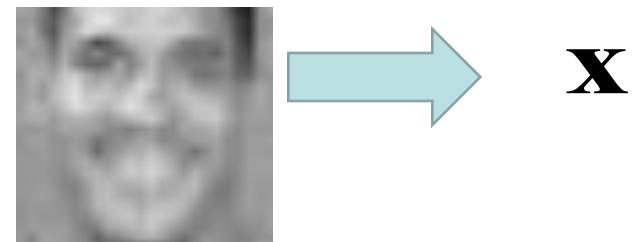
Image feature representation

- Images as a **high dimensional vector**



NN classifier

1. Treat pixels as a vector



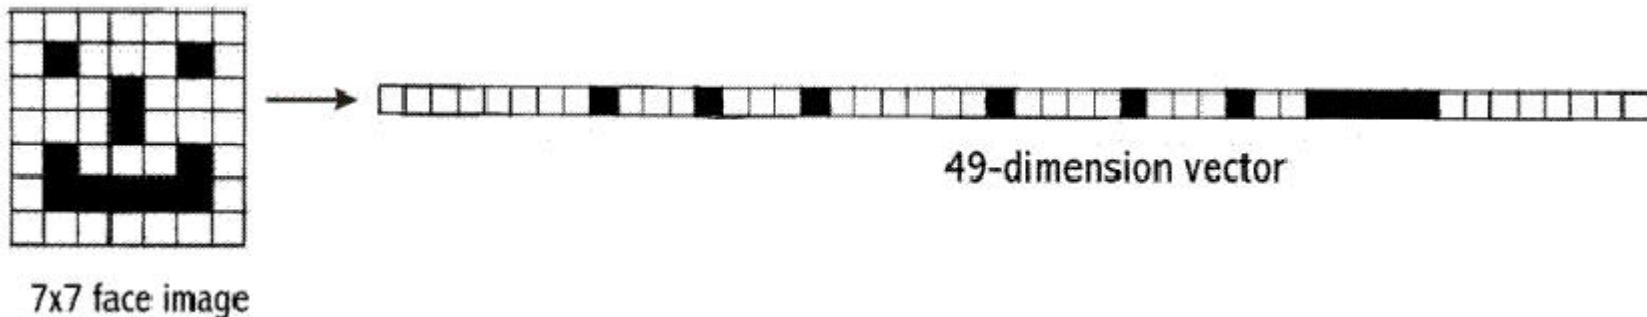
2. Recognize face by nearest neighbor



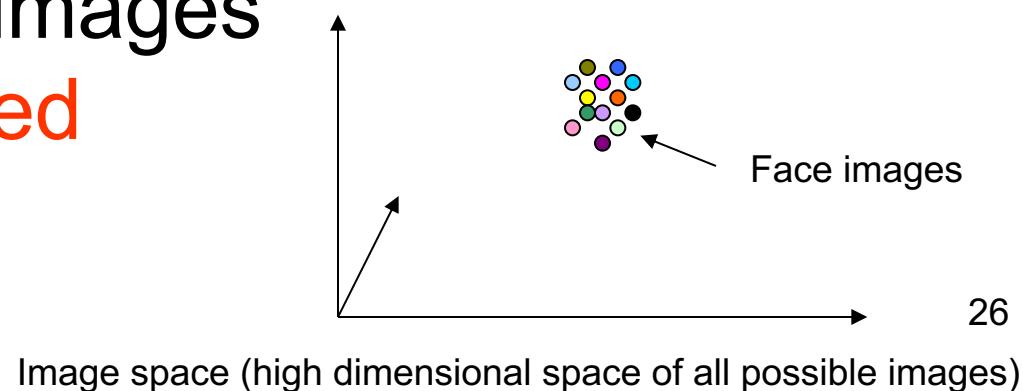
$$k = \operatorname{argmin}_k \|\mathbf{y}_k^T - \mathbf{x}\|$$

High Dimensional Correlated Data

- Images as a **high dimensional vector**



- A typical image used for image processing will be $512 \times 512 = 262144$ dimension vector!
- (Registered) face images are **highly correlated**



The space of all face images

- When viewed as vectors of pixel values, face images are extremely high-dimensional
 - 512x512 image = 262,144 dimensions
 - Slow and lots of storage
- But very few 262,144-dimension long vectors are valid face images; real face vectors are sparse (i.e. sparse distribution) in the data space.
- We want to effectively model the subspace of face images.

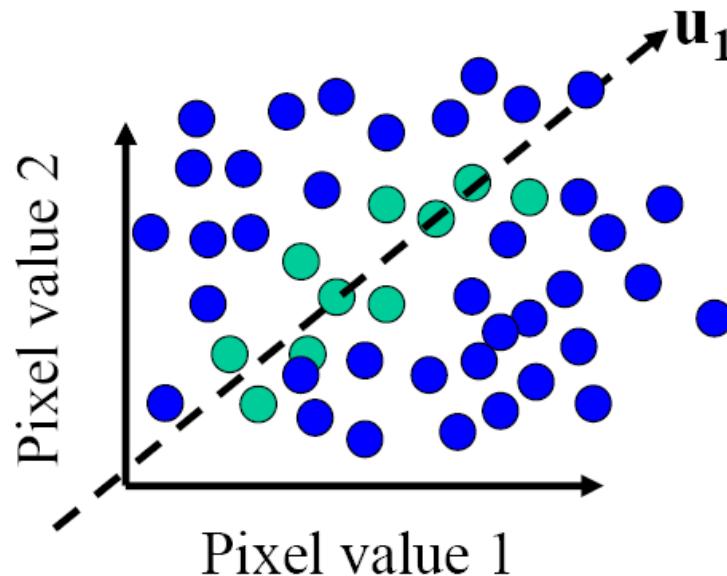


Image feature representation:

Dimensionality reduction

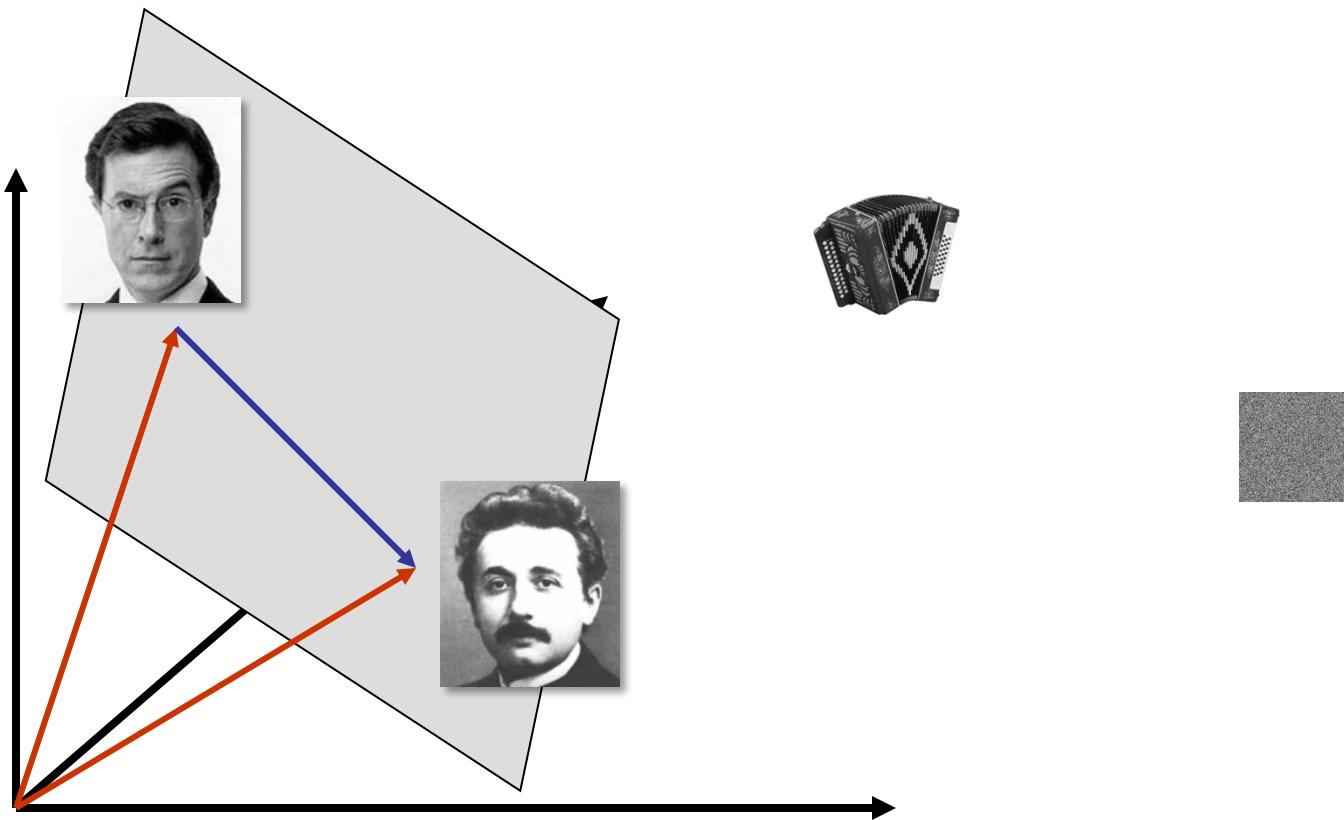
The space of all face images

- Eigenface idea:
 - Construct a low-dimensional linear subspace that best explains the variation in the set of face images



- A face image
- A (non-face) image

Dimensionality reduction



- The set of faces is a “subspace” of the set of all images
 - Suppose it is K dimensional
 - We can find the best subspace using PCA (principal component analysis)
 - This is like fitting a “hyper-plane” to the set of faces
 - spanned by vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_K$
 - any face $\mathbf{x} \approx \bar{\mathbf{x}} + a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_k\mathbf{v}_k$

Principal Component Analysis (PCA)

- General dimensionality reduction technique
- It preserves most of variance with a much more compact representation
 - Lower storage requirements (eigenvectors + a few numbers per face)
 - Faster for matching

Principal Component Analysis (PCA)

- Given: N data points $\mathbf{x}_1, \dots, \mathbf{x}_N$ in \mathbb{R}^d
- We want to find a new set of features that are linear combinations of original ones:

$$w = \mathbf{u}^T(\mathbf{x}_i - \boldsymbol{\mu})$$

($\boldsymbol{\mu}$: mean of data points)

- Choose unit vector \mathbf{u} in \mathbb{R}^d that captures the most data variance

Principal Component Analysis

- Direction that maximizes the variance of the projected data:

$$\begin{array}{ll} \text{Maximize}_{\mathbf{u}} & \frac{1}{N} \sum_{i=1}^N \underbrace{\mathbf{u}^T(\mathbf{x}_i - \mu)(\mathbf{u}^T(\mathbf{x}_i - \mu))^T}_{\text{Projection of data point}} \\ \text{subject to } \|\mathbf{u}\|=1 & \end{array}$$

$$\begin{aligned} &= \mathbf{u}^T \left[\underbrace{\frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^T}_{\text{Covariance matrix of data}} \right] \mathbf{u} \\ &= \mathbf{u}^T \Sigma \mathbf{u} \end{aligned}$$

→ The direction that maximizes the variance is the eigenvector associated with the largest eigenvalue of Σ

Another perspective

- Minimize approximation error

$$\begin{aligned}\|\vec{x}_i - (\vec{w} \cdot \vec{x}_i) \vec{w}\|^2 &= (\vec{x}_i - (\vec{w} \cdot \vec{x}_i) \vec{w}) \cdot (\vec{x}_i - (\vec{w} \cdot \vec{x}_i) \vec{w}) \\ &= \vec{x}_i \cdot \vec{x}_i - \vec{x}_i \cdot (\vec{w} \cdot \vec{x}_i) \vec{w} \\ &\quad - (\vec{w} \cdot \vec{x}_i) \vec{w} \cdot \vec{x}_i + (\vec{w} \cdot \vec{x}_i) \vec{w} \cdot (\vec{w} \cdot \vec{x}_i) \vec{w} \\ &= \|\vec{x}_i\|^2 - 2(\vec{w} \cdot \vec{x}_i)^2 + (\vec{w} \cdot \vec{x}_i)^2 \vec{w} \cdot \vec{w} \\ &= \vec{x}_i \cdot \vec{x}_i - (\vec{w} \cdot \vec{x}_i)^2\end{aligned}$$

Another perspective

- Minimize approximation error

$$\begin{aligned} MSE(\vec{w}) &= \frac{1}{n} \sum_{i=1}^n \|\vec{x}_i\|^2 - (\vec{w} \cdot \vec{x}_i)^2 \\ &= \frac{1}{n} \left(\sum_{i=1}^n \|\vec{x}_i\|^2 - \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i)^2 \right) \end{aligned}$$

Another perspective

- Minimize approximation error

$$\begin{aligned} MSE(\vec{w}) &= \frac{1}{n} \sum_{i=1}^n \|\vec{x}_i\|^2 - (\vec{w} \cdot \vec{x}_i)^2 \\ &= \frac{1}{n} \left(\sum_{i=1}^n \|\vec{x}_i\|^2 - \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i)^2 \right) \end{aligned}$$

Equivalent to Maximize: $\sum_{i=1}^n (\vec{w} \cdot \vec{x}_i)^2$

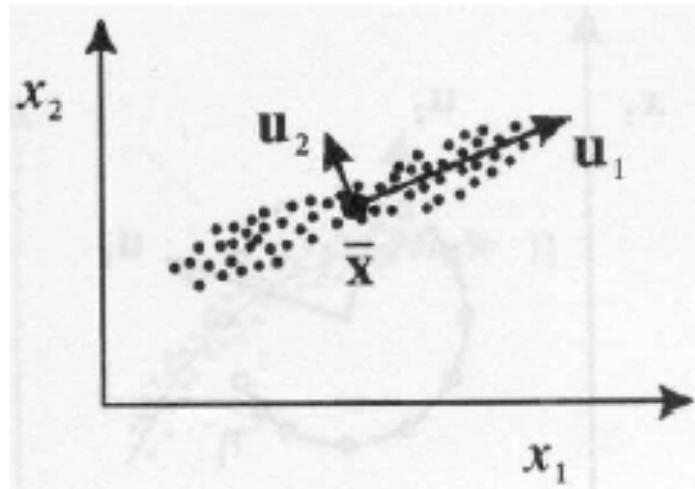
Another perspective

- Maximize variance

$$\begin{aligned}\sigma_{\vec{w}}^2 &= \frac{1}{n} \sum_i (\vec{x}_i \cdot \vec{w})^2 \\ &= \frac{1}{n} (\mathbf{x}\mathbf{w})^T (\mathbf{x}\mathbf{w}) \\ &= \frac{1}{n} \mathbf{w}^T \mathbf{x}^T \mathbf{x}\mathbf{w} \\ &= \mathbf{w}^T \frac{\mathbf{x}^T \mathbf{x}}{n} \mathbf{w} \\ &= \mathbf{w}^T \mathbf{v}\mathbf{w}\end{aligned}$$

Geometric interpretation

- PCA projects the data along the directions where the data varies most.
- These directions are determined by the eigenvectors of the covariance matrix corresponding to the largest eigenvalues.
- The magnitude of the eigenvalues corresponds to the variance of the data along the eigenvector directions.



PCA for dimension reduction

- Lower dimensionality basis
 - Approximate vectors by finding a basis in an appropriate lower dimensional space.
 - (1) Higher-dimensional space representation:

$$x = a_1 v_1 + a_2 v_2 + \cdots + a_N v_N$$

v_1, v_2, \dots, v_N is a basis of the N -dimensional space

- (2) Lower-dimensional space representation:

$$\hat{x} = b_1 u_1 + b_2 u_2 + \cdots + b_K u_K$$

u_1, u_2, \dots, u_K is a basis of the K -dimensional space

- Note: if both bases have the same size ($N = K$), then $x = \hat{x}$)

PCA Algorithm

- Suppose x_1, x_2, \dots, x_M are $N \times 1$ vectors

Step 1: $\bar{x} = \frac{1}{M} \sum_{i=1}^M x_i$

Step 2: subtract the mean: $\Phi_i = x_i - \bar{x}$ (i.e., center at zero)

Step 3: form the matrix $A = [\Phi_1 \ \Phi_2 \ \cdots \ \Phi_M]$ ($N \times M$ matrix), then compute:

$$C = \frac{1}{M} \sum_{n=1}^M \Phi_n \Phi_n^T = \frac{1}{M} A A^T$$

(sample **covariance** matrix, $N \times N$, characterizes the *scatter* of the data)

Step 4: compute the eigenvalues of C : $\lambda_1 > \lambda_2 > \cdots > \lambda_N$

Step 5: compute the eigenvectors of C : u_1, u_2, \dots, u_N

PCA Algorithm

- Since C is symmetric, u_1, u_2, \dots, u_N form a basis, (i.e., any vector x or actually $(x - \bar{x})$, can be written as a linear combination of the eigenvectors):

$$x - \bar{x} = b_1 u_1 + b_2 u_2 + \cdots + b_N u_N = \sum_{i=1}^N b_i u_i \quad b_i = \frac{(x - \bar{x}) \cdot u_i}{(u_i \cdot u_i)}$$

Step 6: (dimensionality reduction step) keep only the terms corresponding to the K largest eigenvalues:

$$\hat{x} - \bar{x} = \sum_{i=1}^K b_i u_i \text{ where } K \ll N$$

- The representation of $\hat{x} - \bar{x}$ into the basis u_1, u_2, \dots, u_K is thus

$$\begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_K \end{bmatrix}$$

Symmetric eigendecomposition

eigenvalues/vectors of a symmetric matrix have important special properties

- all the eigenvalues are real
- the eigenvectors corresponding to different eigenvalues are orthogonal
- a symmetric matrix is diagonalizable by an orthogonal similarity transformation:

$$Q^T A Q = \Lambda, \quad Q^T Q = I$$

in the remainder of the lecture we assume that A is symmetric (and real)

PCA algorithm

- The linear transformation $R^N \rightarrow R^K$ that performs the dimensionality reduction is:

$$\begin{bmatrix} b_1 \\ b_2 \\ \dots \\ b_K \end{bmatrix} = \begin{bmatrix} u_1^T \\ u_2^T \\ \dots \\ u_K^T \end{bmatrix} (x - \bar{x}) = U^T(x - \bar{x})$$

(i.e., simply computing coefficients of linear expansion)

Eigenfaces (PCA on face images)

1. Compute covariance matrix of face images
2. Compute the principal components (“eigenfaces”)
 - K eigenvectors with top K largest eigenvalues
3. Represent all face images in the dataset as linear combinations of eigenfaces
 - Perform nearest neighbor on these coefficients

Eigenfaces example

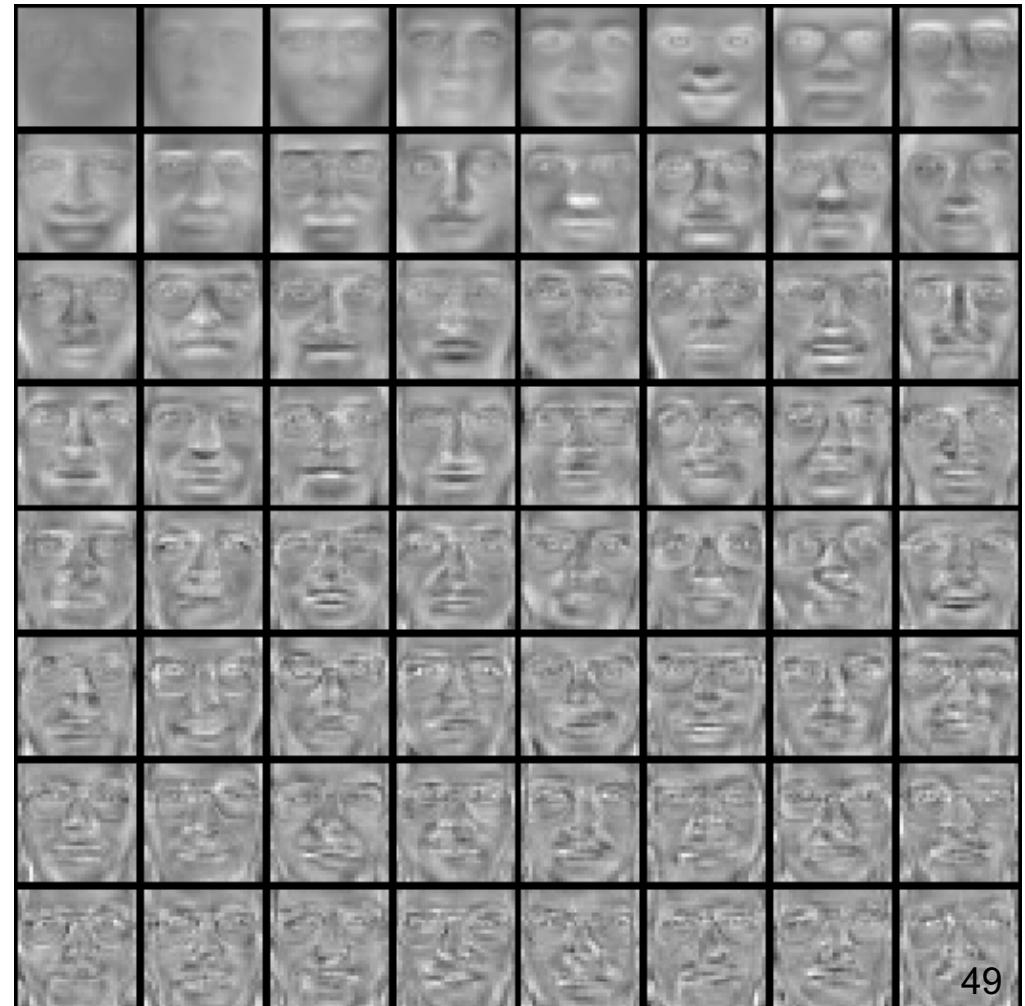
- Training images
- $\mathbf{x}_1, \dots, \mathbf{x}_N$



Eigenfaces example

Top eigenvectors: u_1, \dots, u_k

Mean: μ



Representation and reconstruction

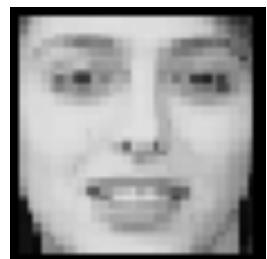
- Face \mathbf{x} in “face space” coordinates:



$$\begin{aligned}\mathbf{x} &\rightarrow [\mathbf{u}_1^T(\mathbf{x} - \mu), \dots, \mathbf{u}_k^T(\mathbf{x} - \mu)] \\ &= [w_1, \dots, w_k]\end{aligned}$$

Representation and reconstruction

- Face \mathbf{x} in “face space” coordinates:



$$\begin{aligned}\mathbf{x} &\rightarrow [\mathbf{u}_1^T(\mathbf{x} - \mu), \dots, \mathbf{u}_k^T(\mathbf{x} - \mu)] \\ &= w_1, \dots, w_k\end{aligned}$$

- Reconstruction:

$$\begin{aligned}\hat{\mathbf{x}} &= \mathbf{\mu} + w_1\mathbf{u}_1 + w_2\mathbf{u}_2 + w_3\mathbf{u}_3 + w_4\mathbf{u}_4 + \dots\end{aligned}$$

The equation shows the reconstruction of a face $\hat{\mathbf{x}}$ from its mean $\mathbf{\mu}$ and coefficients w_i multiplied by basis vectors \mathbf{u}_i . The basis vectors are represented as a horizontal strip of seven smaller face images, each showing a different facial feature or expression.

Reconstruction example

- The visualization of eigenvectors:



These are the top 4 eigenvectors from a training set of 400 images (ORL Face Database). They look like faces, hence called Eigenface.

Reconstruction

$P = 4$



$P = 200$



$P = 400$



After computing eigenfaces using 400 face images from ORL face database

How to choose K?

- Choose K using the following criterion:

$$\frac{\sum_{i=1}^K \lambda_i}{N} > \text{Threshold} \quad (\text{e.g., } 0.9 \text{ or } 0.95)$$
$$\sum_{i=1}^K \lambda_i$$

- In this case, we say that we “preserve” 90% or 95% of the information (variance) in the data.
- If $K=N$, then we “preserve” 100% of the information in the data.

Eigenfaces



Computed using 400 face images from ORL face database

Recognition with eigenfaces

Process labeled training images

- Find mean μ and covariance matrix Σ
- Find k principal components (eigenvectors of Σ) $\mathbf{u}_1, \dots, \mathbf{u}_k$
- Project each training image \mathbf{x}_i onto subspace spanned by principal components:
 $(w_{i1}, \dots, w_{ik}) = (\mathbf{u}_1^T(\mathbf{x}_i - \mu), \dots, \mathbf{u}_k^T(\mathbf{x}_i - \mu))$

Given novel image \mathbf{x}

- Project onto subspace:
 $(w_1, \dots, w_k) = (\mathbf{u}_1^T(\mathbf{x} - \mu), \dots, \mathbf{u}_k^T(\mathbf{x} - \mu))$
- Optional: check reconstruction error $\hat{\mathbf{x}} - \mathbf{x}$ to determine whether image is really a face
- Classify as closest training face in k -dimensional subspace

Recognition Using Eigenfaces

- Given an unknown face image Γ (centered and of the same size like the training faces) follow these steps:

Step 1: normalize Γ : $\Phi = \Gamma - \Psi$

Step 2: project on the eigenspace

$$\hat{\Phi} = \sum_{i=1}^K w_i u_i \quad (w_i = u_i^T \Phi) \quad (\text{where } \|u_i\|=1)$$

Step 3: represent Φ as: $\Omega = \begin{bmatrix} w_1 \\ w_2 \\ \dots \\ w_K \end{bmatrix}$

Step 4: find $e_r = \min_l \|\Omega - \Omega^l\|$ where $\|\Omega - \Omega^l\| = \sum_{i=1}^K (w_i - w_i^l)^2$

Step 5: if $e_r < T_r$, then Γ is recognized as face l from the training set.

The distance e_r is called *distance in face space (difs)*

Reconstruction from partial information

- Robust to partial face occlusion.

Input



Reconstructed



Limitations

Global appearance method: not robust to misalignment, background variation



Issues?

- Image of size: 256×256 , each image as a vector: 65526×1
- Training set size: M is around 1000
- High dimension of covariance matrix Σ :
 65526×65526
- Eigen value decomposition is *intractable* for even small size image
- Solution?

S

S^T

SS^T

$S^T S$

Eigenvalues of $S^T S$ and SS^T are the same

Theorem 1.4. Let \mathbf{S} be an $n \times k$ matrix with $k < n$. Then the non-zero eigenvalues of $\mathbf{S}\mathbf{S}^T$ are equal to the eigenvalues of $\mathbf{S}^T\mathbf{S}$, and if \mathbf{v} is an eigenvector of $\mathbf{S}^T\mathbf{S}$ with eigenvalue λ , then $\mathbf{S}\mathbf{v}$ is an eigenvalue of $\mathbf{S}\mathbf{S}^T$ with the same eigenvalue λ . The corresponding unit eigenvector is therefore $\mathbf{S}\mathbf{v}/\|\mathbf{S}\mathbf{v}\|$.

Proof Suppose that $\mathbf{S}^T\mathbf{S}\mathbf{v} = \lambda\mathbf{v}$. Multiplying by \mathbf{S} on the left gives

$$\begin{aligned}\mathbf{S}(\mathbf{S}^T\mathbf{S})\mathbf{v} &= \mathbf{S}\lambda\mathbf{v} \\ (\mathbf{S}\mathbf{S}^T)(\mathbf{S}\mathbf{v}) &= \lambda(\mathbf{S}\mathbf{v}) .\end{aligned}$$

PCA dimension for high-dimensional ambient space.

- (i) Find the centroid of the points \mathbf{y}_i , given by

$$\bar{\mathbf{y}} = \frac{1}{k} \sum_{i=1}^k \mathbf{y}_i .$$

- (ii) Form the matrix \mathbf{S} whose columns are the vectors $\mathbf{y}_i - \bar{\mathbf{y}}$. This is a matrix of dimension $n \times k$.
- (iii) Find the m largest eigenvalues of $\mathbf{S}^T \mathbf{S}$ and corresponding eigenvectors \mathbf{v}_i .
- (iv) Let \mathbf{A} be the matrix whose columns are the vectors $\mathbf{S}\mathbf{v}_i / \|\mathbf{S}\mathbf{v}_i\|$. These are the eigenvectors of $\mathbf{S}\mathbf{S}^T$.
- (v) The best hyperplane of dimension m fit to the points $\{\mathbf{y}_i\}$ is the affine subspace of points of the form $\mathbf{x}' = \mathbf{Ax} + \bar{\mathbf{y}}$, for $\mathbf{x} \in \mathbb{R}^m$.
- (vi) This can also be written as the points \mathbf{x}' satisfying the equation

$$(\mathbf{A}^\perp)^T (\mathbf{x}' - \bar{\mathbf{y}}) = 0$$

where \mathbf{A}^\perp is the matrix consisting of the columns of \mathbf{Y} corresponding to the $n - m$ **smallest** eigenvalues.

Table 1.2. PCA algorithm when ambient space is high-dimensional

Recall: PCA Algorithm

- Suppose x_1, x_2, \dots, x_M are $N \times 1$ vectors

$$\underline{\text{Step 1:}} \bar{x} = \frac{1}{M} \sum_{i=1}^M x_i$$

Step 2: subtract the mean: $\Phi_i = x_i - \bar{x}$ (i.e., center at zero)

Step 3: form the matrix $A = [\Phi_1 \ \Phi_2 \ \cdots \ \Phi_M]$ ($N \times M$ matrix), then compute:

$$C = \frac{1}{M} \sum_{n=1}^M \Phi_n \Phi_n^T = \frac{1}{M} A A^T$$

(sample **covariance** matrix, $N \times N$, characterizes the *scatter* of the data)

Step 4: compute the eigenvalues of C : $\lambda_1 > \lambda_2 > \cdots > \lambda_N$

Step 5: compute the eigenvectors of C : u_1, u_2, \dots, u_N

Solution

- Follow Step1 and Step2 from former slide.
- Step3: Form the matrix $A = [\varphi_1, \varphi_2, \varphi_3, \varphi_4, \dots, \varphi_M]$, A is of size $N \times M$, and compute $L = A^T A$ where L is of size $M \times M$, and $L_{mn} = \varphi_m^T \varphi_n$
- Step 4: Compute eigen values of L.
- Step 5: Compute eigen vectors of L for the k largest eigen values, defined as q_i
- Step 6: Compute eigen vectors u_i of the original covariance matrix, namely the eigen faces as: $u_i = A q_i$

Step 7: Normalize the eigen vector $u_i = \frac{u_i}{\|u_i\|}$

Another possible solution

PCA using Singular Value Decomposition.

- (i) Find the centroid of the points \mathbf{y}_i , given by

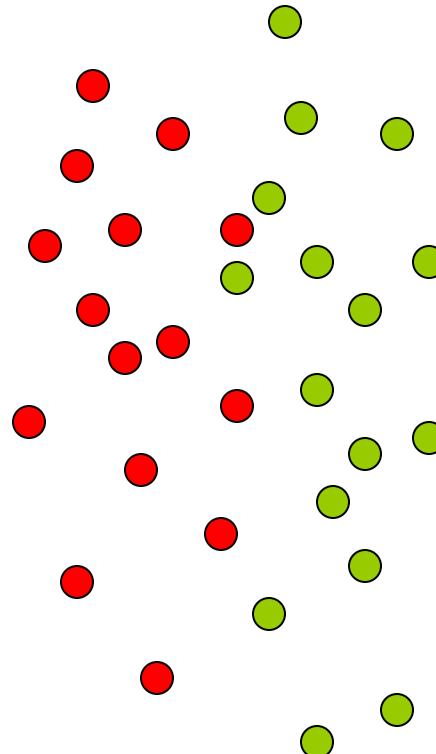
$$\bar{\mathbf{y}} = \frac{1}{k} \sum_{i=1}^k \mathbf{y}_i .$$

- (ii) Form the matrix \mathbf{S} whose columns are the vectors $\mathbf{y}_i - \bar{\mathbf{y}}$. This is a matrix of dimension $n \times k$.
- (iii) Find the SVD $\boxed{\mathbf{S} = \mathbf{U}\mathbf{D}\mathbf{V}^T}$ and let \mathbf{A} be the matrix consisting of the first m columns of \mathbf{U} .
- (iv) The best hyperplane of dimension m fit to the points $\{\mathbf{y}_i\}$ is the affine subspace of points of the form $\mathbf{x}' = \mathbf{Ax} + \bar{\mathbf{y}}$, for $\mathbf{x} \in \mathbb{R}^m$.

Table 1.3. PCA algorithm using the Singular Value Decomposition.

Limitations

- The direction of maximum variance is not always good for classification

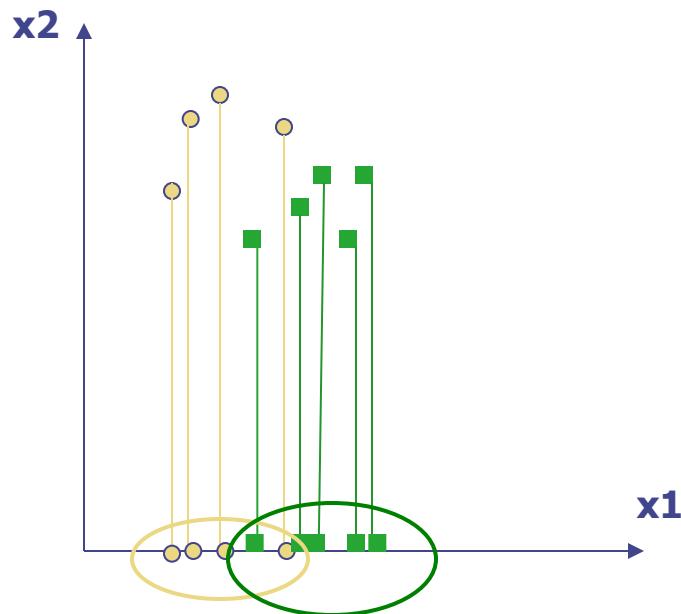


A more discriminative subspace: FLD

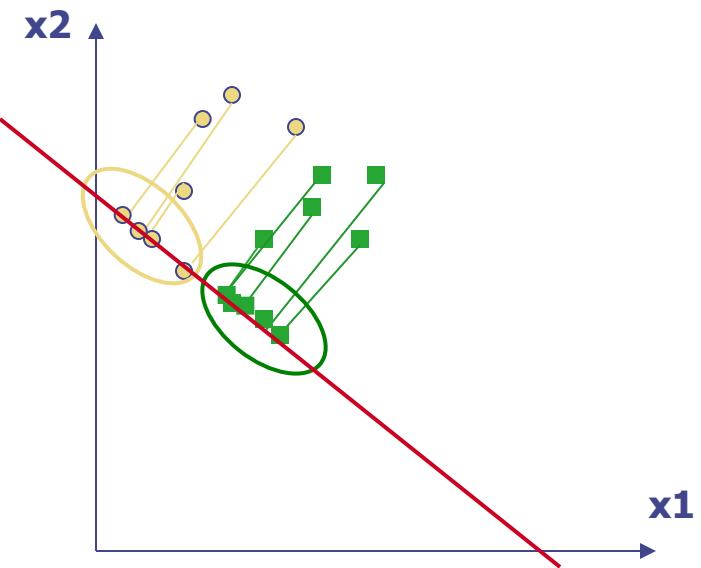
- Fisher Linear Discriminants → “Fisher Faces”
- PCA preserves maximum variance
- FLD preserves discrimination
 - Find projection that maximizes scatter between classes and minimizes scatter within classes

Illustration of the Projection

- ◆ Using two classes as an example:

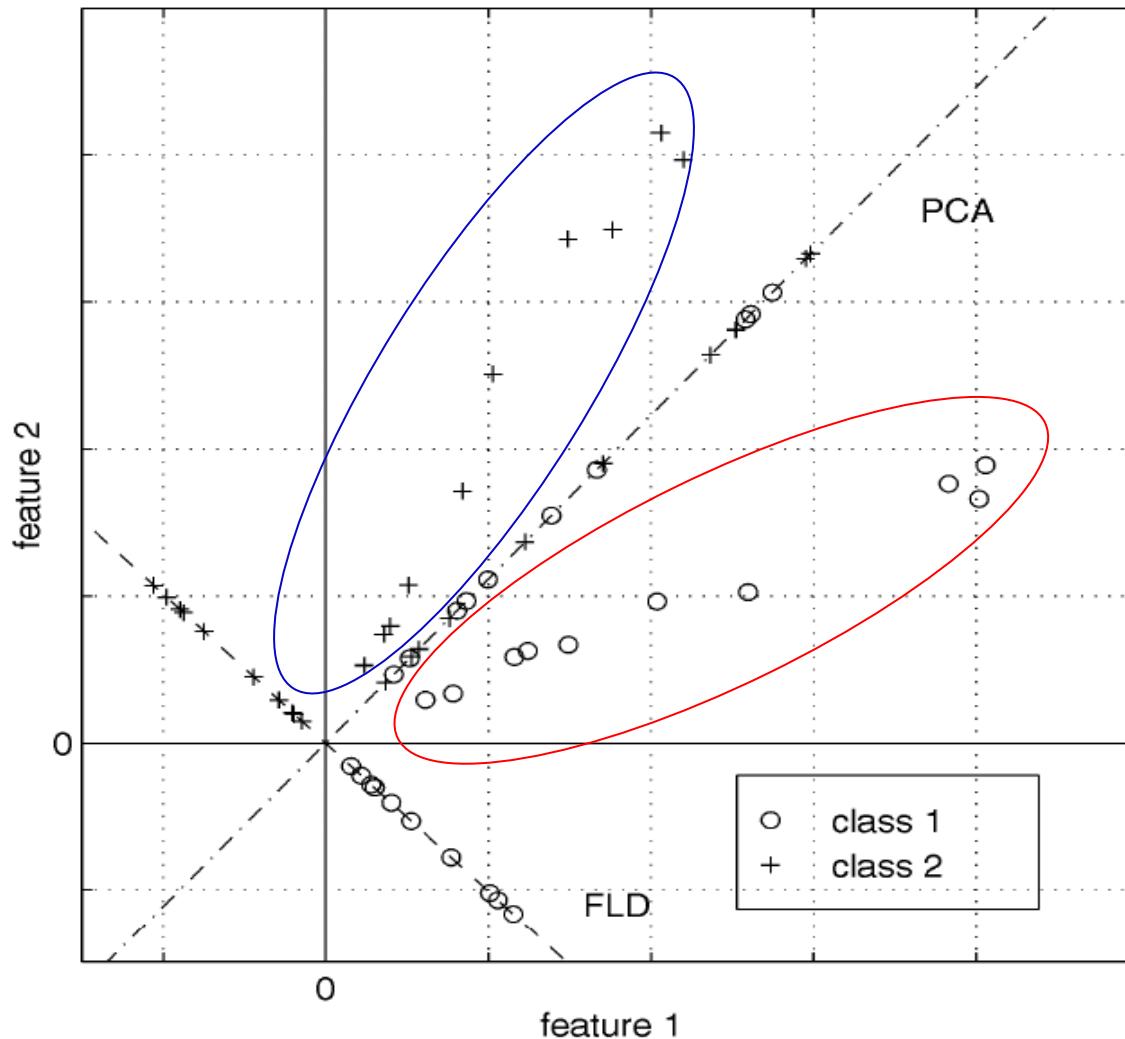


Poor Projection



Good

Comparing with PCA



Variables

- N Sample images:

$$\{x_1, \dots, x_N\}$$

- c classes:

$$\{\chi_1, \dots, \chi_c\}$$

- Average of each class:

$$\mu_i = \frac{1}{N_i} \sum_{x_k \in \chi_i} x_k$$

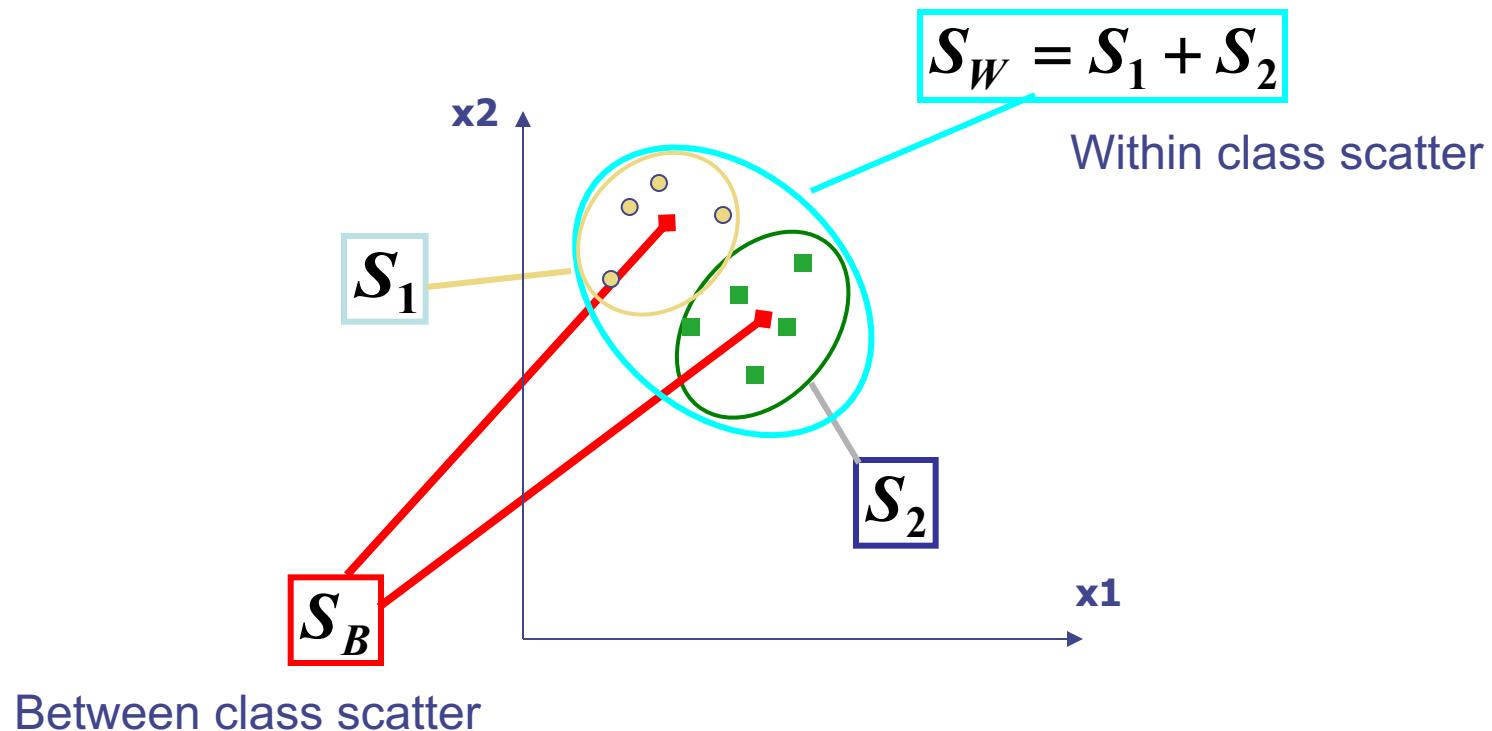
- Average of all data:

$$\mu = \frac{1}{N} \sum_{k=1}^N x_k$$

Scatter Matrices

- Scatter of class i:
$$S_i = \sum_{x_k \in \chi_i} (x_k - \mu_i)(x_k - \mu_i)^T$$
- Within class scatter:
$$S_W = \sum_{i=1}^c S_i$$
- Between class scatter:
$$S_B = \sum_{i=1}^c N_i (\mu_i - \mu)(\mu_i - \mu)^T$$

Illustration



Mathematical Formulation

- After projection
 - Between class scatter $\tilde{S}_B = W^T S_B W$
 - Within class scatter $\tilde{S}_W = W^T S_W W$

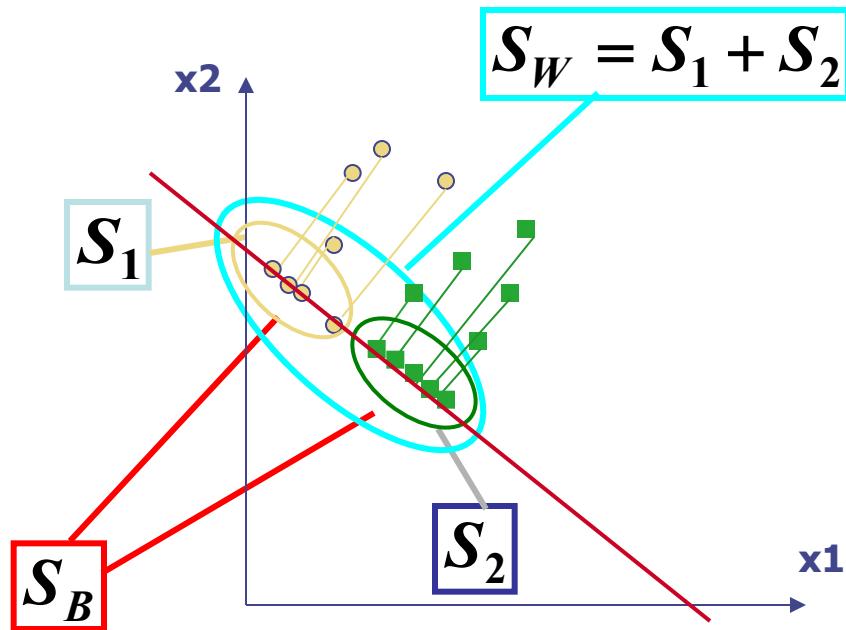
- Objective

$$W_{opt} = \arg \max_w \frac{|\tilde{S}_B|}{|\tilde{S}_W|} = \arg \max_w \frac{|W^T S_B W|}{|W^T S_W W|}$$

- Solution: Generalized Eigenvectors

$$S_B w_i = \lambda_i S_W w_i \quad i = 1, \dots, m$$

Illustration



Recognition with FLD

- Similar to “eigenfaces”
- Compute within-class and between-class scatter matrices

$$S_i = \sum_{x_k \in \chi_i} (x_k - \mu_i)(x_k - \mu_i)^T \quad S_W = \sum_{i=1}^c S_i \quad S_B = \sum_{i=1}^c N_i (\mu_i - \mu)(\mu_i - \mu)^T$$

- Solve generalized eigenvector problem

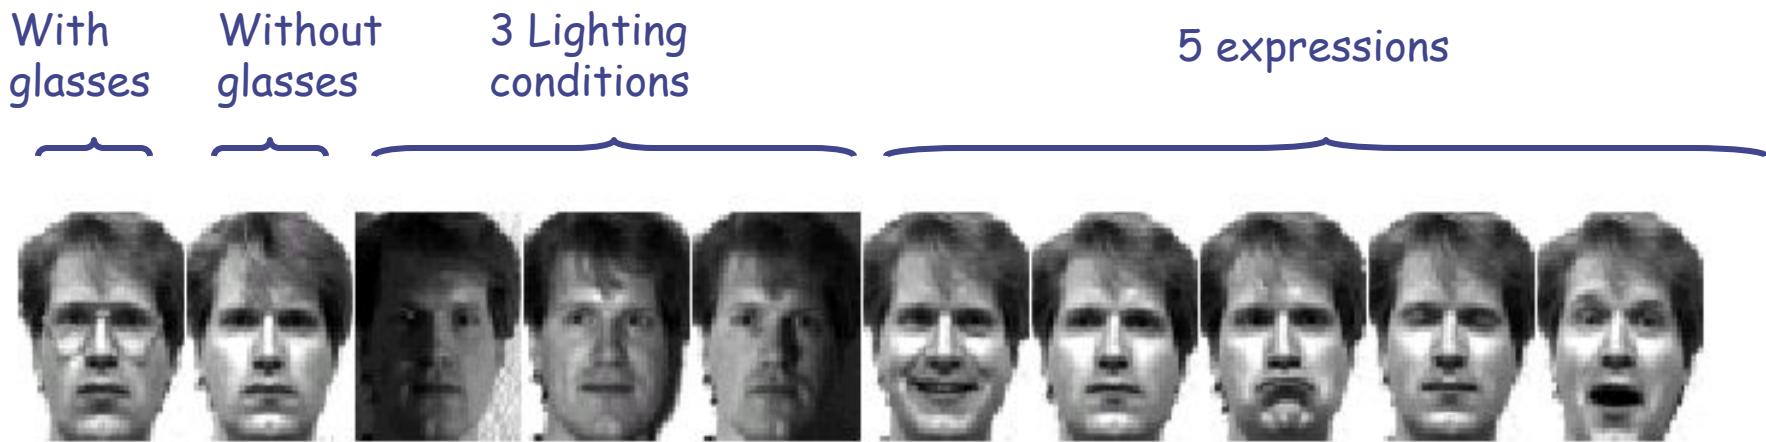
$$W_{opt} = \arg \max_w \frac{|W^T S_B W|}{|W^T S_W W|} \quad S_B w_i = \lambda_i S_W w_i \quad i = 1, \dots, m$$

- Project to FLD subspace and classify by nearest neighbor

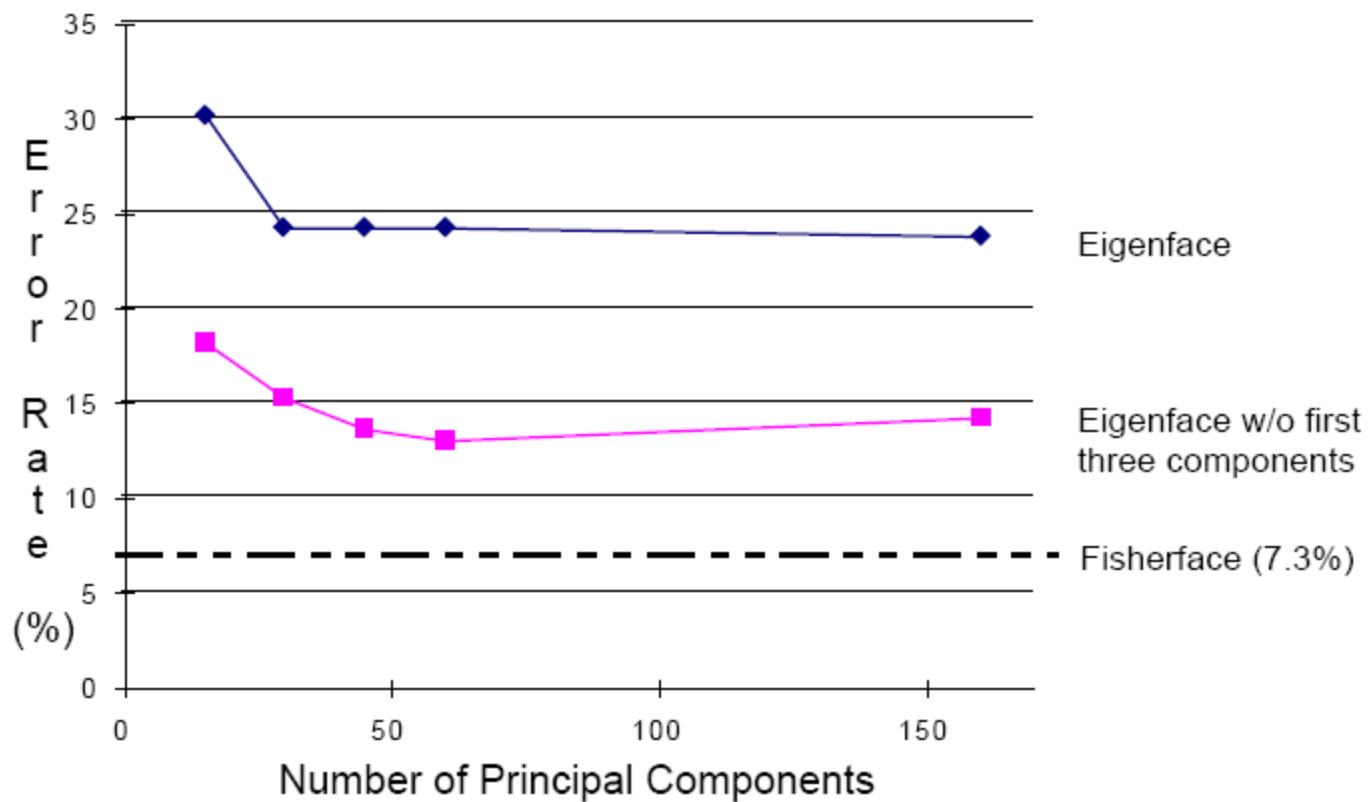
$$\hat{x} = {W_{opt}}^T x$$

Results: Eigenface vs. Fisherface

- Input: 160 images of 16 people
- Train: 159 images
- Test: 1 image
- Variation in Facial Expression, Eyewear, and Lighting



Eigenfaces vs. Fisherfaces

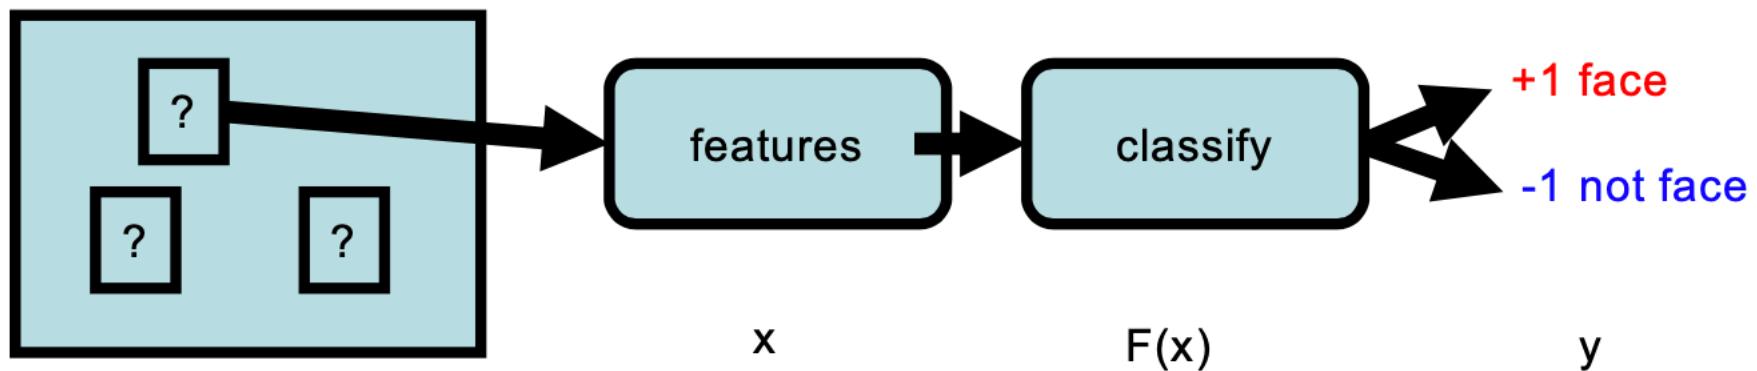


Readings

- Chapter 5.2.3 in the book:

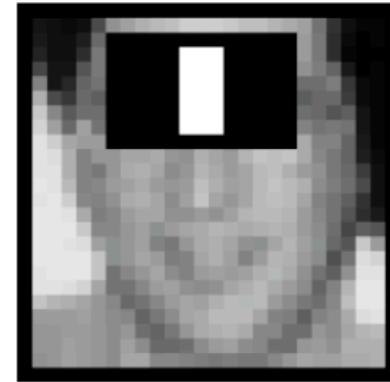
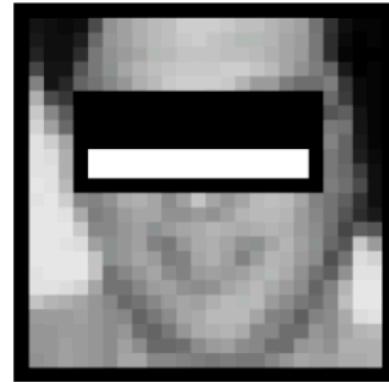
Computer Vision: Algorithms and Applications

Face detection -- basic scheme



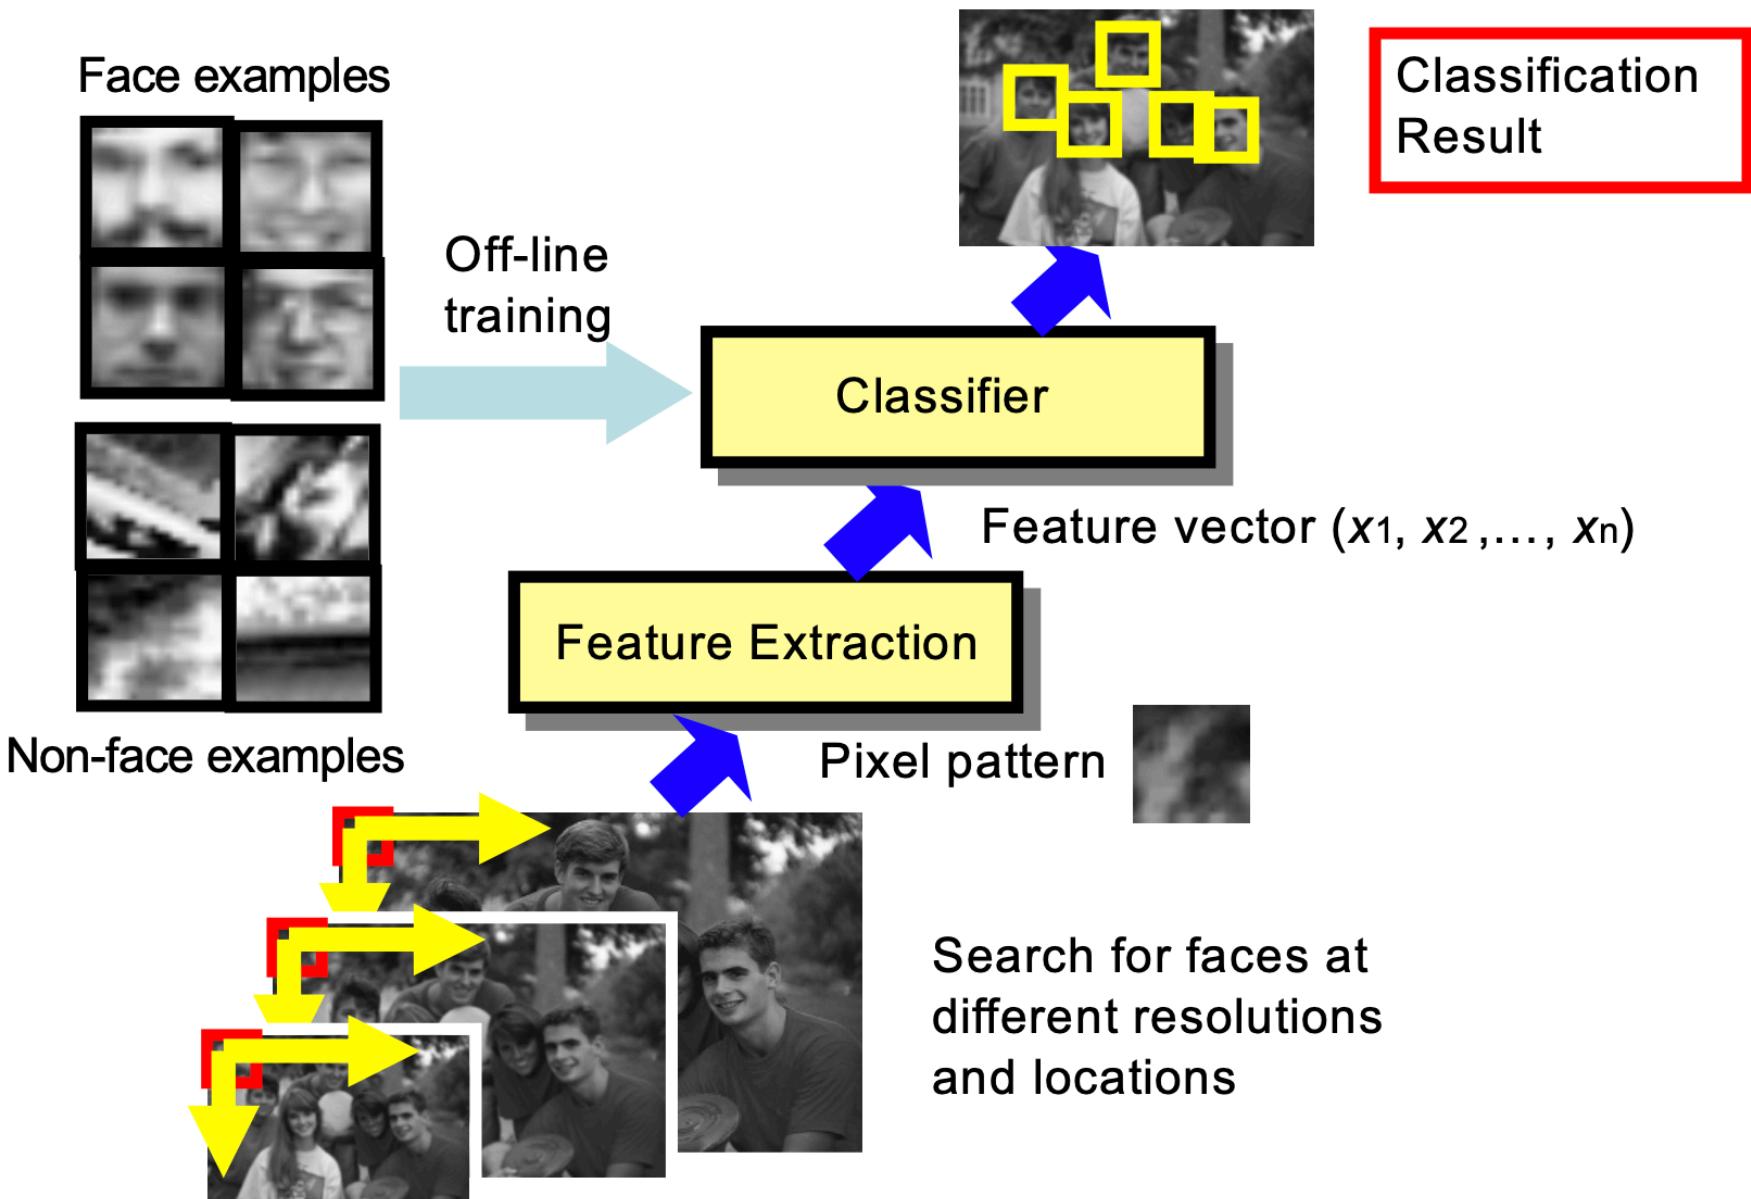
- We slide a window over the image
- Extract features for each window
- Classify each window into face/non-face

What is a face?



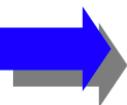
- Eyes are dark (eyebrows+shadows)
- Cheeks and forehead are bright.
- Nose is bright

Face detection – basic scheme



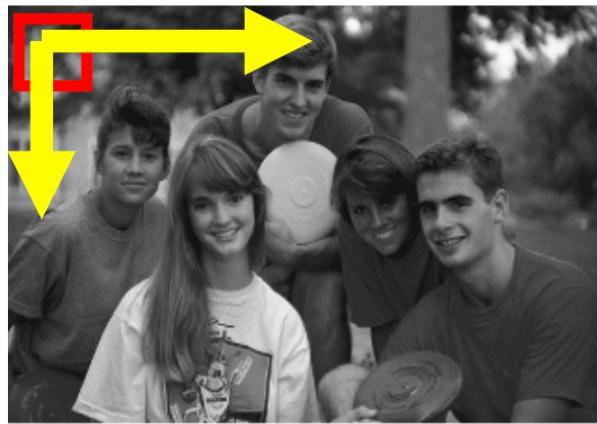
Training and Testing

Training Set



Train Classifier

Labeled Test Set

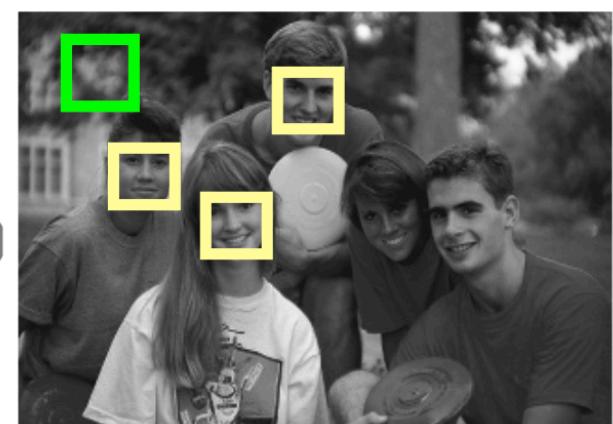


Sensitivity

Classify



False Positive



Correct

Viola-Jones face detection

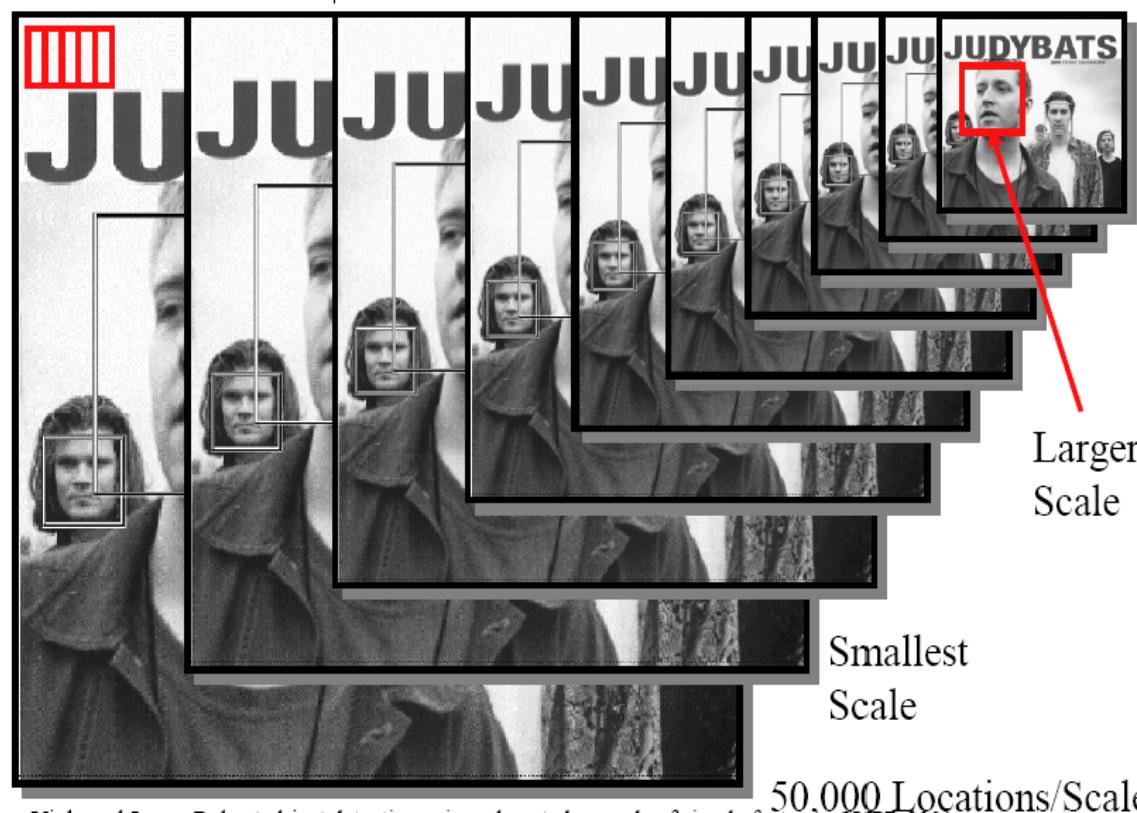
Paul A. Viola and Michael J. Jones

Intl. J. Computer Vision

57(2), 137–154, 2004

Some slides (*slides adapted from Bill Freeman, MIT 6.869, April 2005*)

Scan classifier over locs. & scales



Learn feature & classifier from data

- Training Data
- 5000 faces (frontal)
- 10^8 non faces
- Faces are normalized
 - Scale, translation
- Many variations
- Across individuals
- Illumination
- Pose (rotation both in plane and out)



Characteristics of the algorithm

- Feature set (...is huge about 16M features)
 - Efficient feature selection using AdaBoost
 - New image representation: Integral Image
 - Cascaded Classifier for rapid detection
-
- Fastest known face detector for gray scale images (circa 2001, Marr Prize Winner).

Viola-Jones face detection

- 3 key ideas
 - Use of Haar features
 - Integral Image
 - Cascade classifier
- For more information, read the original paper.
- Also, lots of videos on the web.
 - <https://www.youtube.com/watch?v=uEJ71VIUmMQ&vl=en-GB> (Google: Viola-Jones Computerphile)

Haar features



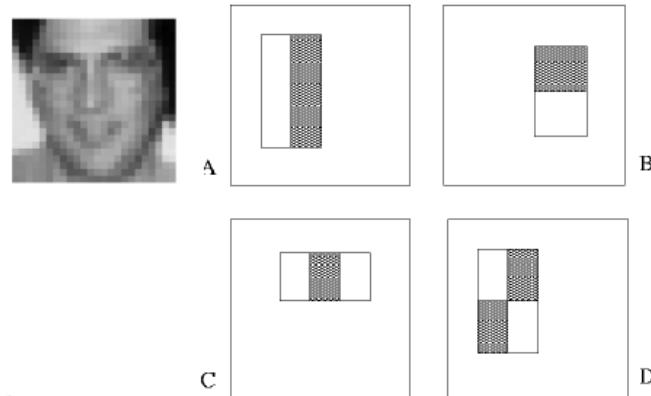
Black = add
White = subtract.



Lots of kernels at
different scales and
shapes.

Image features

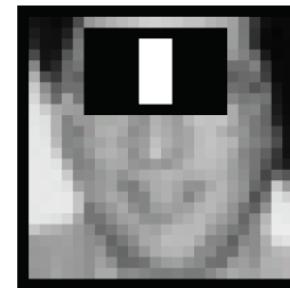
- “Rectangle filters”
- Differences between sums of pixels in adjacent rectangles



$$h_t(x) = \begin{cases} +1 & \text{if } f_t(x) > \theta_t \\ -1 & \text{otherwise} \end{cases}$$

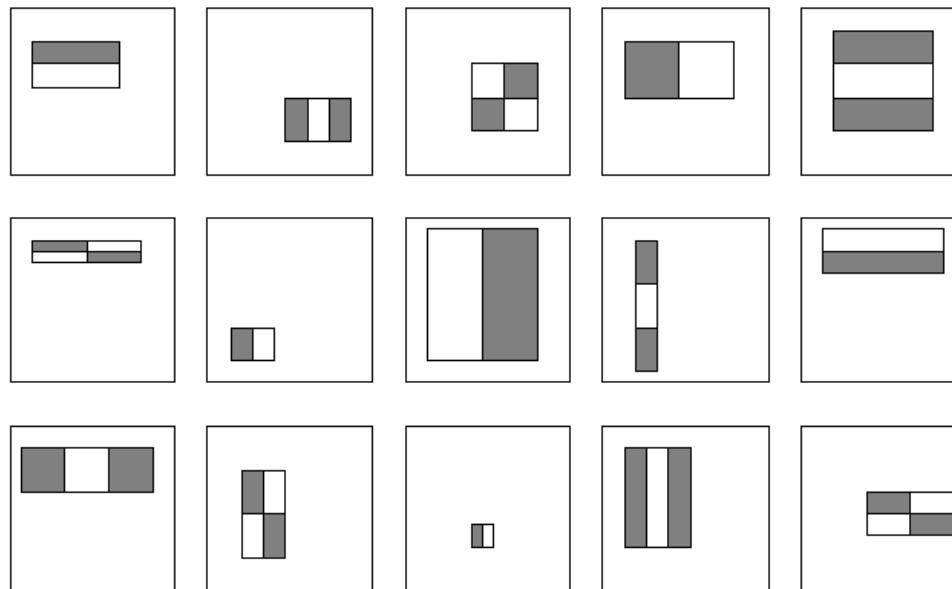
Threshold determined by training

What is a face?



- Eyes are dark (eyebrows+shadows)
- Cheeks and forehead are bright.
- Nose is bright

Huge library of Haar kernels



The integral image

3	2	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Given an image ...

3	2	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

What is the sum of all pixels in a given box?
(here equal to 36)

Construct the “integral image”

Integral image

- Efficient way to add up all pixels
in a given region

3	2	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Integral image

- Efficient way to add up all pixels
in a given region

3	5	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Integral image

- Efficient way to add up all pixels
in a given region

3	5	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Integral image

- Efficient way to add up all pixels
in a given region

3	5	6	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Integral image

- Efficient way to add up all pixels
in a given region

3	5	6	11	15	22
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
9	11	15	16	16	22
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

```
// Start with array A[0:nrows][0:ncols]
// Horizontal pass
for i=0, ..., nrows
    for j=1, ..., ncols
        A[i][j] = A[i][j] + A[i][j-1]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
9	11	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]
// Horizontal pass
for i=0, ..., nrows
    for j=1, ..., ncols
        A[i][j] = A[i][j] + A[i][j-1]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
9	11	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]  
  
// Horizontal pass  
for i=0, ..., nrows  
    for j=1, ..., ncols  
        A[i][j] = A[i][j] + A[i][j-1]  
  
// Vertical pass  
for i=1, ..., nrows  
    for j=0, ..., ncols  
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
9	11	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]  
  
// Horizontal pass  
for i=0, ..., nrows  
    for j=1, ..., ncols  
        A[i][j] = A[i][j] + A[i][j-1]  
  
// Vertical pass  
for i=1, ..., nrows  
    for j=0, ..., ncols  
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
12	11	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]  
  
// Horizontal pass  
for i=0, ..., nrows  
    for j=1, ..., ncols  
        A[i][j] = A[i][j] + A[i][j-1]  
  
// Vertical pass  
for i=1, ..., nrows  
    for j=0, ..., ncols  
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
12	11	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]  
  
// Horizontal pass  
for i=0, ..., nrows  
    for j=1, ..., ncols  
        A[i][j] = A[i][j] + A[i][j-1]  
  
// Vertical pass  
for i=1, ..., nrows  
    for j=0, ..., ncols  
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels in a given region

3	5	6	11	15	22
12	16	15	16	16	22
3	10	18	18	20	24
1	1	3	7	13	14
1	3	7	12	20	26
8	12	14	15	18	23

```
// Start with array A[0:nrows][0:ncols]  
  
// Horizontal pass  
for i=0, ..., nrows  
    for j=1, ..., ncols  
        A[i][j] = A[i][j] + A[i][j-1]  
  
// Vertical pass  
for i=1, ..., nrows  
    for j=0, ..., ncols  
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels
in a given region

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

```
// Start with array A[0:nrows][0:ncols]

// Horizontal pass
for i=0, ..., nrows
    for j=1, ..., ncols
        A[i][j] = A[i][j] + A[i][j-1]

// Vertical pass
for i=1, ..., nrows
    for j=0, ..., ncols
        A[i][j] = A[i][j] + A[i-1][j]
```

Integral image

- Efficient way to add up all pixels in a given region

Each value $A[i][j]$ in the integral image represents the sum of the pixels above and to the right.

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

Eg 45 = Sum of pixel in the original image:

3	2	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Finding the sum of all pixels in any box

3	2	1	5	4	7
9	2	4	1	0	6
3	7	8	0	2	4
1	0	2	4	6	1
1	2	4	5	8	6
8	4	2	1	3	5

Sum = 36

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

$$\text{Sum} = 36 = 64 + 3 - 15 - 16$$

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

Sum = 36 = 64

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

$$\text{Sum} = 36 = 64 - 15$$

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

$$\text{Sum} = 36 = 64 - 15 - 16$$

3	5	6	11	15	22
12	16	21	27	31	44
15	26	39	45	51	68
16	27	42	52	64	82
17	30	49	64	84	108
25	42	63	79	102	131

$$\text{Sum} = 36 = 64 - 15 - 16 + 3$$

Summation of all pixels in any rectangle just requires summing (or subtracting) 4 values.

Readings

- M. Turk and A. Pentland, Face Recognition using Eigenfaces, CVPR 1991
- P. Viola and M. Jones, Rapid Object Detection using a Boosted Cascade of Simple Features, CVPR 2001