

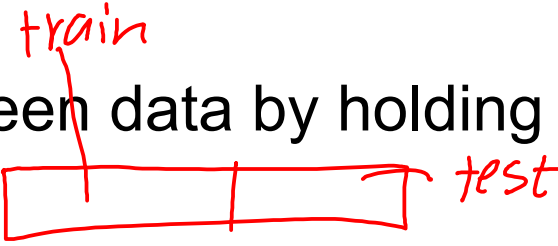
When Models Meet Data 2

Liang Zheng

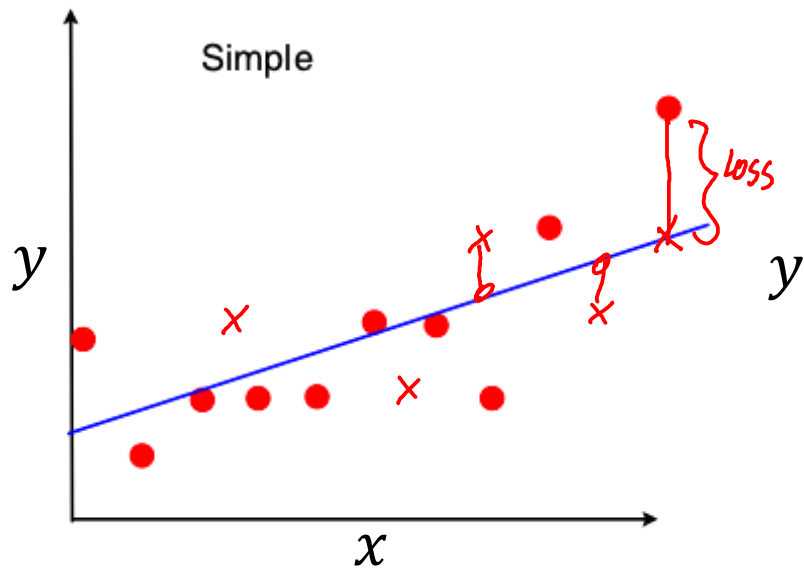
Australian National University

liang.zheng@anu.edu.au

Overfitting

- The aim of a machine learning predictor is to perform well on unseen data.
- We simulate the unseen data by holding out a proportion of the whole dataset.
- This hold out set is called test set.
- In practice, we split data into a training set and a test set.
- **Training set**: fit the model
- **Test set**: not seen during training, used to evaluate generalization performance
- It is important for the user to not cycle back to a new round of training after having observed the test set.

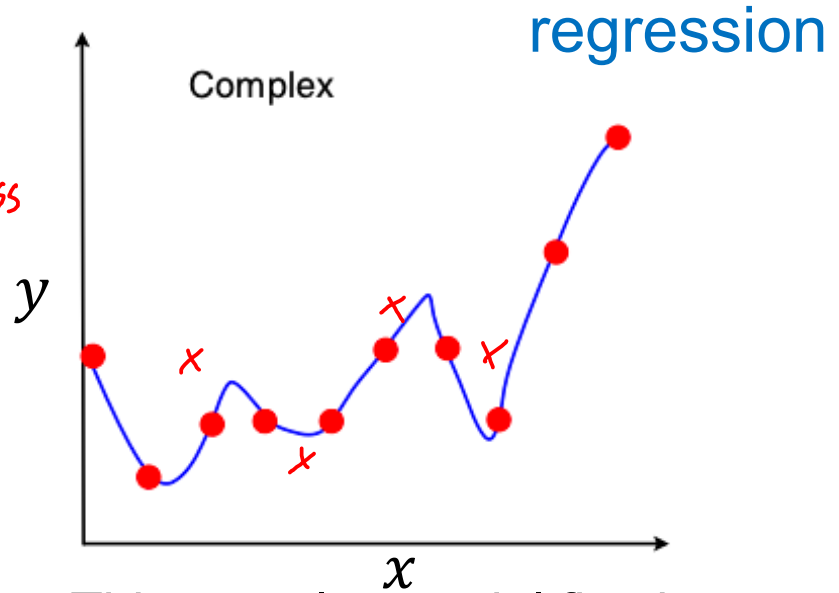
- Empirical risk minimization = average loss can lead to **overfitting**.
- the predictor fits too closely to the training data and does not generalize well to new data



This simple model fits the training data less well.

A larger empirical risk.

A good machine learning model.



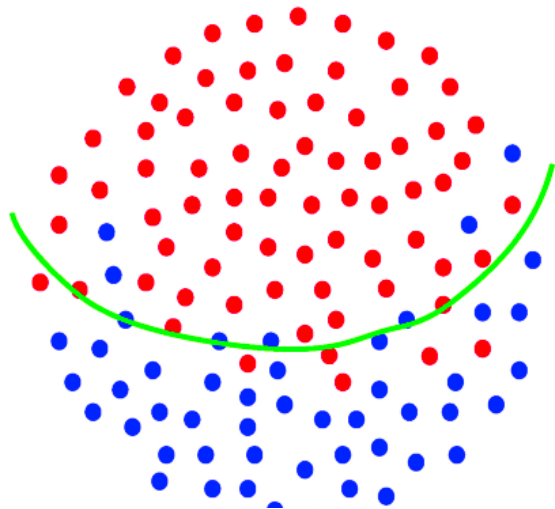
This complex model fits the training data very well.

A very small empirical risk.

A poor machine learning model due to overfitting.

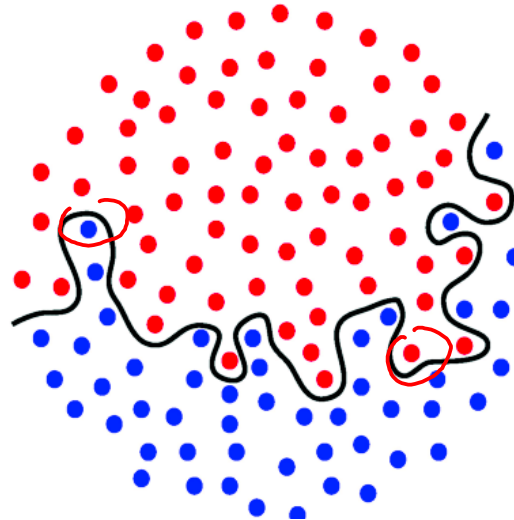
- Empirical risk minimization can lead to **overfitting**.
- the predictor fits too closely to the training data and does not generalize well to new data

A good model



● data, class 1
● data, class 2

A poor model **classification**

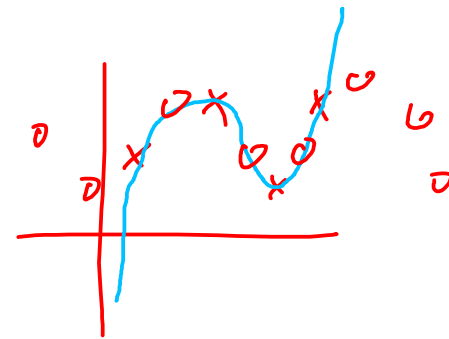


overfitted classification model

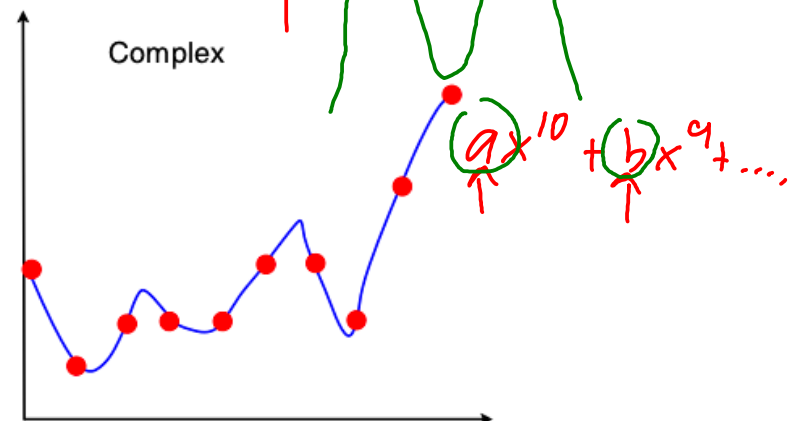
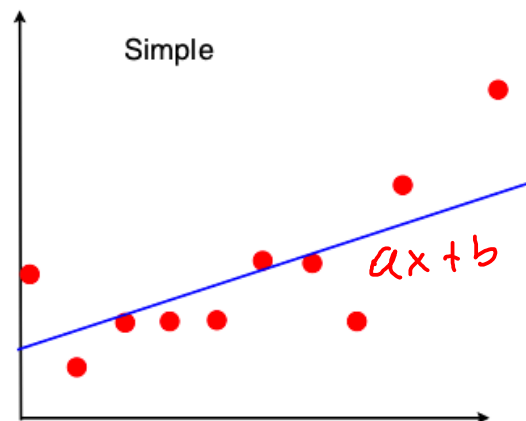
regularised classification model

8.2.3 Regularization to Reduce Overfitting

- When overfitting happens, we have
 - very ^{Zero} small average loss on the training set but large average loss on the test set
- Given a predictor f , overfitting occurs when
 - the risk estimate from the training data $R_{\text{emp}}(f, X_{\text{train}}, y_{\text{train}})$ underestimates the expected risk $R_{\text{true}}(f)$. In other words,
 - $R_{\text{emp}}(f, X_{\text{train}}, y_{\text{train}})$ is much smaller than $R_{\text{true}}(f)$ which is estimated using $R_{\text{emp}}(f, X_{\text{test}}, y_{\text{test}}) \approx R_{\text{true}}(f)$
- Overfitting occurs usually when ^{model} ~~we have little data and a complex hypothesis class~~



- How to prevent overfitting?
- We can bias the search for the minimizer of empirical risk by introducing a penalty term
- The penalty term makes it harder for the optimizer to return an overly flexible predictor
- The penalty term is called regularization.
- Regularization is an approach that discourages complex or extreme solutions to an optimization problem.



$$\begin{matrix} \uparrow & & \uparrow \\ [\lambda_1, \lambda_2, \lambda_3]^T & & \uparrow \\ \text{big} & & \text{small} \end{matrix} \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

$$\cancel{\theta_1} x^2 + \theta_2 x + \theta_3$$

$$\bar{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

$$\begin{matrix} y_1 x_1 \\ \vdots \\ y_n x_n \end{matrix}$$

$$\begin{matrix} \sin x_1 \\ \sin x_2 \\ \vdots \end{matrix}$$

$$X = \begin{bmatrix} 1 & x_1 & x_1^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix}$$

- Example
- Least-squares problem

$$\min_{\theta} \frac{1}{N} \|\mathbf{y} - \mathbf{X}\bar{\theta}\|^2$$

ground truth \hat{y}

- To regularize this formulation, we add a penalty term

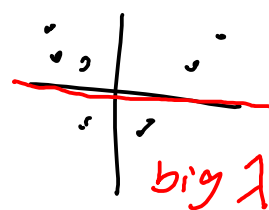
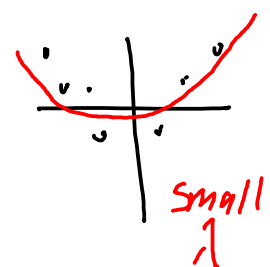
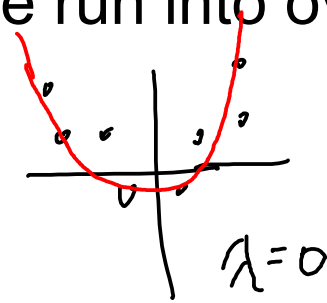
$$\min_{\theta} \left[\underbrace{\frac{1}{N} \|\mathbf{y} - \mathbf{X}\bar{\theta}\|^2}_{\text{Loss}} + \underbrace{\lambda \|\bar{\theta}\|^2}_{\text{regularizer}} \right]$$

- The addition term $\|\bar{\theta}\|^2$ is called the **regularizer** or **penalty term**, and the parameter regularizer λ is the **regularization parameter**.

- λ enables a trade-off between minimizing the loss on the training set and the amplitude of the parameters θ

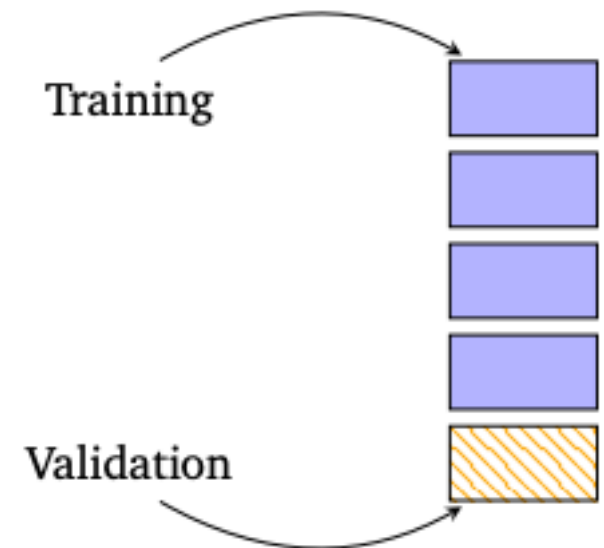
- It often happens that the **amplitude** of the parameters in θ becomes relatively large if we run into overfitting

- λ is a hyperparameter



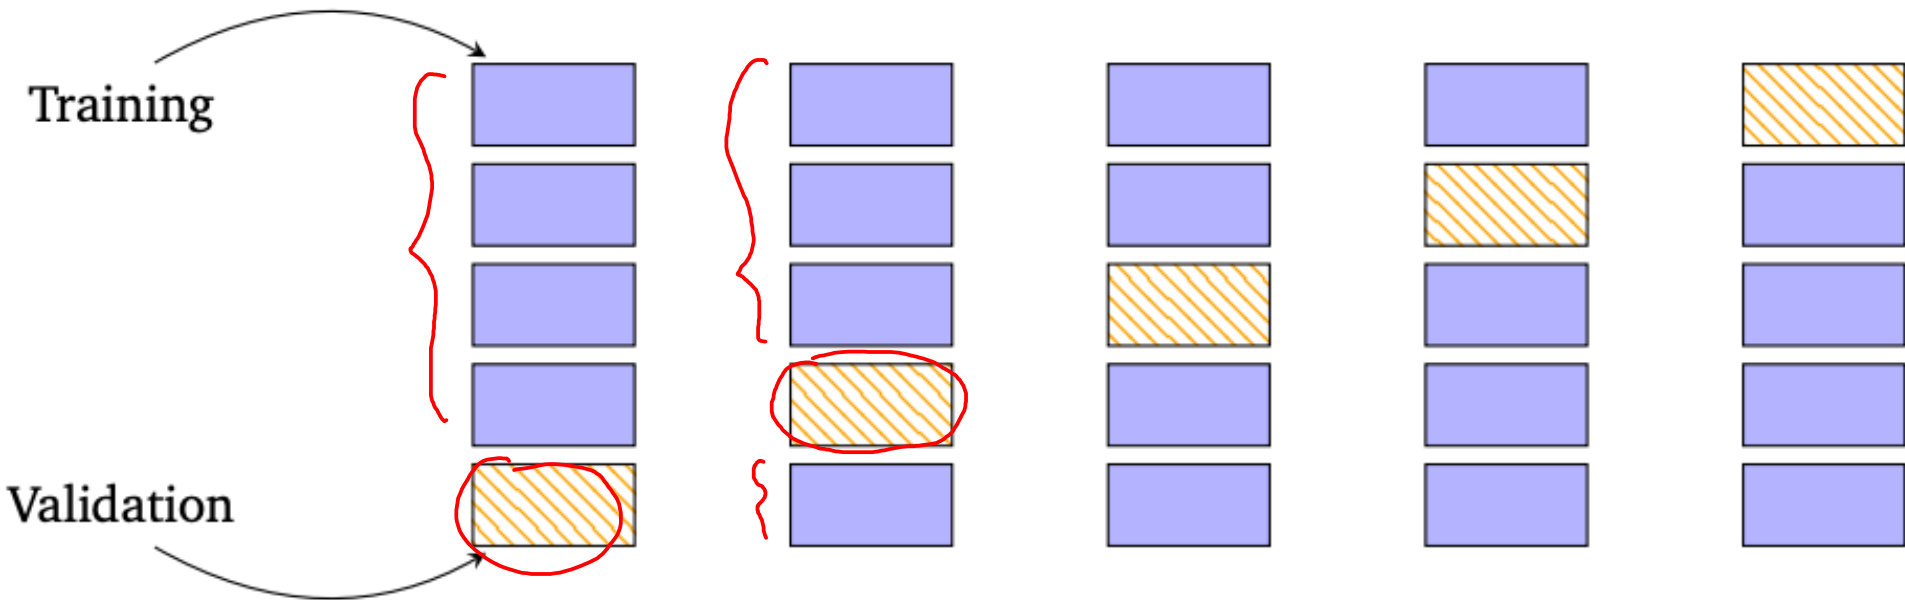
8.2.4 Cross-Validation to Assess the Generalization Performance

- We mentioned that we split a dataset into a **training set** and a **test set**
- we measure **generalization error** by applying the predictor on test data.
- This data is also sometimes referred to as the validation set.
- Validation set is from the entire data, and has no overlap with the training data.
- We want the training set to be **large**
- That leaves the validation set **small**
- A small validation set makes the **result less stable** (large variances)



- Basically, we want the training set to be **large**
- We want the validation to be **large**, too
- How to solve these contradictory objectives?
- **Cross-validation**: (K) -fold cross-validation \equiv *development*

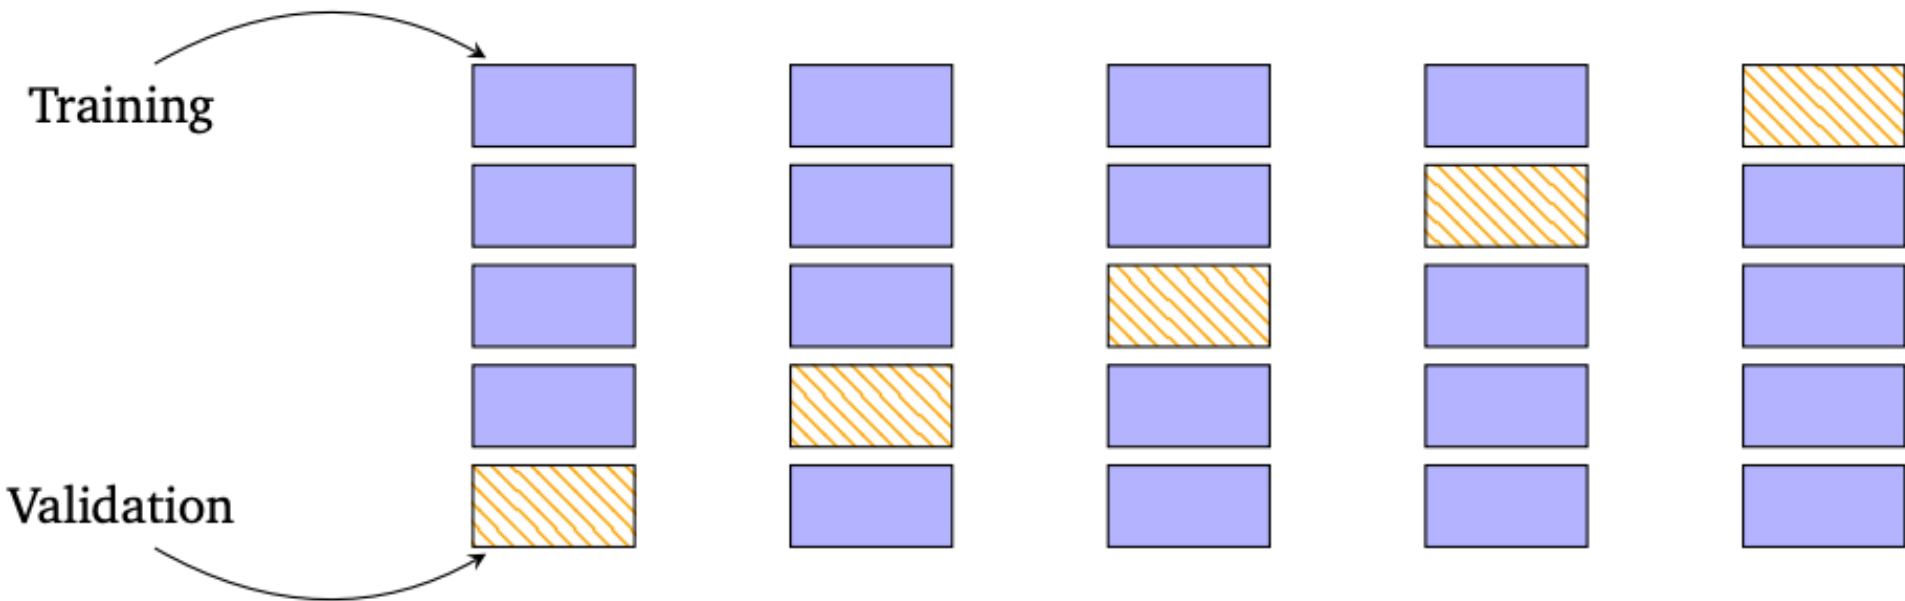
Example: $K = 5$



Cross-validation

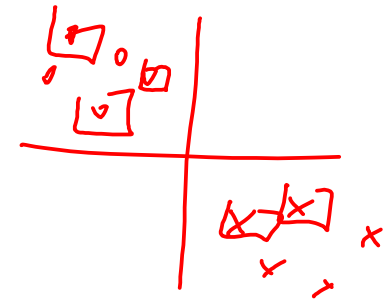
- K -fold cross-validation partitions the data into K chunks
- $K - 1$ trunks form the training set \mathcal{R}
- The last trunk is the validation set \mathcal{V}
- This procedure is repeated for all K choices for the validation set, and the performance of the model from the K runs is averaged

Example: $K = 5$



Cross-validation

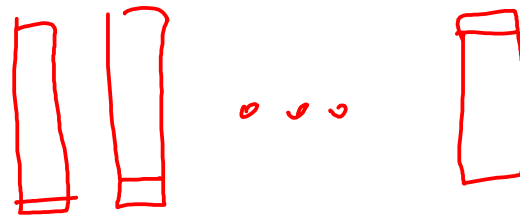
- Formally, we partition our training set into two sets $\mathcal{D} = \mathcal{R} \cup \mathcal{V}$, such that they do not overlap, i.e., $\mathcal{R} \cap \mathcal{V} = \phi$
- We train our model on \mathcal{R} (training set)
- We evaluate our model on \mathcal{V} (validation set)
- We have K partitions. In each partition k :
 - Training set $\mathcal{R}^{(k)}$ produces a predictor $f^{(k)}$
 - $f^{(k)}$ is applied to validation set $\mathcal{V}^{(k)}$ to compute the empirical risk $R(f^{(k)}, \mathcal{V}^{(k)})$
 - All the empirical risks are averaged to approximate the expected generalization error



$$\mathbb{E}_V[R(f, \mathcal{V})] \approx \frac{1}{K} \sum_{k=1}^K R(f^{(k)}, \mathcal{V}^{(k)})$$

Cross-validation – some understandings



- The training set is limited -- not producing the best $f^{(k)}$
- The testing set is limited – ^{may} producing an inaccurate estimation of $R(f^{(k)}, \mathcal{V}^{(k)})$
- After averaging, the results are stable and indicative
- An extreme: leave-one-out cross-validation, where the validation set only contains one example.

- A potential drawback – computation cost
 - The training can be time-consuming
 - If the model has several parameters to tune, it is hard to evaluate those hyperparameters.
- This problem can be solved by parallel computing, given enough computational resources

Check your understanding

- When your model works poorly on the training set, your model will also work poorly on the test set. **most*
- When your model works poorly on the training set, your model may also have overfitting. *by def.*
- Overfitting happens when your model is too complex given your training data. *by def*
- Regularization alleviates overfitting by improving the complexity of your training data. *data is fixed*
- In K -fold cross-validation, we will get more stable test accuracy if K increases.
- In 2-fold cross-validation, you can obtain 2 results from the 2 test sets, and they may differ a lot with each other.









