

ENGN2219/COMP6719

Computer Systems & Organization

Convener: Shoaib Akram

shoaib.akram@anu.edu.au



Australian
National
University

Plan: Week 6

Last week: Instruction Set Architecture (specification)

This Week: Microarchitecture (implementation)

Machine Language

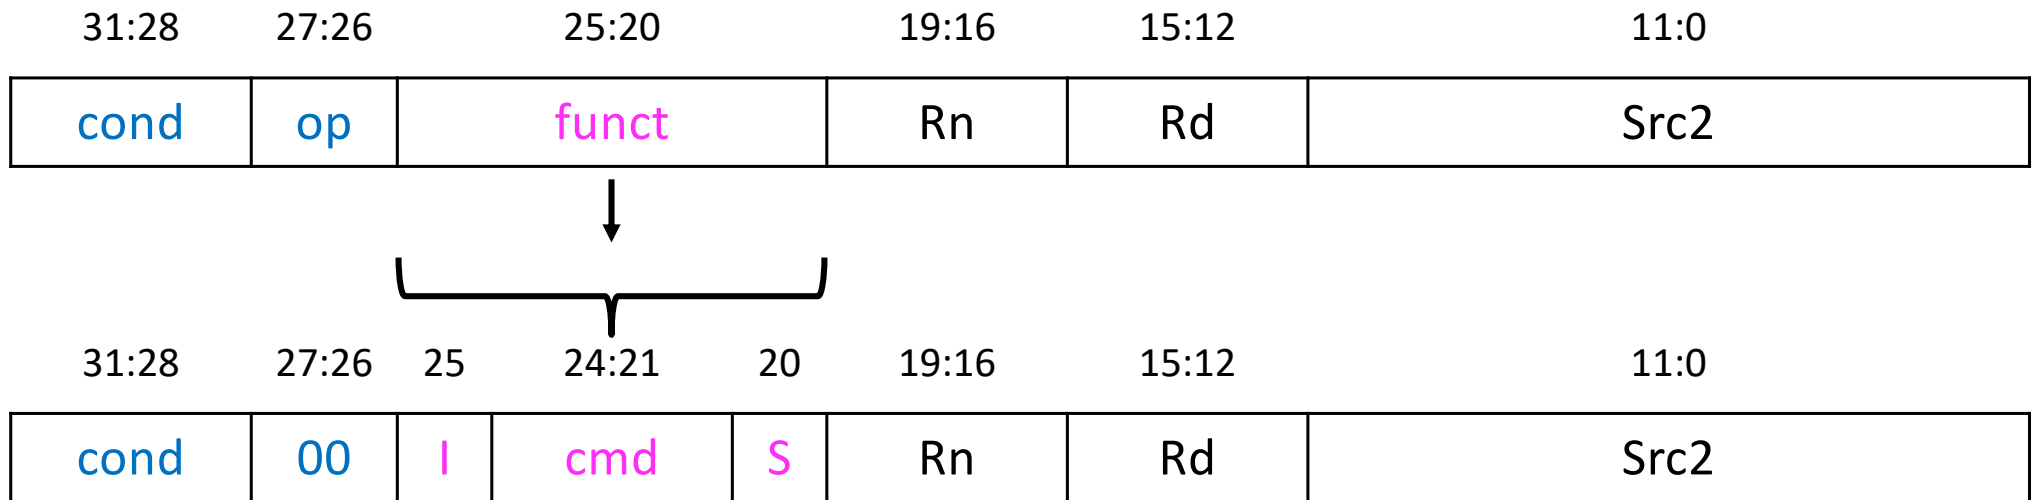
- Instructions are stored in memory as fixed-size words
 - For ARMv4, each instruction is 32 bits wide
- ARM defines three instruction formats
 - Data processing (DP)
 - Memory
 - Branch
- Instruction *fields* encode the instruction operation and its operands
- Understanding binary encoding of instructions is necessary for building the CPU

Instruction Format – 1: DP

31:28	27:26	25:20	19:16	15:12	11:0
cond	op	funct	Rn	Rd	Src2

- Operands
 - Rn: first source operand register (0000, 0001, ..., 1111)
 - Src2: second source register or immediate
 - Rd: destination register
- Control fields
 - cond: specifies conditional execution (1110 for unconditional)
 - op: the operation code or opcode (00)
 - funct: the function/operation to perform

Instruction Format – 1: DP



- op = 00 for DP instructions
- cmd specifies the specific DP instruction (0100 for ADD and 0010 for SUB)
- I-bit
 - I = 0: Src2 is a register
 - I = 1: Src2 is an immediate
- S-bit: 1 if the instruction sets the condition flags

DP with Src2 as Immediate

- Bit 25 (I) informs the CPU how to interpret Src2
 - I = 1, CPU interprets Src2[7:0] as an unsigned 8-bit constant
- Format (Src2 = immediate)

ADD R0, R1, #16

ADD Rd, Rn, #imm8

31:28		27:26		25	24:21		20	19:16		15:12		11:8		7:0			
cond		00		1	cmd		S	Rn		Rd		0	0	0	0	imm8	

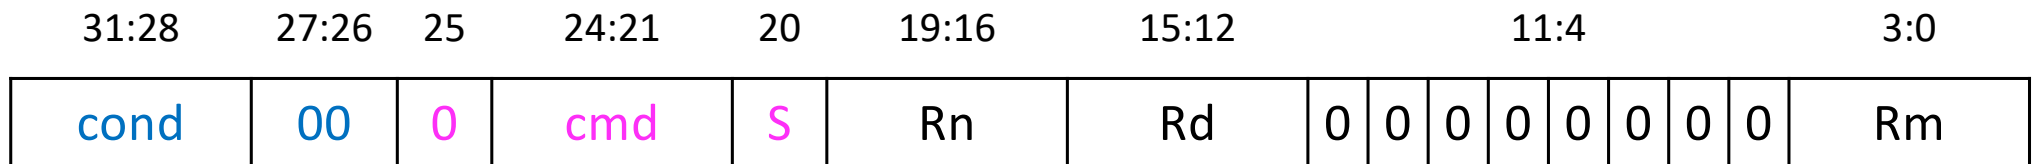
DP with Src2 as Register

- Bit 25 informs the CPU how to interpret Src2
 - I = 0, CPU interprets Src2[3:0] as a register
- Format (Src2 = Register)

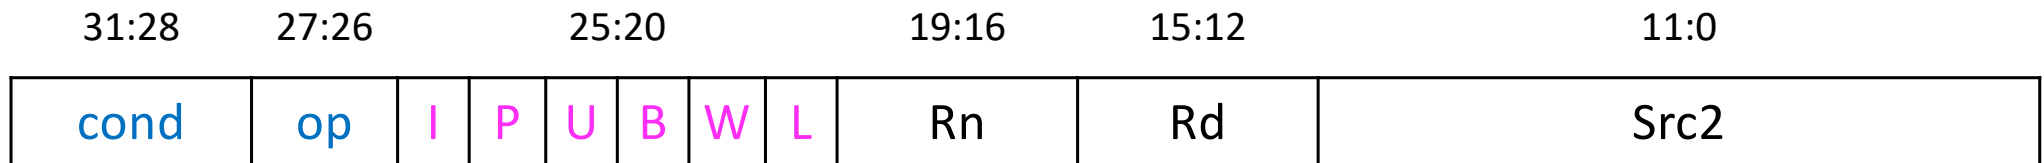
ADD R0, R1, R3



ADD Rd, Rn, Rm



Instruction Format – 2: Memory



- op = 01
- Rn = base register (base address)
- Rd = destination (load), source (store)
- Src2 = offset (register, shifted register, immediate)
- funct = 6 control bits
 - I (Bit 25): Encoding of Src2
 - L (Bit 20): Load or Store
 - Remaining bits (ignore)

LDR with Src2 as Immediate

- I (Bit 25) = 1: Src2 = imm12 where imm2 is a 12-bit unsigned offset added to the value in the base register (Rn)
- Format of Load Register instruction

LDR R0, [R1, #12]

↓ ↓ ↓

LDR Rd, [Rn, #imm12]
- L (Bit 20) = 1: CPU performs an LDR

31:28	27:26	25:20						19:16	15:12	11:0
cond	01	1	1	1	0	0	1	Rn	Rd	imm12

STR with Src2 as Immediate

- I (Bit 25) = 1: Src2 = imm12 where imm2 is a 12-bit unsigned offset added to the value in the base register (Rn)
- Format of STore Register instruction

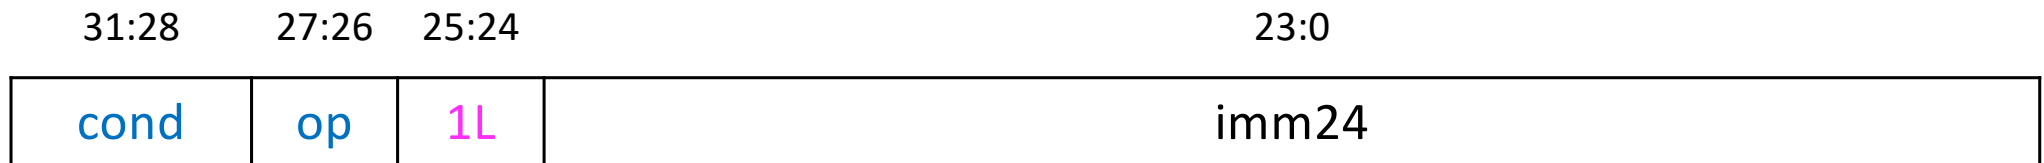
STR R0, [R1, #12]

↓ ↓ ↓

STR Rd, [Rn, #imm12]
- L (Bit 20) = 0: CPU performs an STR

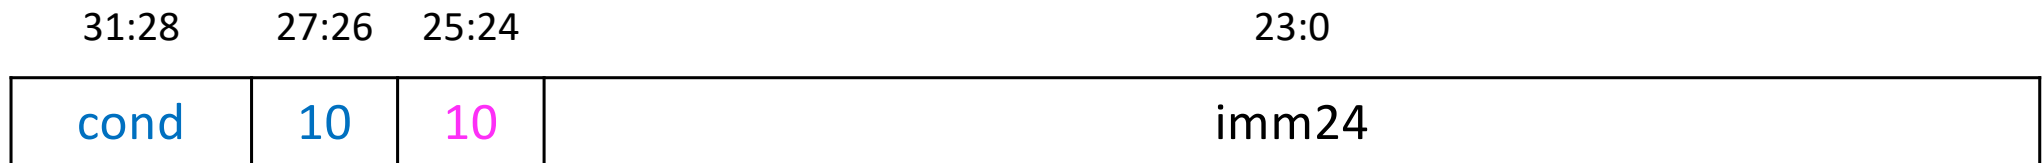
31:28	27:26	25:20						19:16	15:12	11:0
cond	01	1	1	1	0	0	1	Rn	Rd	imm12

Instruction Format – 3: Branch



- op = 10
- imm24 = 24-bit **signed** immediate
- The two bits 25:24 form the **funct** field
 - Bit 25 is always 1
 - L bit: L = 0 for B (Branch)
 - L bit: L = 1 for BL (Branch and Link, ignore for now)

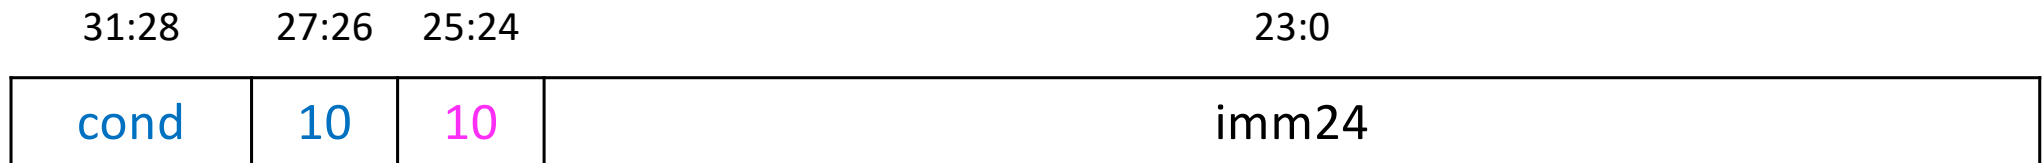
Branch with L = 0



- op = 10
- imm24 = 24-bit **signed** immediate
- The two bits 25:24 form the **funct** field
 - Bit 25 is always 1
 - L bit: L = 0 for B (Branch)

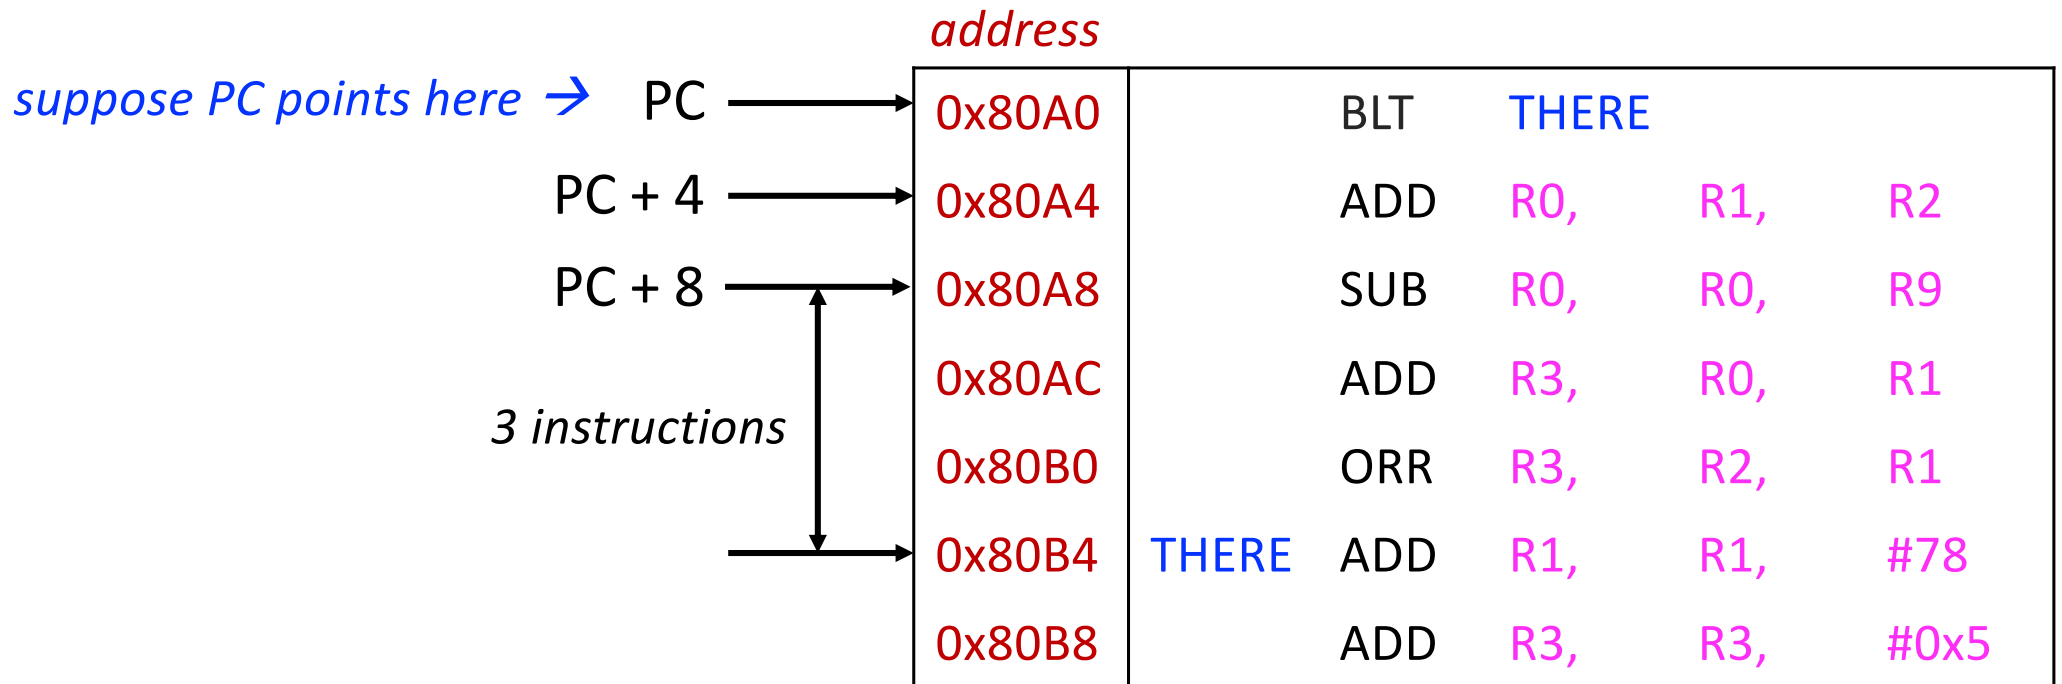
- Format
B TARGET
↓
B imm24

Branch Target Address (BTA)



- The 24-bit two's complement **imm24** field specifies the instruction address relative to **PC + 8**
 - Why **PC + 8**? (historical reasons)
- **BTA**: The address of the next instruction to execute if the branch is taken
- The imm8 field is the number of instructions between the BTA and PC + 8 (two instructions past the branch)

BTA Calculation Example

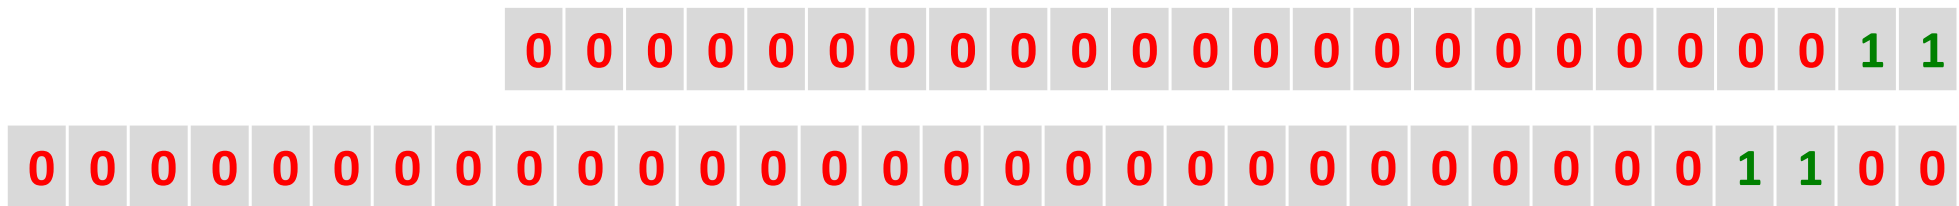


31:28	27:26	25:24	23:0
cond	10	10	imm24 = 3 (00000000000000000000000011)

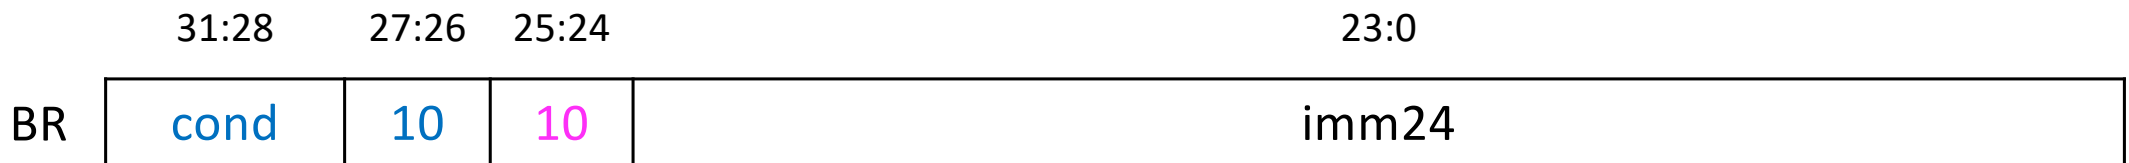
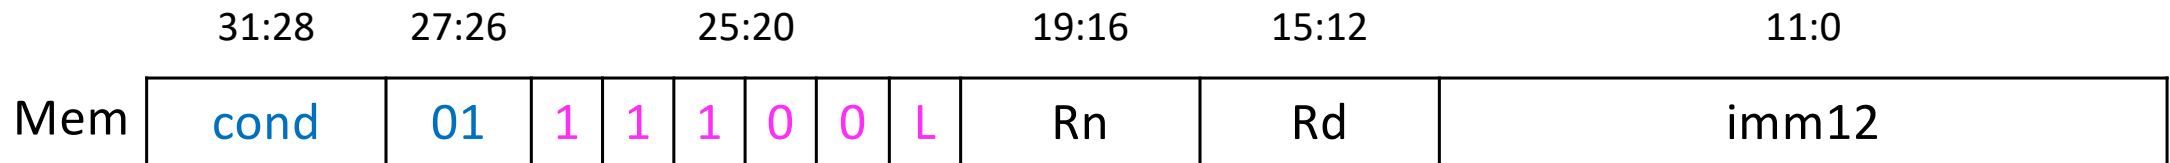
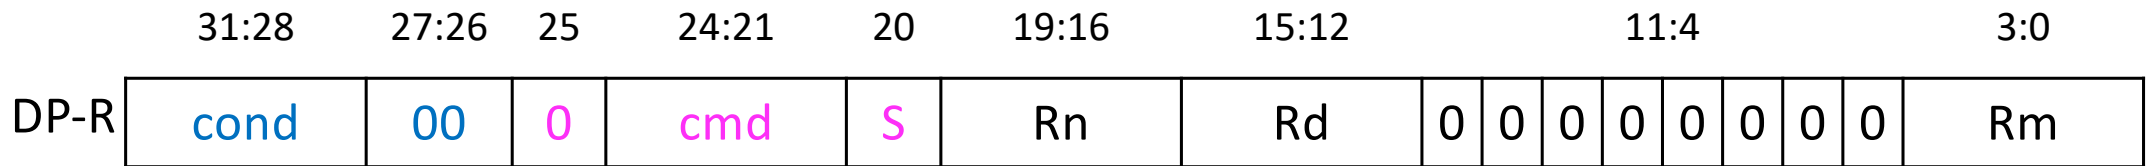
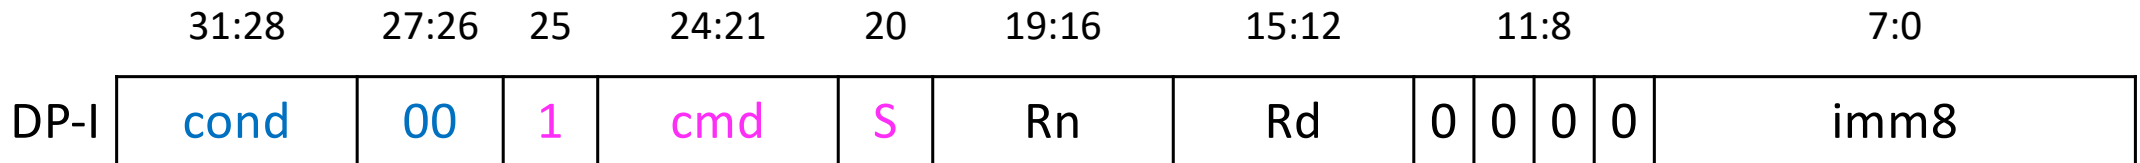
BTA Calculation

The processor calculates the BTA in three steps

1. Shift left imm24 by 2 (to convert words to bytes)
2. Sign-extend (copy Instr₂₅ into Instr_{31:26})
3. Add PC + 8

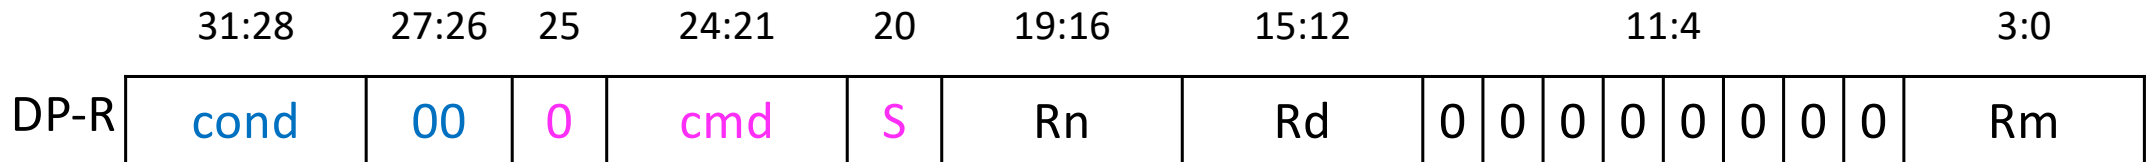


Summary of Formats



Exercise – 1

- Write the 32-bit ARM machine language representation for the following instruction?
 - SUB R8, R9, R10
 - For SUB, **cmd** is **0010**; for unconditional exec, **cond** = **1110**



Exercise – 1

- Write the 32-bit ARM machine language representation for the following instruction?
 - SUB R8, R9, R10
 - For SUB, **cmd** is **0010**; for unconditional exec, **cond** = **1110**

	31:28	27:26	25	24:21	20	19:16	15:12	11:4								3:0
DP-R	cond	00	0	cmd	S	Rn	Rd	0	0	0	0	0	0	0	0	Rm

	31:28	27:26	25	24:21	20	19:16	15:12	11:4								3:0
SUB	1110	00	0	0010	0	1001	1000	0	0	0	0	0	0	0	0	1010

HEX: E0 49 80 0A

Exercise – 2

- Write the 32-bit ARM machine language representation for the following instruction?
 - ADD R0, R1, #42
 - For ADD, **cmd** is **0100**; for unconditional exec, **cond** = **1110**

	31:28	27:26	25	24:21	20	19:16	15:12	11:8	7:0
DP-I	cond	00	1	cmd	S	Rn	Rd	0 0 0 0	imm8

Exercise – 2

- Write the 32-bit ARM machine language representation for the following instruction?
 - ADD R0, R1, #42
 - For ADD, **cmd** is 0100; for unconditional exec, **cond** = 1110

	31:28	27:26	25	24:21	20	19:16	15:12	11:8	7:0
DP-I	cond	00	1	cmd	S	Rn	Rd	0 0 0 0	imm8
	31:28	27:26	25	24:21	20	19:16	15:12	11:8	7:0
ADD	1110	00	1	0100	0	0001	0000	0 0 0 0	00101010

HEX: E2 81 00 2A

Exercise – 3

- Write the 32-bit ARM machine language representation for the following instructions?
 - LDR R7, [R0, #16]

RISC vs. CISC Architectures

- Reduced Instruction Set Computer (RISC)
 - Small # of simple instructions
 - Simple hardware (e.g., easy to decode)
 - ARM, RISC-V, QuAC, MIPS
- Complex Instruction Set Computer (CISC)
 - Many more instructions, and some instructions are very complex
 - x86 DP/ALU instructions can have memory operands
 - One CISC instruction = Many RISC instructions
 - Complex instructions are costly to implement

RISC Principles

Simplicity favors regularity

- Consistent number of operands in DP instructions

Make the common case fast

- Include a few frequently used instructions

Smaller is faster

- Use a small and fast register file
- Large # registers = Large decoding circuitry

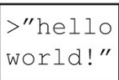


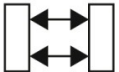
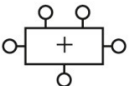

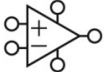


Good design demands good compromises

- Simplicity encourages a single instruction format, but that is too restrictive
- Define multiple formats, but have some regularity among instructions to simplify decoding



*Patterson & Hennessy
2018 Turing Award*

The Big Picture

Application Software		Programs
Operating Systems		Device Drivers
Architecture		Instructions Registers
Micro-architecture		Datapaths Controllers
Logic		Adders Memories
Digital Circuits		AND Gates NOT Gates
Analog Circuits		Amplifiers Filters
Devices		Transistors Diodes
Physics		Electrons

- The *compiler* translates high-level programs to assembly
- The *assembler* translates assembly code to machine code

Software

Hardware

ISA is the Sw/Hw boundary
(Contract/agreement)



- Machine code resides in main memory
- CPU fetches/executes instructions
- Microarchitecture is the specific arrangement of adders, memories, registers, etc., to implement an ISA

Plan for Weeks 6 – 7

Week 5

- We simplified the ISA by setting some control bits to 0 or 1

Week 6

- Build a CPU for the simplified ARM 32-bit instruction set

Week 7

- Study the remaining ISA features

Microarchitecture

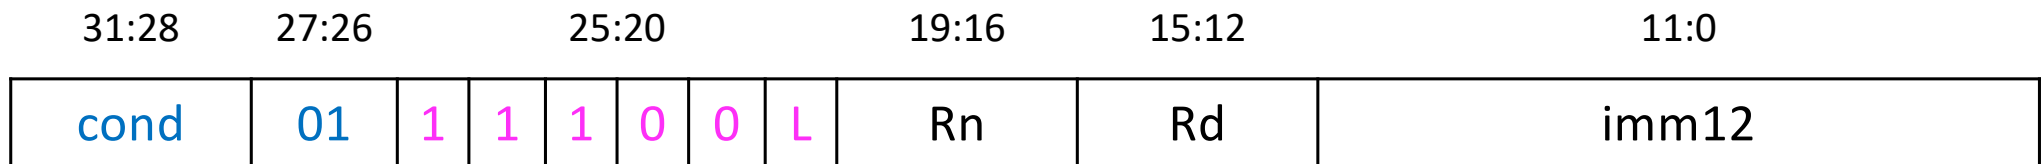
- Two interacting parts
 - Datapath
 - Control unit
- Datapath *operate* on words of data
 - *Register file, ALU, memories*
- Control unit *informs* the datapath how to execute an instruction
 - *Reads the instruction and generate multiplexer selects, register enable, and memory write signals*

LDR Instruction

- We will add the logic for one instruction at a time beginning with the LDR instruction
- Format of Load Register instruction

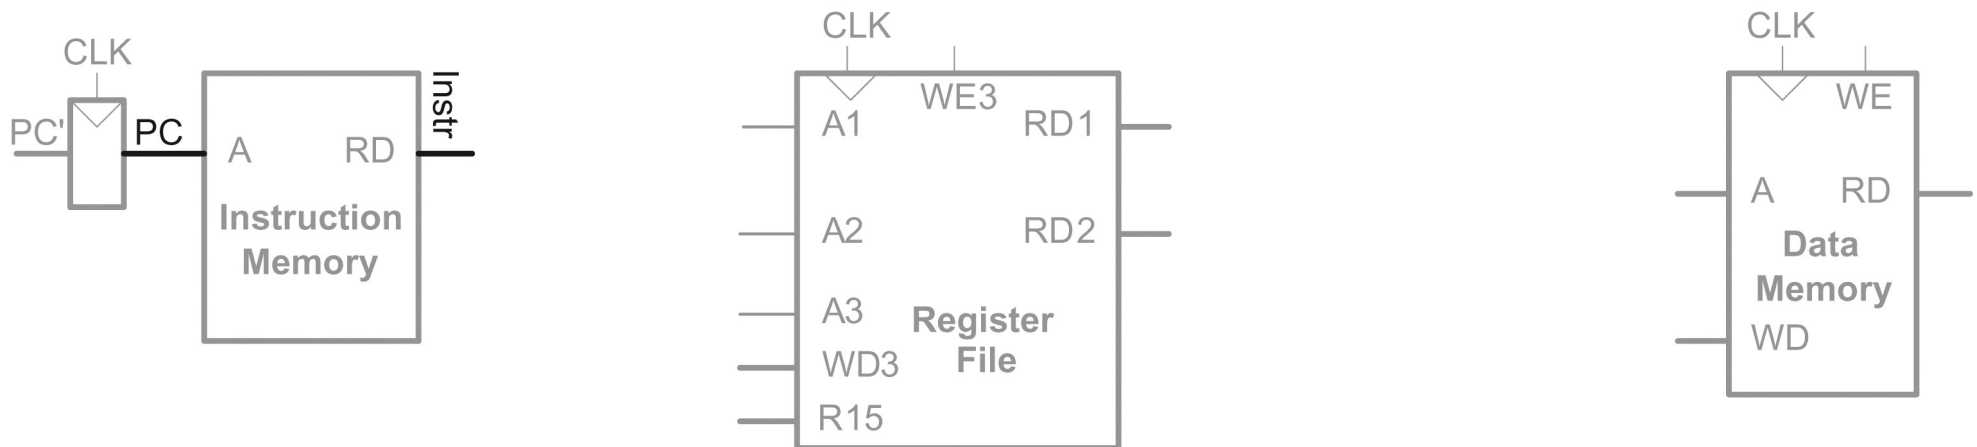
LDR R0, [R1, #12]

LDR Rd, [Rn, #imm12]



The LDR Datapath

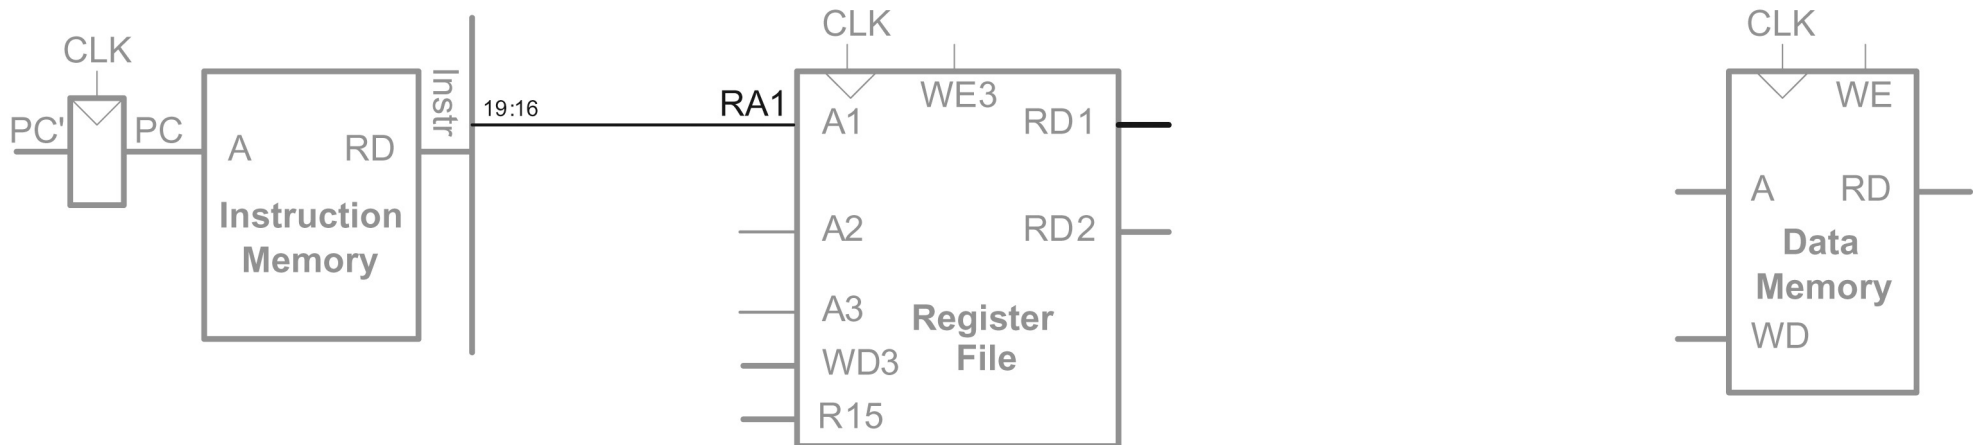
Step 1: Read (Fetch) instruction from memory



31:28	27:26	25:20						19:16	15:12	11:0			
cond	01	1	1	1	0	0	L	Rn	Rd	imm12			

The LDR Datapath

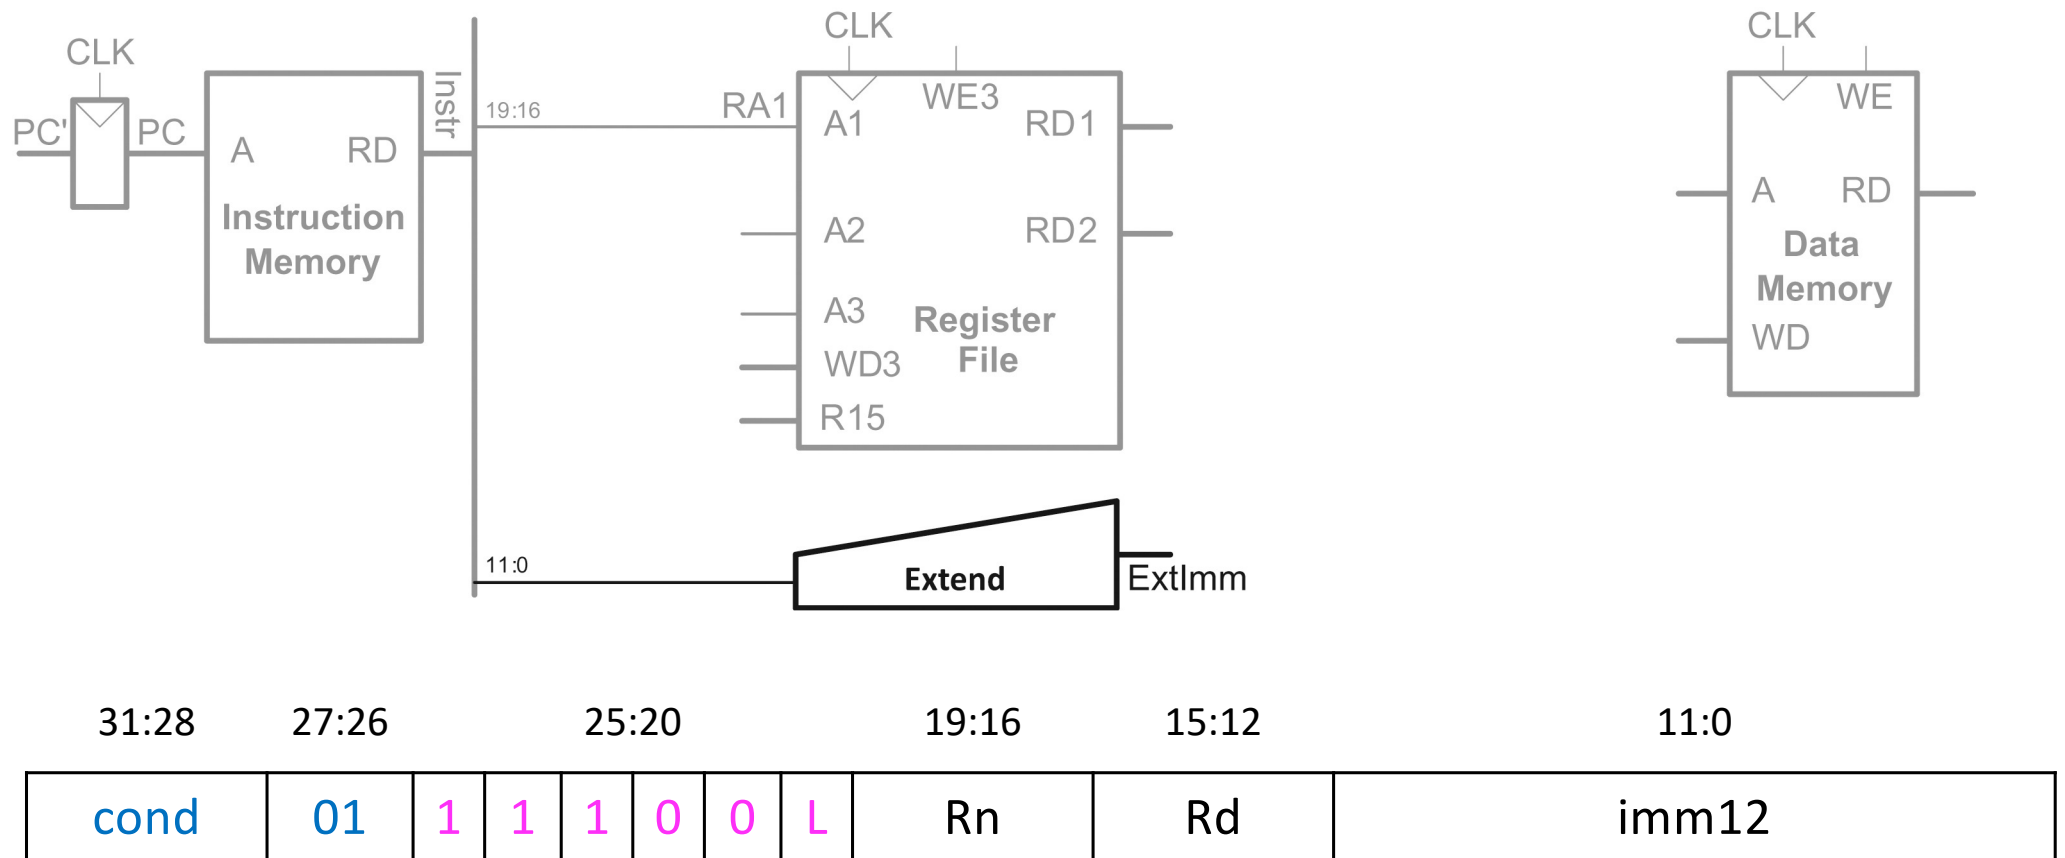
Step 2: Read source operand (base register) from register file



31:28	27:26	25:20						19:16	15:12	11:0			
cond	01	1	1	1	0	0	L	Rn	Rd	imm12			

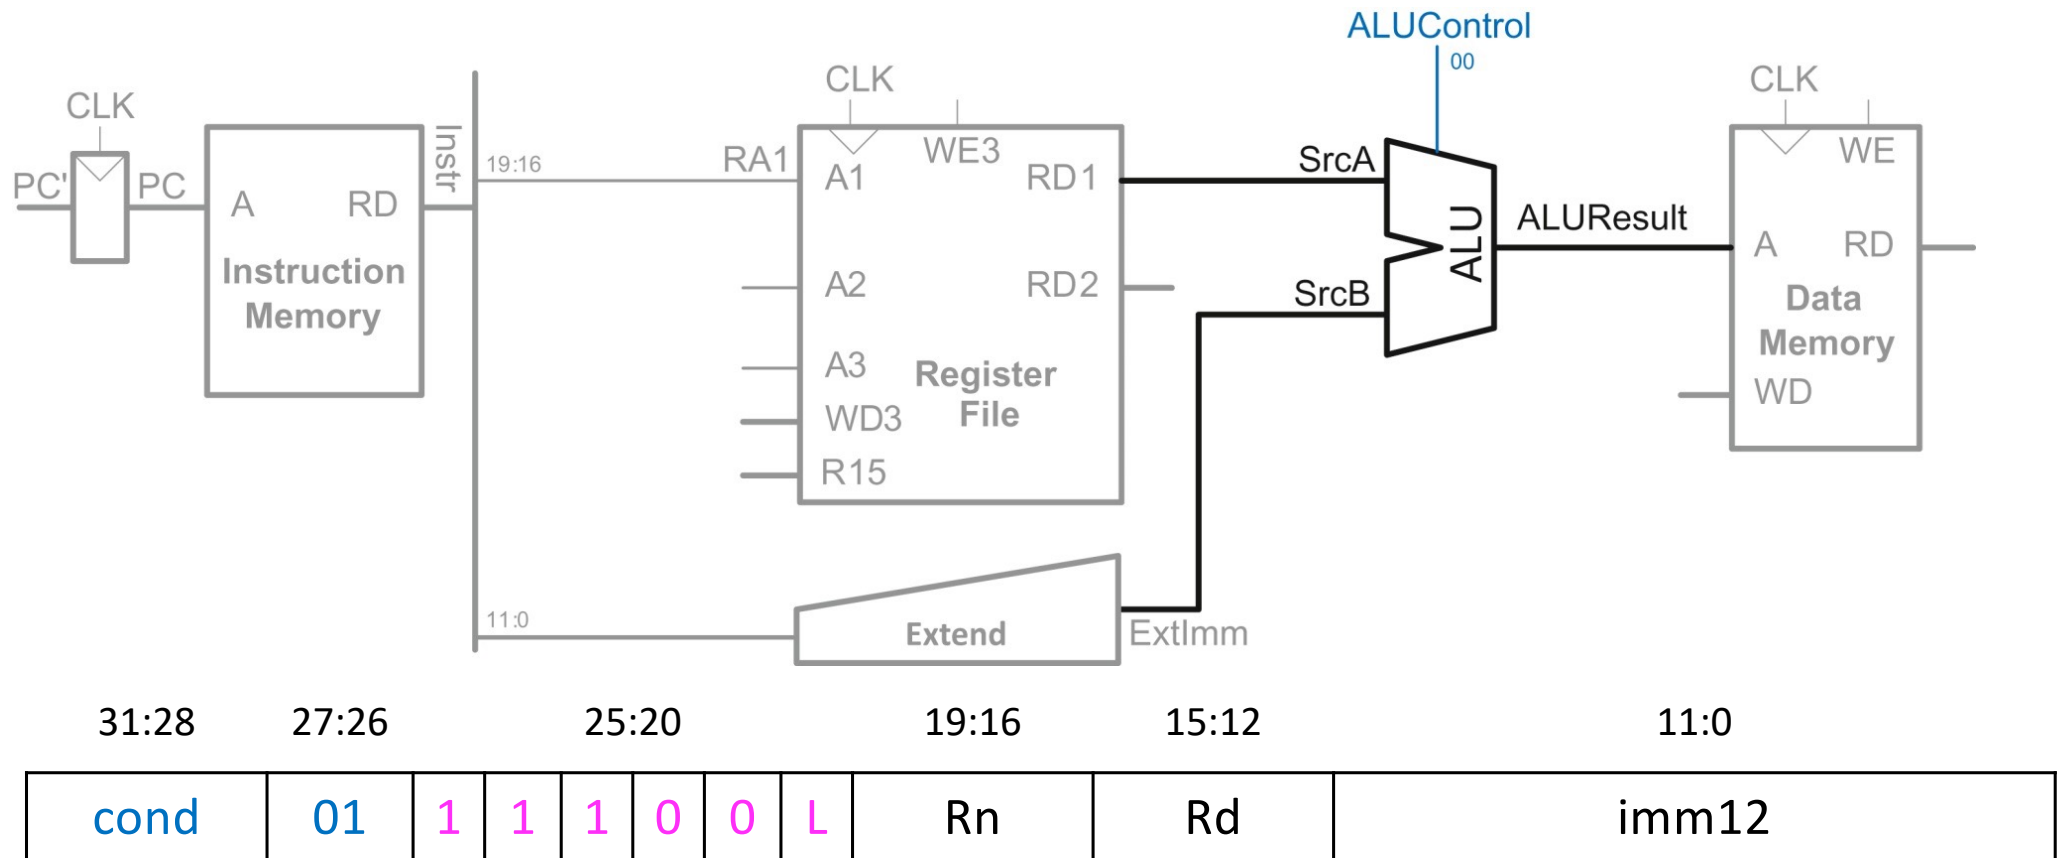
The LDR Datapath

Step 3: Zero-extend the immediate field stored in Instr_{11:0}



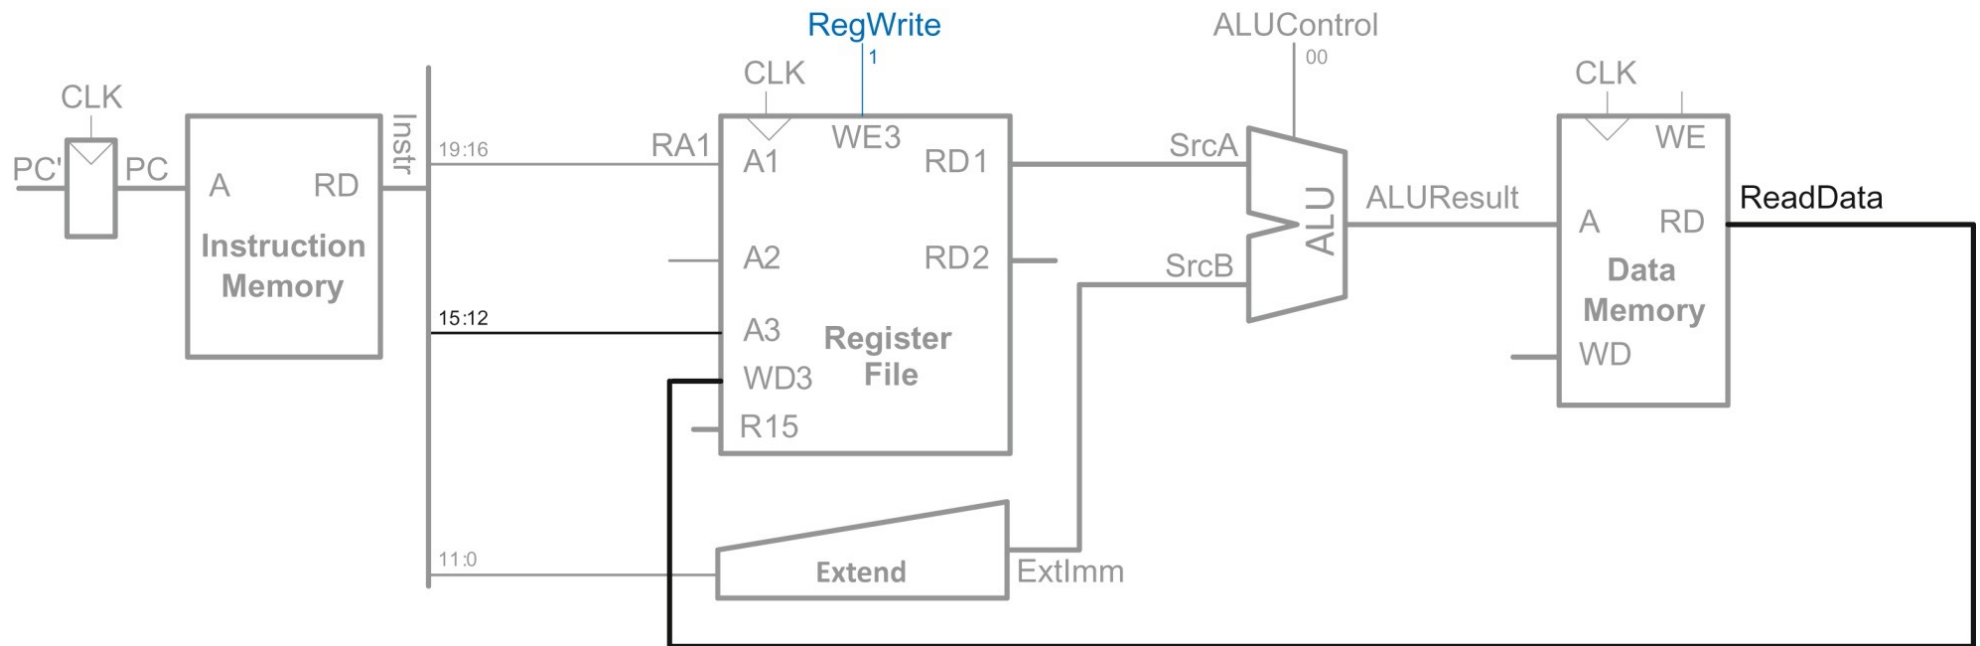
The LDR Datapath

Step 4: Compute memory address (**ALUControl** = 00)



The LDR Datapath

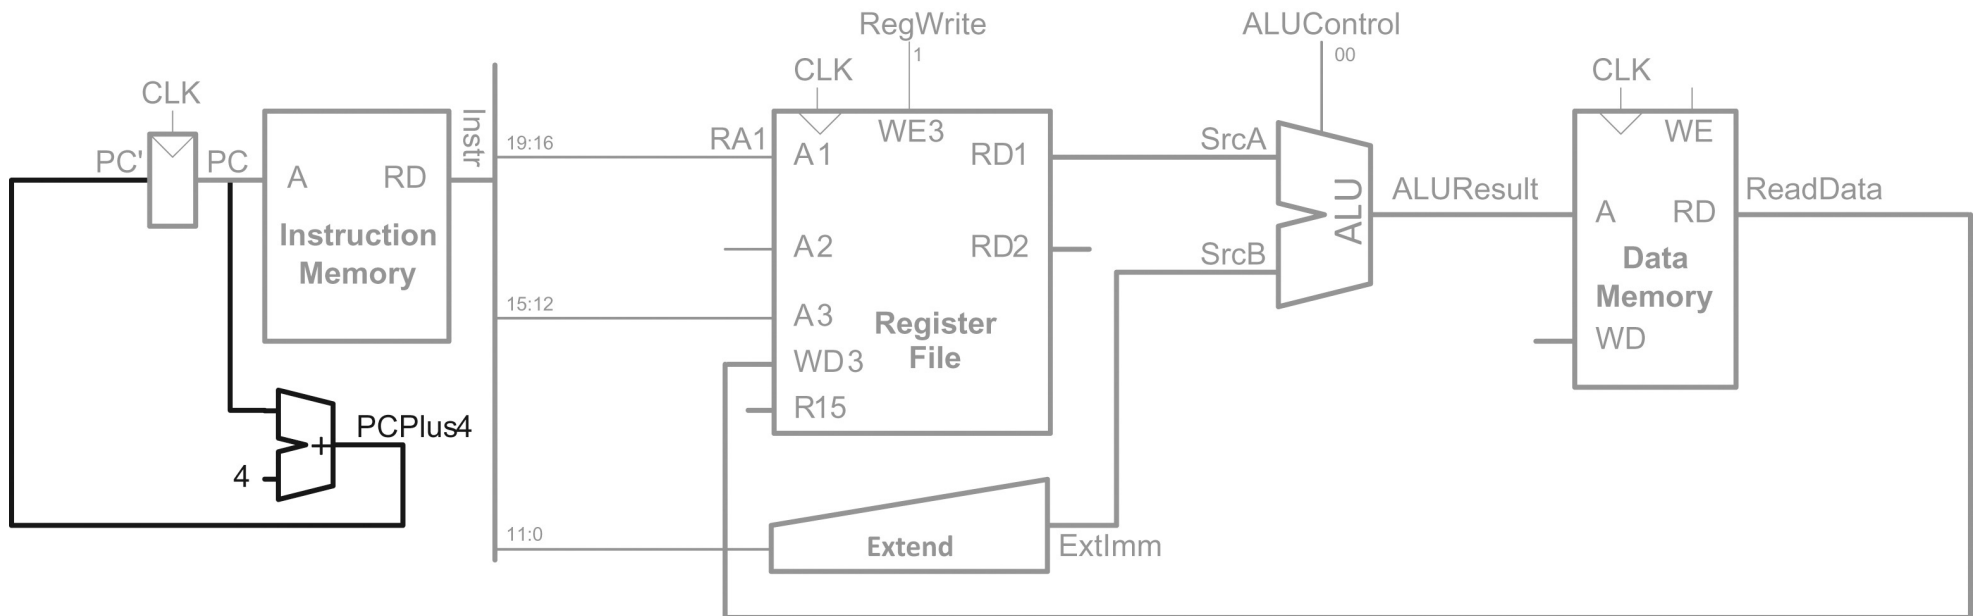
Step 5: Write back data from read by data memory to Rd in Reg File



31:28	27:26	25:20	19:16	15:12	11:0
cond	01	11100L	Rn	Rd	imm12

The LDR Datapath

Step 6: Compute address of next instruction ($PC' = PC + 4$)



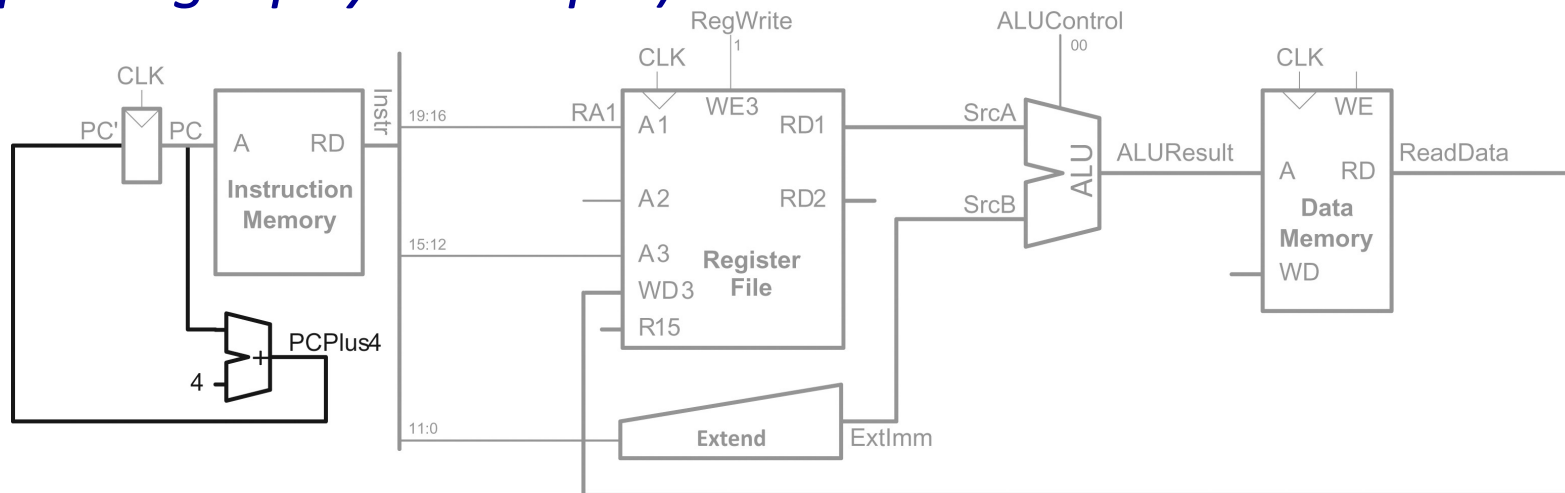
PC will become PC' the following cycle (recall photography example)

31:28	27:26	25:20	19:16	15:12	11:0
cond	01	1 1 1 0 0 L	Rn	Rd	imm12

The LDR Datapath

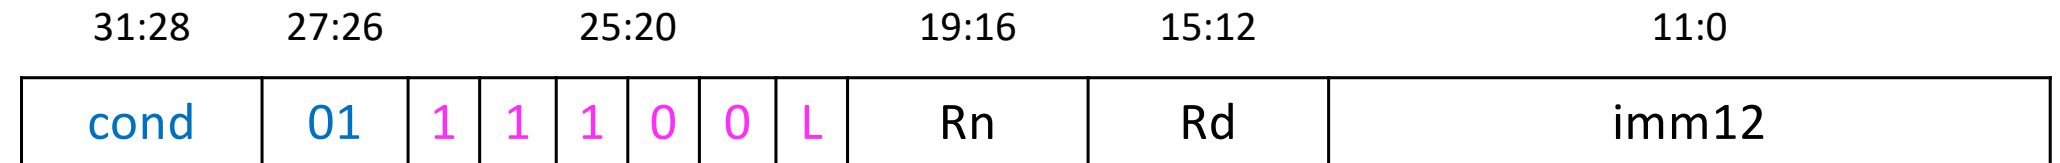
Step 6: Compute address of next instruction ($PC' = PC + 4$)

*Important: PC will become PC' the following cycle
(recall photography example)*

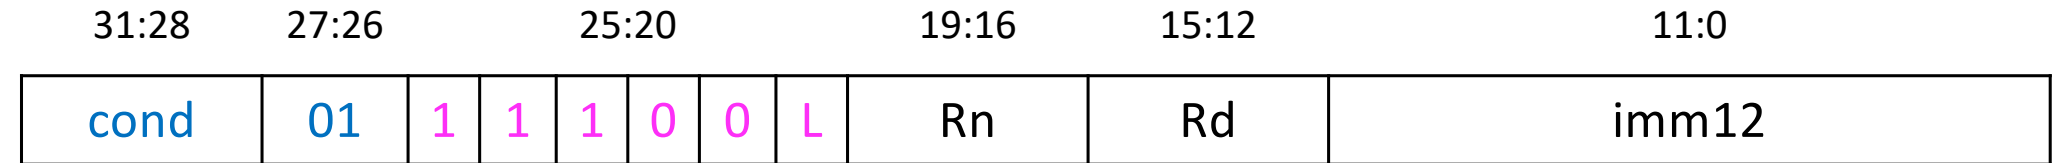


31:28	27:26	25:20					19:16	15:12	11:0		
cond	01	1	1	1	0	0	L	Rn	Rd	imm12	

Step 7/a: Reading register R15 returns PC + 8



Step 7/b: Writing register R15 (PC may be an instruction's result)

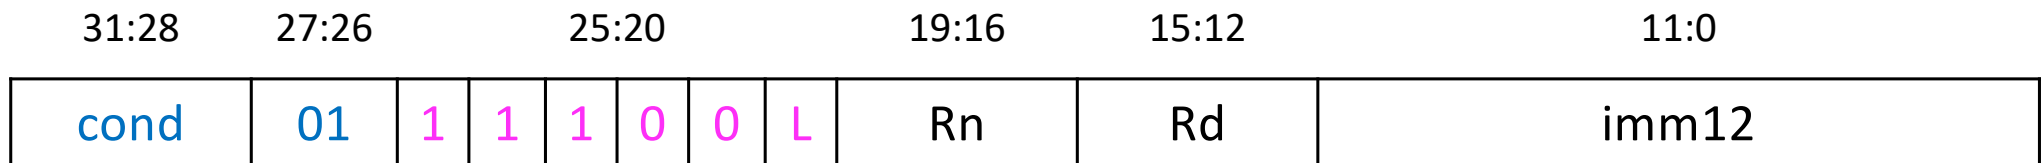


STR Instruction

- STR instruction uses the same instruction format
- LDR and STR behave differently at the machine level
- Rd is a source operand (specifies the register to store to mem)
- Format of **ST**ore **R**egister instruction

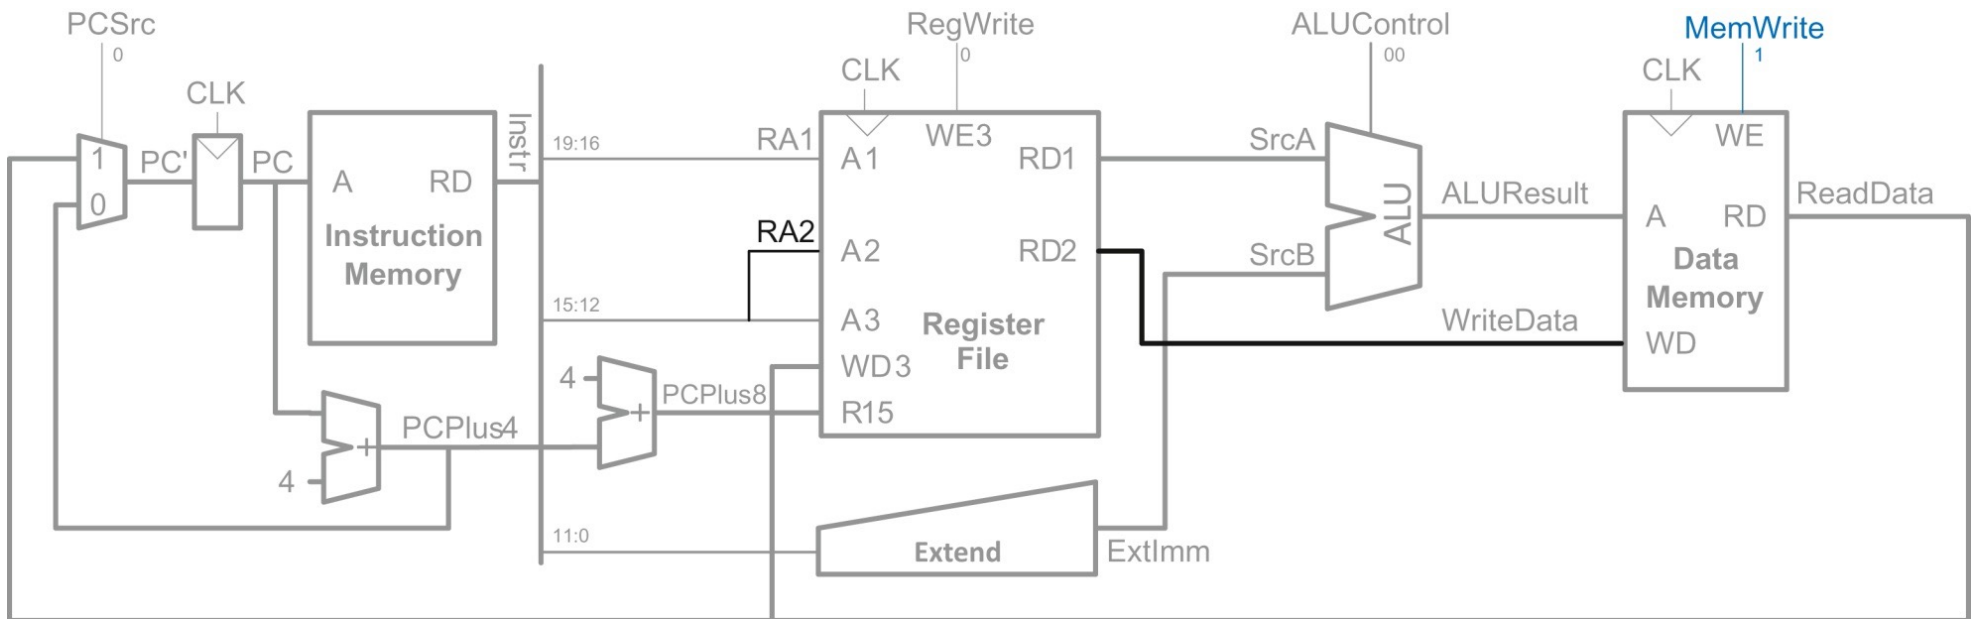
STR R0, [R1, #12]

STR Rd, [Rn, #imm12]



The STR Datapath

Step 8: Read a second register (Rd) and write its value to memory



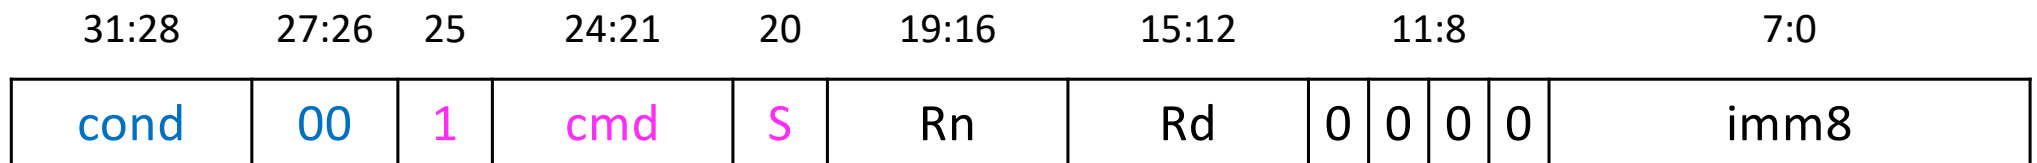
31:28	27:26	25:20					19:16	15:12	11:0			
cond	01	1	1	1	0	0	L	Rn	Rd	imm12		

DP Instructions: Immediate

- Like the LDR instruction, but two important differences
 - imm8 instead of imm12
 - The destination register stores the result of the ALU operation instead of memory access
- Format

ADD R0, R1, #16

ADD Rd, Rn, #imm8

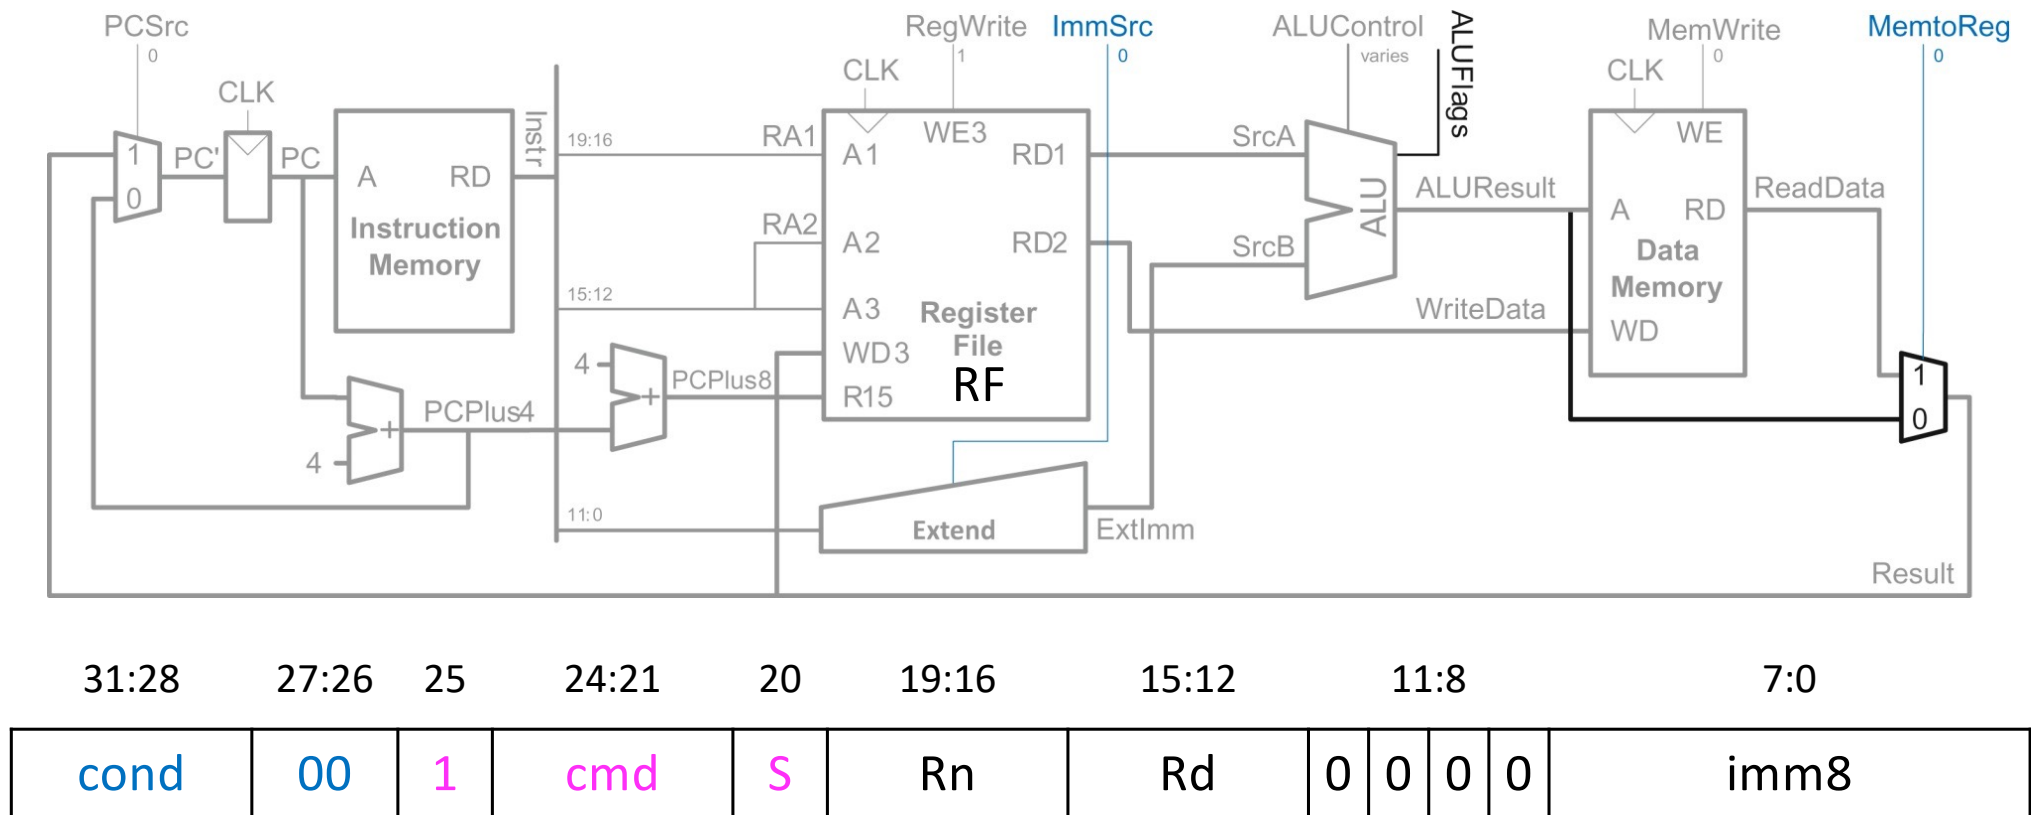


Recall the ALU Functions

ALUControl	Function
00	ADD
01	SUB
10	AND
11	ORR

DP-Immediate Datapath

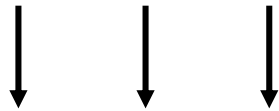
Step 9: Change extend block, and add signal to write ALU result to RF



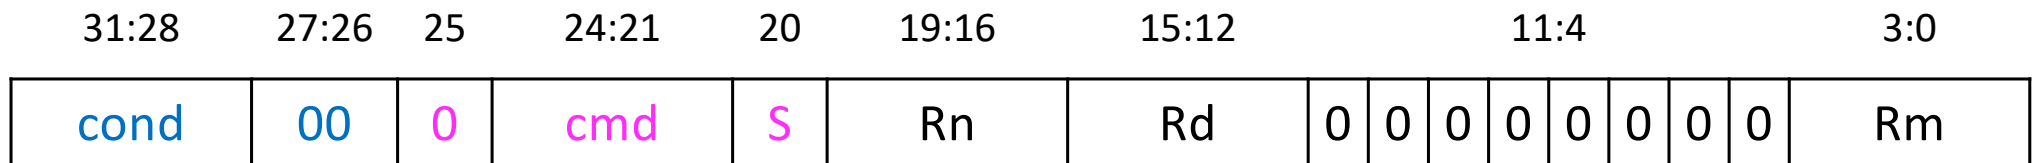
DP Instructions: Register

- The second source operand is Rm instead of an immediate
- Place Rm on the A2 port of the register file for DP instructions with register as the second operand
- Format

ADD R0, R1, R3

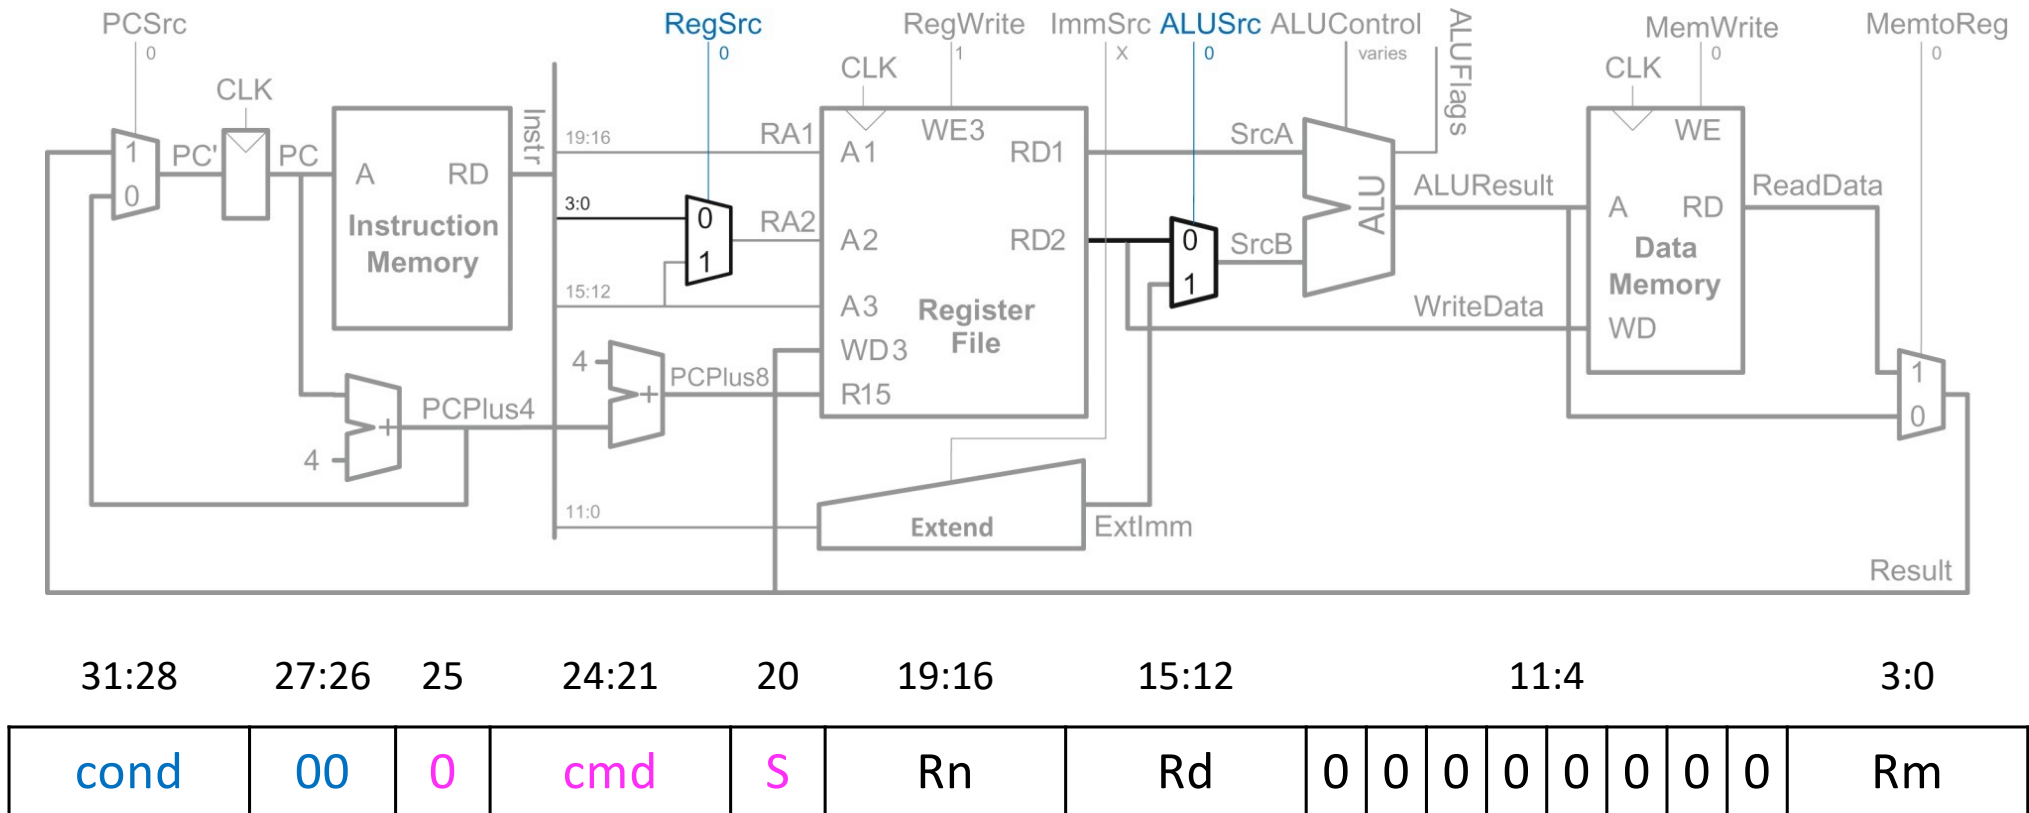


ADD Rd, Rn, Rm



DP-Register Datapath

Step 10: Read 2nd register (Rm) from Reg File and send RD2 to ALU



Branch Instruction: Unconditional

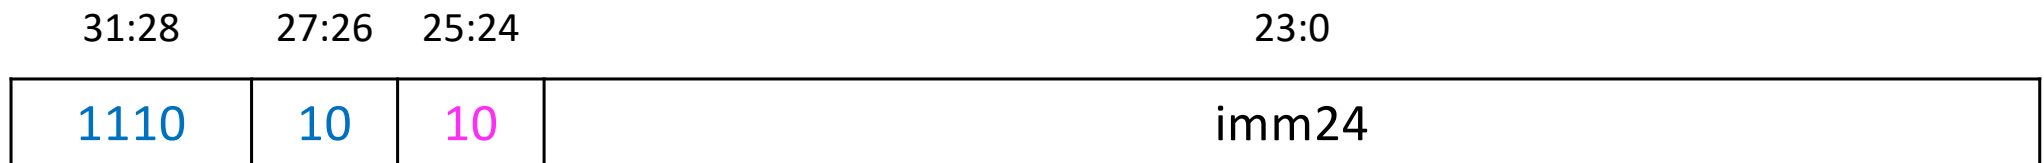
- The second source operand is Rm instead of an immediate
- Place Rm on the A2 port of the register file for DP instructions with register as the second operand

- Format

B TARGET

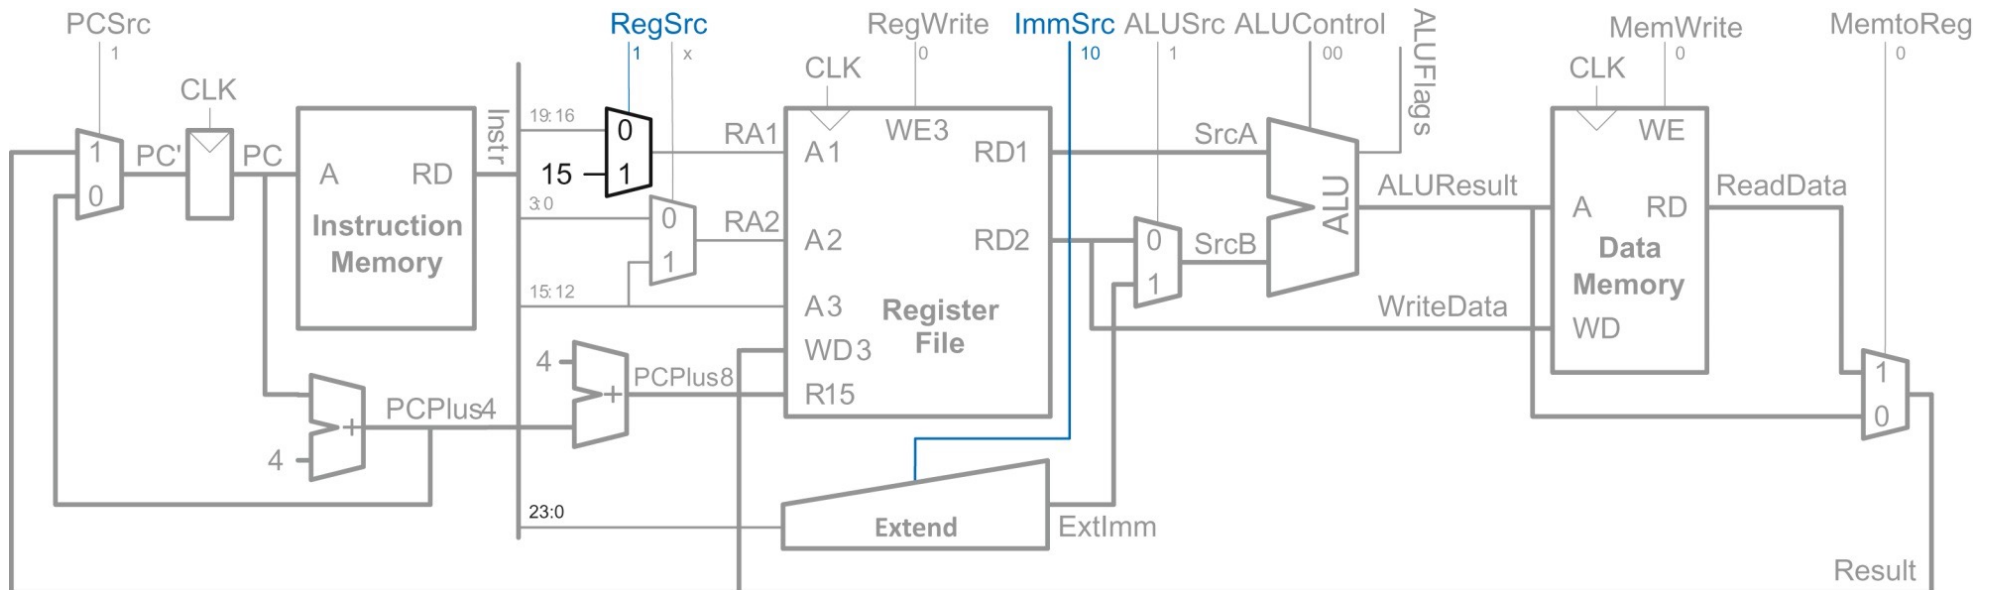


B imm24



Branch Datapath

Step 11: Change extend block, and add a bit to **RegSrc** for branch



31:28

27:26

25:24

23:0

1110	10	10	imm24
------	----	----	-------

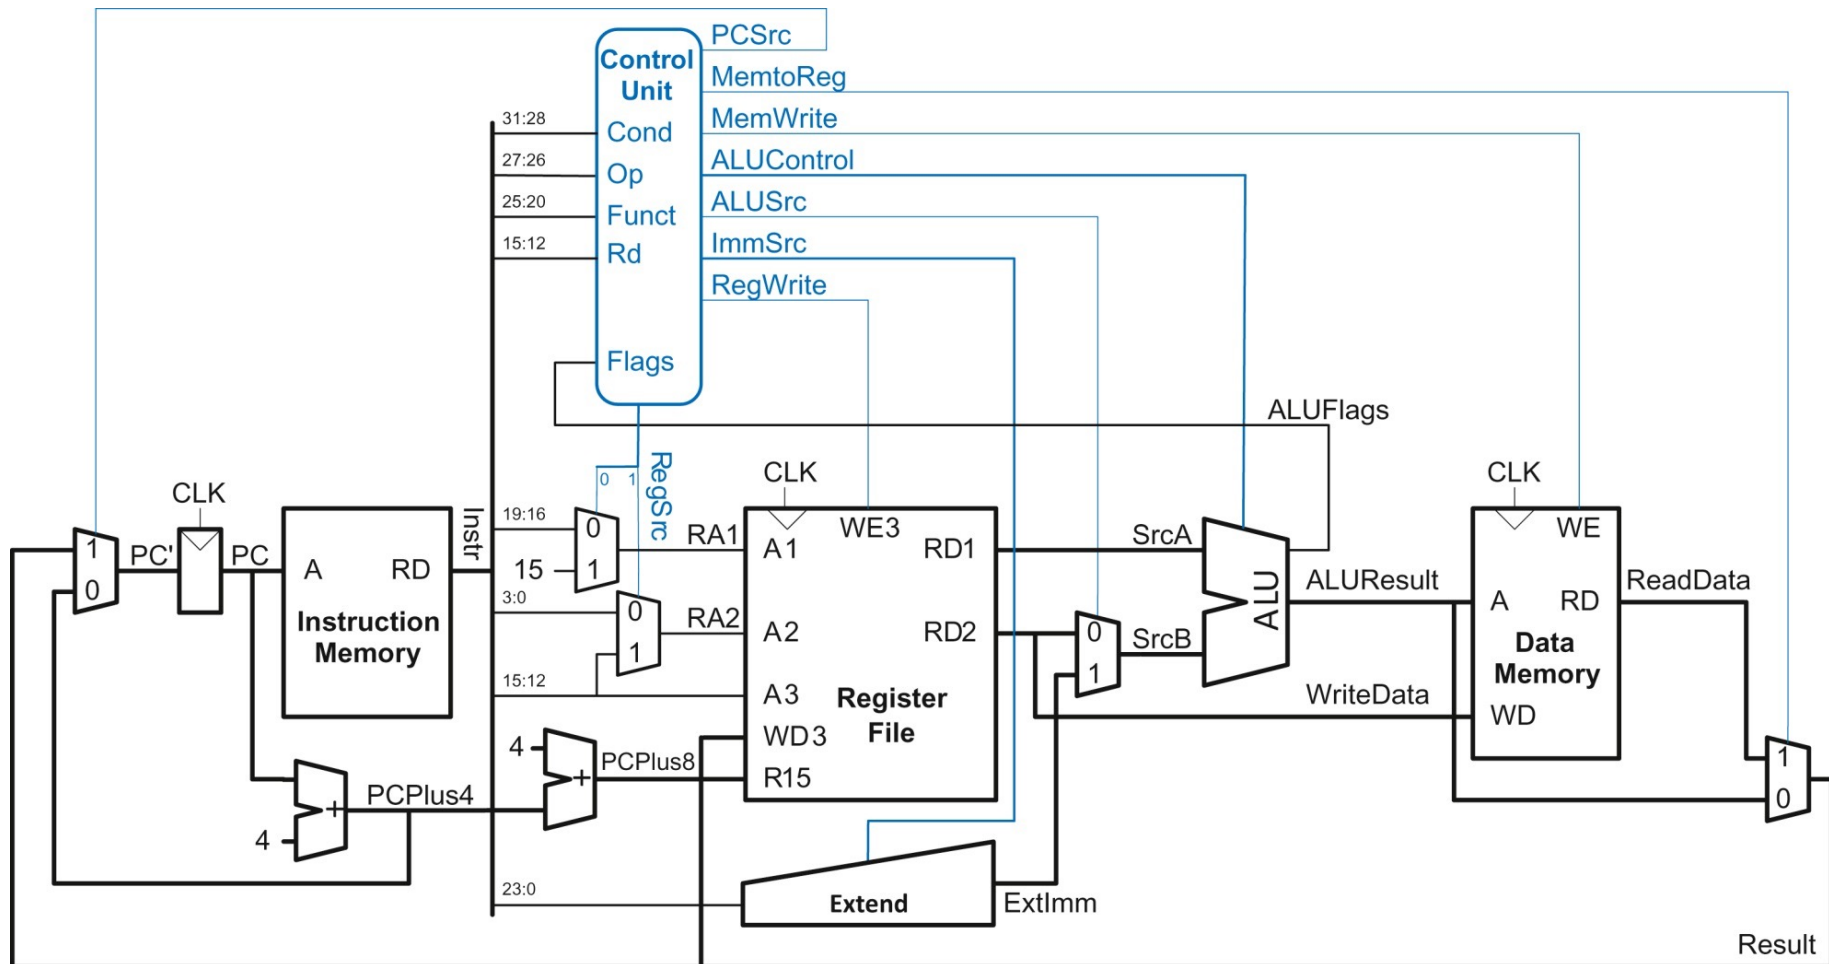
Operation of the Extend Block

Each of the three instruction formats interpret the immediate field differently

- $\text{ImmSrc}_{1:0}$ is the 2-bit control signal input to the extend block

$\text{ImmSrc}_{1:0}$	ExtImm	Description
00	$\{24'b0, \text{Instr}_{7:0}\}$	Zero-extended <i>imm8</i>
01	$\{20'b0, \text{Instr}_{11:0}\}$	Zero-extended <i>imm12</i>
10	$\{6\{\text{Instr}_{23}\}, \text{Instr}_{23:0}\}$	Sign-extended <i>imm24</i>

Datapath w/t Control



Control Unit

- Generate control signals based on instruction fields
 - Instr_{31:20} (cond)
 - Instr_{27:26} (op)
 - Instr_{25:20} (funct)
 - Flags
 - Destination register (*why does the controller needs this?*)
- Controller for “*our microarchitecture*” is purely combinational
- Things to ponder
 - Ensure correct PC update (branch or Rd == R15)
 - Update state (Reg file, Data memory) based on conditional execution

Decoder Truth Table

- Only selected signals are shown in the truth table

Op	Funct ₅	Funct ₀	Type	Branch	MemtoReg	MemW	ALUSrc	ImmSrc	RegW	RegSrc	ALUOp
00	0	X	DP Reg	0	0	0	0	XX	1	00	1
00	1	X	DP Imm	0	0	0	1	00	1	X0	1
01	X	0	STR	0	X	1	1	01	0	10	0
01	X	1	LDR	0	1	0	1	01	1	X0	0
11	X	X	B	1	0	0	1	10	0	X1	0