

# ENGN2219/COMP6719

## Computer Systems & Organization

Convener: Shoaib Akram

[shoaib.akram@anu.edu.au](mailto:shoaib.akram@anu.edu.au)



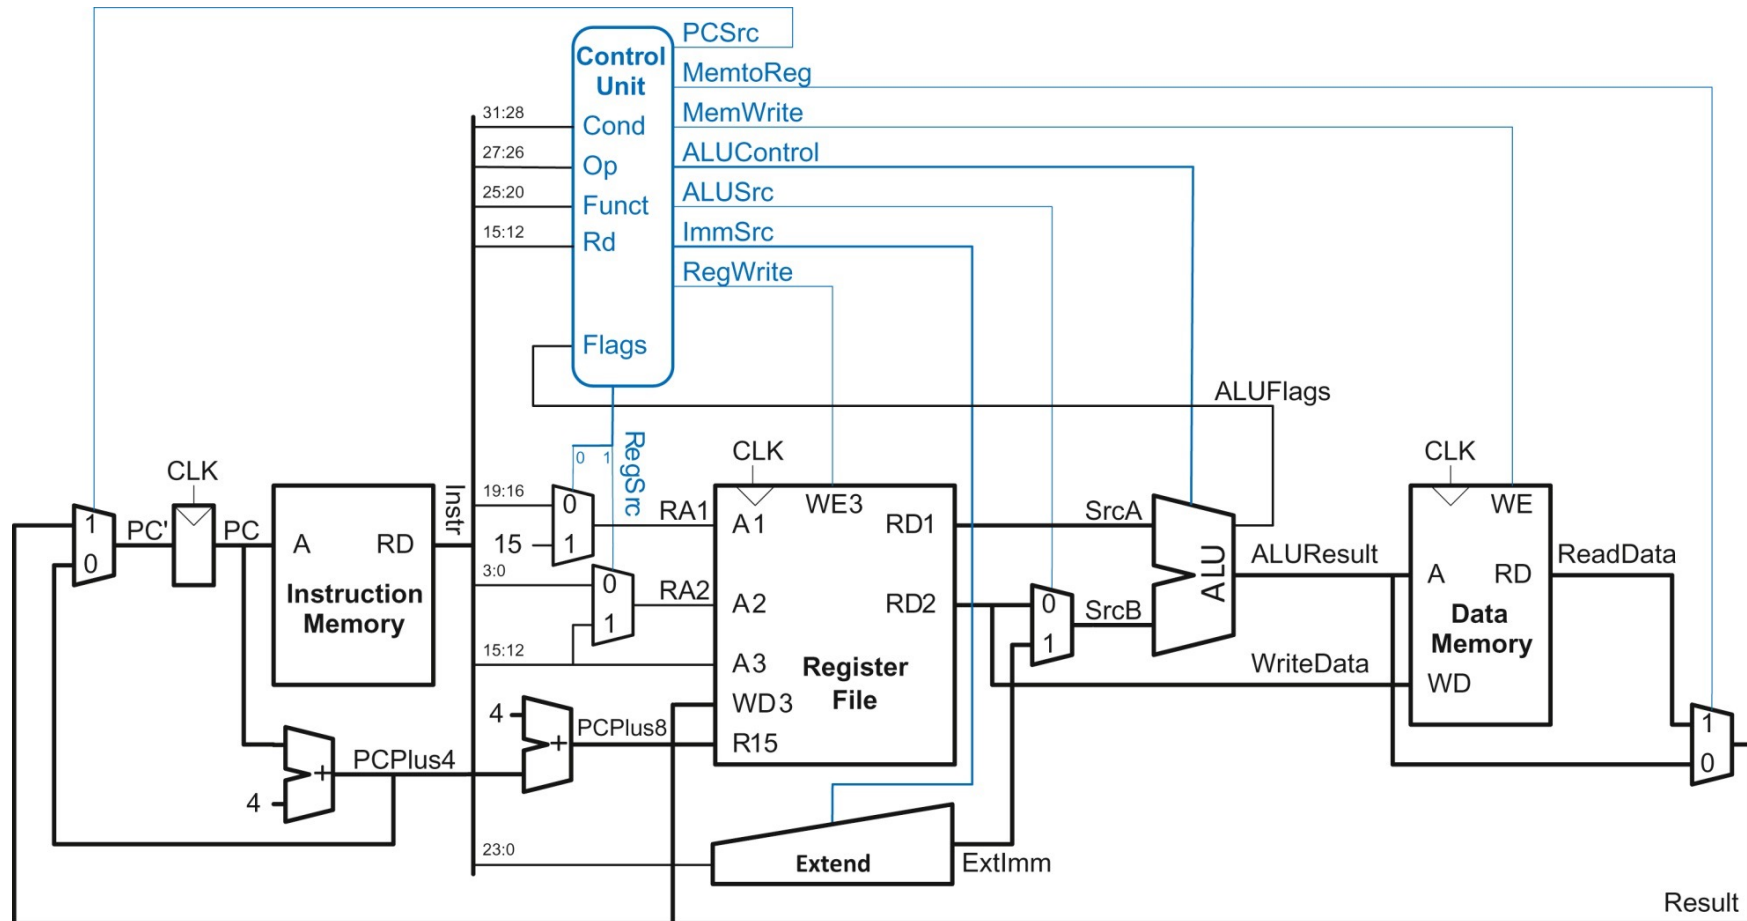
Australian  
National  
University

# Plan: Week 6

*Last week: Instruction Set Architecture (specification)*

*This Week: Microarchitecture (implementation)*

# Processor Operation: ORR

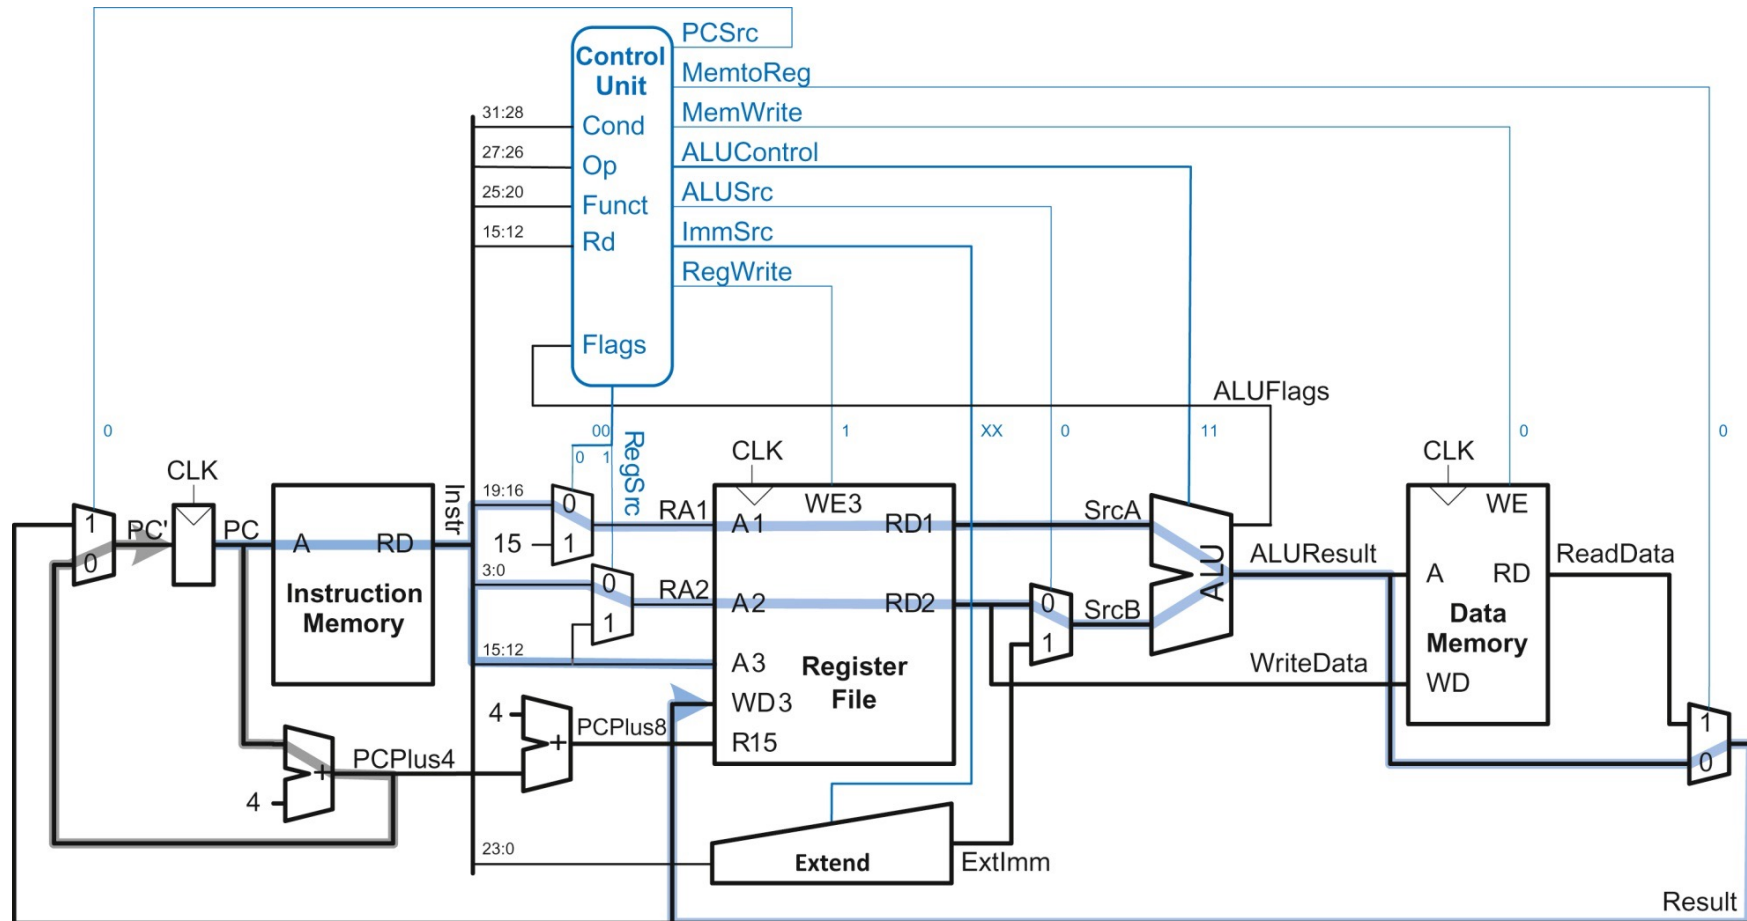


# Processor Operation: ORR


PCSrc	0
MemtoReg	0
MemWrite	0
ALUControl	11
ALUSrc	0
ImmSrc <sub>0:1</sub>	XX
RegWrite	1
RegSrc <sub>0:1</sub>	00

ALUControl	Function
00	ADD
01	SUB
10	AND
11	ORR

# Processor Operation: ORR



# Processor Operation: LDR



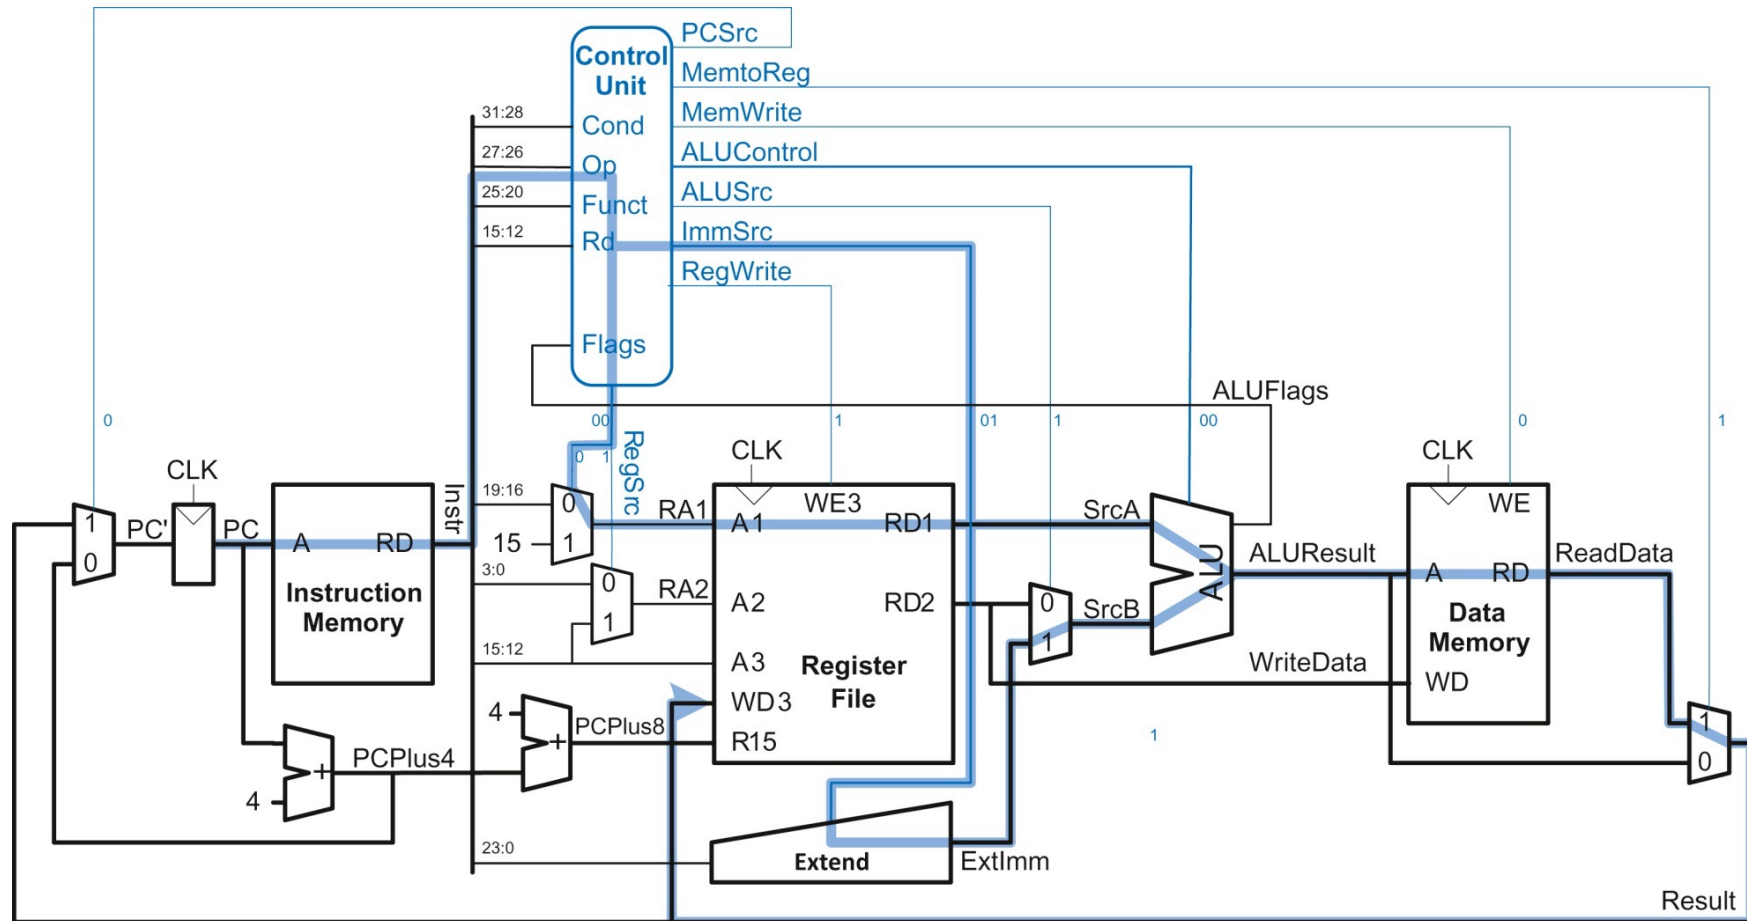
# Processor Operation: LDR

PCSrc	0
MemtoReg	1
MemWrite	0
ALUControl	00
ALUSrc	1
ImmSrc <sub>0:1</sub>	01
RegWrite	1
RegSrc <sub>0:1</sub>	00

ALUControl	Function
00	ADD
01	SUB
10	AND
11	ORR

ImmSrc <sub>1:0</sub>	ExtImm	Description
00	{24'b0, Instr <sub>7:0</sub> }	Zero-extended <i>imm8</i>
01	{20'b0, Instr <sub>11:0</sub> }	Zero-extended <i>imm12</i>
10	{6{Instr <sub>23</sub> }, Instr <sub>23:0</sub> }	Sign-extended <i>imm24</i>

# Processor Operation: LDR





# Whiteboard

# Exercise: Generating PCSrc Signal

- PCSrc is 1 when
  - Destination register (Rd) is R15
  - RegW is 1 (ADD/SUB or LDR)
  - Instruction is a branch
- $PCSrc = ((Rd == 15) \& RegW) \mid Branch$ 
  - Assuming the control unit generates a signal called Branch when opcode is 10 (B or BL)
- **Important:** Be careful to take conditional execution into account in the lab task and assignment!

# Critical Path Analysis

- Each instruction in our CPU takes one clock cycle
- To determine the clock cycle time requires us to find the critical path
- Different instructions use different resources
  - LDR uses instruction and data memory
  - ADD does not use data memory
  - STR does not write anything back to the register file
- Which instruction is the slowest?
  - Let's go back to the figures and find out!

# Elements of Critical Path

Parameter	Description
$t_{pcq\_PC}$	PC clock-to-Q delay
$t_{mem}$	Memory read
$t_{dec}$	Decoder propagation delay
$t_{mux}$	Multiplexer delay
$t_{RFread}$	Register file read
$t_{ext}$	Extension block delay
$t_{ALU}$	ALU delay
$t_{RFsetup}$	Set up RF for write (next cycle)

# Critical Path: LDR

$$T_c = t_{pcq\_PC} + t_{mem} + t_{dec} + \max[t_{mux} + t_{RFread}, t_{ext} + t_{mux}] + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- Memories & register files slower than combinational logic
  - Therefore,  $t_{mux} + t_{RFread} \gg t_{ext} + t_{mux}$

Final Equation

$$T_c = t_{pcq\_PC} + 2t_{mem} + t_{dec} + t_{RFread} + t_{ALU} + 2t_{mux} + t_{RFsetup}$$

# Critical Path: DP-R

$$T_c = t_{pcq\_PC} + t_{mem} + t_{dec} + t_{mux} + t_{RFread} + t_{mux} + t_{ALU} + t_{mux} + t_{RFsetup}$$

Final Equation

$$T_c = t_{pcq\_PC} + t_{mem} + t_{dec} + t_{RFread} + t_{ALU} + 3t_{mux} + t_{RFsetup}$$

# Critical Path Analysis

- Different instructions have different critical paths
  - LDR is the slowest instruction
  - DP-R and B have shorter critical paths because they do not need to access data memory
- Single-cycle processor is a synchronous sequential circuit
  - Clock period is constant and long enough to accommodate the slowest instruction
- The numerical values of different variables in the critical path equation depend on the specific technology

# Single-Cycle Microarchitecture

- Single-cycle microarchitecture
  - Execute the entire instruction in a single cycle
- Drawbacks
  - Separate memories for instructions and data
  - Cycle time is determined by the slowest instruction
  - Requires three adders (expensive)

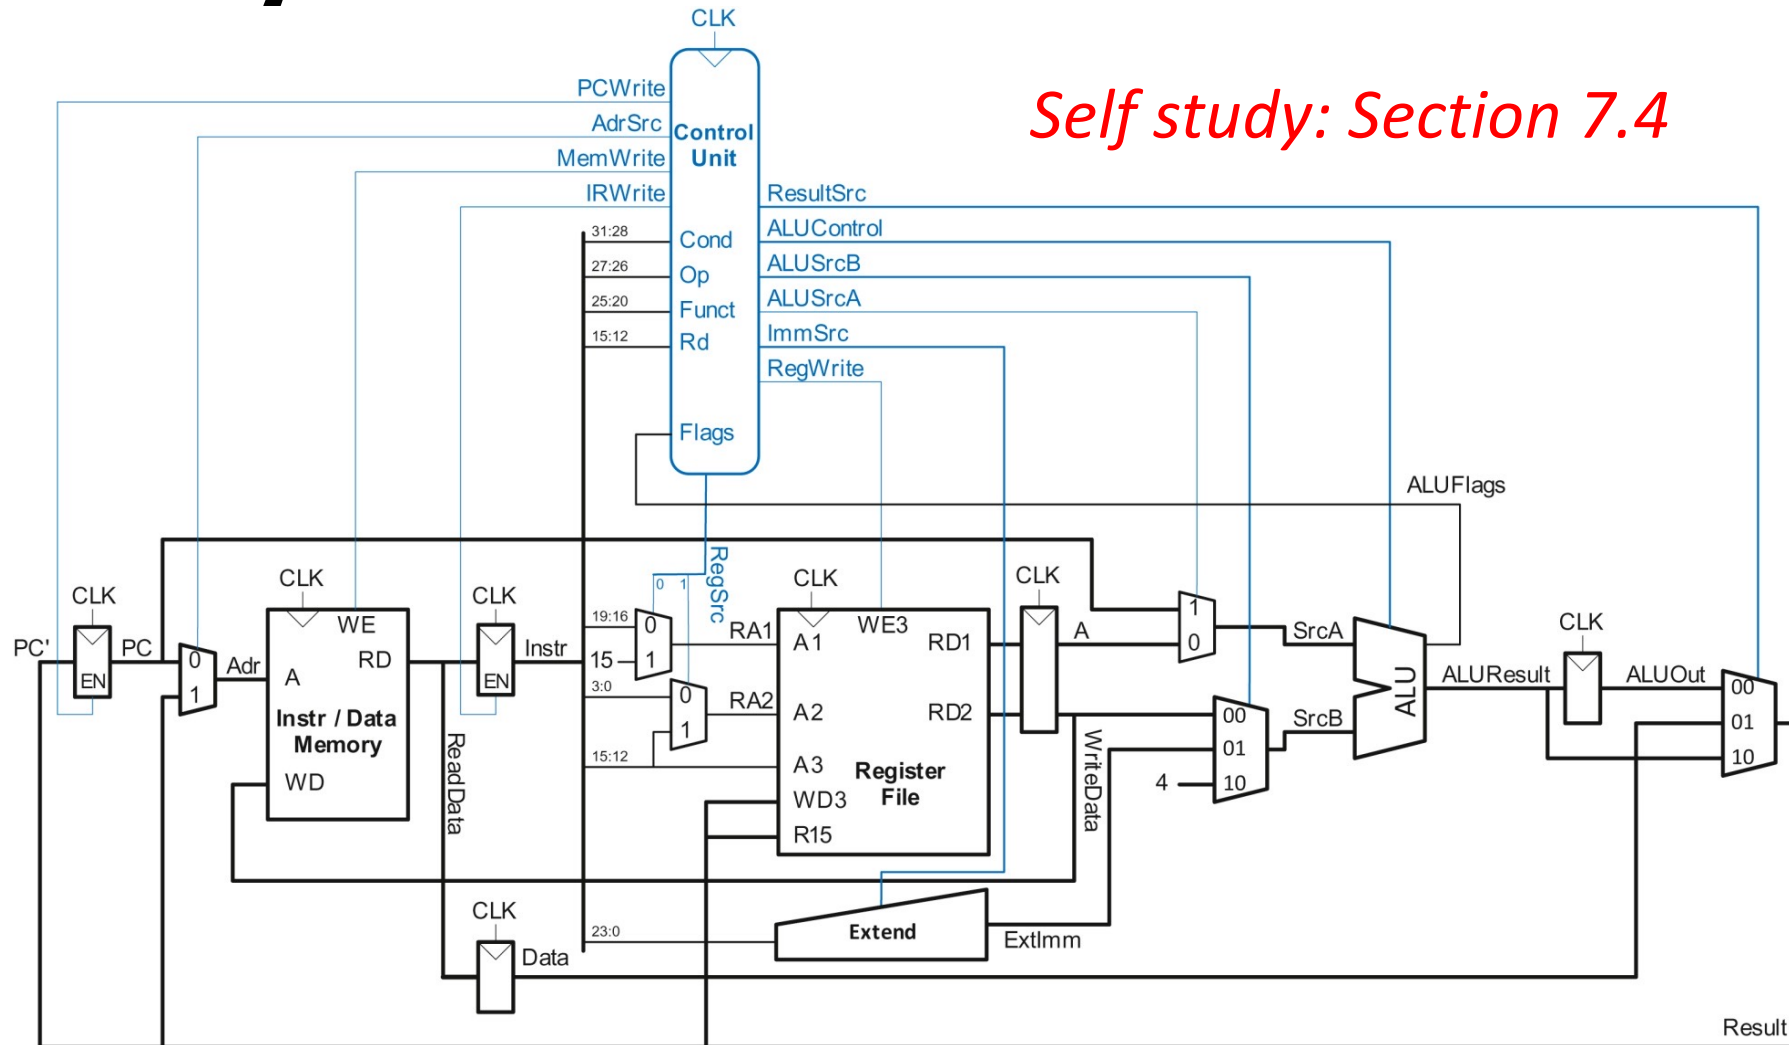


# Multicycle Microarchitecture

- Break the instruction into shorter steps
  - Access memory or use ALU in one step
  - Read instruction in one step and data in another step
  - Different instructions take different # steps or cycles
- Advantages
  - Some instructions finish faster than others
  - Needs one adder and one memory
- Drawbacks
  - Decoder complexity (decoder is an FSM)

# Multicycle uarch

*Self study: Section 7.4*



# Architectural State

- Modern systems run many applications at a time
- The operating system *interrupts* a program and runs another one, before restarting the interrupted program
- **Architectural state:** State needed to safely restart the program if it is interrupted by the operating system
- What state do we need to save and restore in ARM architecture?
  - 16 registers
  - Status register

# Microarchitectural State

- Any other state in the CPU is called *microarchitectural* state or *non-architectural* state
  - This state is not visible to assembly programs
  - Programmers are unaware of microarchitectural state
- **Note:** Memory is part of the architectural state but typically we do not consider memory part of CPU
  - Memory and storage is part of the computer system

# Pipelined Microarchitecture

- Multicycle microarchitecture
  - At any time, only one instruction is in execution
- Pipelined microarchitecture
- Pipelining divides the single-cycle CPU into stages
  - Each instruction goes through many stages
  - In each clock cycle, multiple instructions are active
  - One instruction can be in Fetch stage, another one can be in decode stage, ....
- Pipelining results in a shorter clock cycle time

# Pipeline Stages

- What steps/stages do an LDR instruction goes through?
  - Fetch
  - Decode (Register Read)
  - Execute (ALU)
  - Memory Access
  - Writeback

# Pipeline Stages

- What steps/stages do an ADD instruction goes through?
  - Fetch
  - Decode (Register Read)
  - Execute (ALU)
  - ~~Memory Access~~
  - Writeback
- More details: We will return to pipelining details after a few lectures!

# Performance Analysis

- We want the fastest (*best performing*) computer for a task
  - How should we measure and report performance?
- Which metric is *fair* for comparing two computers?
  - # Instructions
  - Clock frequency
  - Cycles per Instruction (**CPI**)
  - # Cores

*Important to understand the **true, gimmick-free** measure of computer performance*





# # Instructions

- RISC computer
  - Many more simple instructions
  - Simple hardware means smaller clock cycle
- CISC computer
  - Small number of instructions
  - Complex hardware means larger clock cycle

**Bottomline:** *Number of instructions alone is not a good metric for quantifying the performance of application A*

# Clock Frequency

- Consider the two scenarios
  - Computer A has a faster clock than computer B, but A executes many more instructions than B
  - Computer A has a faster clock than computer B, but takes multiple cycles to finish/execute a single instruction
- Is A faster than B?

**Bottomline:** *Clock frequency alone is not a good metric for quantifying the performance of application A*

# CPI

- Cycles per instruction
  - Ratio of # cycles to # instructions
- A program has 10 instructions. Each instruction takes one cycle
  - Instructions = 10, Cycles = 10, CPI = 1
- A program has 10 instructions. Two out of 10 instructions take two cycles
  - Instructions = 10, Cycles = 14, CPI = 1.4
- The inverse of CPI is called IPC
  - Instructions per cycle
  - For the above examples, IPC is 1 and 0.7

# CPI

- How can each instruction take multiple cycles?
  - Multi-cycle CPU
  - Memory accesses take more than one cycle
- On most computers, LDR takes variable # cycles because
  - Data may be present in faster (SRAM) memory called CPU cache
  - Or it may be present in main memory which is much slower

# CPI

- Two computers A and B have the same CPI for a specific program. Is there performance equivalent?
  - We need to know the cycle time
  - We need to know the # instructions

**Bottomline:** *CPI alone is not a good metric for quantifying the performance of application A*

# Execution Time

- The time it takes for a program to execute from start to finish is the only true measure of performance

$$\text{Execution time} = (\# \text{instructions}) \left( \frac{\text{cycles}}{\text{instruction}} \right) \left( \frac{\text{seconds}}{\text{cycle}} \right)$$

- seconds per cycle = cycle time
- Execution time is measured in seconds
- Golden metric for quantifying computer performance!

# Execution Time

$$\text{Execution time} = (\text{\#instructions}) \left( \frac{\text{cycles}}{\text{instruction}} \right) \left( \frac{\text{seconds}}{\text{cycle}} \right)$$

## # instructions

- Depends on the ISA, skill of programmer, compiler, algorithm

## cycles per instruction

- Depends on the microarchitecture esp. memory system

## seconds per cycle

- critical path, circuit technology, type of adders, gate-level details

# Exercise: Perf Analysis

- Find the time it takes to execute a program with 100 billion instructions on a single-cycle CPU in 16 nm CMOS manufacturing process. See the table for delays of logic elements.

Parameter	Delay (ps)
$t_{pcq\_PC}$	40
$t_{mem}$	200
$t_{dec}$	70
$t_{mux}$	25
$t_{RFread}$	100
$t_{ALU}$	120
$t_{RFsetup}$	60

$$T_c = t_{pcq\_PC} + 2t_{mem} + t_{dec} + t_{RFread} + t_{ALU} + 2t_{mux} + t_{RFsetup}$$



# Measurement Methodology

- The **good practice**: Take a program of interest and measure its execution time
- The **better practice**: Take a collection of programs like the programs of interest, and measure their performance
  - You do not have the program yet
  - Some one else is measuring the performance independently
- This collection of programs is called a *benchmark suite*
  - Dhrystone and CoreMark (embedded systems)
  - SPEC (Standard Performance Evaluation Corporation)
  - SPEC is standard suite for high-performance processors

# For Loop in C

C code:

```
int i;  
int sum = 0;  
  
for (i = 0; i < 10; i = i + 1) {  
    sum = sum + i;  
}
```

- The variable “i” is called the loop index
- $i = 0$  : index initialization
- $i < 10$  : loop termination condition
- $i = i + 1$  : loop advancement

# For Loop: C to Assembly

## C code:

```
int i;  
int sum = 0;  
  
for (i = 0; i < 10; i = i + 1) {  
    sum = sum + i;  
}
```

*check termination condition  
to break out of the loop if  
condition is met*

*keep iterating by  
branching back*

## ARM Assembly code

*; R0 = i, R1 = sum*

	MOV	R1,	#0
	MOV	R0,	#0
FOR			
	CMP	R0,	#10
	BGE	DONE	
	ADD	R1,	R1, R0
	ADD	R0,	R0, #1
	B	FOR	
DONE			

# For Loop: Perf Analysis

ARM Assembly code

*; R0 = i, R1 = sum*

```
FOR
    MOV    R1,    #0
    MOV    R0,    #0
    CMP    R0,    #10
    BGE    DONE
    ADD    R1,    R1,    R0
    ADD    R0,    R0,    #1
    B      FOR
DONE
```

- The clock cycle time  $T_c$  is 840 ps
- Find the time it takes to execute the for loop
  - # instructions = ?
  - CPI = 1
  - Execution time = ?

# Alternative For Loop

	MOV	R1,	#0	
	MOV	R0,	#0	
COND				
	CMP	R0,	#10	
	BLT	LOOP		
	B	DONE		
LOOP				
	ADD	R1,	R1,	R0
	ADD	R0,	R0,	#1
	B	COND		
DONE				

- We can implement the for loop in a different way
- Find the time it takes to execute the for loop again
  - # instructions = ?
  - Execution time = ?

**Bottom line:** Execution time depends on how we write code and microarchitecture details

# While Loop in C

C code:

```
int pow = 1;
int x = 0;

while (pow != 128) {
    pow = pow * 2;
    x = x + 1;
}
```

- For loop: iterate N times
- While loop: Iterate until a condition is not met

# While Loop in C

## C code:

```
int pow = 1;
int x = 0;

while (pow != 128) {
    pow = pow * 2;
    x = x + 1;
}
```

## ARM Assembly code

*; R0 = pow, R1 = x*

	MOV	R0,	#1
	MOV	R1,	#0
WHILE			
	CMP	R0,	#128
	BEQ	DONE	
	LSL	R0,	R0, #1
	ADD	R1,	R1, #1
	B	WHILE	
DONE			

# Exercise

- The clock cycle time  $T_c$  is 840 ps
- Find the time it takes to execute the While loop
  - # instructions = ?
  - CPI = 1
  - Execution time = ?



# Shift Instructions

- Shift the value in a register left or right, drop bits off the end
  - Logical shift left (**LSL**)
  - Logical shift right (**LSR**)
  - Arithmetic shift right (**ASR**)
  - Rotate right (**ROR**)
- Logical shift: shifts the number to the left or right and fills the empty slots with zero
- Arithmetic shift: on right shifts fill the most significant bits with zero
- Rotate: rotates number in a circle such that empty spots are filled with bits shifted off the other end

# Example: Shift Operations

- **Immediate** shift amount (5-bit immediate)
- Shift amount: 0-31

Source register


R5	1111 1111	0001 1100	0001 0000	1110 0111
----	-----------	-----------	-----------	-----------

Assembly Code


Result

LSL R0, R5, #7	R0	1000 1110	0000 1000	0111 0011	1000 0000
LSR R1, R5, #17	R1	0000 0000	0000 0000	0111 1111	1000 1110
ASR R2, R5, #3	R2	1111 1111	1110 0011	1000 0010	0001 1100
ROR R3, R5, #21	R3	1110 0000	1000 0111	0011 1111	1111 1000

# Shifts: Machine Representation




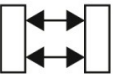
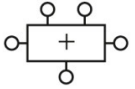

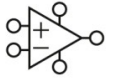

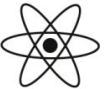


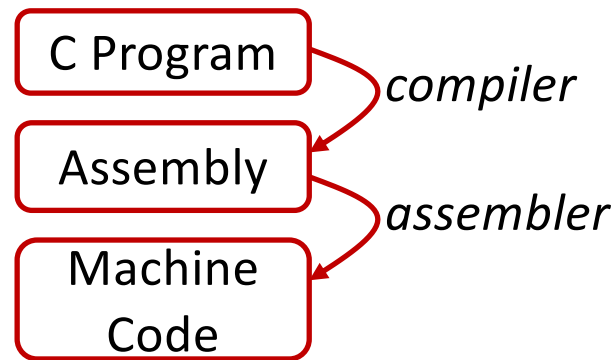
## Shift Instructions



- cmd = 1101
- sh = 00 (LSL), 01 (LSR), 10 (ASR), 11 (ROR)
- Rn = 0
- shamt5 = 5-bit shift amount

# Big Picture

Application Software		Programs
Operating Systems		Device Drivers
Architecture		Instructions Registers
Micro-architecture		Datapaths Controllers
Logic		Adders Memories
Digital Circuits		AND Gates NOT Gates
Analog Circuits		Amplifiers Filters
Devices		Transistors Diodes
Physics		Electrons



*Software*

*Hardware*

*ISA is the boundary  
(Contract)*



01010010  
10101010  
10101001  
10000011  
*Memory*

*Instructions stored as 0's and 1's*



*CPU*

*Fetch, decode, execute  
instructions*

Happy Teaching Break!

Remember the Assignment (30%)

Keep the transformation hierarchy  
fresh in your memory!

See you in two weeks 😊