

# ENGN2219/COMP6719

## Computer Systems & Organization

Convener: Shoaib Akram

[shoaib.akram@anu.edu.au](mailto:shoaib.akram@anu.edu.au)



Australian  
National  
University

# Plan: Lectures

- C and assembly
  - *Hardware/software interaction*
- Memory and storage devices
  - *How do the devices work?*
  - *How are they exposed to C programs?*
- Hardware optimizations
  - *Caches and virtual memory (memory-side)*
  - *Pipelining (CPU-side)*

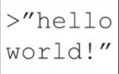


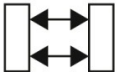
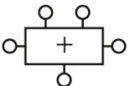

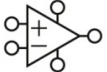

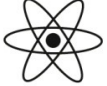
# Plan: Labs

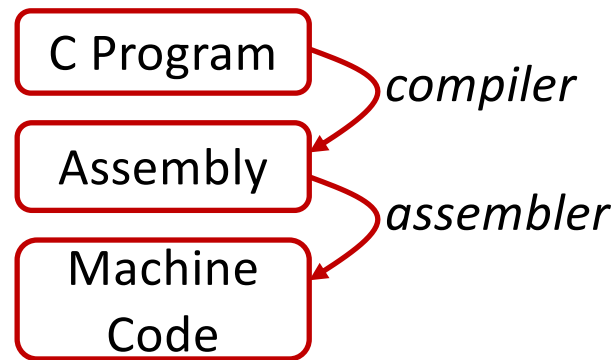
- Introduction to C – Part 1 ✓
- Introduction to C – Part 2
  - Control flow, bitwise operations, more pointers, strings
- Data Structures (beyond arrays)
  - Structs, unions, linked lists, read/write file I/O
- Dynamic Memory Allocation (rich topic)
- Assignment 2
  - Problem specification (QuAC CPU model, memory allocator)
  - Your task: Solve the problem in C (knowing assembly helps)
  - If you do Labs 1 – 3 diligently, you will (mostly) nail it!

# Hardware/Software Interaction

- *Predominantly Assembly, some C (2 lectures)*
- *Exclusively C (2 – 3 lectures)*

# Big Picture

Application Software		Programs
Operating Systems		Device Drivers
Architecture		Instructions Registers
Micro-architecture		Datapaths Controllers
Logic		Adders Memories
Digital Circuits		AND Gates NOT Gates
Analog Circuits		Amplifiers Filters
Devices		Transistors Diodes
Physics		Electrons



*Software*

*Hardware*

*ISA is the boundary  
(Contract)*



01010010  
10101010  
10101001  
10000011  
*Memory*

*Instructions stored as 0's and 1's*



*CPU*

*Fetch, decode, execute  
instructions*


# Shift Instructions

- Shift the value in a register left or right, drop bits off the end
  - Logical shift left (**LSL**)
  - Logical shift right (**LSR**)
  - Arithmetic shift right (**ASR**)
  - Rotate right (**ROR**)
- Logical shift: shifts the number to the left or right and fills the empty slots with zero
- Arithmetic shift: on right shifts fill the most significant bits with the sign bit
- Rotate: rotates number in a circle such that empty spots are filled with bits shifted off the other end

# Logical Shift Left (LSL)


ARM Instruction

LSL R0, R5, #3




- Shift all bits left 3 positions, insert three 0's from the left
- Drop the 3 bits from the right

*Result*




# Logical Shift Right (LSR)

ARM Instruction LSR    R0,    R5,    #3



- Shift all bits right 3 positions, insert three 0's from the right
- Drop the 3 bits from the left


*Result*






# Arithmetic Shift Right (LSR)

ARM Instruction ASR    R0,    R5,    #3



- Shift all bits right 3 positions, insert three 0's from the right
- Drop the 3 bits from the left


*Result*




# Rotate Right (ROR)

ARM Instruction


ROR R0, R5, #21




- Do a circular shift
- Right shift by 21 and put back bits that fall off at left end



# Shifts: Machine Representation



## Shift Instructions



- cmd = 1101
- sh = 00 (LSL), 01 (LSR), 10 (ASR), 11 (ROR)
- Rn = 0
- shamt5 = 5-bit shift amount

# Shifts: Machine Representation

- Format (Src2 = Register)

LSL R0, R5, #3

LSL Rd, Rm, shamt5

31:28	27:26	25	24:21	20	19:16	15:12	11:7	6:5	4	3:0
cond	00	0	cmd	S	Rn	Rd	shamt5	sh	0	Rm

# Shift Instructions

- **Immediate** shift amount (5-bit immediate)
- Shift amount: 0-31
- Rn is not used
- sh encodes the type of shift
- ARM also has instructions with shift amount held in a register

LSL	R4,	R8,	R6
-----	-----	-----	----

ROR	R5,	R8,	R6
-----	-----	-----	----

31:28	27:26	25	24:21	20	19:16	15:12	11:8	7	6:5	4	3:0
cond	00	0	cmd	S	Rn	Rd	Rs	0	sh	1	Rm

# Shift Instructions

- Having dedicated instructions for shift operations is useful for systems programming
- Bit masks are a common requirement in low-level hardware resource management
- Code that manages network protocols or file formats
- Anything related to compression/decompression or packing/unpacking of information
- Rotation is used in cyclic codes (cryptography, compression)

# Control Flow

- In real programs, the *order* in which statements execute is not always sequential (one after the other)
- Decisions and iterating the same task repeatedly is common
  - `if/else` ✓
  - `switch`
  - `for` and `while`
- How can we write these statements in assembly?
  - Performance analysis
  - Evaluate alternatives

# For Loop in C: Sum

C code:

```
int i;  
int sum = 0;  
  
for (i = 0; i < 10; i = i + 1) {  
    sum = sum + i;  
}
```

C code:

```
int i;  
int sum = 0;  
  
for (i = 9; i >= 0; i = i - 1) {  
    sum = sum + i;  
}
```

*Decremental version*

- The variable “i” is called the loop index/counter
- $i = 0$  : index initialization
- $i < 10$  : loop termination condition
- $i = i + 1$  : loop advancement



# Sum: Assembly

## C code:

```
int i;  
int sum = 0;  
  
for (i = 0; i < 10; i = i + 1) {  
    sum = sum + i;  
}
```

*check termination condition  
to break out of the loop if  
condition is met*

*keep iterating by  
branching back*

## ARM Assembly code

*; R0 = i, R1 = sum*

	MOV	R1,	#0
	MOV	R0,	#0
FOR			
	CMP	R0,	#10
	BGE	DONE	
	ADD	R1,	R1, R0
	ADD	R0,	R0, #1
	B	FOR	
DONE			

# Sum: Perf Analysis

ARM Assembly code

*; R0 = i, R1 = sum*

	MOV	R1,	#0
	MOV	R0,	#0
FOR			
	CMP	R0,	#10
	BGE	DONE	
	ADD	R1,	R1, R0
	ADD	R0,	R0, #1
	B	FOR	
DONE			

How long does it take to execute the loop (*frequency = 1 GHz, CPI = 1*)

- # instructions = ?
- Clock cycle time,  $T_c$  = ?
- Execution time = ?

# Sum: Alternative Approach

	MOV	R1,	#0	
	MOV	R0,	#0	
COND				
	CMP	R0,	#10	
	BLT	FOR		
	B	DONE		
FOR				
	ADD	R1,	R1,	R0
	ADD	R0,	R0,	#1
	B	COND		
DONE				

- More faithfully follow the for loop semantics in C
- Use BLT instead of BGE

How long does it take to execute the loop (*frequency = 1 GHz, CPI = 1*)

- Instruction count = ?
- Clock cycle time,  $T_c = ?$
- Execution time = ?

# Sum: Decremental Version

## C code:

```
int i;  
int sum = 0;  
  
for (i = 9; i != 0; i = i - 1) {  
    sum = sum + i;  
}
```

## ARM Assembly code

*; R0 = i, R1 = sum*

```
MOV    R1,    #0  
MOV    R0,    #9  
FOR  
ADD     R1,    R1,    R0  
SUBS    R0,    R0,    #1  
BNE     FOR
```

DONE

*sum = sum + 1*

*i = i - 1 and set flags*

## Saves 2 instructions per iteration:

- Decrement loop variable & compare: SUBS R0, R0, #1
- Only 1 branch – instead of 2

# Exercise

- Find the time it takes to execute the loop now if the clock cycle time is 1 ns

# Lessons

Execution time depends on how we write code and microarchitecture details

Always make the common case fast!

**Note:** *Eliminating a branch is always desirable in pipelined processors because the CPU needs to wait for the branch to finish execution in order to fetch the next instruction*

# While Loop in C

- For loop iterate N times
  - Used when N is known in advance
- While loop
  - Iterate until the *controlling condition* is false
- Determine **x** such that  $2^x = 128$

C code:

```
int pow = 1;
int x = 0;

while (pow != 128) {
    pow = pow * 2;
    x = x + 1;
}
```

# Two Interesting While Loops

C code:

```
while (1) {  
    // iterates forever  
}
```

```
while (0) {  
    // iterates 0 times  
}
```



# While Loop: C and Assembly

## C code:

```
int pow = 1;
int x = 0;

while (pow != 128) {
    pow = pow * 2;
    x = x + 1;
}
```

Determine  $x$  such that  
 $2^x = 128$

## ARM Assembly code

*; R0 = pow, R1 = x*

```
MOV    R0,    #1
MOV    R1,    #0
WHILE
CMP     R0,    #128
BEQ     DONE
LSL     R0,    R0,    #1
ADD     R1,    R1,    #1
B       WHILE
DONE
```

# Exercise

- Find the time it takes to execute the While loop if the clock cycle time is 1 ns.
- Write `sum` as a while loop and find the time it takes to execute the resulting loop if all CPU parameters are the same as before.

# switch/case Statement

C code:

```
switch(button) {  
    case 1: atm = 20; break;  
    case 2: atm = 50; break;  
    case 3: atm = 100; break;  
    default; atm = 0;  
}
```



C code for if...else ladder:

```
if (button == 1)  
    atm = 20;  
else if (button == 2)  
    atm = 50;  
else if (button == 3)  
    atm = 100;  
else  
    atm = 0;
```

- Execute one of several blocks of code depending on the condition and *break* out of the entire switch block
- If no conditions are met, the *default* block is executed

# switch/case Statement

C code:

```
switch(button) {  
    case 1: atm = 20; break;  
    case 2: atm = 50; break;  
    case 3: atm = 100; break;  
    default; atm = 0;  
}
```

ARM Assembly code

*; R0 = button, R1 = atm*

CMP	R0,	#1
MOVEQ	R1,	#20
BEQ	DONE	
CMP	R0,	#2
MOVEQ	R1,	#50
BEQ	DONE	
CMP	R0,	#3
MOVEQ	R1,	#50
BEQ	DONE	
MOV	R1,	#0

DONE

# Arrays

- Arrays contain a collection of similarly typed elements
- Elements are stored contiguously in memory

**int** is 4 bytes on most architectures

C code:

```
int marks[5] = {19, 10, 8, 17, 9};
```

```
int a = marks[0];
```

```
int b = 2;
```

```
marks[3] = b;
```

Address	Data	Index	Element
⋮	⋮	⋮	⋮
00000010	9	4	marks[4]
0000000C	17	3	marks[3]
00000008	8	2	marks[2]
00000004	10	1	marks[1]
00000000	19	0	marks[0]

← 4 Bytes →

# Array of Characters

- Array of **characters** (**char** is a data type in C)
- **char** is used for representing characters

**char** is always 1 byte

C code:

```
char alphas[5] = {'a', 'b', 'c', 'd', 'e'};
```

Address	Data	Index	Element
⋮	⋮	⋮	⋮
00000004	'e'	4	alphas[4]
00000003	'd'	3	alphas[3]
00000002	'c'	2	alphas[2]
00000001	'b'	1	alphas[1]
00000000	'a'	0	alphas[0]

← 1 Byte →

# Array of Characters

- Array of **characters** (char is a data type in C)
- char is used for representing characters

**char** is always 1 byte

C code:

```
char alphas[5] = {'a', 'b', 'c', 'd', 'e'};
```

Address

•  
•  
•

00000010

0000000C


00000008

00000004

00000000

Data

•  
•  
•



4 Bytes

# Endianness

For large objects (greater than 1 byte), byte order in memory matters: *Which byte of a 4-byte int is stored at the lowest address?*

- **Little Endian:** Little end (**LSB**) stored first (at lowest address)
  - Intel x86
- **Big Endian:** Big end (**MSB**) stored first,
  - SPARC, Motorola processors
- ARM is bi-endian (supports both)

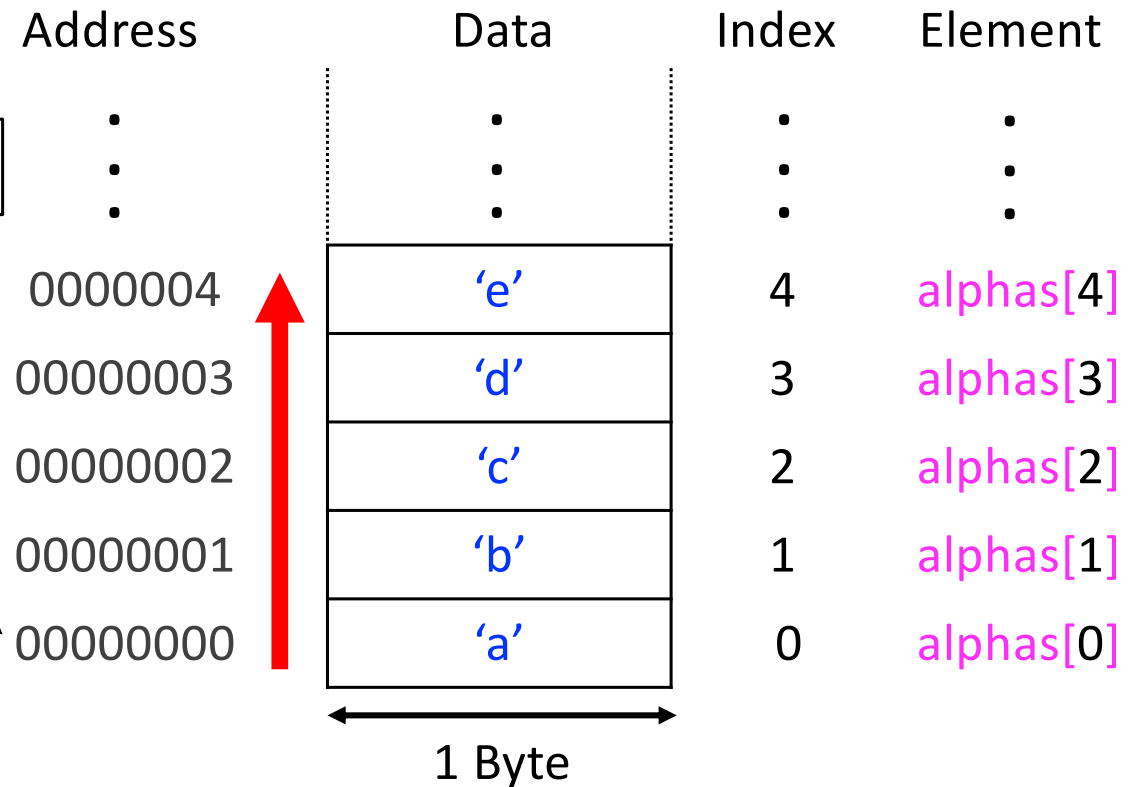


# Characters and Endianness

- Characters are 1-byte each
- There is no ambiguity in which byte to store first
  - Always Big-Endian


```
char alphas[5] = {'a', 'b', 'c', 'd', 'e'};
```

`alphas[0]` is stored in memory  
at address 0x00000000



# Integers and Endianness

- A 32-bit integer is stored at memory address 100
- Stored in locations: 100, 101, 102, 103
- Which part of the 32-bit value is stored first?



00010010001101000101011001111000

Hex 12 34 56 78

# Integers and Endianness

- A 32-bit integer is stored at memory address 100
- Stored in locations: 100, 101, 102, 103
- Which part of the 32-bit value is stored first?

104	
103	12
102	34
101	56
100	78

00010010001101000101011001111000

Hex 12 34 56 78

Little Endian: 78 at location 100

# Integers and Endianness

- A 32-bit integer is stored at memory address 100
- Stored in locations: 100, 101, 102, 103
- Which part of the 32-bit value is stored first?

Address	Data
104	
103	78
102	56
101	34
100	12

00010010001101000101011001111000

Hex 12 34 56 78

Big Endian: 12 at location 100

# Endianness: Pros

- Little Endian
  - Easy to create small values from large values
  - Previous example: Read byte at address 100
  - On Big Endian, add 4 to 100, then read byte to find out
- Big Endian
  - Easy to test sign and range of a value

QuAC architecture in labs evades the entire issue of endianness with word-addressable memory

# Array Sum

- Example to illustrate how instructions are picked for ISAs
- Add a constant 10 to each element of the scores array

## C code:

```
int i;  
int scores[200];  
// initialization code  
...  
for (i = 0; i < 200; i++)  
    scores[i] = scores[i] + 10;
```

## Assembly code:

```
; R0 = array base address, R1 = i  
MOV    R0, #0x14000000  
MOV    R1, 0  
LOOP  
    CMP    R1, #200  
    BGE    L3  
    LSL    R2, R1, #2  
    LDR    R3, [R0, R2]  
    ADD    R3, R3, #10  
    STR    R3, [R0, R2]  
    ADD    R1, R1, #1  
    B      LOOP  
L3
```

- R0 = base address
- i = 0
- i < 200?
- if not, exit loop
- R2 = i\*4
- R3 = scores[i]
- R3 = scores[i] + 10
- scores[i] += 10
- i = i + 1
- repeat loop

# LDR with register as offset

- New LDR variant (LDR with register as offset)

LDR R3, [R0, R2]

dest base offset

The diagram illustrates the instruction `LDR R3, [R0, R2]`. Below the instruction, three labels are positioned: `dest` under `R3`, `base` under `R0`, and `offset` under `R2`. Arrows point from each label to its corresponding register in the instruction: a vertical arrow from `dest` to `R3`, a vertical arrow from `base` to `R0`, and a diagonal arrow from `offset` to `R2`.

- Very common to load from memory with base + offset addressing, so there is an instruction for that
- R2 is called the index register

# Condensing Array Sum – 1

- LSL/LDR combo often used in tandem in array traversals
  - There is support for that in the ISA
- Eliminates LSL instruction

*ARM has an instruction that scales index reg. R1*

LDR R3, [R0, R1, LSL #2]

Left shift is a multiply by 2

Memory address =  $R0 + (R1 * 4)$

Assembly code:

; R0 = array base address, R1 = i

MOV R0, #0x14000000

MOV R1, 0

LOOP

CMP R1, #200

BGE L3

LSL R2, R1, #2

LDR R3, [R0, R2]

ADD R3, R3, #10

STR R3, [R0, R2]

ADD R1, R1, #1

B LOOP

L3

■ R0 = base address

■ i = 0

■ i < 200?

■ if not, exit loop

■ R2 = i\*4

■ R3 = scores[i]

■ R3 = scores[i] + 10

■ scores[i] += 10

■ i = i + 1

■ repeat loop