

# ENGN2219/COMP6719

## Computer Systems & Organization

Convener: Shoaib Akram

[shoaib.akram@anu.edu.au](mailto:shoaib.akram@anu.edu.au)



Australian  
National  
University

# Recap: Array Sum

Add 10 to each element of the 200-element scores array

C code:

```
int i;  
int scores[200];  
// initialization code not  
//shown  
...  
for (i = 0; i<200; i++)  
    scores[i] = scores[i] + 10;
```

# Array Sum

Add 10 to each element of the 200-element scores array

C code:

```
int i;  
int scores[200];  
// initialization code not  
//shown  
...  
for (i = 0; i<200; i++)  
    scores[i] = scores[i] + 10;
```

<i>address</i>	<i>data</i>	
0x14000010	90	scores[4]
0x1400000C	76	...
0x14000008	80	scores[2]
0x14000004	40	scores[1]
base → 0x14000000	100	scores[0]

← 4 bytes →

Showing the scores array in memory

# Array Sum

Add 10 to each element of the 200-element scores array

C code:

```
int i;
int scores[200];
// initialization code not
//shown
...
for (i = 0; i<200; i++)
    scores[i] = scores[i] + 10;
```

address      data

0x14000010	90	scores[4]
0x1400000C	76	...
0x14000008	80	scores[2]
0x14000004	40	scores[1]
base → 0x14000000	100	scores[0]

4 bytes

Showing the scores array in memory

Assembly code:

```
; R0 = array base address
; R1 = i
```

```
MOV R0, #0x14000000
```

```
MOV R1, 0
```

LOOP

```
CMP R1, #200
```

```
BGE L3
```

```
LSL R2, R1, #2
```

```
LDR R3, [R0, R2]
```

```
ADD R3, R3, #10
```

```
STR R3, [R0, R2]
```

```
ADD R1, R1, #1
```

```
B LOOP
```

L3

- R0 = base addr
- i = 0
- i < 200?
- no? exit loop
- i = i + 1
- R3 = scores[i]
- R3 = R3 + 10
- scores[i] += 10
- i = i + 1
- repeat

# Array Sum

Add 10 to each element of the 200-element scores array

C code:

```
int i;
int scores[200];
// initialization code not
//shown
...
for (i = 0; i<200; i++)
    scores[i] = scores[i] + 10;
```

address      data

0x14000010	90	scores[4]
0x1400000C	76	...
0x14000008	80	scores[2]
0x14000004	40	scores[1]
base → 0x14000000	100	scores[0]

4 bytes

Showing the scores array in memory

Assembly code:

```
; R0 = array base address
; R1 = i
```

```
MOV R0, #0x14000000
```

```
MOV R1, 0
```

LOOP

```
CMP R1, #200
```

```
BGE L3
```

```
LSL R2, R1, #2
```

```
LDR R3, [R0, R2]
```

```
ADD R3, R3, #10
```

```
STR R3, [R0, R2]
```

```
ADD R1, R1, #1
```

```
B LOOP
```

L3

▪ R0 = base addr

▪ i = 0

▪ i < 200?

▪ no? exit loop

▪ i = i + 1

▪ R3 = scores[i]

▪ R3 = R3 + 10

▪ scores[i] += 10

▪ i = i + 1

▪ repeat

# Another LDR Variant

- We have seen two ways of specifying the offset so far

LDR R3, [R0, #16]

dest base offset

LDR R3, [R0, R2]

dest base offset in  
index/offset register

- LSL + LDR combo often used in tandem in array traversals
- ISA supports eliminating the extra LSL instruction

LDR R3, [R0, R1, LSL #2]

Left shift is the same  
as multiplying by 2

- Memory address
  - Left shift R1 by 2 (scale R1)
  - Add to R0
  - Address =  $R0 + (R1 * 4)$

# Condensing Array Sum – 1

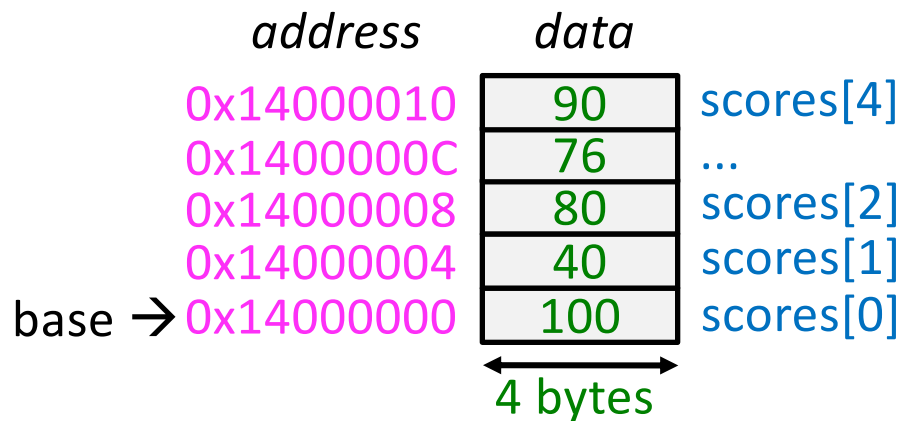
Add 10 to each element of the 200-element scores array

C code:

```
int i;
int scores[200];
// initialization code not
//shown
...
for (i = 0; i<200; i++)
    scores[i] = scores[i] + 10;
```

Assembly code:

```
; R0 = array base address
; R1 = i
MOV    R0,    #0x14000000
MOV    R1,    0
LOOP
CMP    R1,    #200
BGE    L3
LDR    R3,    [R0, R1, LSL, #2]
ADD    R3,    R3,    #10
STR    R3,    [R0, R2]
ADD    R1,    R1,    #1
B      LOOP
L3
```



Showing the scores array in memory

# ARM Indexing Modes

- **Offset Addressing** `LDR R0, [R1, R2]`
  - *Address is the sum of base register + offset (#20, #-20, -R2)*
  - *Base register is unchanged*
- **Pre-indexed Addressing** `LDR R0, [R1, R2]!`
  - *Address is the sum of base register + offset*
  - *Base register is updated with the address*
- **Post-index Addressing** `LDR R0, [R1], R2`
  - *Address is the base register*
  - *Base register is updated with the new address only after the memory access*



# ARM Indexing Modes

- Offset Addressing

LDR R0, [R1, R2]

- $Address = R1 + R2$

- $R1 = \text{Unchanged}$

- Pre-indexed Addressing

LDR R0, [R1, R2]!

- $Address = R1 + R2$

- $R1 = R1 + R2$

- Post-index Addressing

LDR R0, [R1], R2

- $Address = R1$

- $R1 = R1 + R2$

- *In all cases, offset can be an immediate*

# Condensing Array Sum – 2

Add 10 to each element of the 200-element scores array

C code:

```
int i;
int scores[200];
// initialization code not
//shown
...
for (i = 0; i<200; i++)
    scores[i] = scores[i] + 10;
```

Assembly code:

```
; R0 = array base address
; R1 = i
```

```
MOV R0, #0x14000000
```

```
MOV R1, R0, #800
```

LOOP

```
CMP R0, R1
```

```
BGE L3
```

```
LDR R2, [R0]
```

```
ADD R2, R2, #10
```

```
STR R2, [R0], #4
```

```
B LOOP
```

L3

- R0 = base addr
- R1 = base + 800
- end of array?
- yes? exit loop
- R2 = scores[i]
- scores[i] + 10
- store scores[i]
- and R0 = R0 + 4
- repeat loop

address data

0x14000010	90	scores[4]
0x1400000C	76	...
0x14000008	80	scores[2]
0x14000004	40	scores[1]
base → 0x14000000	100	scores[0]

4 bytes

Showing the scores array in memory

# Condensing Array Sum – 2

Add 10 to each element of the 200-element scores array

Assembly code:

```
; R0 = array base address
; R1 = i
MOV    R0,    #0x14000000
MOV    R1,    R0,    #800
LOOP
  CMP    R0,    R1
  BGE    L3
  LDR    R2,    [R0]
  ADD    R2,    R2,    #10
  STR    R2,    [R0], #4

  B      LOOP
L3
```

- This version of Array Sum first computes the address of the last byte of the array (**#0x14000800**)
- Each iteration of **LOOP** checks if **R0** is greater than or equal to **#0x14000800**
- If so, we are done, so step out of **LOOP**
- **STR R2, [R0], #4**
  - Stores **R2** at **[R0]**, and after that, adds 4 to **R0**

# Explaining

1. `CMP R0, R1`
2. `BGE LOOP`

- When CPU encounters: `CMP R0, R1`
  - It subtracts `R1` from `R0`
  - It sets the `flags` in the CSPR register
  - No register is updated with the result (no side-effects)
- When CPU encounters: `BGE LOOP`
  - CPU checks the `flags` to establish if `R0` is greater than or equal to `R1`
- Some conditions are easy to establish, and others harder
  - `EQ` is easily established by looking at the `zero` flag
  - If the `zero` flag is `1`, it means  $(R0 - R1)$  is `0`, meaning `R0` and `R1` are equal

# Conditional Execution

- Week 5, part 2 lecture
  - Two ways of setting the flags in the CPSR register
  - Conditional execution in general

# Bytes and Characters

- Characters on the English keyboard can be encoded in a single byte
- **char** in C is an 8-bit integer under the hood
  - C operators close to the hardware (basic types)
  - No string, list, or other composite types
- ASCII standard is for mapping characters to integer codes
  - Other standards such as Unicode are ASCII supersets
- Need instructions to manipulate bytes!

## Decimal - Binary - Octal - Hex – ASCII Conversion Chart

Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII
0	00000000	000	00	NUL	32	00100000	040	20	SP	64	01000000	100	40	@	96	01100000	140	60	`
1	00000001	001	01	SOH	33	00100001	041	21	!	65	01000001	101	41	A	97	01100001	141	61	a
2	00000010	002	02	STX	34	00100010	042	22	"	66	01000010	102	42	B	98	01100010	142	62	b
3	00000011	003	03	ETX	35	00100011	043	23	#	67	01000011	103	43	C	99	01100011	143	63	c
4	00000100	004	04	EOT	36	00100100	044	24	\$	68	01000100	104	44	D	100	01100100	144	64	d
5	00000101	005	05	ENQ	37	00100101	045	25	%	69	01000101	105	45	E	101	01100101	145	65	e
6	00000110	006	06	ACK	38	00100110	046	26	&	70	01000110	106	46	F	102	01100110	146	66	f
7	00000111	007	07	BEL	39	00100111	047	27	'	71	01000111	107	47	G	103	01100111	147	67	g
8	00001000	010	08	BS	40	00101000	050	28	(	72	01001000	110	48	H	104	01101000	150	68	h
9	00001001	011	09	HT	41	00101001	051	29	)	73	01001001	111	49	I	105	01101001	151	69	i
10	00001010	012	0A	LF	42	00101010	052	2A	*	74	01001010	112	4A	J	106	01101010	152	6A	j
11	00001011	013	0B	VT	43	00101011	053	2B	+	75	01001011	113	4B	K	107	01101011	153	6B	k
12	00001100	014	0C	FF	44	00101100	054	2C	,	76	01001100	114	4C	L	108	01101100	154	6C	l
13	00001101	015	0D	CR	45	00101101	055	2D	-	77	01001101	115	4D	M	109	01101101	155	6D	m
14	00001110	016	0E	SO	46	00101110	056	2E	.	78	01001110	116	4E	N	110	01101110	156	6E	n
15	00001111	017	0F	SI	47	00101111	057	2F	/	79	01001111	117	4F	O	111	01101111	157	6F	o
16	00010000	020	10	DLE	48	00110000	060	30	0	80	01010000	120	50	P	112	01110000	160	70	p
17	00010001	021	11	DC1	49	00110001	061	31	1	81	01010001	121	51	Q	113	01110001	161	71	q
18	00010010	022	12	DC2	50	00110010	062	32	2	82	01010010	122	52	R	114	01110010	162	72	r
19	00010011	023	13	DC3	51	00110011	063	33	3	83	01010011	123	53	S	115	01110011	163	73	s
20	00010100	024	14	DC4	52	00110100	064	34	4	84	01010100	124	54	T	116	01110100	164	74	t
21	00010101	025	15	NAK	53	00110101	065	35	5	85	01010101	125	55	U	117	01110101	165	75	u
22	00010110	026	16	SYN	54	00110110	066	36	6	86	01010110	126	56	V	118	01110110	166	76	v
23	00010111	027	17	ETB	55	00110111	067	37	7	87	01010111	127	57	W	119	01110111	167	77	w
24	00011000	030	18	CAN	56	00111000	070	38	8	88	01011000	130	58	X	120	01111000	170	78	x
25	00011001	031	19	EM	57	00111001	071	39	9	89	01011001	131	59	Y	121	01111001	171	79	y
26	00011010	032	1A	SUB	58	00111010	072	3A	:	90	01011010	132	5A	Z	122	01111010	172	7A	z
27	00011011	033	1B	ESC	59	00111011	073	3B	;	91	01011011	133	5B	[	123	01111011	173	7B	{
28	00011100	034	1C	FS	60	00111100	074	3C	<	92	01011100	134	5C	\	124	01111100	174	7C	
29	00011101	035	1D	GS	61	00111101	075	3D	=	93	01011101	135	5D	]	125	01111101	175	7D	}
30	00011110	036	1E	RS	62	00111110	076	3E	>	94	01011110	136	5E	^	126	01111110	176	7E	~
31	00011111	037	1F	US	63	00111111	077	3F	?	95	01011111	137	5F	_	127	01111111	177	7F	DEL

# Loading/Storing Bytes

- **LDRB**

- *Load byte in register, and zero extend to fill the 32 bits*

- **LDRSB**

- *Load byte in register, and sign-extend to fill the 32 bits*

- **STRB**

- *Store the LSB of the 32-bit integer into the specified byte in memory*
  - *More significant bits of the register are ignored*



# Loading/Storing Bytes

- What is in `R1`, `R2`, and `memory` after each of the instruction has executed?
- Assume `R4 = 0`

Byte Address	Data
4	...
3	F7
2	8C
1	42
0	03

## Registers

XX	XX	XX	XX	LDRB	R1, [R4, #2]
XX	XX	XX	XX	LDRSB	R2, [R4, #2]
11	10	A1	9B	STRB	R3, [R4, #3]

# Loading/Storing Bytes

- What is in `R1`, `R2`, and `memory` after each of the instruction has executed?
- Assume `R4 = 0`

Byte Address	Data
4	...
3	9B
2	8C
1	42
0	03

## Registers

00	00	00	8C	LDRB	R1, [R4, #2]
FF	FF	FF	8C	LDRSB	R2, [R4, #2]
XX	XX	XX	9B	STRB	R3, [R4, #3]

# Strings in C

- A series of characters is a string
    - `char welcome[6] = {'H', 'E', 'L', 'L', 'O', '\0'};`
    - `char welcome[] = "HELLO";`
  - Compiler figures out the length
    - 5 + 1 for `'\0'`
    - Manually track length (unlike Python)
  - Compiler inserts a null terminator `'\0'` automatically
  - Need a way to know the end of the string
  - C strings are null-terminated
- 
- ```
graph TD; A["char welcome[6] = {'H', 'E', 'L', 'L', 'O', '\0'};"] --> B["Compiler figures out the length"]; A --> C["Compiler inserts a null terminator '\0' automatically"]; A --> D["Need a way to know the end of the string"]; B --> E["5 + 1 for '\0'"]; B --> F["Manually track length (unlike Python)"]; D --> G["C strings are null-terminated"];
```

# Ex: Manipulating Char Array

C code:

```
char array[10] = "ENGN2219!";  
int i;  
  
for (i = 0; i < 10; i = i + 1)  
    array[i] = array[i] - 32;
```

# Ex: Manipulating Char Array

- Transform the 10-character ASCII string, namely array, from lower case to upper case

## C code:

```
char array[10] = "finalexam";
int i;

for (i = 0; i < 10; i = i + 1)
    array[i] = array[i] - 32;
```

## Assembly code:

```
; R0 = base addr, R1 = i
MOV    R0,    #0
LOOP
CMP     R1,    #10
BGE     DONE
LDRB    R2,    [R0, R1]
SUB     R2,    R2,    #32
STRB    R2,    [R0, R1]
ADD     R1,    R1,    #1
B       LOOP
DONE
```


- $i = 0$
- $i < 10?$
- if  $i \geq 10$ , exit
- $R2 = \text{array}[i]$
- subtract 32
- store  $\text{array}[i]$
- $i = i + 1$
- repeat loop

# Exercise: Strings in Memory

- Show how “HELLO!” is stored in memory below at address 0x1522FFF0.

## *ASCII Encoding*

|      |      |
|------|------|
| H    | 0x48 |
| E    | 0x65 |
| L    | 0x6C |
| O    | 0x6F |
| !    | 0x21 |
| Null | 0x00 |




# Exercise: Strings in Memory

- Show how “HELLO!” is stored in memory below at address 0x1522FFF0.

## *ASCII Encoding*

|      |      |
|------|------|
| H    | 0x48 |
| E    | 0x65 |
| L    | 0x6C |
| O    | 0x6F |
| !    | 0x21 |
| Null | 0x00 |



# Practice

## C Code

```
int array[5];  
array[0] = array[0] * 8;  
array[1] = array[1] * 8;
```

## ARM Assembly Code

; R0 = array base address

```
MOV R0, #0x60000000 ; R0 = 0x60000000
```

```
LDR R1, [R0] ; R1 = array[0]
```

```
LSL R1, R1, #3 ; R1 = R1 << 3 = R1*8
```

```
STR R1, [R0] ; array[0] = R1
```

```
LDR R1, [R0, #4] ; R1 = array[1]
```

```
LSL R1, R1, #3 ; R1 = R1 << 3 = R1*8
```

```
STR R1, [R0, #4] ; array[1] = R1
```



# Exercise

## C Code

```
int array[200];
int i;
for (i=199; i >= 0; i = i - 1)
    array[i] = array[i] * 8;
```


## ARM Assembly Code

```
; R0 = array base address, R1 = i
MOV R0, 0x60000000
MOV R1, #199

FOR
    LDR R2, [R0, R1, LSL #2] ; R2 = array(i)
    LSL R2, R2, #3           ; R2 = R2<<3 = R2*8
    STR R2, [R0, R1, LSL #2] ; array(i) = R2
    SUBS R0, R0, #1          ; i = i - 1
                                ; and set flags
    BPL FOR                  ; if (i>=0) repeat
loop
```

# Functions

- High-level languages offer functions to enable
  - Abstraction & Modularity
  - Code reuse
  - Readability
  - Testability & validation
  - Maintainability
- Functions are also called *procedures* or *subroutines*
- Functions are ubiquitous, encouraging **ISA** support
  - Special jump instructions
  - Special scratch space to store temporary variables
  - Ways to reduce interference b/w functions



# Functions: Our Goal

- Architectural support for functions
  - Branch and Link instruction (BL)
  - Stack Pointer (SP)
  - Link Register (LR)
- Microarchitecture-level impacts of programming styles (Iteration vs. Recursion)
- Provides a deeper understanding of hardware/software interaction and tradeoffs

# Functions in C

## C Code

```
void main()  
{  
    int y;  
    y = sum(42, 7);  
    ...  
}
```

*42 and 7 are function arguments provided by caller, i.e., main()*

```
int sum(int a, int b)  
{  
    return (a + b);  
}
```

- **main**( ) is caller (calling someone else)
  - Returns nothing (**void**)
  - No input arguments
- **sum**( ) is *callee* (being called by someone)
  - Two input arguments of type int: **a** and **b**
  - Return type: integer
  - Returns the sum of **a** and **b**

# Leaf and Non-Leaf Functions

- `sum( )` is a leaf function
  - It does not call another function
- `main( )` is a non-leaf function
  - It calls another function
- Non-leaf functions are more complicated especially at the assembly level
- `sum( )` can be called from many different functions
  - Code reuse

# Functions as Detectives

- Secret mission
- Acquire necessary resources
- Perform the mission
- Leave no trace
- Return safely



# Functions as Detectives

- Caller stores *arguments* in specific registers
- Caller transfers *flow control* to the callee (**call**)
- Callee acquires/allocates memory for doing work
- Callee executes the function body
- Callee stores the result in a specific register
- Callee *returns* control to the caller (**return**)

# ARM Function Calls

- Instruction for calling the function
  - **BL** (**B**ranch and **L**ink)
  - CPU branches to the label specified by BL
  - CPU stores the *return address* in the link register (LR)
- Return address is the address of the next instruction after the function call
- Returning from function
  - Move the link register into PC
  - MOV PC, LR
- Passing arguments (convention)
  - R0, R1, R2, R3
- Returning value (convention)
  - R0



# Function Calls

## C Code

```
int main() {  
    simple();  
    a = b + c;  
}  
  
void simple() {  
    return;  
}
```

## ARM Assembly Code

```
0x00000200 MAIN      BL    SIMPLE  
0x00000204          ADD    R4, R5, R6  
...  
  
0x00401020 SIMPLE    MOV    PC, LR
```

- **BL** branches to **SIMPLE**  
 $LR = PC + 4 = 0x00000204$
- **MOV PC, LR** makes  $PC = LR$   
(the next instruction executed is at **0x00000200**)

- **MAIN** and **SIMPLE** are labels (memory addresses) in assembly
- **BL** transfers flow to **SIMPLE** and stores the *return address* in **LR**
- The function returns after **MOV**, and the next instruction (**ADD**) is executed

# Example: Difference of Sums

C code:

```
int main() {  
    int y;  
    ...  
    y = diffofsums(2, 3, 4, 5);  
    ...  
}  
  
int diffofsums(int f, int g, int h, int i) {  
    int result;  
    result = (f + g) - (h + i);  
    return result;  
}
```

## ARM Assembly Code

```
; R4 = y
```

```
MAIN
```

```
...
```

```
MOV R0, #2      ; argument 0 = 2
MOV R1, #3      ; argument 1 = 3
MOV R2, #4      ; argument 2 = 4
MOV R3, #5      ; argument 3 = 5
BL  DIFFOFSUMS  ; call function
MOV R4, R0      ; y = returned value
```

```
...
```

```
; R4 = result
```

```
DIFFOFSUMS
```

```
ADD R8, R0, R1   ; R8 = f + g
ADD R9, R2, R3   ; R9 = h + i
SUB R4, R8, R9   ; result = (f + g) - (h + i)
MOV R0, R4       ; put return value in R0
MOV PC, LR       ; return to caller
```

# Questions

- How can we pass more than 4 function arguments?
- How can we ensure that registers in use by the caller are not corrupted?
  - `DIFFOFSUMS` overwrites `R4`, `R8`, `R9`
  - `MAIN` may need these registers after return
- The Stack
  - A special area in memory used across function calls
  - Preserving registers, passing arguments, scratch space

# The Stack

- Abstract view
  - Last In First Out (LIFO) Queue
- Stored in memory at some arbitrary address in memory
- Caller and callee can *push* things onto the stack and *pop* things off the stack
- Stack expands and contracts over time as function call and return

