

ENGN2219/COMP6719

Computer Systems & Organization

Convener: Shoaib Akram

shoaib.akram@anu.edu.au



Australian
National
University

Word Count (**wordcount.c**)

Problem statement: Count the number of words input by the user via keyboard

*A **word** contains a sequence of characters (a-z, A-Z) without a blank space or a new line*

✓ ENGN2219

✓ 2219ENGN

✓ ENGN

✗ 2219

Input/Output

Input: *Read from keyboard character-wise until end of file is encountered. Blank space and new-line starts a new word.*

Output: *# of words*

Flow of our solution

1. Read characters until the **end of file** is encountered
EOF is a symbolic constant in C (equivalent to pressing Ctrl-d)
2. For each character, check if it is in the range: *a to z or A to Z*
3. Keep count of the number of words seen so far

Reading characters

```
char c;  
while ( (c = getchar()) != EOF ) {
```

Checking for characters in range

In C, the *char* type or the character variable holds an ASCII value between 0 and 127

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	NUL	(null)	32	20	040	 Space	64	40	100	@ @	96	60	140	` `		
1	1	001	SOH	(start of heading)	33	21	041	! !	65	41	101	A A	97	61	141	a a		
2	2	002	STX	(start of text)	34	22	042	" "	66	42	102	B B	98	62	142	b b		
3	3	003	ETX	(end of text)	35	23	043	# #	67	43	103	C C	99	63	143	c c		
4	4	004	END	(end of transmission)	36	24	044	$ \$	68	44	104	D D	100	64	144	d d		
5	5	005	ENQ	(enquiry)	37	25	045	% %	69	45	105	E E	101	65	145	e e		
6	6	006	ACK	(acknowledge)	38	26	046	& &	70	46	106	F F	102	66	146	f f		
7	7	007	BEL	(bell)	39	27	047	' '	71	47	107	G G	103	67	147	g g		
8	8	010	BS	(backspace)	40	28	050	((72	48	110	H H	104	68	150	h h		
9	9	011	TAB	(horizontal tab)	41	29	051))	73	49	111	I I	105	69	151	i i		
10	A	012	LF	(NL line feed, new line)	42	2A	052	* *	74	4A	112	J J	106	6A	152	j j		
11	B	013	VT	(vertical tab)	43	2B	053	+ +	75	4B	113	K K	107	6B	153	k k		
12	C	014	FF	(NP form feed, new page)	44	2C	054	, ,	76	4C	114	L L	108	6C	154	l l		
13	D	015	CR	(carriage return)	45	2D	055	- -	77	4D	115	M M	109	6D	155	m m		
14	E	016	SO	(shift out)	46	2E	056	. .	78	4E	116	N N	110	6E	156	n n		
15	F	017	SI	(shift in)	47	2F	057	/ /	79	4F	117	O O	111	6F	157	o o		
16	10	020	DLE	(data link escape)	48	30	060	0 0	80	50	120	P P	112	70	160	p p		
17	11	021	DC1	(device control 1)	49	31	061	1 1	81	51	121	Q Q	113	71	161	q q		
18	12	022	DC2	(device control 2)	50	32	062	2 2	82	52	122	R R	114	72	162	r r		
19	13	023	DC3	(device control 3)	51	33	063	3 3	83	53	123	S S	115	73	163	s s		
20	14	024	DC4	(device control 4)	52	34	064	4 4	84	54	124	T T	116	74	164	t t		
21	15	025	NAK	(negative acknowledge)	53	35	065	5 5	85	55	125	U U	117	75	165	u u		
22	16	026	SYN	(synchronous idle)	54	36	066	6 6	86	56	126	V V	118	76	166	v v		
23	17	027	ETB	(end of trans. block)	55	37	067	7 7	87	57	127	W W	119	77	167	w w		
24	18	030	CAN	(cancel)	56	38	070	8 8	88	58	130	X X	120	78	170	x x		
25	19	031	EM	(end of medium)	57	39	071	9 9	89	59	131	Y Y	121	79	171	y y		
26	1A	032	SUB	(substitute)	58	3A	072	: :	90	5A	132	Z Z	122	7A	172	z z		
27	1B	033	ESC	(escape)	59	3B	073	; ;	91	5B	133	[[123	7B	173	{ {		
28	1C	034	FS	(file separator)	60	3C	074	< <	92	5C	134	\ \	124	7C	174	|		
29	1D	035	GS	(group separator)	61	3D	075	= =	93	5D	135]]	125	7D	175	} }		
30	1E	036	RS	(record separator)	62	3E	076	> >	94	5E	136	^ ^	126	7E	176	~ ~		
31	1F	037	US	(unit separator)	63	3F	077	? ?	95	5F	137	_ _	127	7F	177	 DEL		

Source: www.LookupTables.com

A note on characters in C

```
if (((c >= 'a') && (c <= 'z')) || ((c >= 'A') && (c <= 'Z'))) {
```

Counting words

First, we need a variable to store the *state* we are in

Inside a word (**IN**)

Outside a word (**OUT**)

If the initial state is **OUT**

State transitions to **IN** if the user types a valid character

State transitions to **OUT** if the user inputs blank space or \n

OUT → IN

We have a new word (words++)


Question?

What happens if we are in state **IN** and the user inputs a valid character?

Printing the #words

```
printf("%d\n", words);
```

State machine diagram



wordcount.c

```
1  #include <stdio.h>

2  #define IN  1  // inside a word
3  #define OUT 0  // outside a word

4  int main(void) {

5      char c;
6      int inorout = OUT;
7      int words   = 0;

8      while ((c = getchar()) != EOF) {
9          if (((c >= 'a') && (c <= 'z')) || ((c >= 'A') && (c <= 'Z'))) {
10             if (inorout == OUT) {
11                 inorout = IN;
12                 words++;
13             }
14         } else if ((c == ' ') || (c == '\n')) {
15             inorout = OUT;
16         } else {
17             // ignore
18         }
19     }

20     printf("%d\n", words);

21     return 0;
22 }
```

Address Space

- **Address range**

- A 32-bit (ARM) CPU generates addresses in the range 0 to 0xFFFFFFFF (4294967292)
- With a 4×10^9 address range, the CPU can access 4 billion individual bytes

- **Address space**

- The address space of a 32-bit CPU is 2^{32} bytes which equals 4 Gigabytes (GB)

0xFFFFFFFFFC

Address Space

- Each word is 32 bits or 4 bytes. Address of first & last word is shown
- The address space is empty as shown here
 - Let's populate with stack and code and data

0x00000000



Questions


- Where is the code, data, and the stack in the address space?
- **Memory map**
 - Defines where code, data, and stack memory are in the program address space
 - Differs from architecture to architecture
 - The subsequent discussion pertains to ARM


ARM 32-bit Memory Map

- Five parts or segments
 - text
 - global data
 - dynamic data
 - OS & I/O
 - Exception handlers

0xFFFFFFFFFC


0x00000000





- *Data in this segment is dynamically allocated and deallocated during program execution*
- *Heap data is allocated by the program at runtime*
 - `malloc()` and `new`
- *Heap grows upward, stack grows downward*
- *Global variables visible to all functions (contrasted with local variables that are only visible to a function)*
- *Machine language program*
- *Also called read-only (**RO**) segment*
- *Literals (constants) such as "Hello"*

Translating/Starting Programs

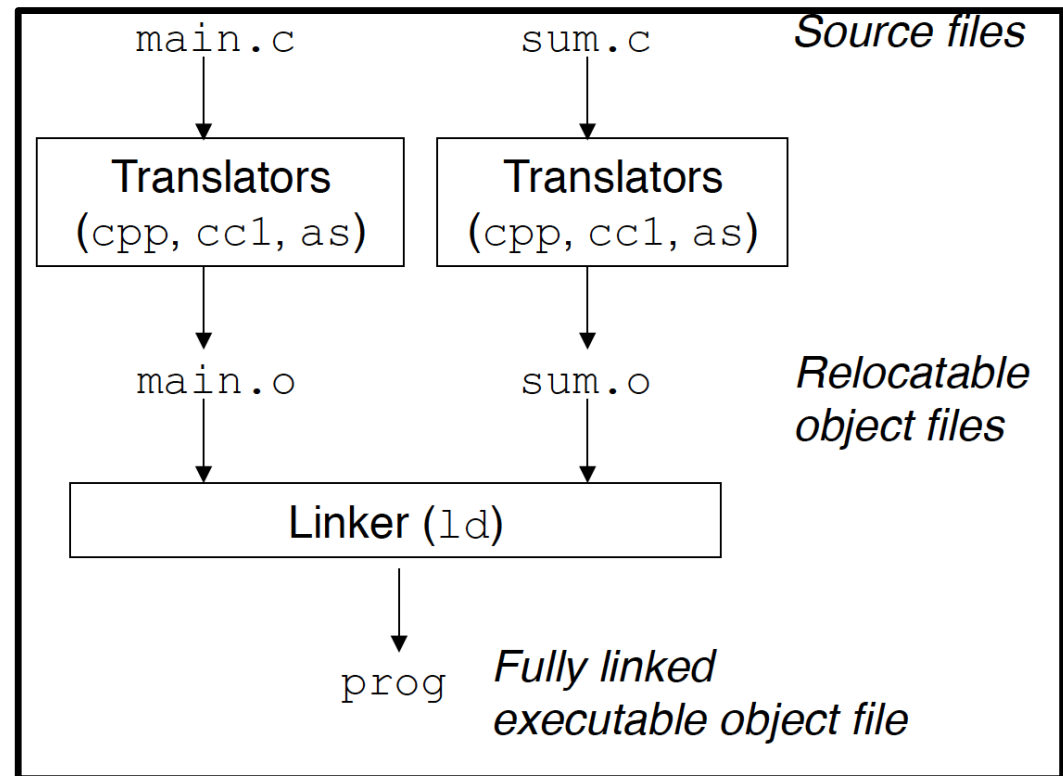


Translating/Starting Programs

- GNU compilation system & Linux specific
 - `gcc -o prog main.c sum.c`
- Invokes the GCC driver
 - GCC performs # steps

Translating/Starting Programs

- GNU compilation system & Linux specific
 - `gcc -o prog main.c sum.c`
- Invokes the GCC driver
 - GCC performs # steps



Example

sum.h:

```
int sum(int a, int b);
```

sum.c:

```
int sum(int a, int b) {  
    return a + b;  
}
```

main.c:

```
#include <stdio.h>  
#include "sum.h"  
int main() {  
    int c = sum(a,b);  
    printf("c = %i \n", c);  
    return 0;  
}
```

cpp, cc1, and as

- **cpp** is the C preprocessor: handles header file inclusion and macro expansion and generates an intermediate (.i) file
 - `#include <stdio.h>` copies the contents of `stdio.h`
 - `#define LEN 100` replaces `LEN` with `100` everywhere in the code
- **cc1** is the C compiler that generates an assembly (.s) file from the intermediate format (internal command)
- **as** is the assembler, which translates the assembly file into a binary relocatable (.o) object file

Relocatable Object File

- Contains binary code and data that can be combined with other relocatable object files at compile time to create an executable object file
- Symbols are not resolved
 - Variables with the same name declared in multiple sources
 - **Symbol resolution:** removes ambiguity by creating a single *linked executable file*
- Functions and variables are not bound to any specific address
- Addresses are still symbols (i.e., starting from 0 in each object file), and not properly assigned to an address in the memory map

Useful gcc commands

- Can break down the steps to generate the final executable file
 - `gcc -c main.c` (generates `main.o`)
 - `gcc -c sum.c` (generates `sum.o`)
 - `gcc -o prog main.o sum.o`
- Can generate the assembly (.s) file to view the assembly
 - `gcc -S main.c` (generates `main.s`)

Program vs. Process

- **Program**
 - A compiled executable binary lies dormant on a storage medium such as disk
- A **process** is a running program
 - The loader loads the executable file (image) in memory and fills up the memory map with code and data
 - Process has an execution environment (stack and heap)

Scope

- C identifiers (variables, functions, macros) have scope that delimits the regions where they can be accessed
- Four type of scope
 - File
 - Block
 - Function prototype
 - Function
- Scope is determined by where a variable is declared in the program

Example: Scope

```
int j;    // file scope of j begins

void f(int i) { // block scope of i begins
    int j = 1;  // block scope of j begins; hides file-scope j
    i++;        // i refers to function parameter
    for (int i = 0; i < 2; i++) { // block scope of loop-local i begins
        int j = 2; // block scope of inner j begins, hides outer j
        printf("%d\n", j); // inner j is in scope, prints ?
    }
    printf("%d\n", j); // outer j is in scope, prints ?
} // the block scope of i and j ends

void g(int j); // j has function prototype scope; hides file-scope j
```

Storage Duration

- Variables (objects) have a storage duration that determines their lifetime
 - There is a physical memory location reserved for the variable
- **Example:** Variables declared inside a code block or function definition are alive during the execution of the block or function
- **Example:** The loop index (declared inside the `for` statement) dies when loop terminates
- Four durations available in C
 - `automatic`
 - `static`
 - `allocated`
 - `thread` (related to concurrency, won't cover)

Storage Duration: automatic

- Consider the function definition below
 - Variable `life_and_death` has automatic storage duration
 - Implicit, no need to specify explicitly
 - It is alive only during the execution of function in which it is declared

```
void function(int i) {  
    int life_and_death = 1;  
    printf("%i\n", life_and_death);  
}
```

```
int main() {  
    int life_and_death = 2;  
    printf("%i\n", life_and_death);  
    return 0;  
}
```

Storage Duration: automatic

- More generally, variables declared in a code block demarcated by `{ . . . }` have automatic storage duration

```
int main() {  
    for (int i = 0; i < 100; i++)  i has automatic storage duration  
        printf("%i\n", i);  
    return 0;  
}
```

- The compiler (or assembly programmer) can reclaim the register in which **i** is stored after the `for` loop terminates

Storage Duration: `static`

- Objects declared in file scope have `static` storage duration
- Array `big_array` has `static` storage duration
 - Lifetime → Entire execution of program
- Static storage duration is implicit for variables declared in file scope

```
int big_array[1L<<24]; //static
int huge_array[1L<<31]; // static

int main() {
    printf("%d\n", life_and_death);
    return 0;
}
```

Storage Duration: `static`

- We can explicitly use `static` storage duration for variables inside functions
- These variables persist after the function exits

```
int big_array[1L<<24];
int huge_array[1L<<31];

void increment() {
    // ctr is only initialized once
    static unsigned int ctr = 0;
    ctr++;
    printf("ctr = %d\n", ctr);
}

int main() {
    for (int i=0; i<5; i++)
        increment();
    return 0;
}
```


Storage Duration: `static`

- We can explicitly use `static` storage duration for variables inside functions
- These variables persist after the function exits
- Output
 - `ctr = 1`
 - `ctr = 2`
 - `ctr = 3`
 - `ctr = 4`
 - `ctr = 5`

```
int big_array[1L<<24];
int huge_array[1L<<31];

void increment() {
    // ctr is only initialized once
    static unsigned int ctr = 0;
    ctr++;
    printf("ctr = %d\n", ctr);
}

int main() {
    for (int i=0; i<5; i++)
        increment();
    return 0;
}
```

Storage Duration: `static`

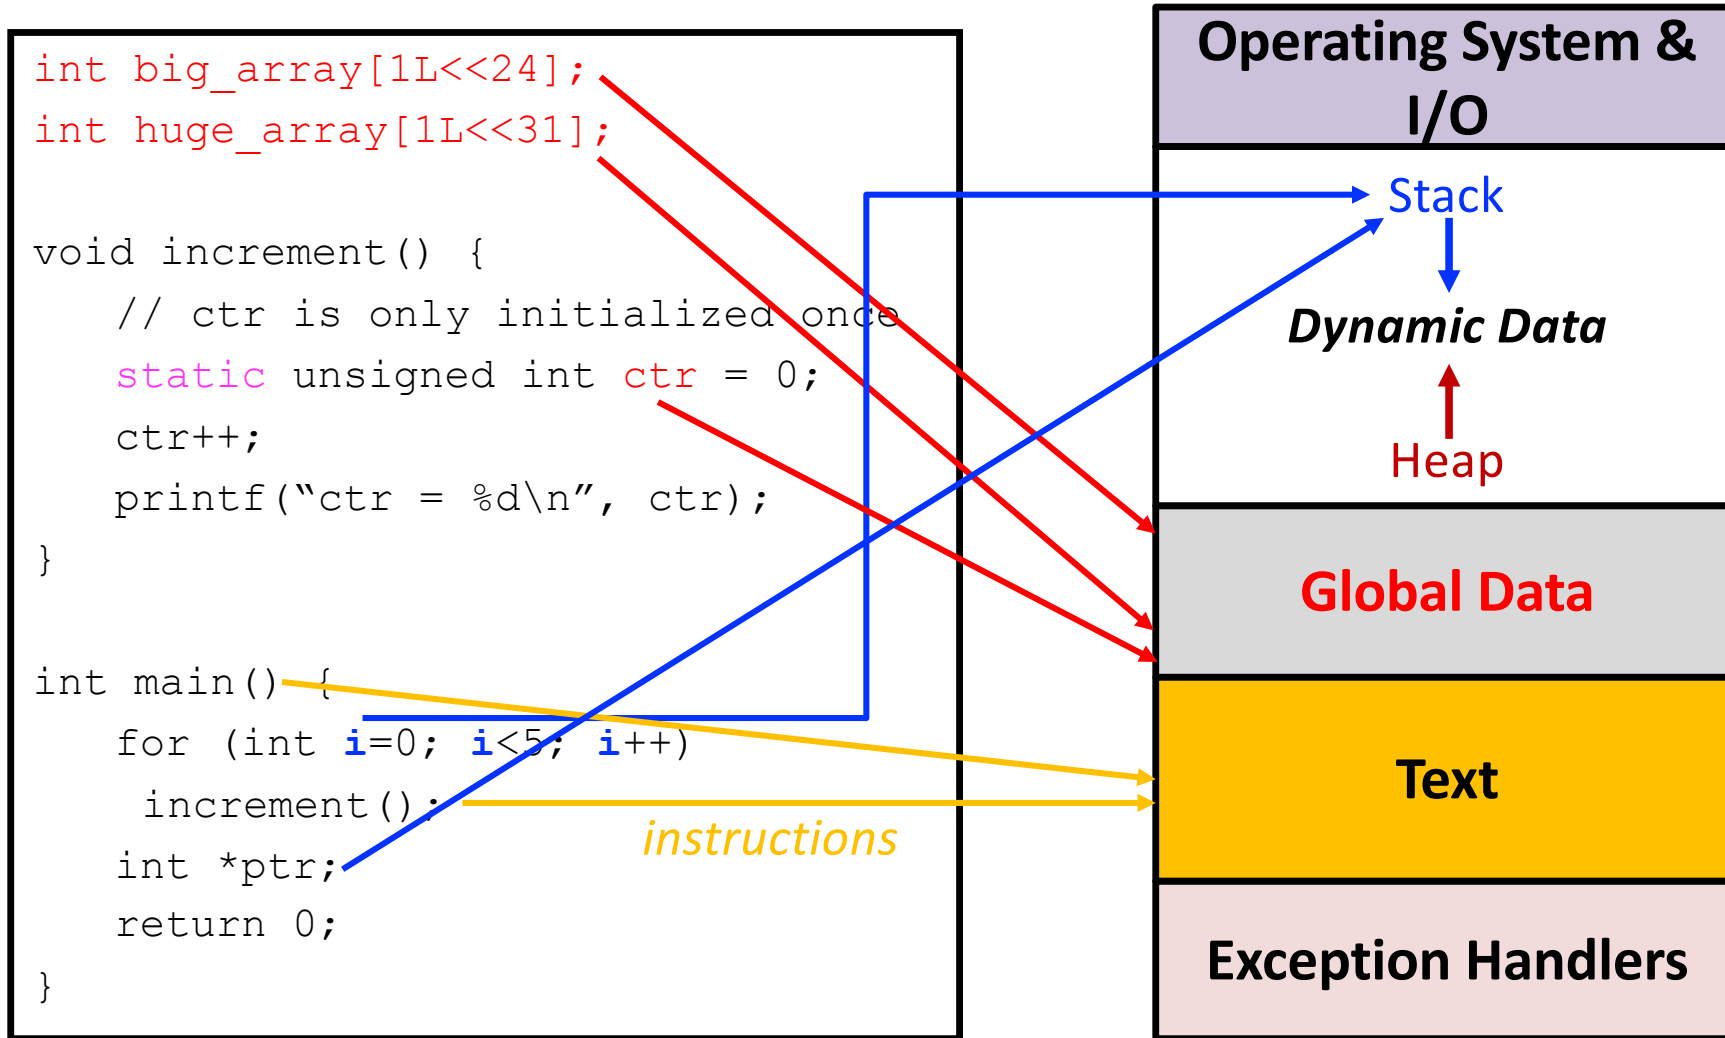
- We can declare `ctr` outside the function `increment()`
- Good software engineering practice to limit the scope of a variable whenever possible
- Static variables cannot be initialized to the symbol/name of another variable
- Initialized only to constants such as `0`, `1`, `"Hello"`

```
int big_array[1L<<24];
int huge_array[1L<<31];

void increment() {
    // ctr is only initialized once
    static unsigned int ctr = 0;
    ctr++;
    printf("ctr = %d\n", ctr);
}

int main() {
    for (int i=0; i<5; i++)
        increment();
    return 0;
}
```

Where is everything mapped?



Statically Allocated Memory

```
int big_array[1L<<24];  
int huge_array[1L<<31];
```

- The above arrays are statically allocated
 - Size must be known at compile time (*limitation of static*)
- We need to specify the size of the statically-allocated arrays

```
int N;  
int big_array[N]; ❌
```

```
int main() {  
    scanf("%i", &N); ✓  
    int array[N];  
}
```


→ *This array is allocated on the stack*

Limitations of Stack Allocation

- Memory is allocated and deallocated (freed) in a specific order
 - Last In First Out (LIFO)
- Memory is retained on the stack even if it is not needed
 - Memory is a precious resource
- No way for programmer to inform the compiler/OS to free unused memory
 - Deallocated only when function returns
 - And in a specific order

Example: Mem Waste

```
int caller_useless_func() {  
    int sum = 0;  
    int array[5] = {1, 2, 3, 4, 5};  
    for (int i = 0; i < 5; i++)  
        sum += array[i];  
    int x = useless_func(sum);  
    return x * sum;  
}  
  
int useless_func(int var) {  
    int multiplied = var * 5;  
    return multiplied;  
}
```



`array` is no longer needed after this statement, but it is still retained on the stack

Another Problem: Stack Allocation

- Sharing data across functions

```
int caller_useless_func() {  
    int sum = 0;  
    int array[5] = {1, 2, 3, 4, 5};  
    for (int i = 0; i < 5; i++)  
        sum += array[i];  
    int *y = useless_func(sum);  
    printf("Now printing .... \n");  
    printf("%d\n", *y * sum);  
}
```

```
int *useless_func(int var) {  
    int multiplied = var * 5;  
    return &multiplied;  
}
```

→ **printf's** stack frame
likely corrupts the
useless_func's
stack frame

→ **multiplied** is dead
after the function returns


Problem: Stack Allocation

- We can solve the problem by allocating **multiplied** as a global variable
- What if we want to share an array b/w functions?
 - Statically allocate the array
 - Ok, but what if we only find out the exact length of the array at run-time? What if we want to resize the array?
 - E.g., number of records in a database not known in advance
 - Student numbers grow dynamically during the semester
- We use the heap for such “dynamically allocated” memory

Pointers: Revision

- A pointer is a variable that contains the address of another variable


```
int A = 19;  
int B = 10;  
int C = 8;  
int D = 17;  
....  
int *P = &B;  
//unary operator & gives  
    the address of a  
    variable
```



Pointers: Revision

- Can use the pointer to access the value stored in a memory location


```
int A = 19;
int B = 10;
int C = 8;
int D = 17;
....
int *P = &B;
*P = 1;
//dereferencing or
// indirection operator
//that accesses the value
// stored at address in P
```



Pointers: Example

- A pointer is 4-bytes on a 32-bit system and 8-bytes on a 64-bits system & it can be stored on the stack or data segment like ordinary variables

```
int A = 19;
int B = 1;
int C = 8;
int D = 17;
....
int *P = &B;
char *Q = &B;
// Both P and Q contain
// 00000004
printf("%i\n", *P); ??
printf("%i\n", *Q); ??
```



Pointers: Bottomline

- A pointer points to a memory location
- Its content is a memory address
- It wears “datatype glasses”
 - Wherever it points, it sees through these glasses
- The variable stored at some memory address can be interpreted (via the dereferencing operator `*`) as character or integer or float, depending on the type of the pointer




Heap

- Heap is for dynamically allocated memory
 - Contrast with statically allocated memory
 - Large subdividable block of memory
 - Programmers explicitly allocate and deallocate memory on the heap
- Lifetime of heap variables/arrays extend from allocation until deallocation
- Memory managers are libraries that manage heap for the programmer (we will refine this view)

Heap Organization


- Programs dynamically allocate objects (arrays, structures) of different types and sizes on the heap over time



- Heap is now full!**

Heap Organization


- C programmers need to track lifetime of heap variables
- Suppose we do not need anymore



- **Heap is now full!**


Heap Organization

- Let's free some space on the heap




Heap Organization

- Let's free some space on the heap




Heap Organization

- Now we can allocate new objects



Heap Organization

- Now we can allocate new objects



Heap Management

- We can deallocate (free) objects in any order (unlike the stack)
- Programmers need to deallocate objects that are no longer needed
 - Otherwise, heap will remain full even if we can have some free space
- Not returning unused heap back to the memory manager is called a *memory leak*

malloc and free

- C library provides **malloc** (short for memory allocate) and **free** to allocate and deallocate heap memory, respectively

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array[i] = i * i;
    int sum = array[0] + array[9];
    free(array);
    printf("%i\n", sum);
    return;
}
```

malloc and free

- **malloc**

- Declaration in C library: `void *malloc(size_t size)`
- Takes input as size (# bytes)
- Returns a void pointer that can be casted to any pointer type

- **free**

- Declaration in C library: `void free(void *ptr)`
- Memory manager knows how many bytes to free, all it needs is the starting address

Quiz: Bug?

- Is there a memory-related bug in the following program?

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array1 = malloc(10 * sizeof(int));
    int *array2 = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++) {
        array1[i] = i * i;
        array2[i] = i + i;
    }
    int sum = array1[0] + array2[9];
    free(array1);
    free(array2);
    printf("%i\n", sum + array1[8]);
    return;
}
```

Quiz: Bug?

- Is there a memory-related bug in the following program?

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array1 = malloc(10 * sizeof(int));
    int *array2 = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++) {
        array1[i] = i * i;
        array2[i] = i + i;
    }
    int sum = array1[0] + array2[9];
    free(array1);
    free(array2);
    printf("%i\n", sum + array1[8]);
    return;
}
```

use after free

Quiz: Bug?

- Is there a memory-related bug in the following program?

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array1 = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    array1 = malloc(5 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    free(array1);
    printf("Done ....\n");
    return;
}
```

Quiz: Bug?

- Is there a memory-related bug in the following program?

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array1 = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    array1 = malloc(5 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    free(array1);
    printf("Done ....\n");
    return;
}
```

memory leak
Original 10-
element array
still on heap
(address is
gone, not saved)

out of bounds