

ENGN2219/COMP6719

Computer Systems & Organization

Convener: Shoaib Akram

shoaib.akram@anu.edu.au



Australian
National
University

Pointers: Example

- Guess the output of the `printf ()` statements

```
int A = 19;
int B = 1;
int C = 8;
int D = 17;
....
int *P = &B;
char *Q = &B;
// Both P and Q contain
    00000004
printf("%i\n", *P); ??
printf("%i\n", *Q); ??
```


Address	Data	Variable
•	•	•
•	•	•
•	•	•
00000010	00000004	P
0000000C	17	D
00000008	8	C
00000004	1	B
00000000	19	A

4 Bytes

Answer

- `printf("%i\n", *P);` Output is always 1
- `printf("%i\n", *Q);` Big Endian: 0, Little Endian: 1

```
int A = 19;
int B = 1;
int C = 8;
int D = 17;
....
int *P = &B;
char *Q = &B;
// Both P and Q contain
    00000004
printf("%i\n", *P); ??
printf("%i\n", *Q); ??
```



malloc and free

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
void useless_func() {
```

```
    int *array1 = malloc(10 * sizeof(int));
```

```
    for (int i = 0; i < 10; i++)
```

```
        array1[i] = i * i;
```

```
    printf("%i\n", sizeof(array1));
```

```
    free(array1);
```

```
    printf("Done ..... \n");
```

```
    return;
```

```
}
```

include for
malloc()

use sizeof(int)
operator rather
than guessing
how big is an
int on a machine

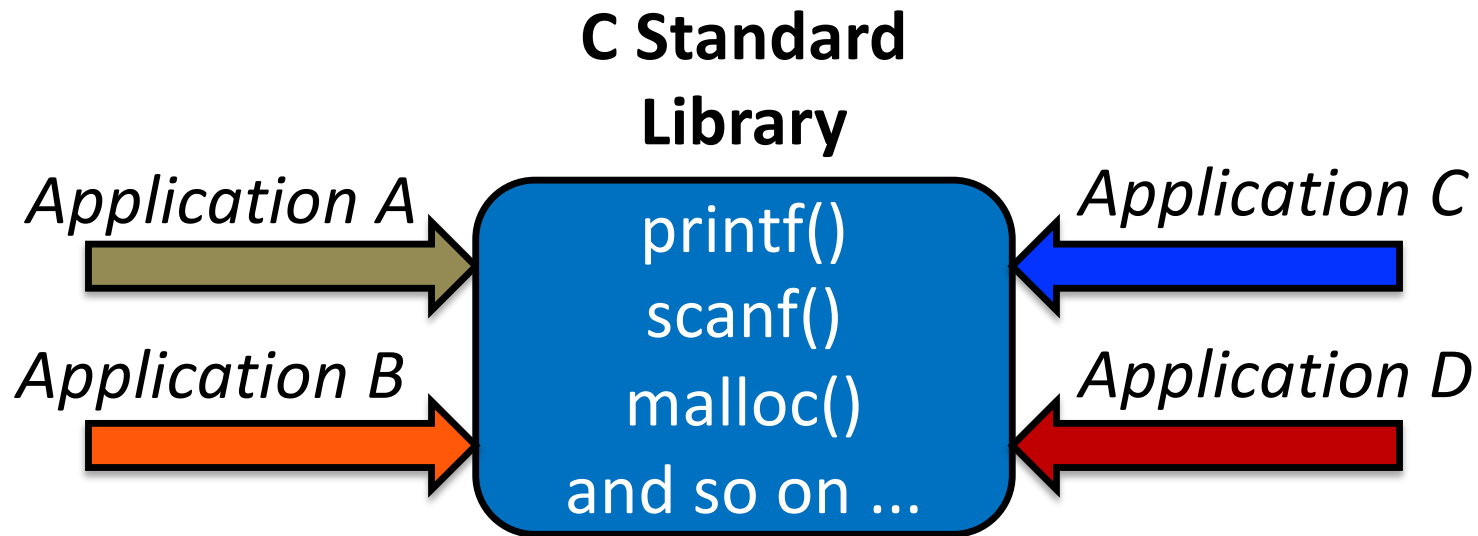
16 or 32 or 64
depending on
the machine
architecture

API




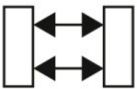
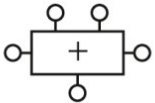
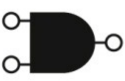
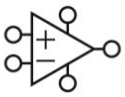


- *Application Programming Interface (API)*
 - Defines the interfaces by which one software program communicates with another at the source code level
 - **Abstraction:** API provides a standard set of interfaces that many different users (writing programs) can invoke
- *API defines the interface only*
 - The user of the API can ignore the implementation
 - Many implementations of an API can exist

API: Example

- The C standard library (`libc`) defines a family of basic and essential functions: memory mgmt. and string manipulation



- `libc` hides a lot of operating system details by interacting with the OS on behalf of the programmer

Application Software		Programs
Operating Systems		Device Drivers
Architecture		Instructions Registers
Micro-architecture		Datapaths Controllers
Logic		Adders Memories
Digital Circuits		AND Gates NOT Gates
Analog Circuits		Amplifiers Filters
Devices		Transistors Diodes
Physics		Electrons

ABI

- *Application Binary Interface*
 - Defines the binary interface b/w two or more pieces of software on a particular architecture
- *ABI ensure binary compatibility*
 - Calling convention, byte ordering, register use, linking, binary object format
 - Enforced by the toolchain (not a programmer's worry)
- ABI is a function of both the architecture (x86, ARM) and operating system (Linux, Windows)

Common Memory Errors

- *Forgetting to allocate memory*

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    char *src = "Hello!";
    char *dst;
    strcpy(dst, src);
}
```

- *Causes a segmentation fault*
 - *System is angry at you, because you did something wrong with memory*

Description

The C library function `char *strcpy(char *dest, const char *src)` copies the string pointed to, by **src** to **dest**.

Declaration

Following is the declaration for strcpy() function.

```
char *strcpy(char *dest, const char *src)
```

Parameters

- **dest** – This is the pointer to the destination array where the content is to be copied.
- **src** – This is the string to be copied.

Common Memory Errors

- *Forgetting to allocate memory*

```
//Fixed program

#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    char *src = "Hello!";
    char *dst = (char*) malloc(strlen(src) + 1);
    strcpy(dst, src);
}
```

Common Memory Errors

- *Not allocating sufficient memory
(a.k.a. buffer overflow)*

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    char *src = "Hello!";
    char *dst = (char*) malloc(strlen(src));
    strcpy(dst, src);
}
```

Common Memory Errors

- *Not allocating sufficient memory (a.k.a. buffer overflow)*
 - There maybe an unused variable at the end of allocated memory (no harm)
 - Memory manager might have allocated a little extra space
 - Fault and crash and security vulnerability is likely

■ *Just because a program runs, does not mean it is correct*

Common Memory Errors

- *Forgetting to initialize allocated memory*
 - Zero initialization is important for avoiding hard to debug problems

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *n = (int*) malloc(sizeof(int));
    ...
    cond = *n;
    ...
    if (cond == 0) {
        ...;
    }
}
```

Common Memory Errors

- *Forgetting to free memory (memory leak)*
 - Huge problem in long-running programs
- When the process exits, the operating system automatically reclaims all allocated memory by the process
 - Still a bad habit to not call `free ()` in short-running programs

Example

- Memory leak example (from last lecture)

```
#include <stdlib.h>
#include <stdio.h>

void useless_func() {
    int *array1 = malloc(10 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    array1 = malloc(5 * sizeof(int));
    for (int i = 0; i < 10; i++)
        array1[i] = i * i;
    free(array1);
    printf("Done ....\n");
    return;
}
```

memory leak
Original 10-
element array
still on heap
(address is
gone, not saved)

out of bounds

Common Memory Errors

- *Freeing memory before the program is done with it*
 - Known as a dangling pointer
 - The subsequent use can crash the program

```
// example 1
void useless_func() {
    int n = 100;
    int *ptr = (int*) malloc(sizeof(int));
    if (n == 100) {
        int v = 10;
        ptr = &v;
    }
    // ptr is now a dangling pointer
}
```


Common Memory Errors

- *Freeing memory before the program is done with it*
 - Known as a dangling pointer
 - The subsequent use can crash the program

```
// example 2
void useless_func() {
    int *ptr = (int*) malloc(sizeof(int));
    ...
    ...
    free(ptr);
    ...
    ...
    *ptr = 2; // ptr is now a dangling pointer
}
```

Common Memory Errors

- *Freeing memory before the program is done with it*
 - Known as a dangling pointer
 - The subsequent use can crash the program

```
//example 2 fix, you will now get a runtime error
void useless_func() {
    int *ptr = (int*) malloc(sizeof(int));
    ...
    ...
    free(ptr);
    ptr = NULL; // good practice
    ...
    ...
    *ptr = 2; // ptr is no longer a dangling pointer
}
```

Common Memory Errors

- *Freeing memory before the program is done with it*
 - Known as a dangling pointer
 - The subsequent use can crash the program

```
// example 3
int *useless_func() {
    int x = 100;
    return &x;
}

// when the function returns, it's stack is
// deallocated, and we must not use the address
// returned by useless_func()
```

Common Memory Errors

- *Freeing memory multiple times*
 - double free (typical name)

```
void useless_func() {  
    int *x = (int *) malloc(100 * sizeof(int));  
    free(x);  
    ...  
    free(x); // confuses the memory manager  
}
```

Common Memory Errors

- *Incorrectly calling `free()`*
 - Avoid these so-called **invalid frees**

```
void useless_func() {  
    int *x = (int *) malloc(100 * sizeof(int));  
    ...  
    x = x + 4;  
    ...  
    free(x);  
}
```

Multi-Dimensional Arrays

- *Statically allocated arrays and stack-allocated 2-dimensional arrays are simple*

```
#define R 5
#define C 4
int matrix[R][C];

for (int i = 0; i < R; i++) {
    for (int j = 0; j < C; j++) {
        matrix[i][j] = i + j;
    }
}
```

[0][0]	[0][1]	[0][2]	[0][3]
[1][0]	[1][1]	[1][2]	[1][3]
[2][0]	[2][1]	[2][2]	[2][3]
[3][0]	[3][1]	[3][2]	[3][3]
[4][0]	[4][1]	[4][2]	[4][3]

Multi-Dimensional Arrays

- *Dynamically allocated 2-dimensional arrays*

```
#define R 5
#define C 4

int **matrix;

void useless_func() {
    matrix = (int **) malloc(R * sizeof(int*));
    for (int i = 0; i < R; i++) {
        matrix[i] = malloc(C * sizeof(int));
    }
}
```

Multi-Dimensional Arrays

- *You can have jagged arrays, where each row has a different # columns*

```
#define R 3

int **matrix;

void useless_func() {
    matrix = (int **) malloc(R * sizeof(int*));
    matrix[0] = malloc(2 * sizeof(int));
    matrix[1] = malloc(5 * sizeof(int));
    matrix[2] = malloc(3 * sizeof(int));
}
```


Structs

- Lab 10 handout introduces structs in C
- How can we create a dynamically allocated array of structs?

```
#define TOTAL 100


struct student {
    char name[16];
    int id;
};

typedef struct student student_t;

student_t *students;

void useless_func() {
    students = (student_t *) malloc(TOTAL * sizeof(student_t));
    students[0].name = "Shane";
    students[0].id = 10;
}
```

can access
members with .
operator



Structs and \rightarrow Operator

- If we have a pointer to a struct, we can use the arrow operator (\rightarrow) to access the members of a struct

```
struct student {  
    char name[16];  
    int id;  
};  
  
typedef struct student student_t;  
  
student_t *student;  
void useless_func() {  
    student = (student_t *) malloc(sizeof(student_t));  
    student->name = "Shane";  
    students->id = 10;  
}
```

Enumerations in C

- Enumeration (or enum) is a user defined data type in C
 - Used to assign names to integral constants (code readability)

```
// An example program to demonstrate working
// of enum in C
#include<stdio.h>

enum week{Mon, Tue, Wed, Thur, Fri, Sat,
Sun};

int main()
{
    enum week day;
    day = Wed;
    printf("%i",day);
    return 0;
}
```

Principle of Locality

- Programs tend to reference (access) data items (words) that:
 - Reside near other recently accessed items
 - Were recently accessed themselves
- Temporal Locality
 - A memory location that is referenced once is likely to be referenced again multiple times soon
- Spatial Locality
 - If a memory location is referenced once, then a nearby location is likely to be referenced soon
- Generally, programs with good locality are likely to run faster than programs with poor locality

Locality is Everywhere

- Computer designers speed up main memory accesses by keeping most recently accessed data items in small fast memories called cache memories
 - Main memory is slow but high capacity (DRAM technology uses capacitors)
 - Cache is fast but low capacity (SRAM uses cross-coupled inverters)
- Web browsers exploit temporal locality by caching recently referenced documents on disks
- One of the enduring ideas in computer systems

Example

- Does the program below exhibit good locality?

```
int sumarray(int a[n]) {  
    int i;  
    int sum = 0;  
    for (i = 0; i < n; i++)  
        sum = sum + a[i];  
    return sum;  
}
```

Address	0	4	8	12	16	20	24	28
Contents	a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]
Access Order	1	2	3	4	5	6	7	8

Example

- Does the program below exhibit good locality?

```
int sumvec(int a[n]) {  
    int i;  
    int sum = 0;  
    for (i = 0; i < n; i++)  
        sum = sum + a[i];  
    return sum;  
}
```

Variable	Spatial	Temporal
sum	Poor	Good
a	Good	Poor

Address	0	4	8	12	16	20	24	28
Contents	a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]
Access Order	1	2	3	4	5	6	7	8

Explanation


- `array a`
 - Each element is accessed only once
 - Neighboring element is accessed soon in the next loop iteration
- `sum`
 - `sum` is a scalar, so no spatial locality
 - `sum` has good temporal locality, as it is accessed in each iteration

Stride

- Visiting each element of an array sequentially is an example of a *stride-1* reference pattern
 - Stride-1 reference pattern is called sequential access pattern
- Visiting every k-th element of a contiguous array is a *stride-k* reference pattern
- As stride increases, the spatial locality decreases
- Sequential accesses are generally highly desirable

Row Major Order

- C arrays are laid out in memory row-wise
 - The entire row is stored contiguously (elements are next to each other, 0, 4, 80)
- a_{mn} : m is row, and n is column

Row 1


Address	0	4	8	12	16	20
Contents	a_{00}	a_{01}	a_{02}	a_{10}	a_{11}	a_{12}
Access Order	1	2	3	4	5	6

Example

- Sum the elements of the array in row-major order

```
int sumarrayrows(int a[M][N]) {  
    int i, j, sum = 0;  
  
    for (i = 0; i < M; i++)  
        for (j = 0; j < N; j++)  
            sum = sum + a[i][j];  
    return sum;  
}
```

good spatial
locality



Address	0	4	8	12	16	20
Contents	a_{00}	a_{01}	a_{02}	a_{10}	a_{11}	a_{12}
Access Order	1	2	3	4	5	6

Example

- What is the impact of interchanging i and j loops?

```
int sumarrayrows(int a[M][N]) {  
    int i, j, sum = 0;  
  
    for (j = 0; j < N; j++)  
        for (i = 0; i < M; i++)  
            sum = sum + a[i][j];  
    return sum;  
}
```

poor spatial
locality




Address	0	4	8	12	16	20
Contents	a_{00}	a_{01}	a_{02}	a_{10}	a_{11}	a_{12}
Access Order	1	3	5	2	4	6

Practice Problem

Permute the loops in the following function so that it scans the 3-dimensional array **a** with a stride-1 reference pattern.


```
int sumarray3d(int a[N][N][N]) {  
    int i, j, k, sum = 0;  
  
    for (i = 0; i < N; i++)  
        for (j = 0; j < N; j++)  
            for (k = 0; k < N; k++)  
                sum = sum + a[k][i][j];  
    return sum;  
}
```



Locality of Instructions

- Program instructions are stored in memory and must be fetched by CPU
 - Locality is also relevant to instruction accesses
- Instructions in the loop body have high locality
 - Good spatial locality because instructions next to each other are executed in sequential order
 - Good temporal locality because the loop body is executed multiple times

Assignment 1 Grade Summary



Something was submitted by the deadline

Issues I Observed ☹️

- Work in a group, but do not submit the group survey
- Dump the submission in the lecturer's inbox
 - Which gets ~100 emails on a slow-paced day
- Do not fork the assignment
- Submit only the top-level circuit and nothing else
- Do not submit report
- Do not respect the deadline, don't ask for extension
- Do not submit Statement of Originality

Good News

- Many reports were a pleasure to read
- Highest: 99.5
- Constantly surprised with the ambitious extensions
 - Multi-Cycle
 - Pipelining
 - New instructions for condensing the code
- On average, neatness and clarity was clearly seen in submissions that were pushed on time